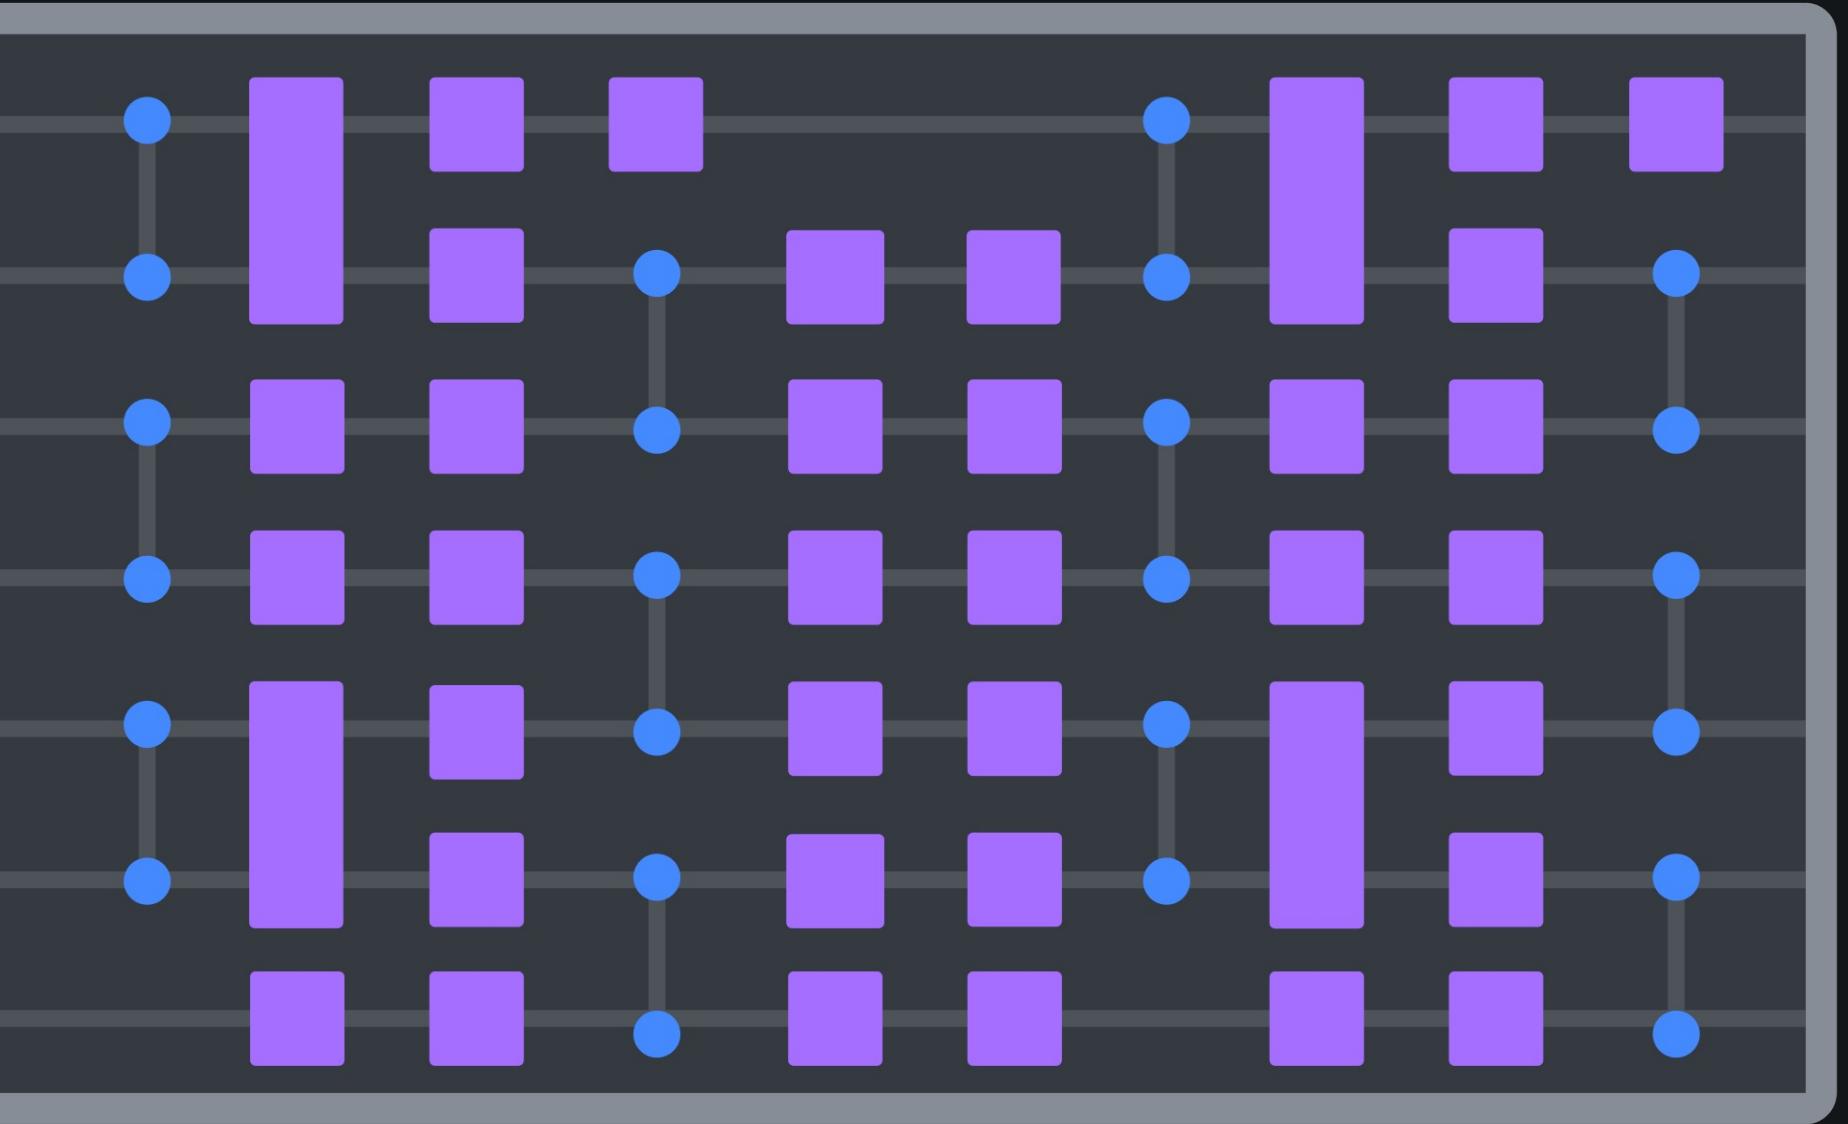


Optimizing quantum circuits with Qiskit

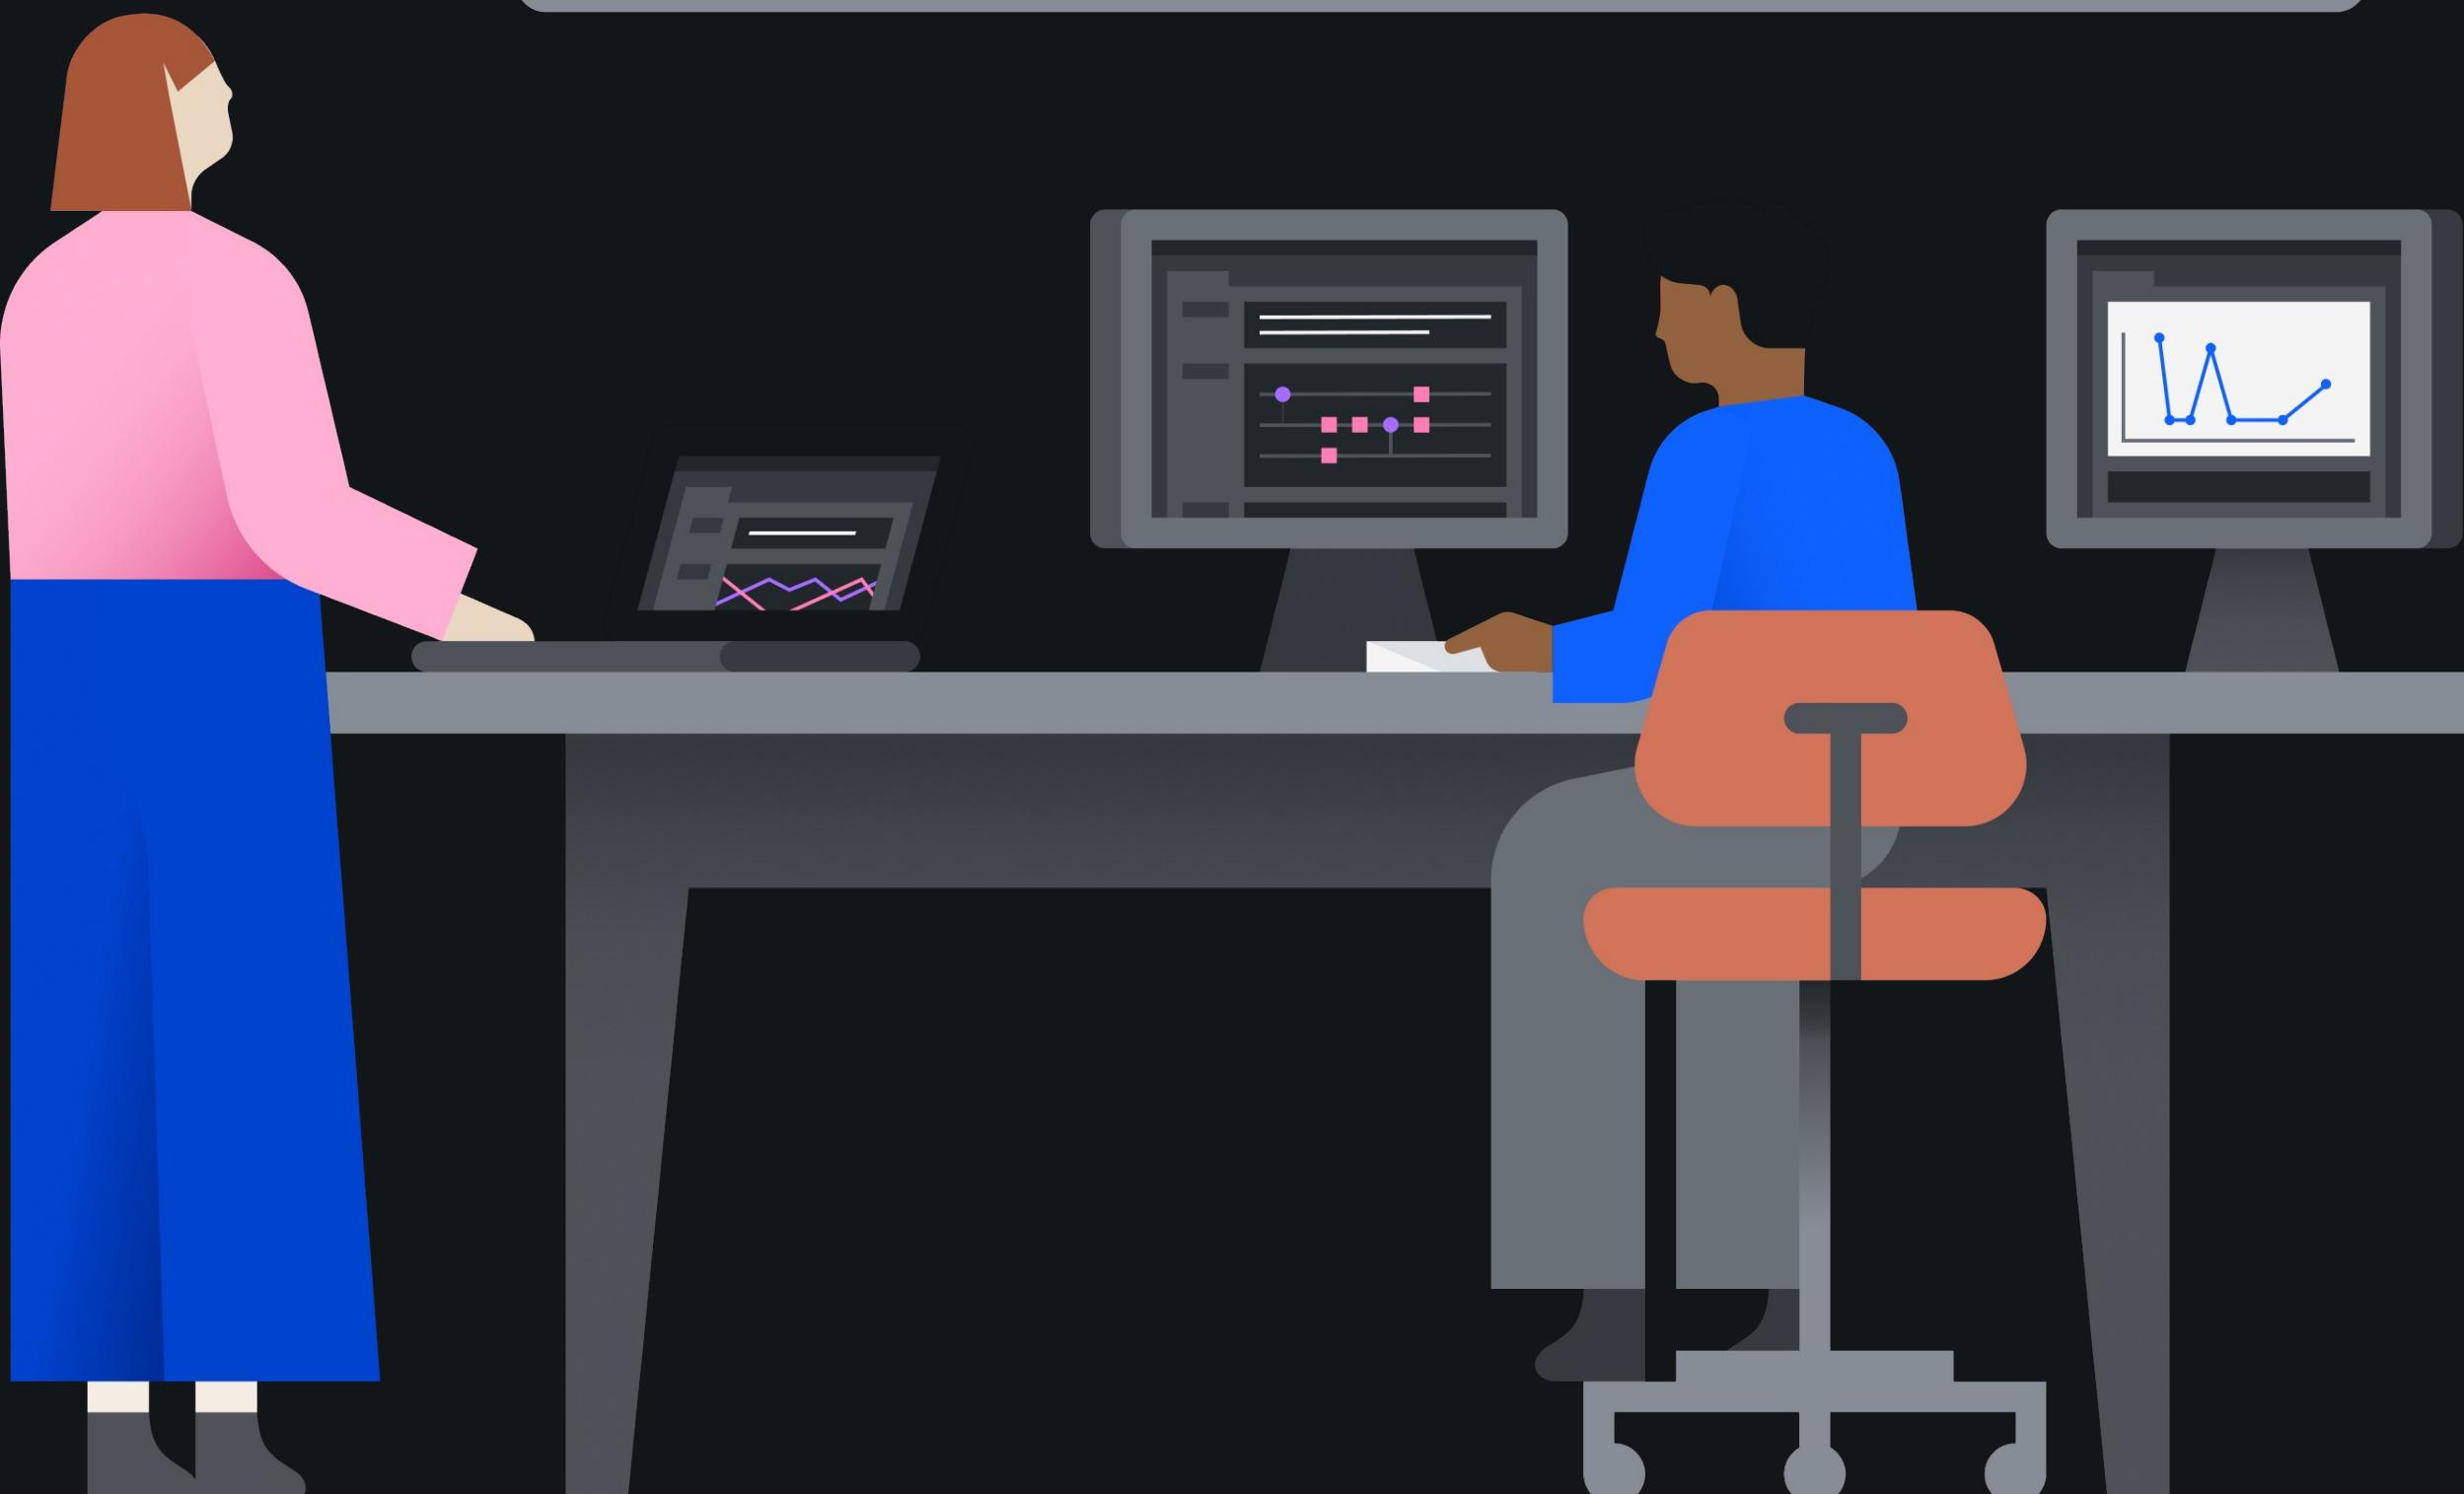


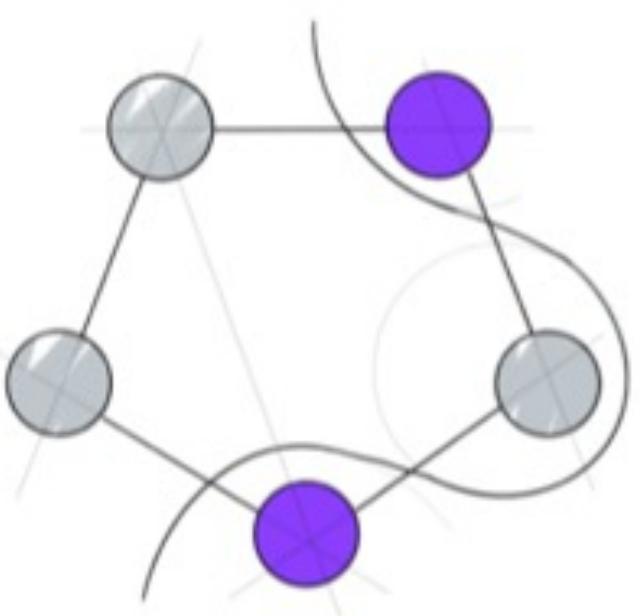
Luciano Bello
Qiskit Developer
IBM Quantum

 @microluciano

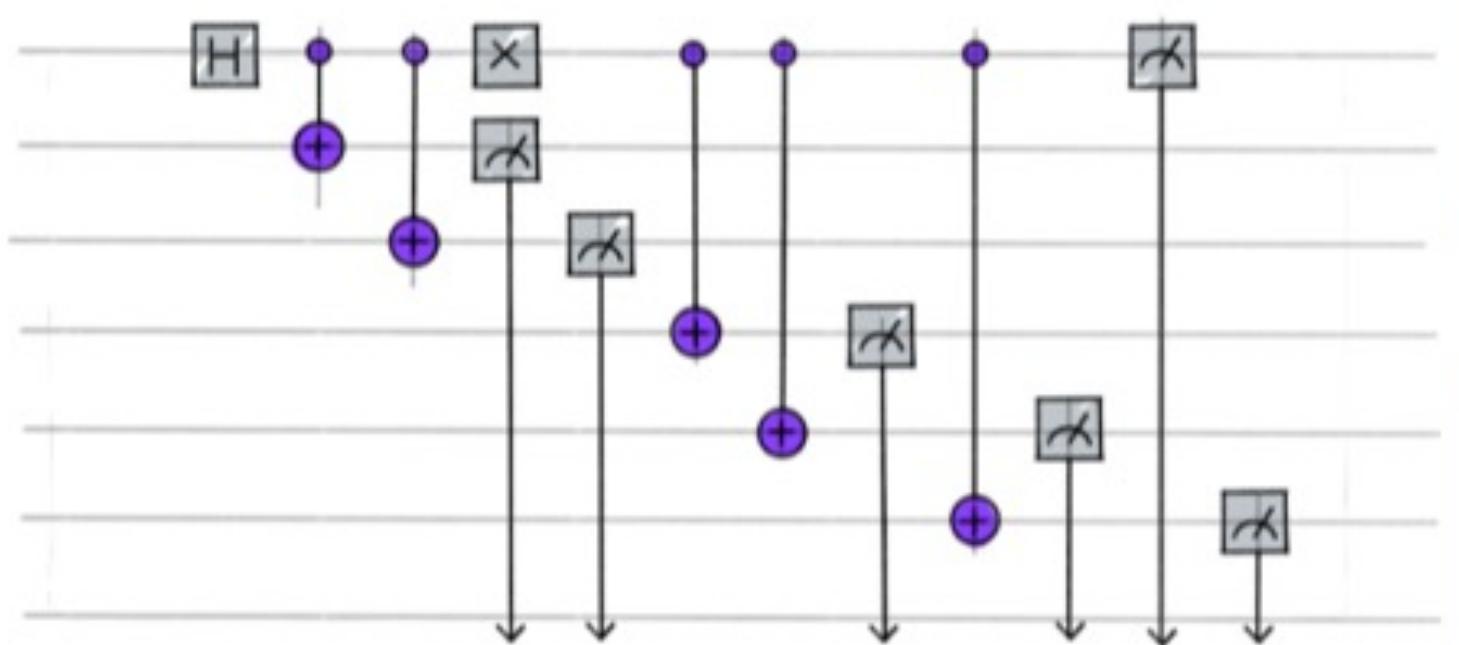
 1ucian0

 0x1ucian0

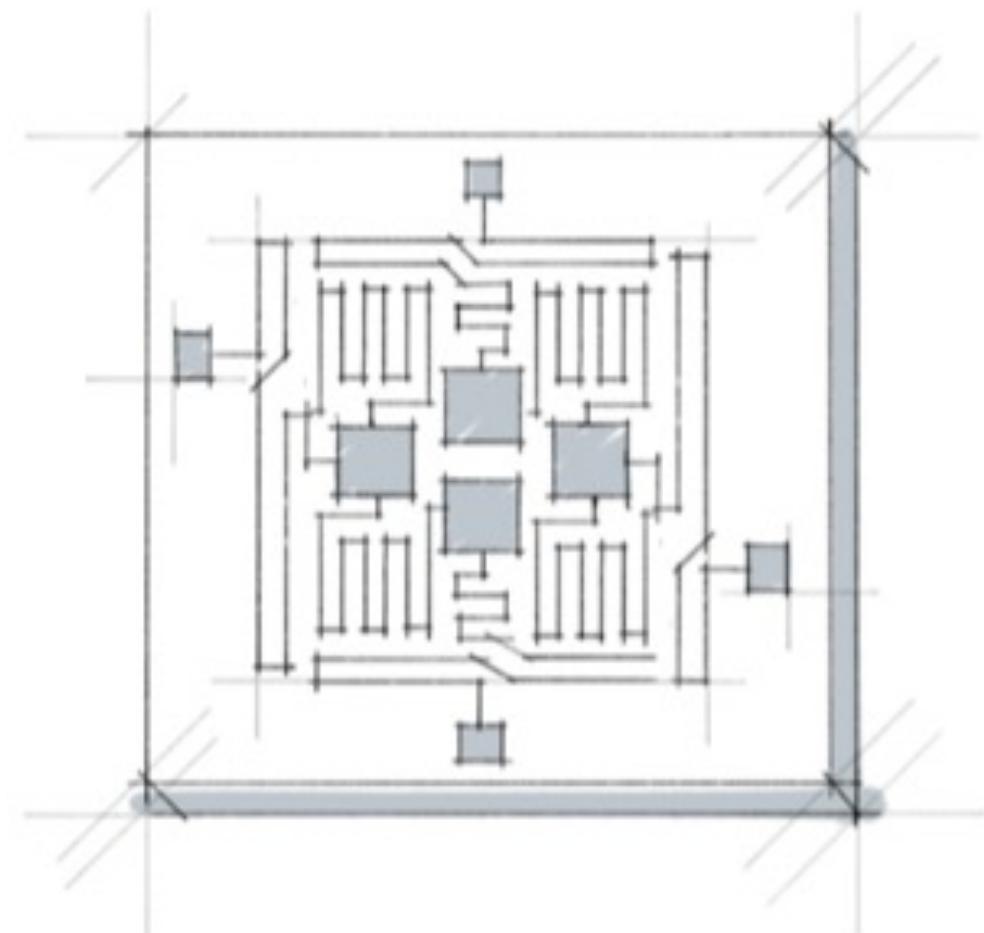




Problem



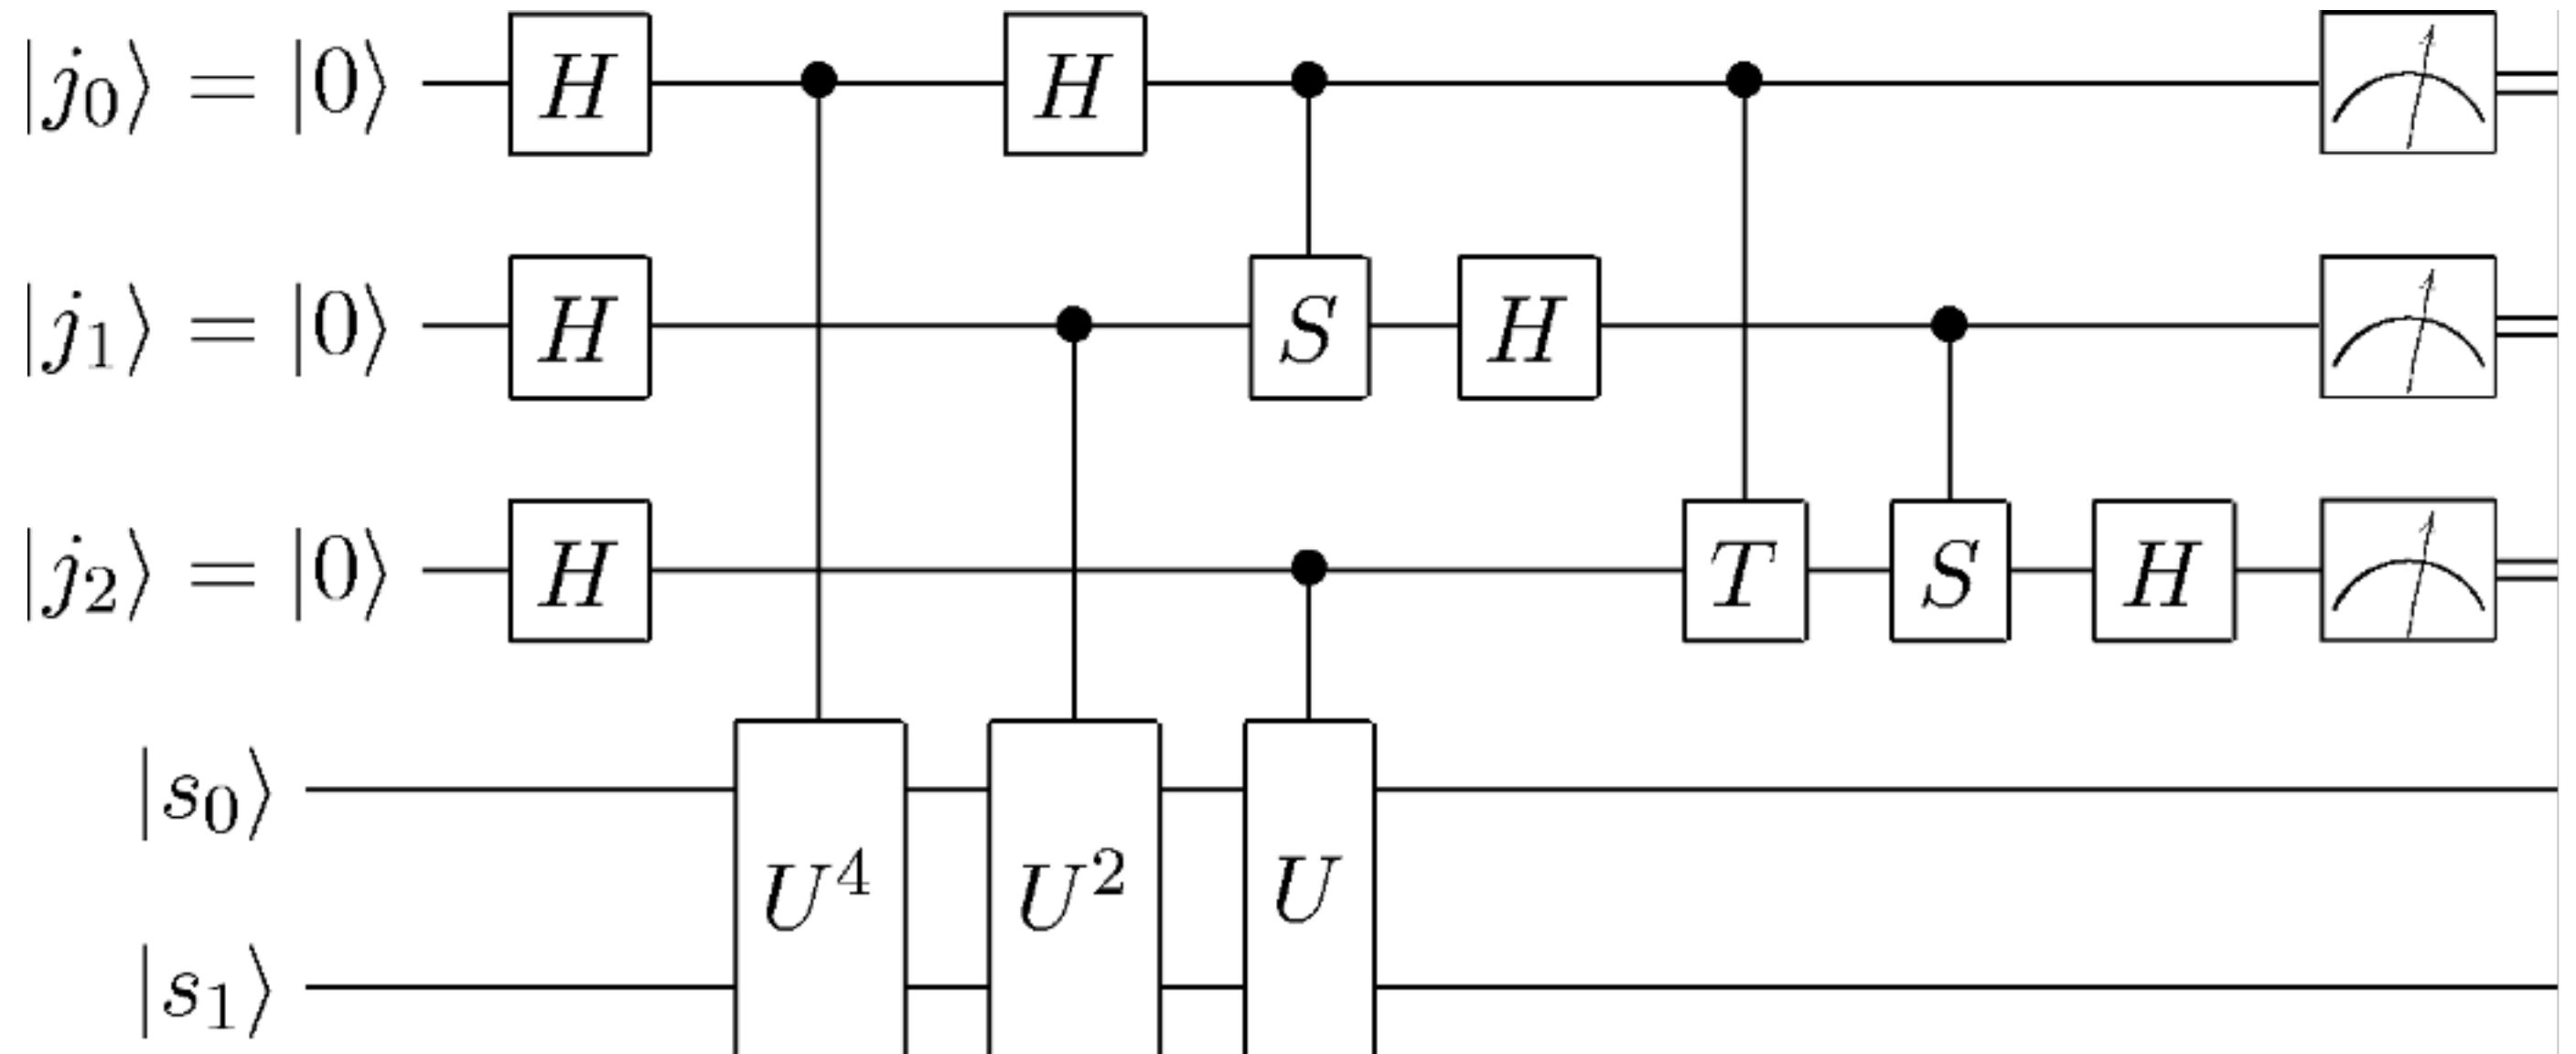
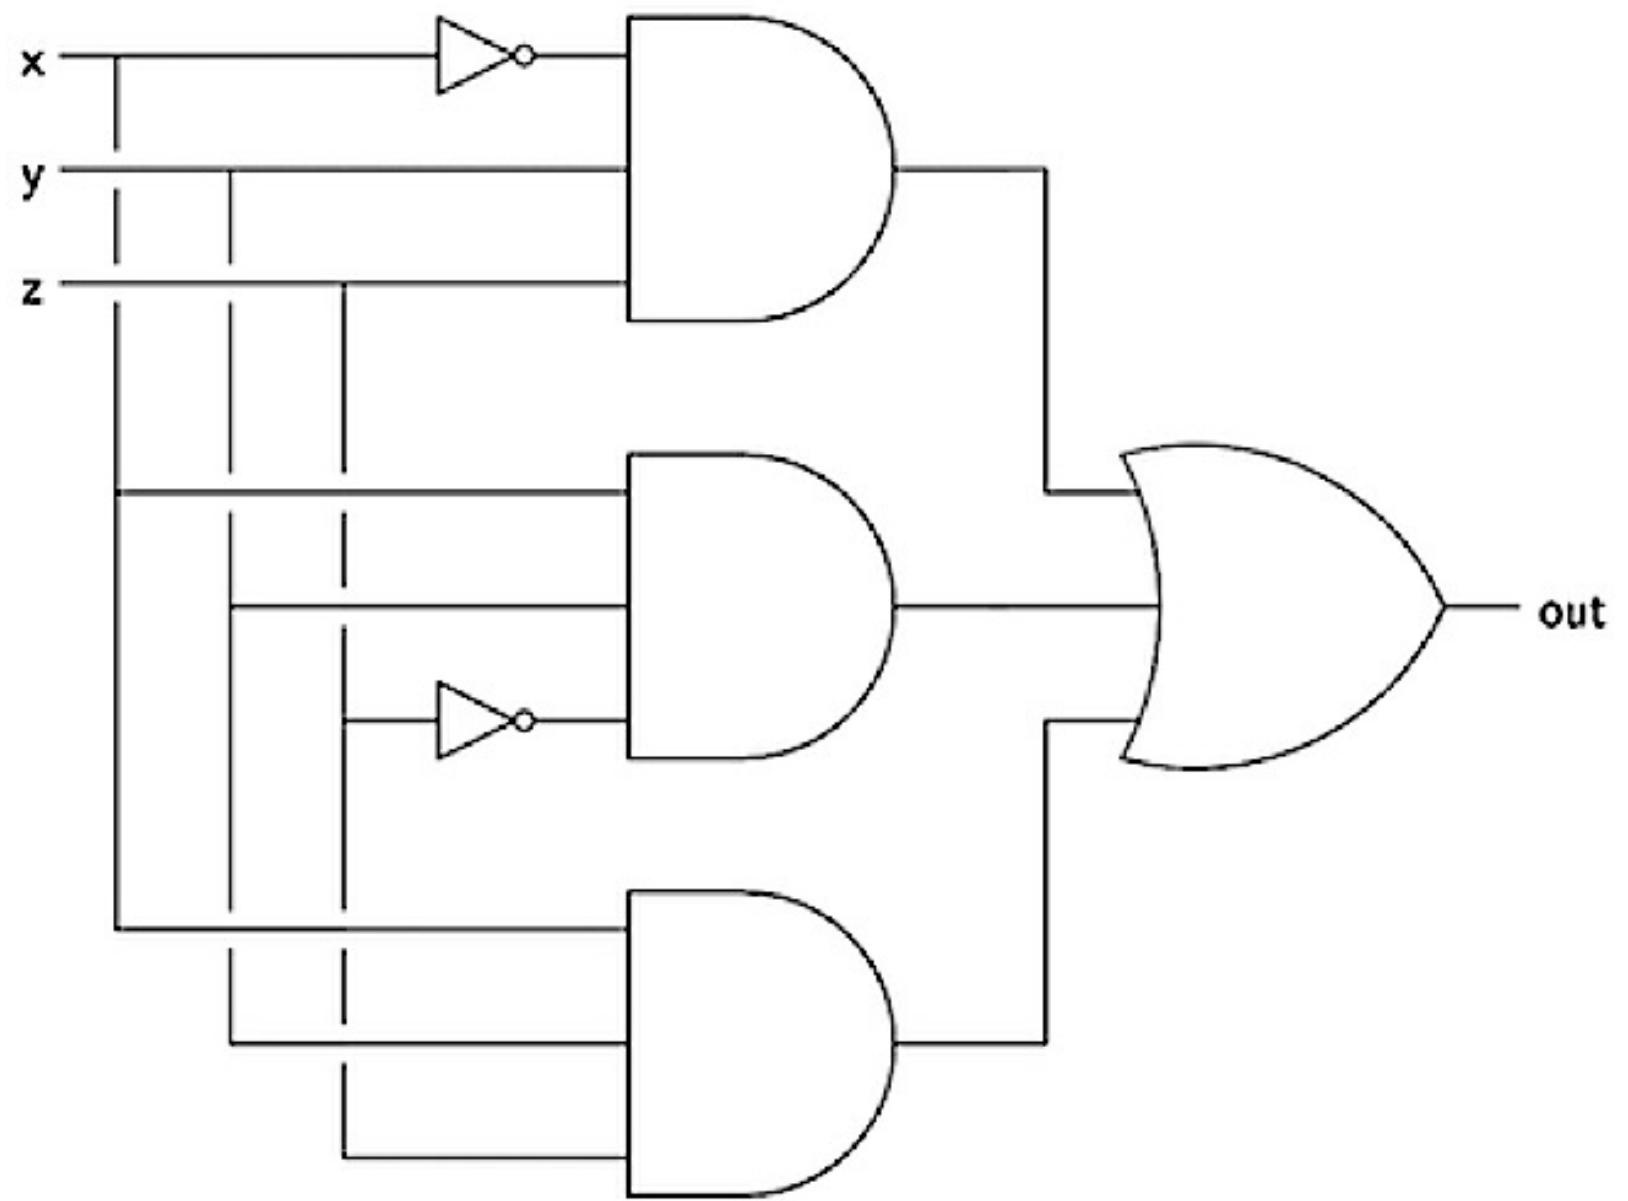
Quantum Algorithm

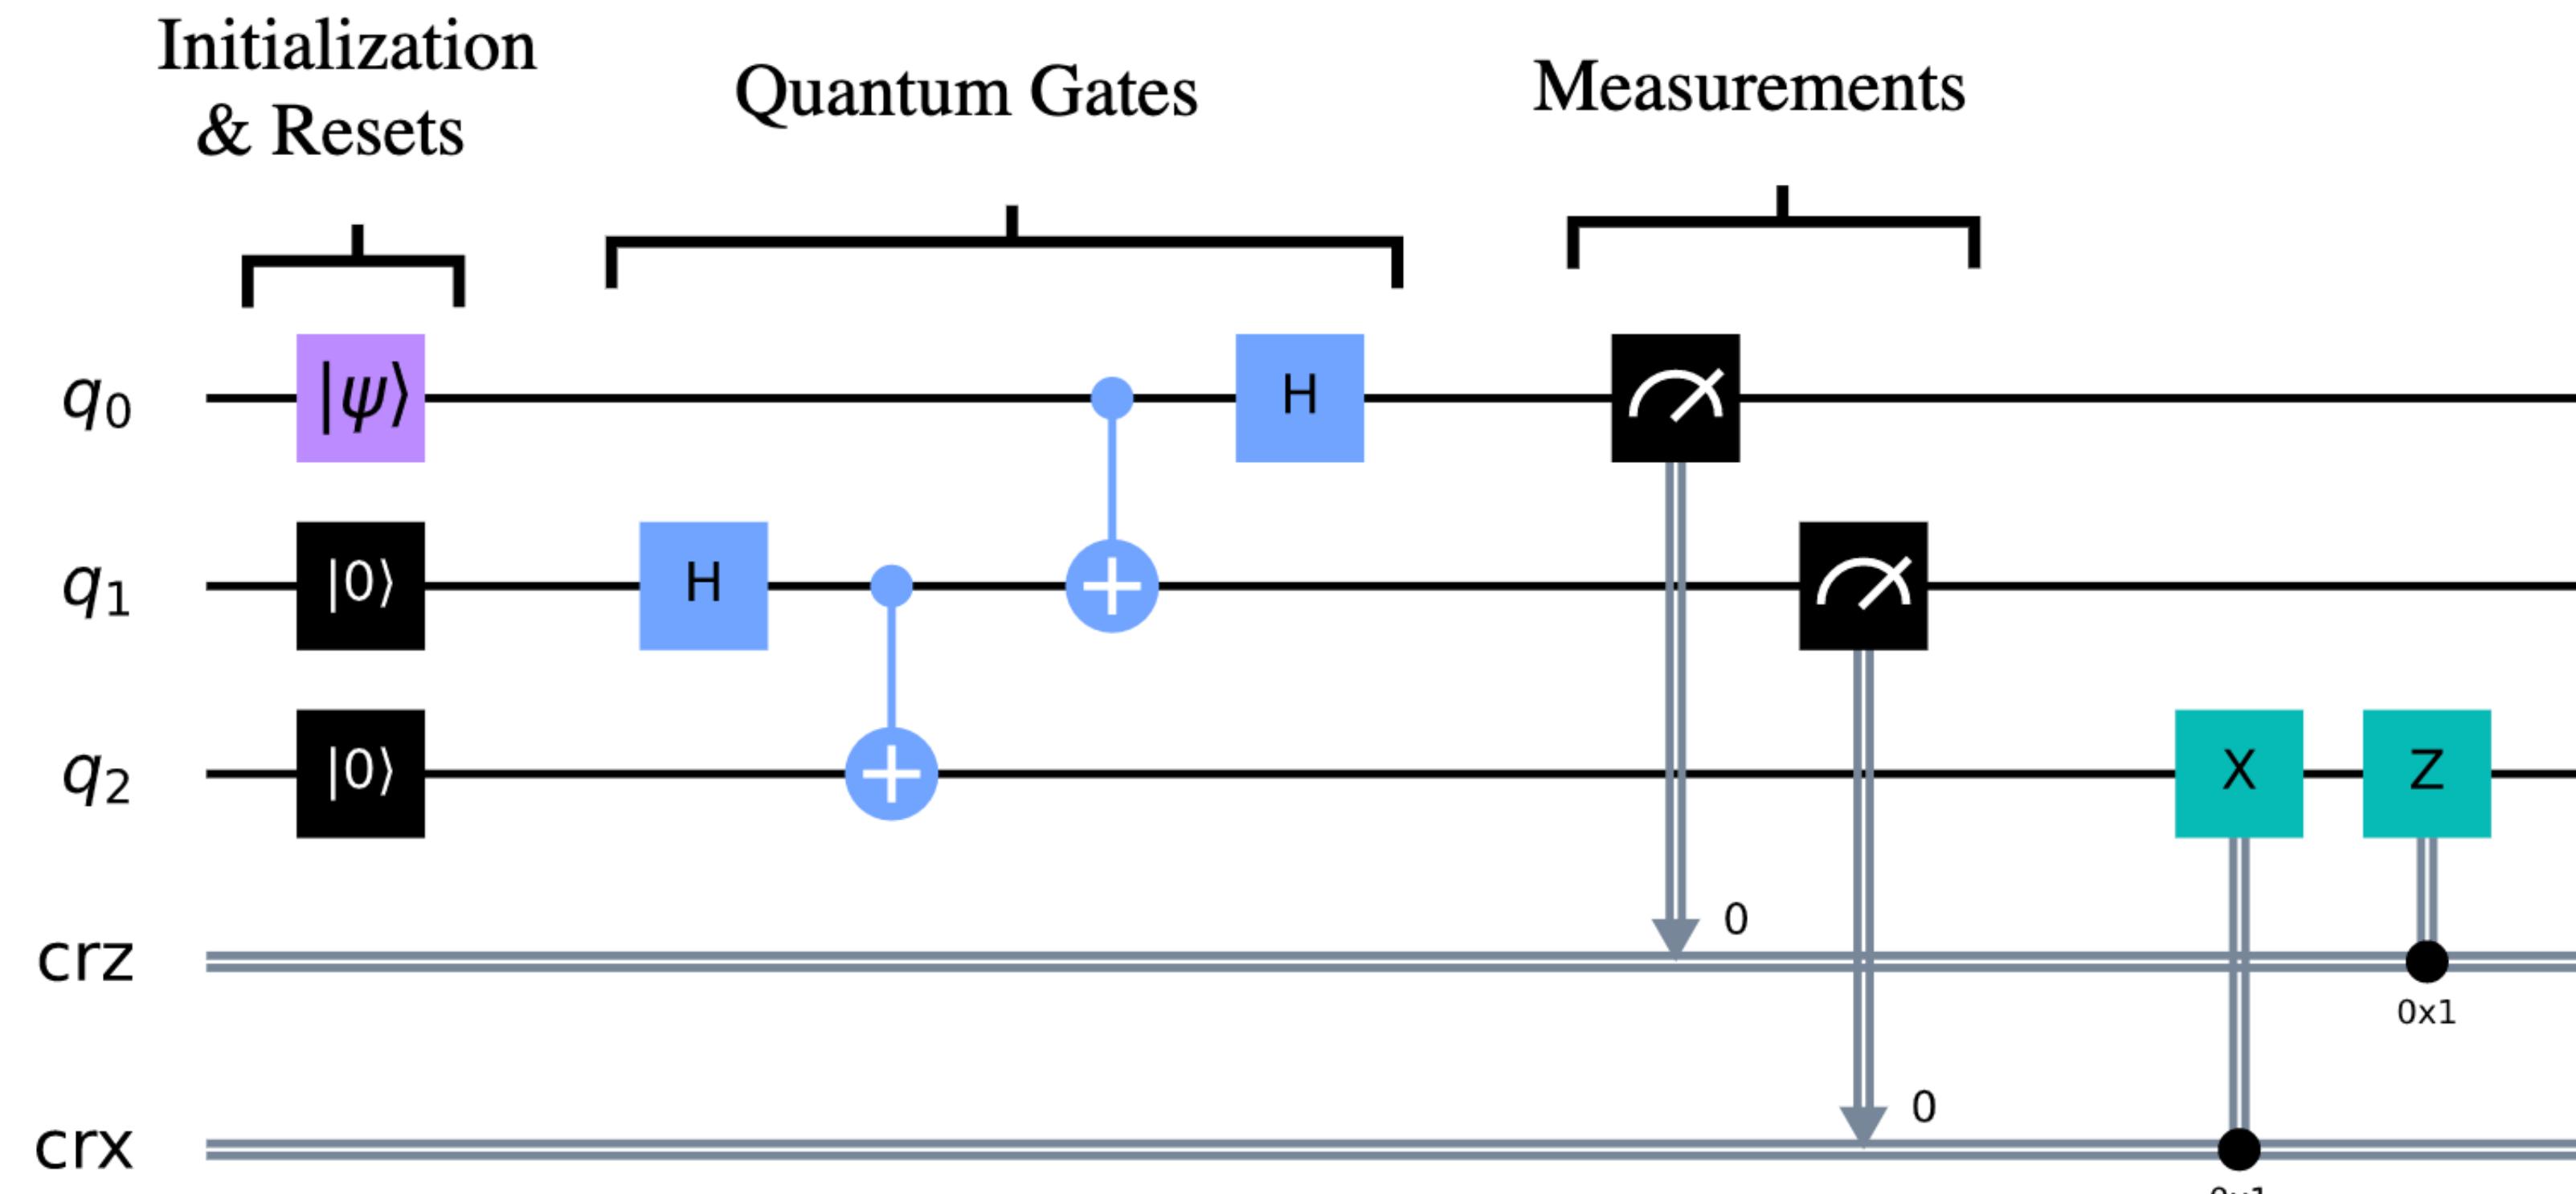


Compiled Quantum Circuit

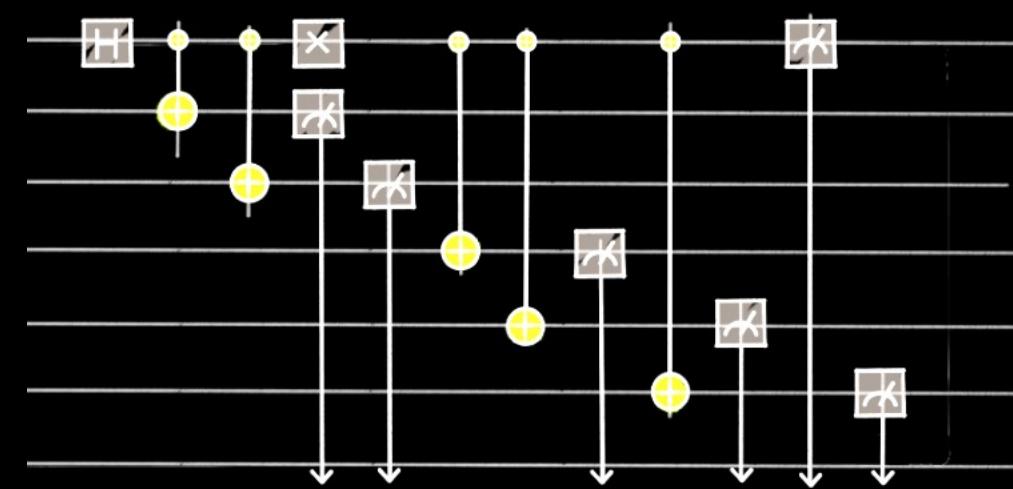
Hardware

Simulator





Classically Controlled
Quantum Gates

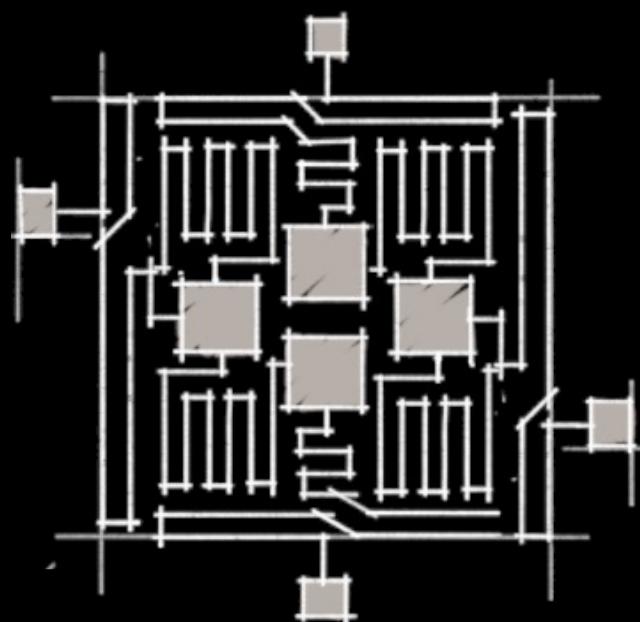


quantum circuit

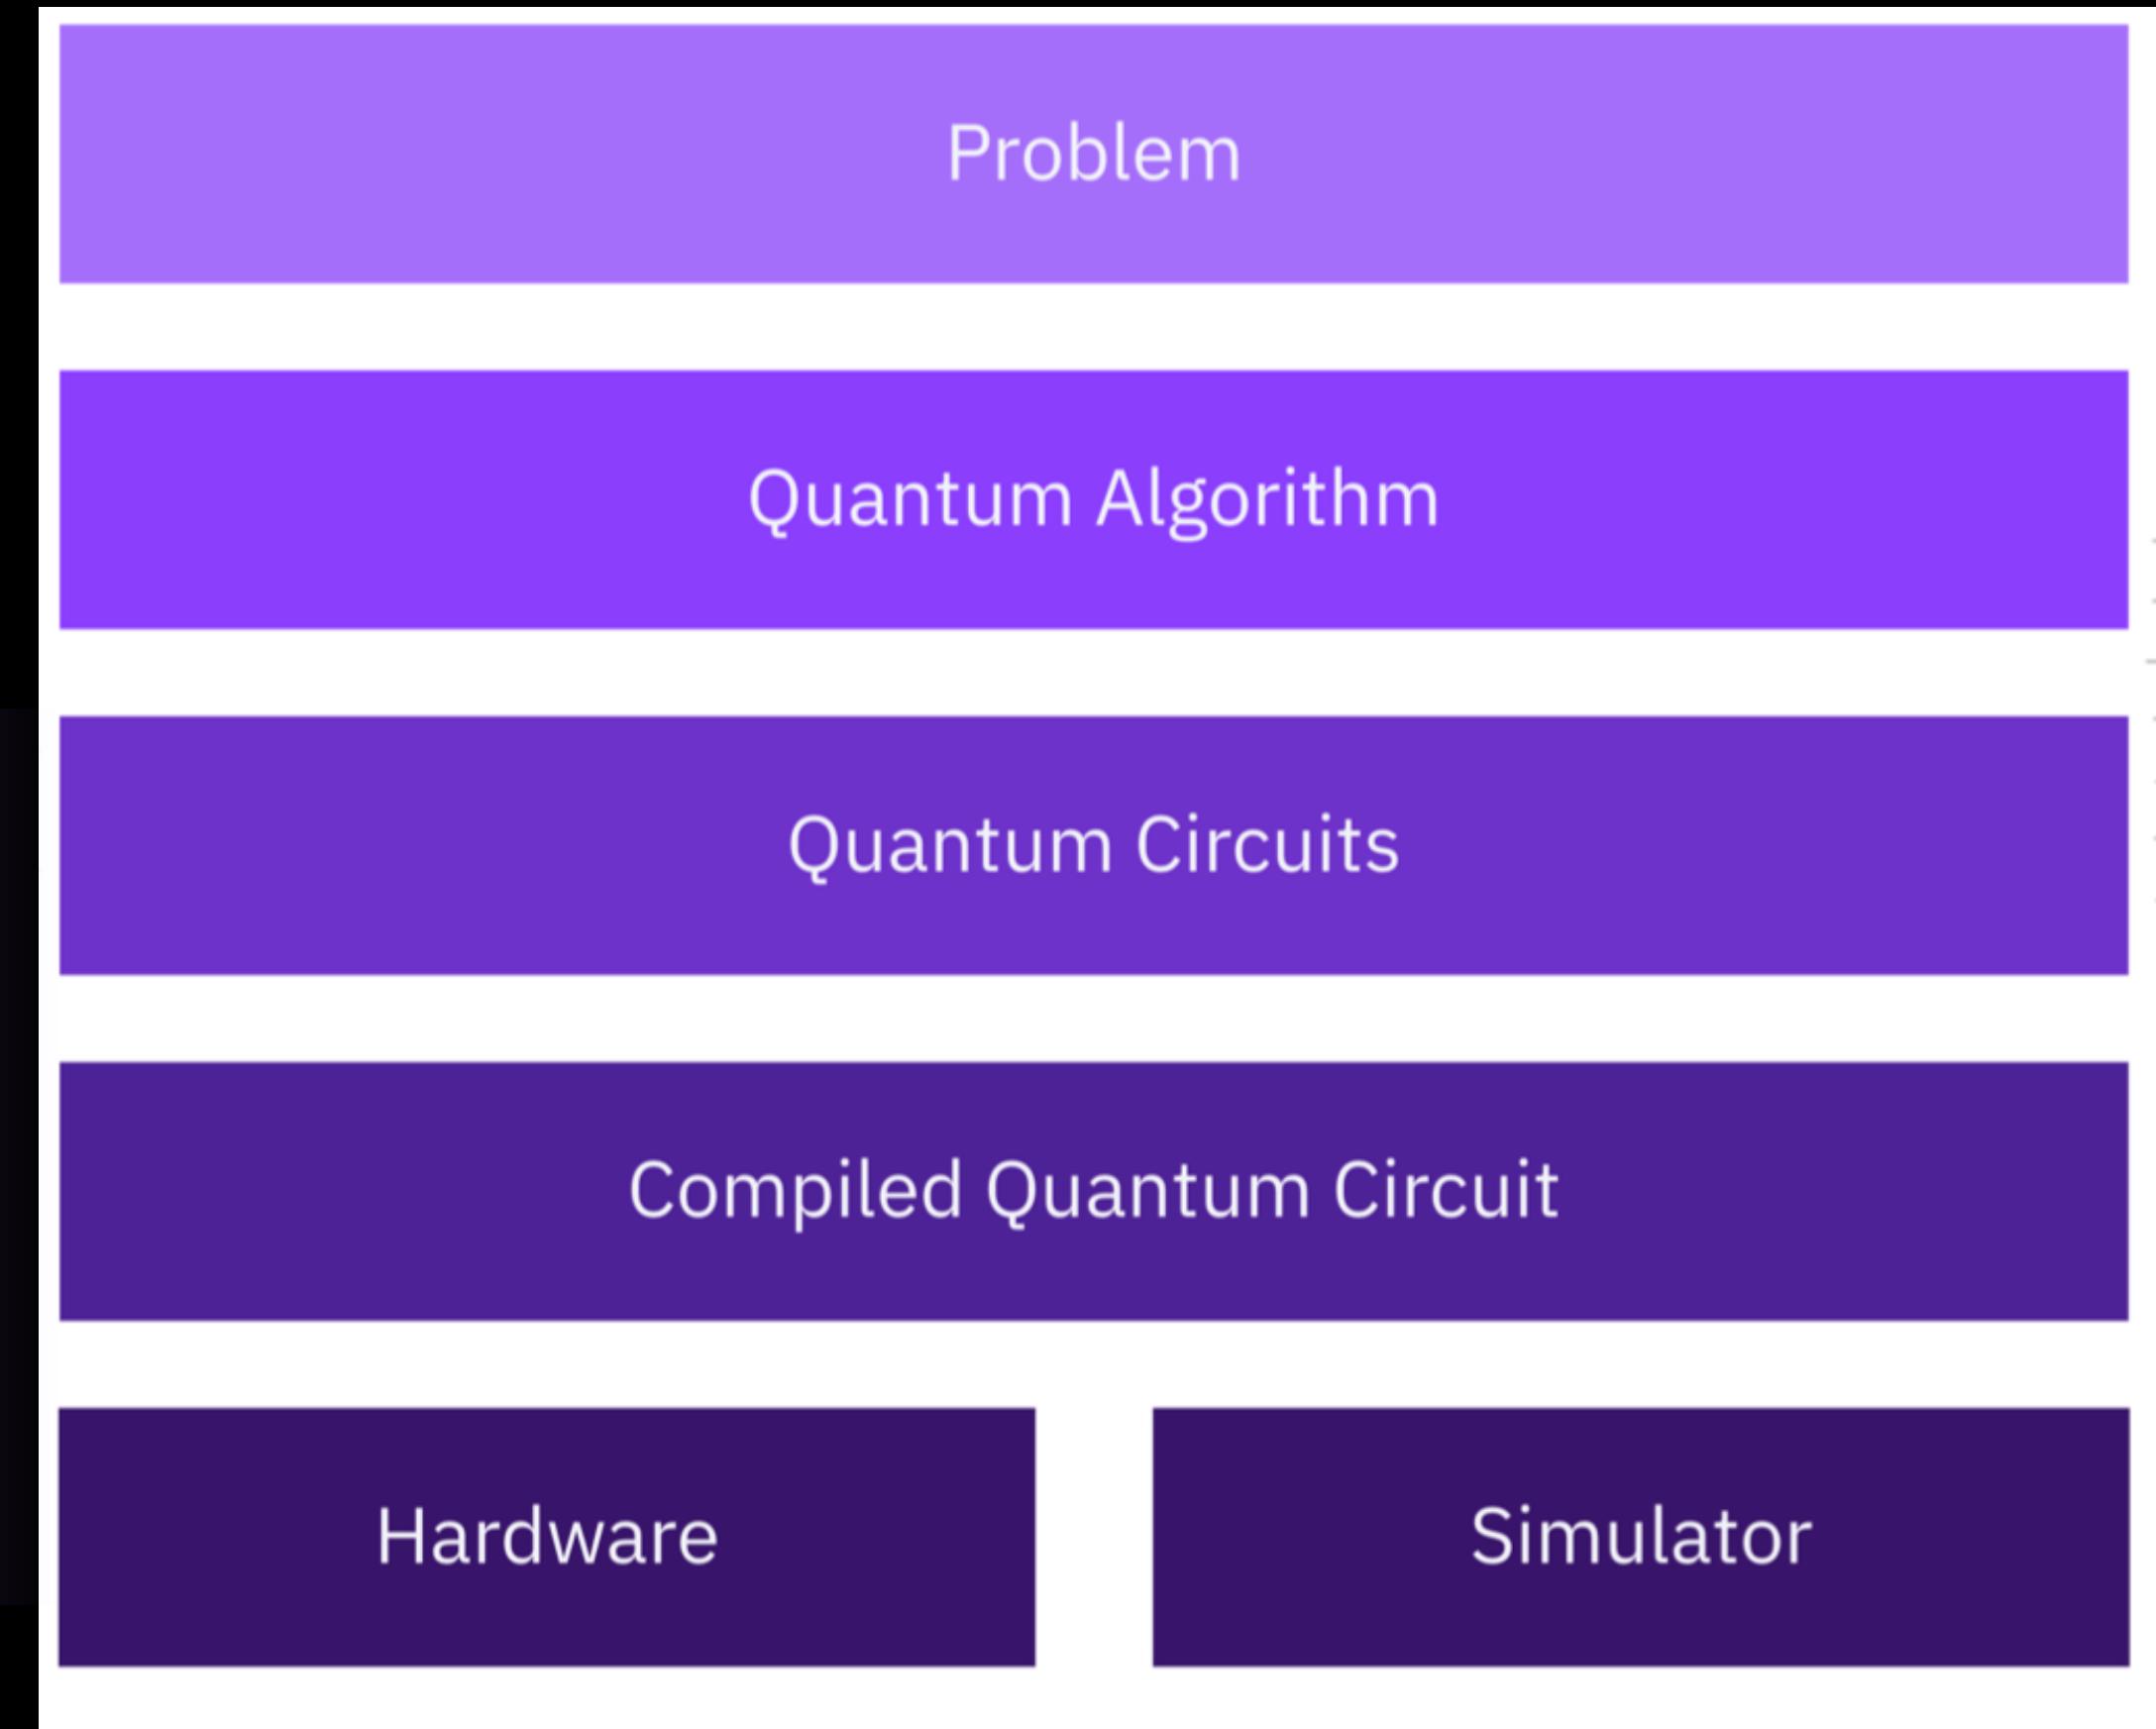
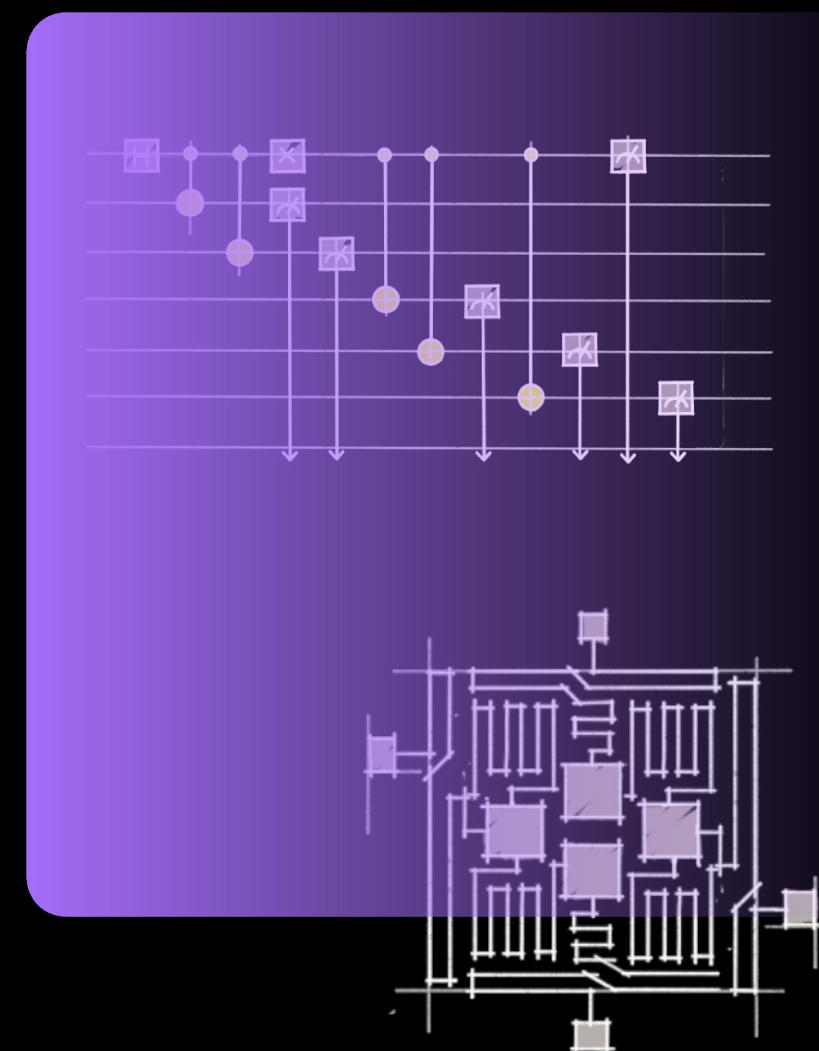


target description

compiled quantum circuit



- Why compiling?
- Transpiler architecture
- Performance



Why compiling?

Optimizations

```
X = 0  
Y = 0  
Z = input()
```

```
if (X == 1) {  
    Y = 5^20  
} else {  
    Y = 20^20  
}
```

```
while (X >= -1)  
{ X = X - 1  
    Z = Z + 1  
}
```

```
return(Z+X)
```

```
X = 0  
Y = 0  
Z = input()
```

```
Y = 20^20
```

```
while (X >= -1)  
{ X = X - 1  
    Z = Z + 1  
}
```

```
return(Z+X)
```

```
X = 0  
Y = 0  
Z = input()
```

```
Y = 20^20
```

```
X = X - 1  
Z = Z + 1  
X = X - 1  
Z = Z + 1
```

```
return(Z+X)
```

```
X = 0  
Z = input()
```

```
X = X - 1  
Z = Z + 1  
X = X - 1  
Z = Z + 1
```

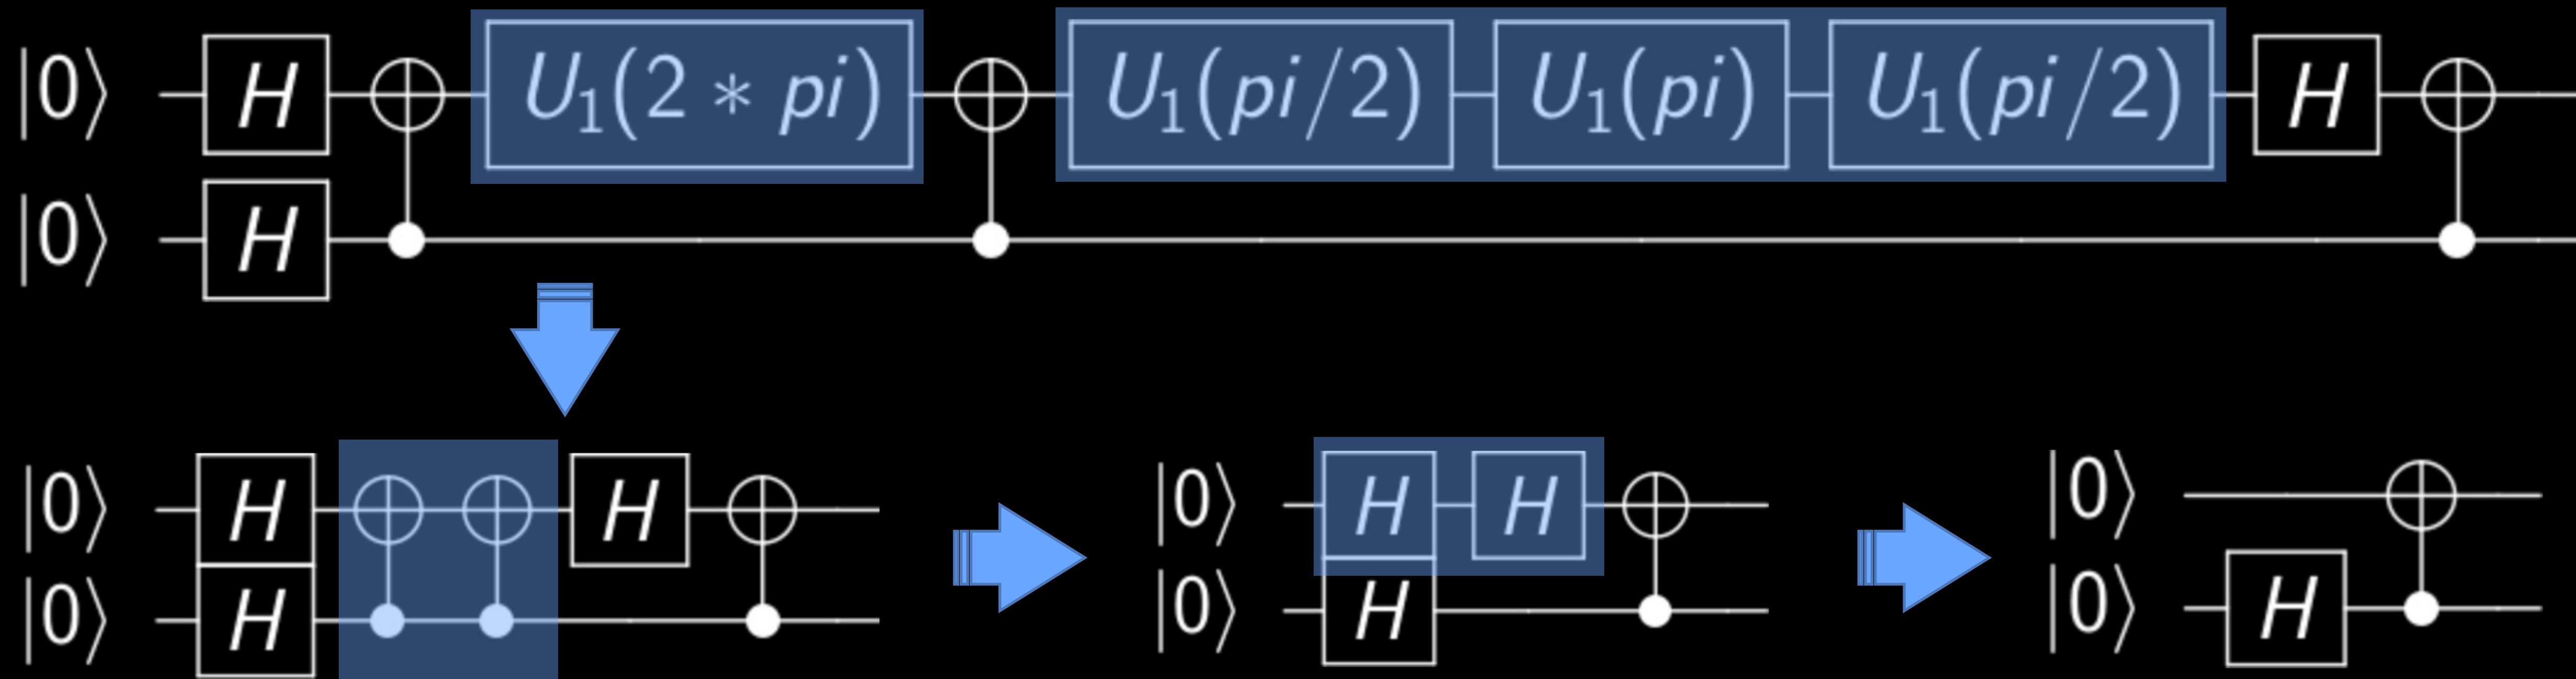
```
return(Z+X)
```

```
Z = input()
```

```
return(Z)
```

Why compiling?

Optimizations



Why compiling?

Optimizations

Instruction set

$c = a + b$

```
push a
push b add
pop c
```

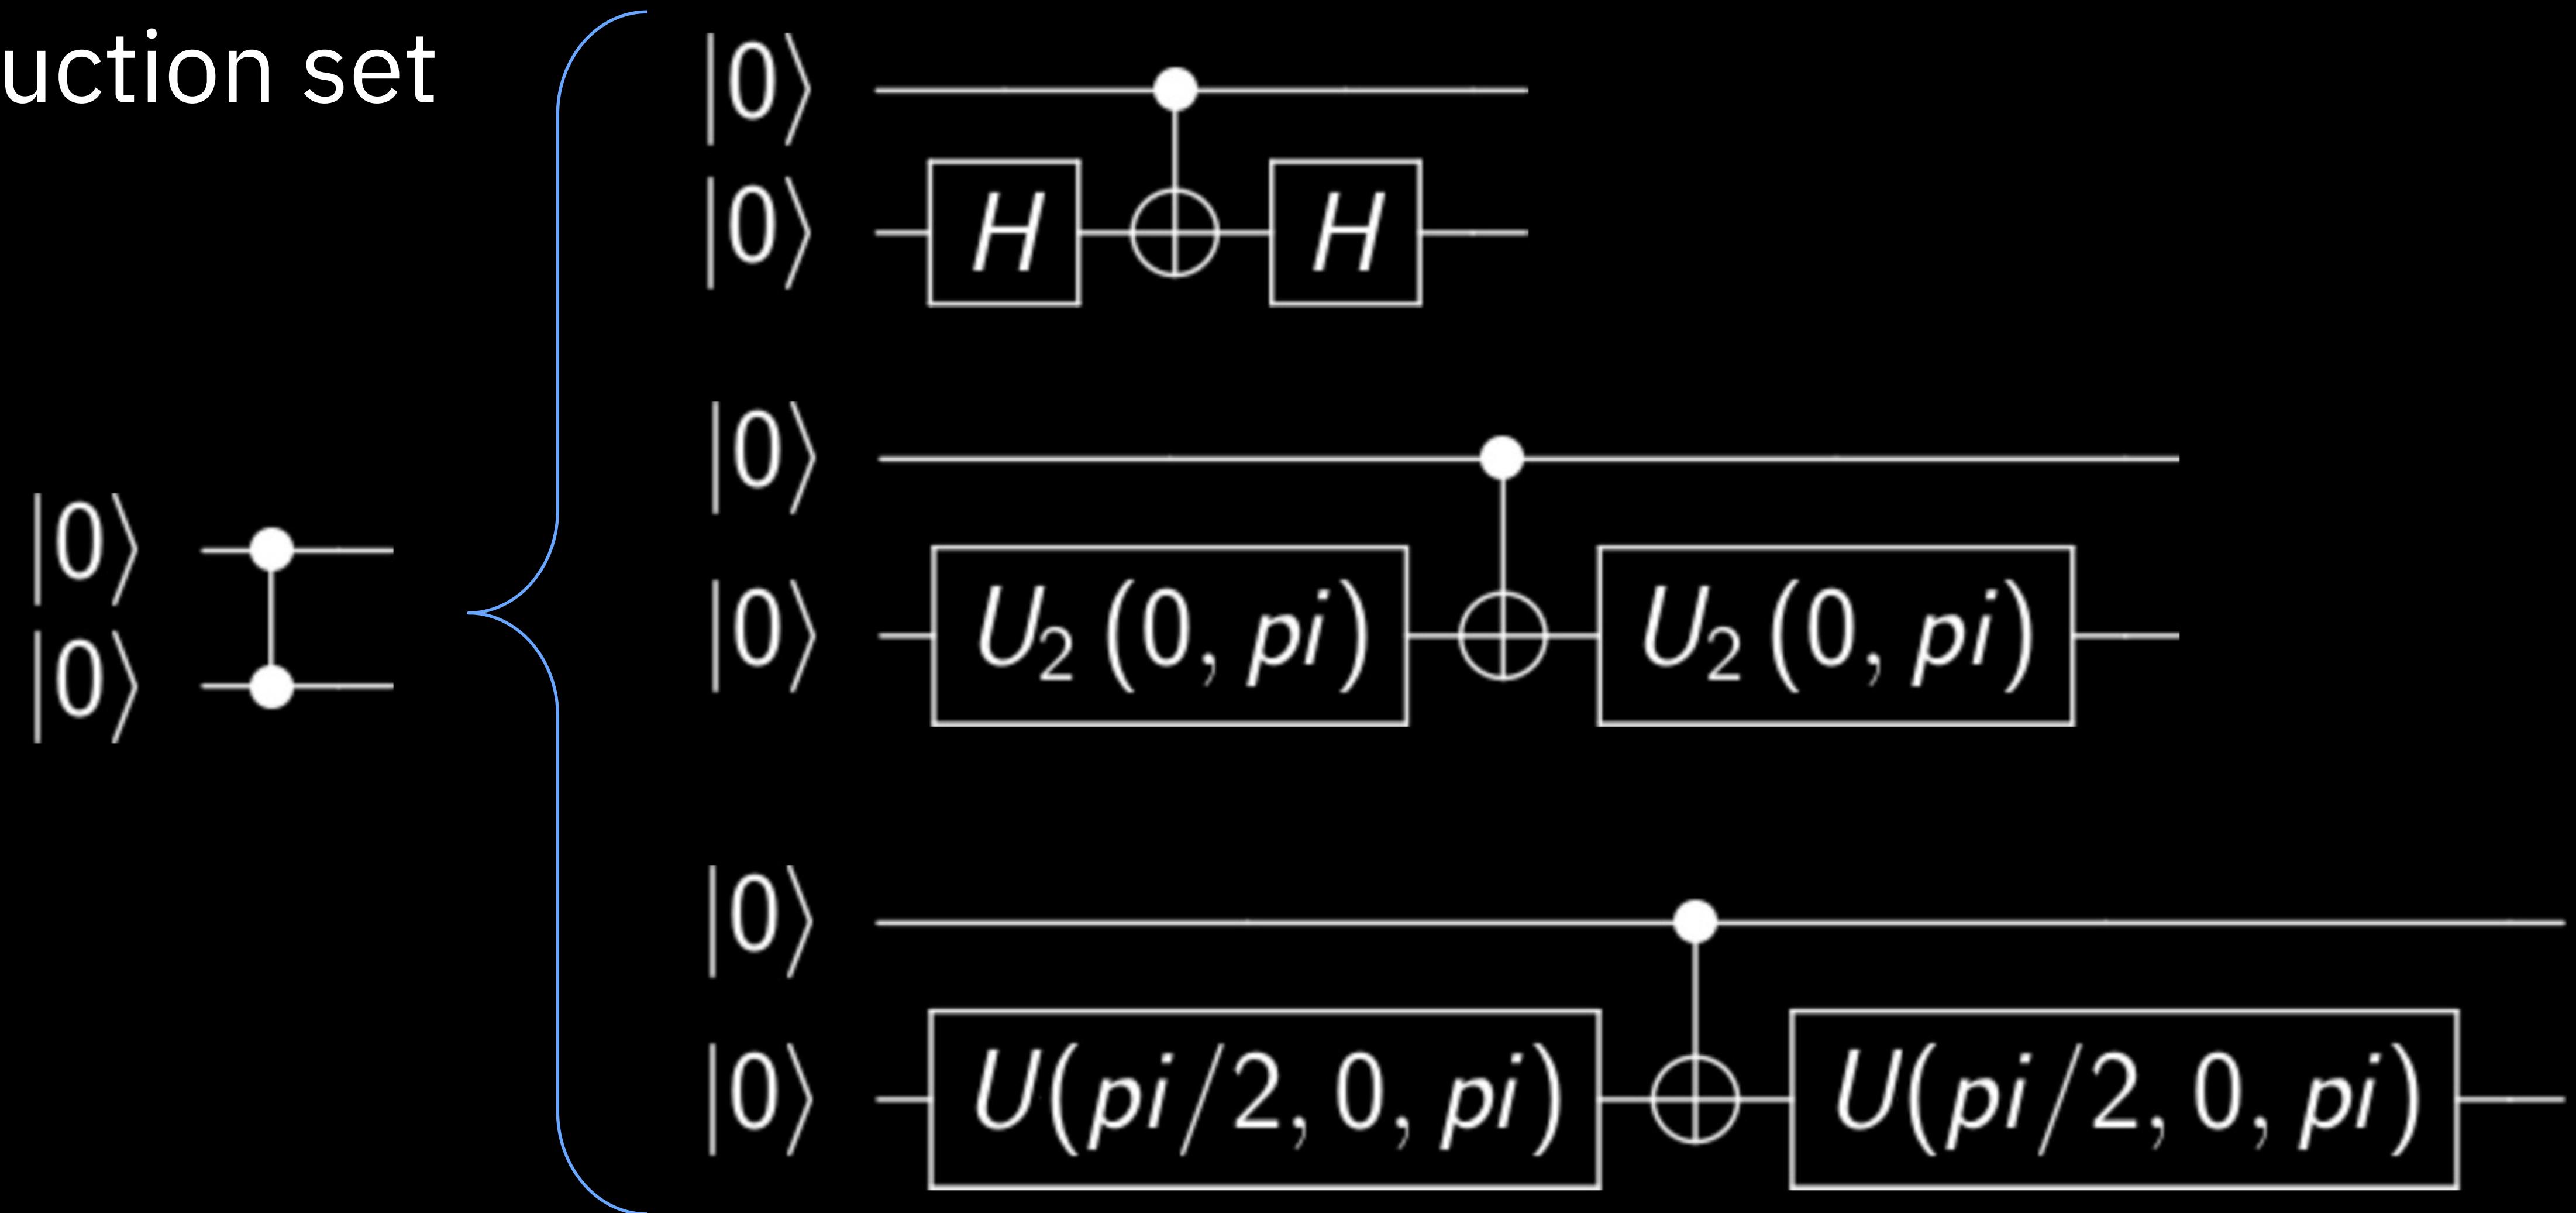
```
load a add
b store c
```

add a, b, c

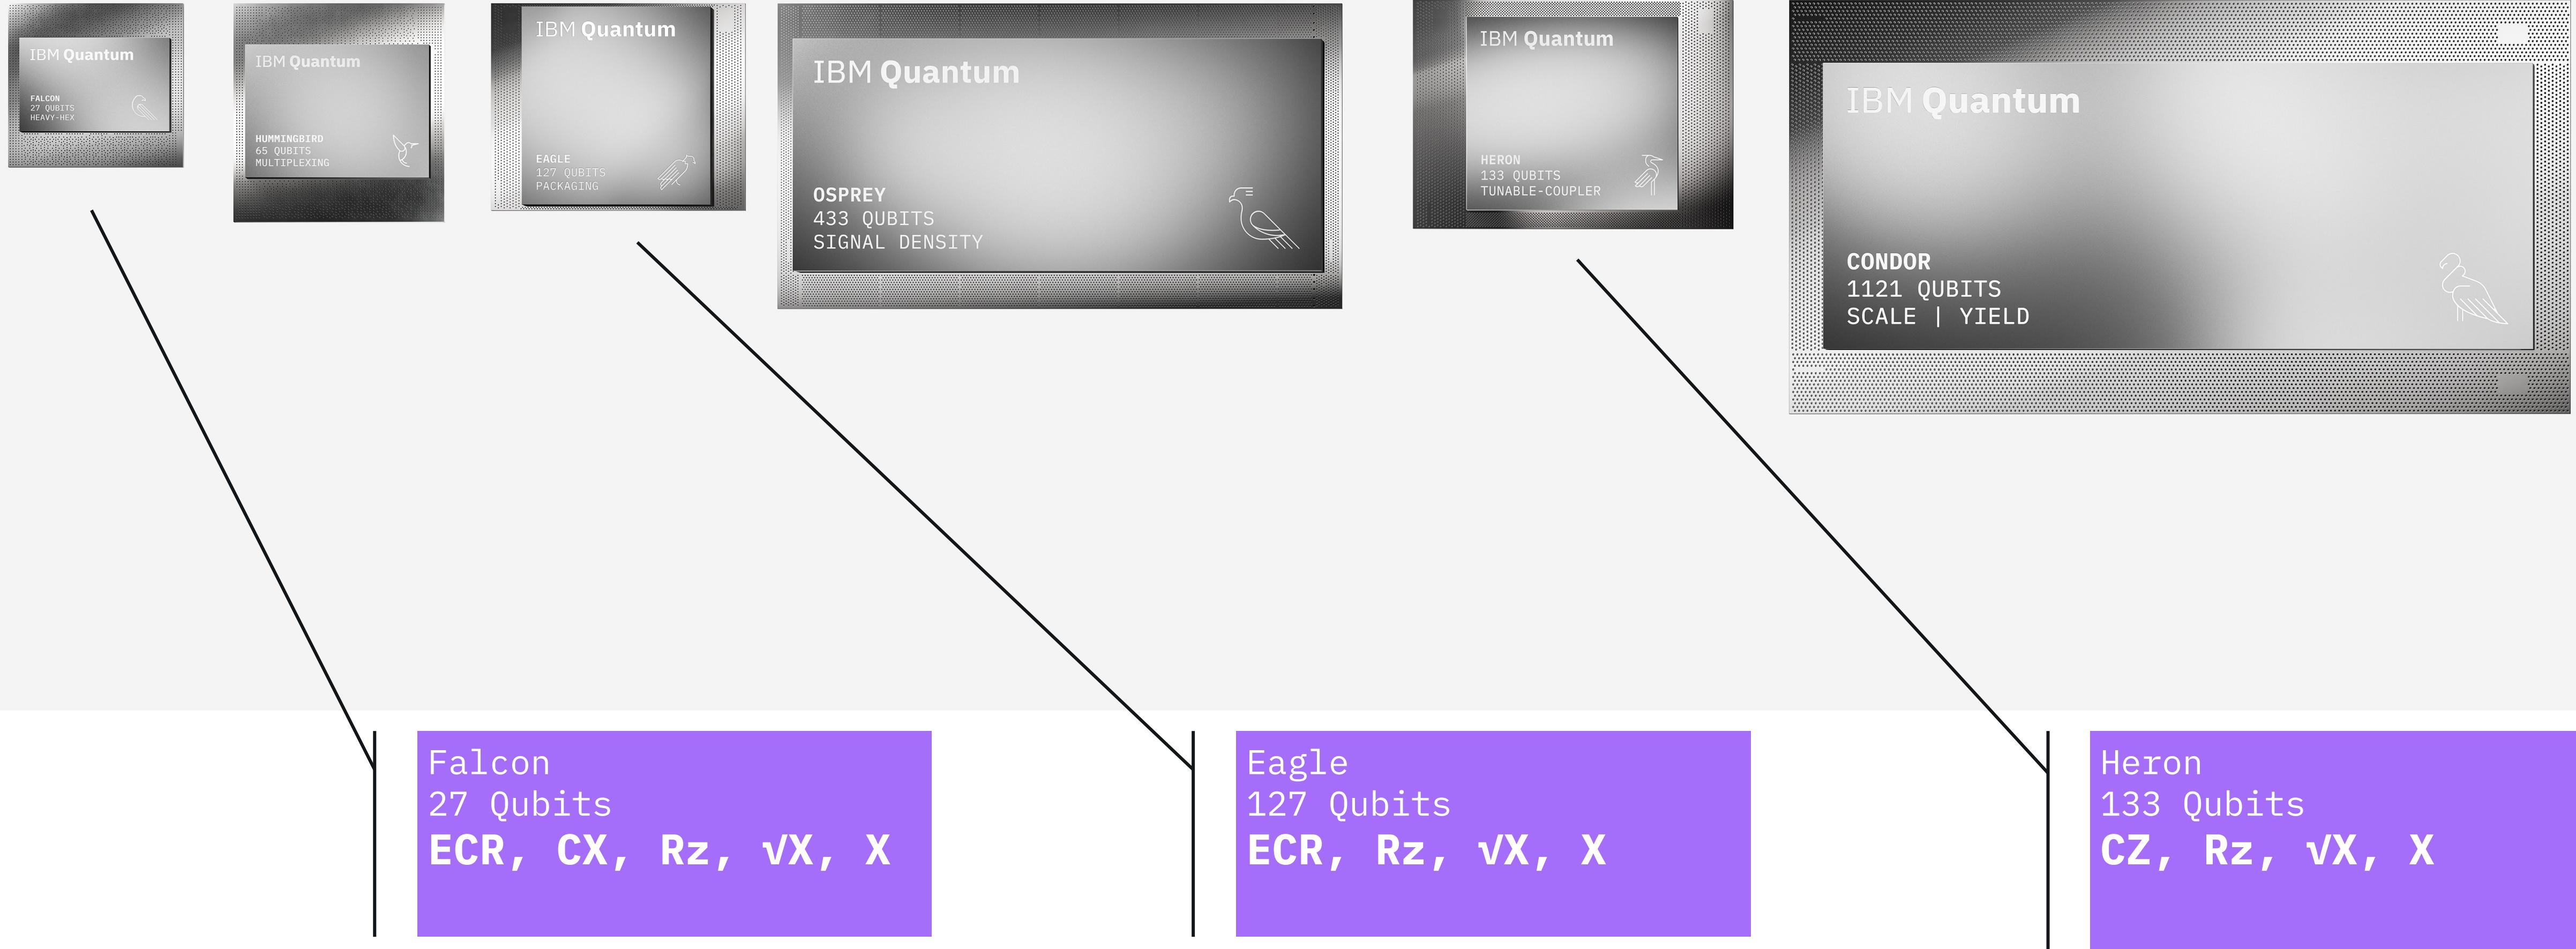
Why compiling?

Optimizations

Instruction set



Why compiling?

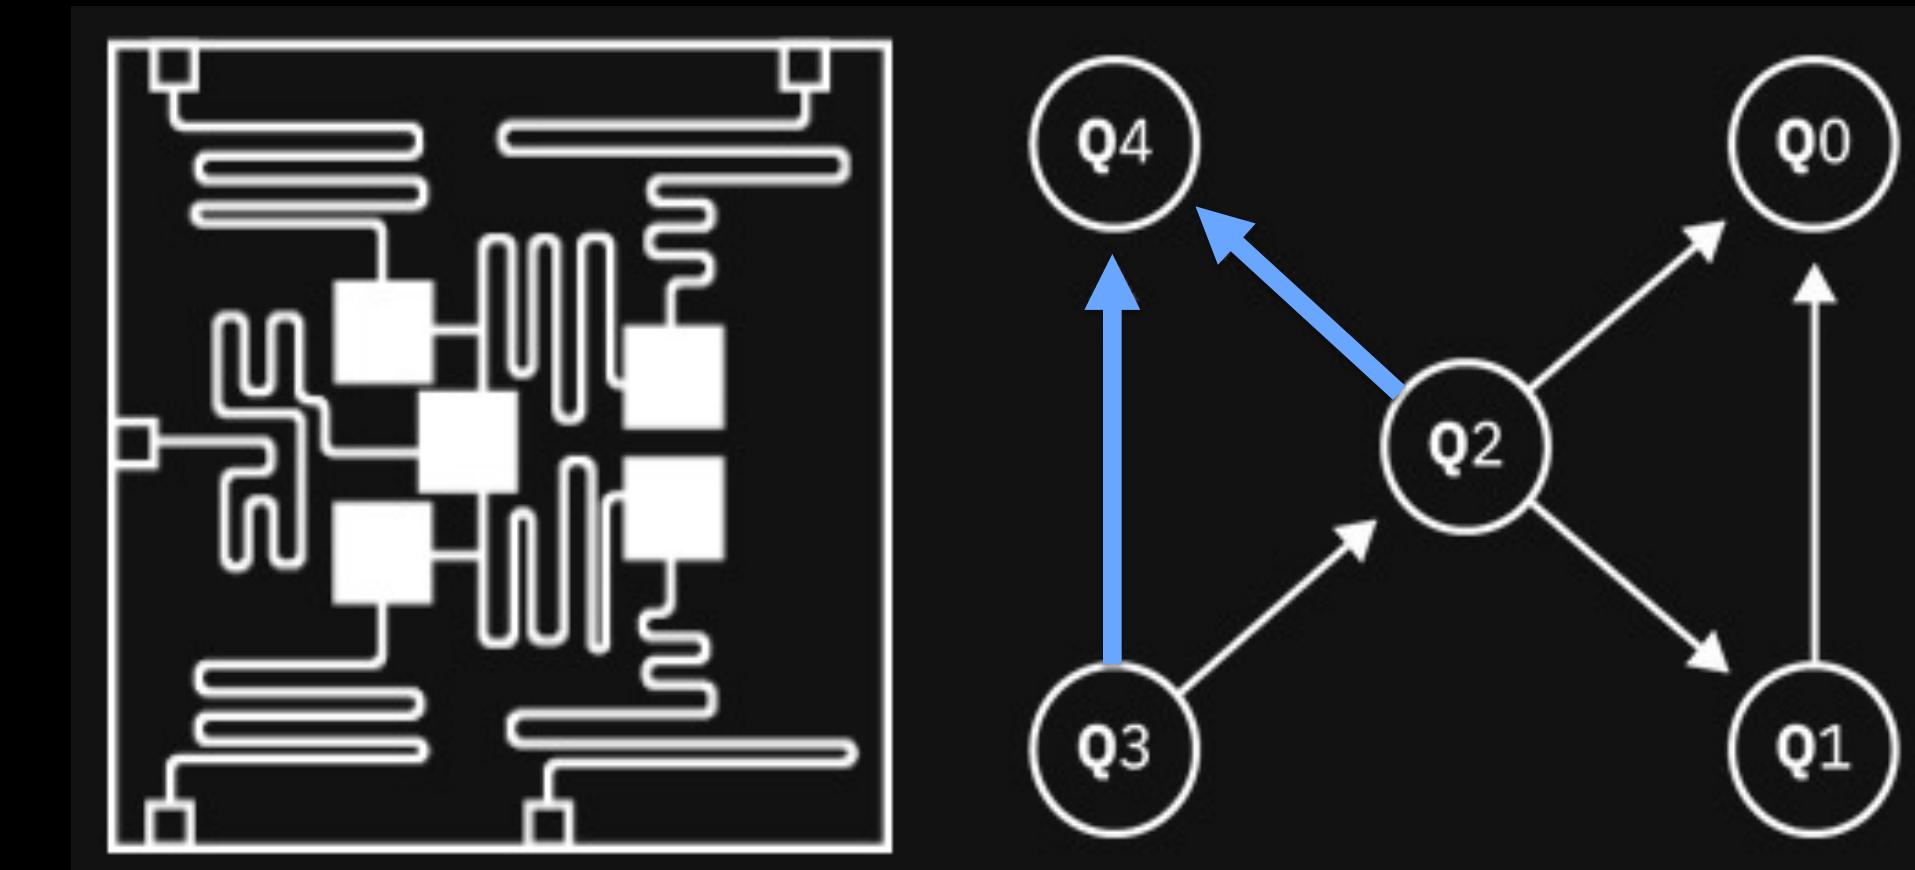
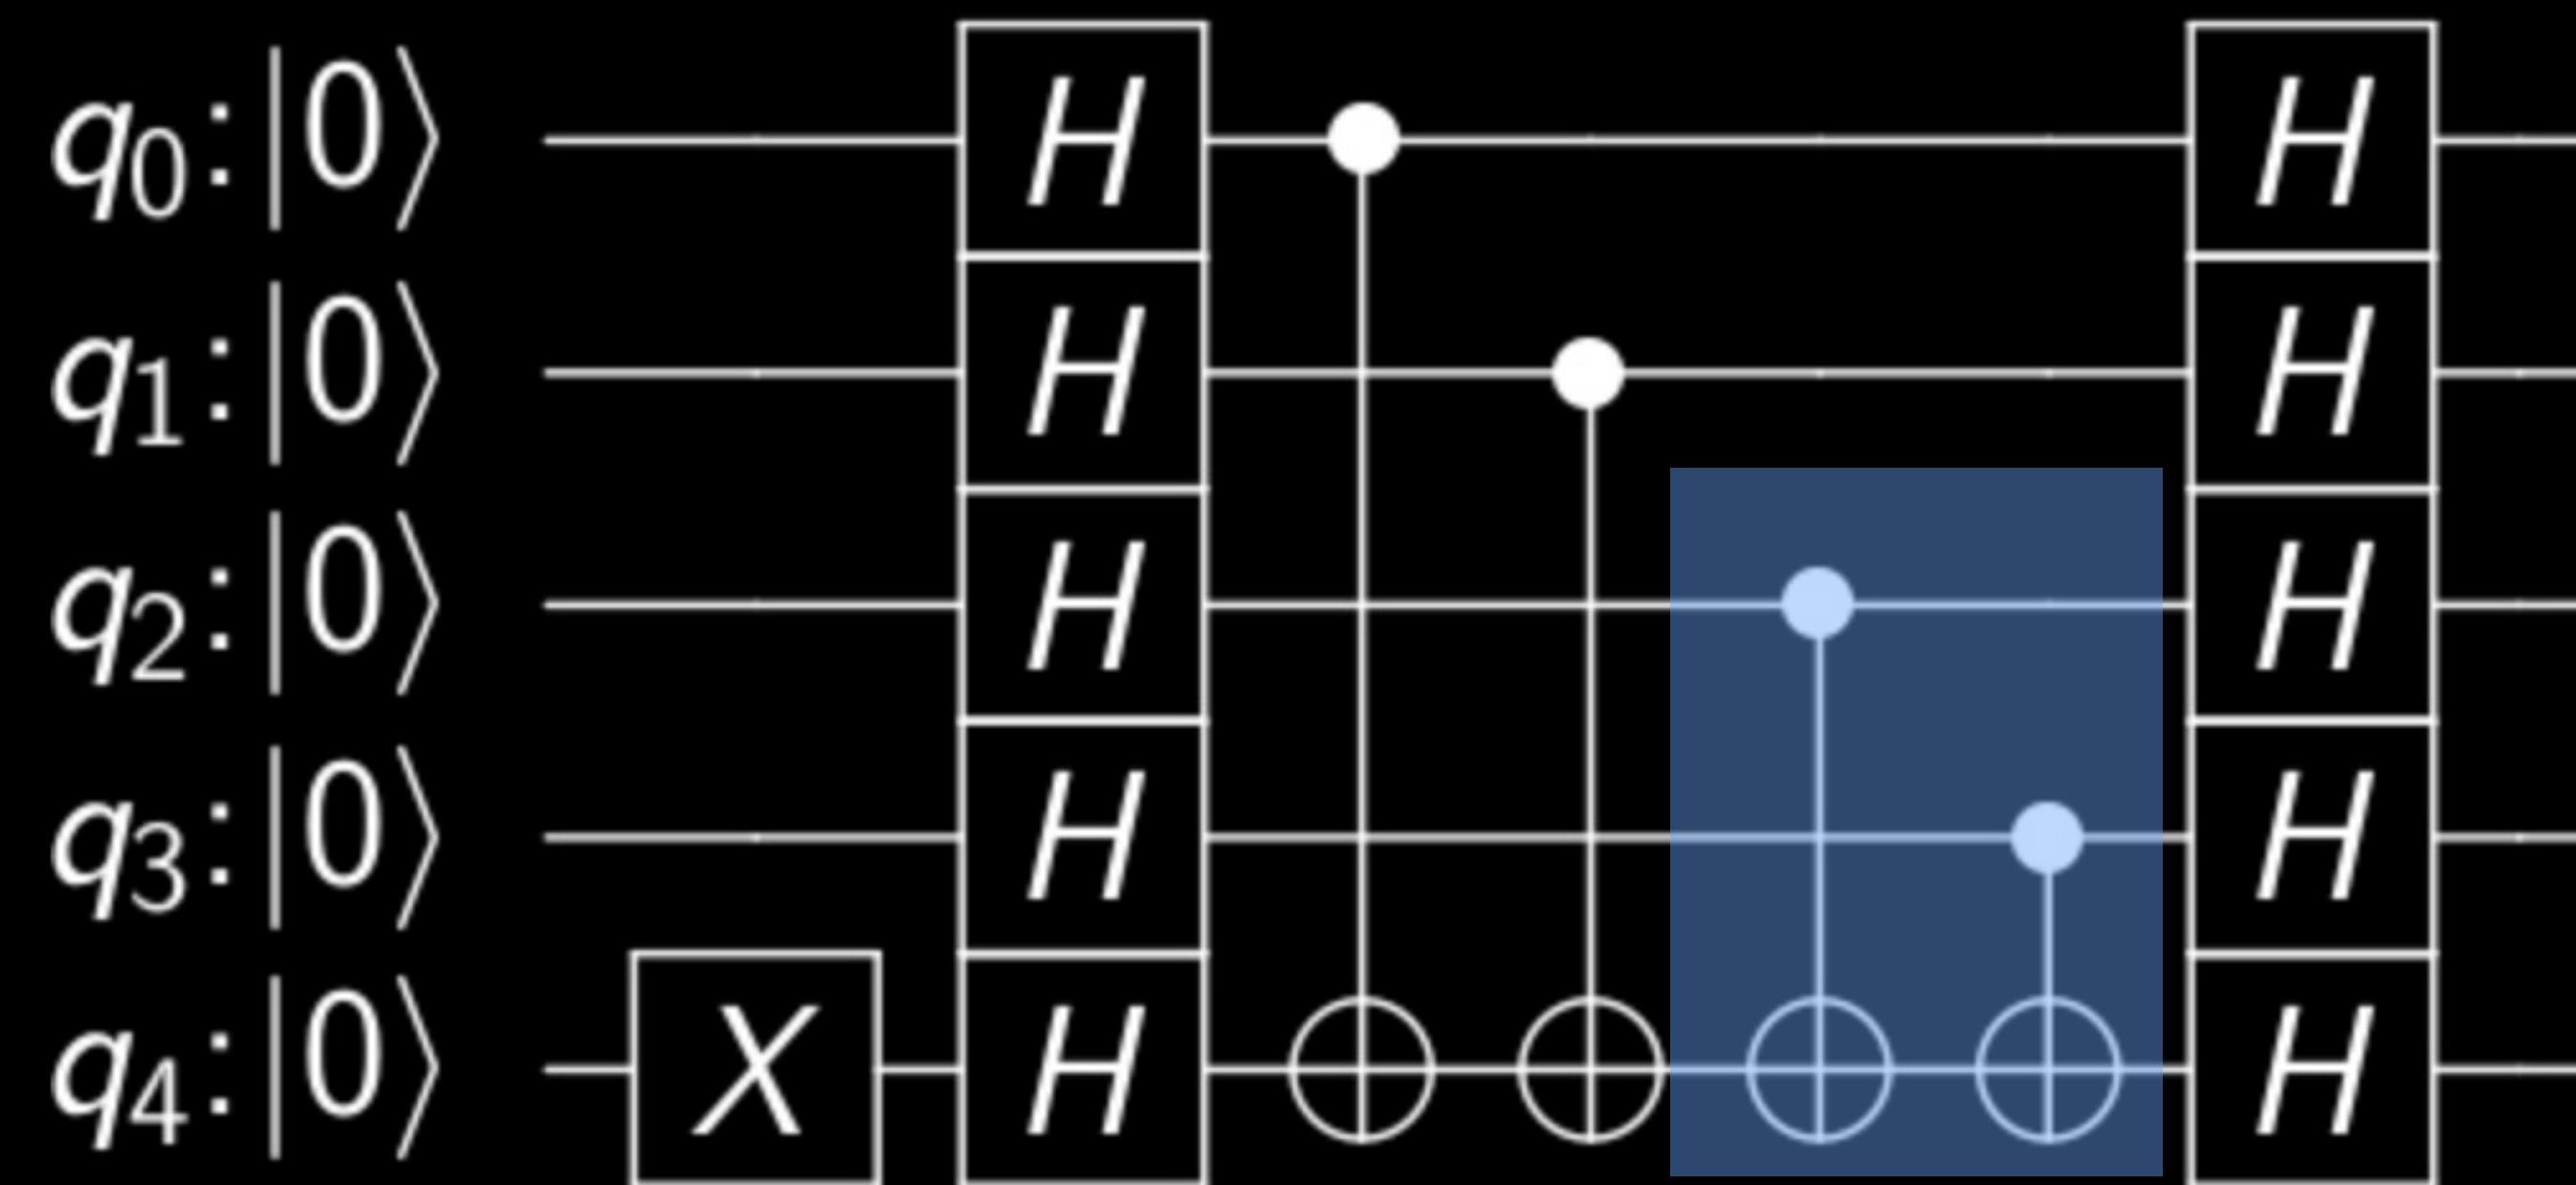


Why compiling?

Optimizations

Instruction set

Hardware connectivity



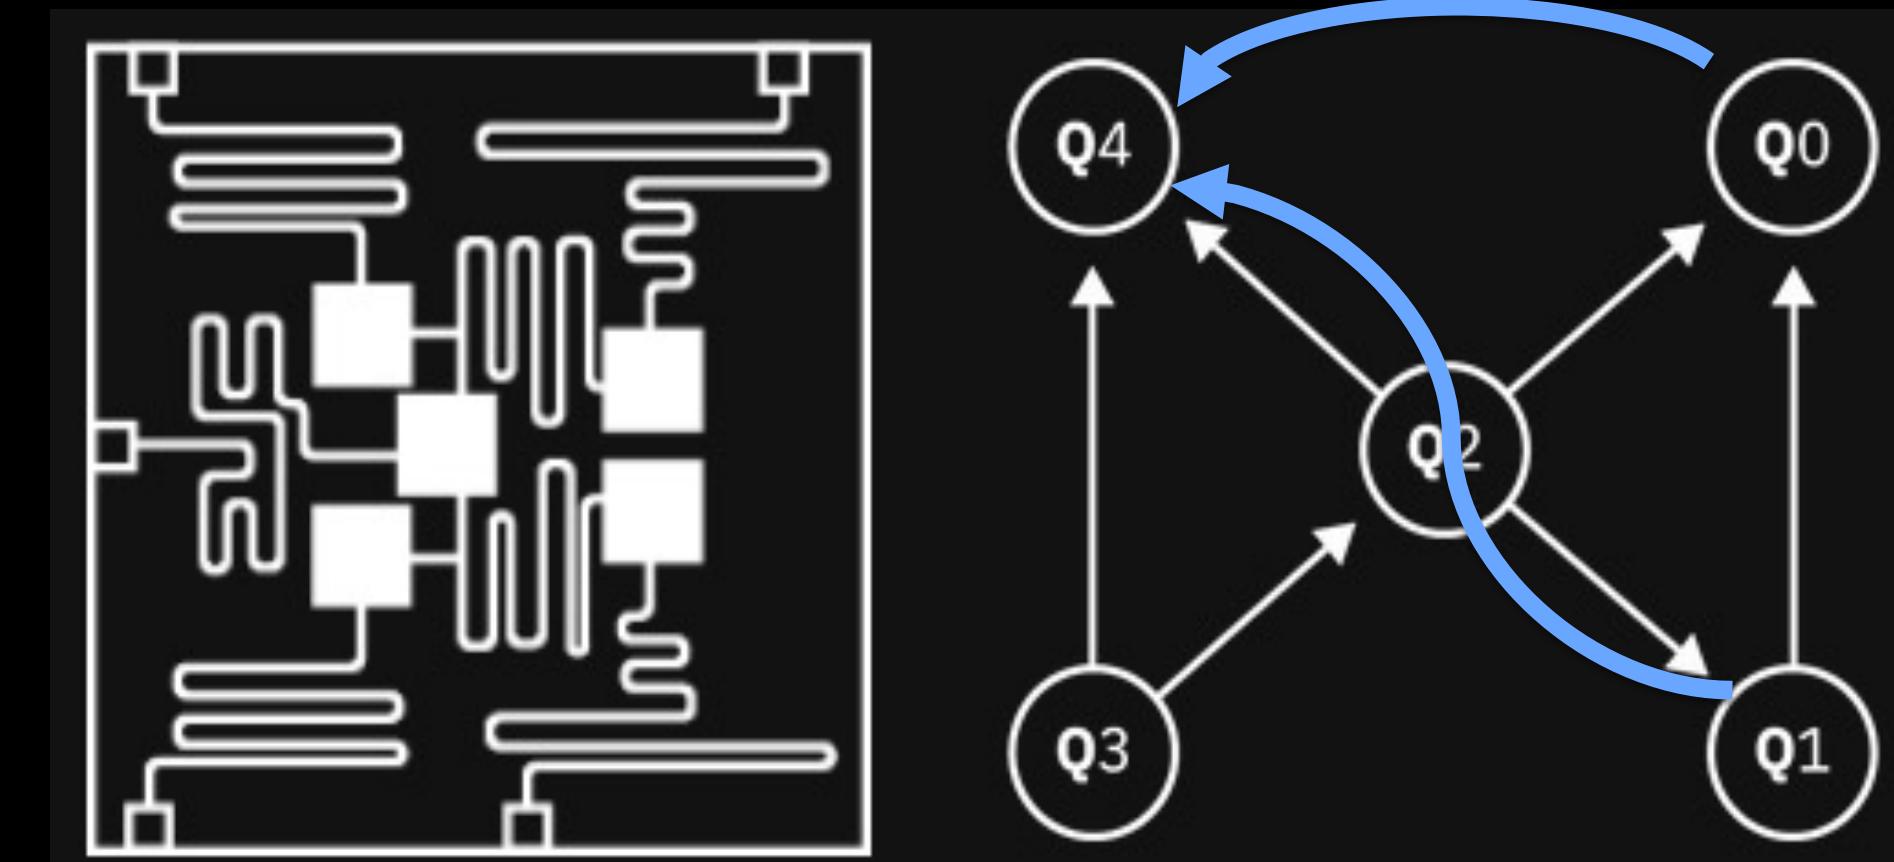
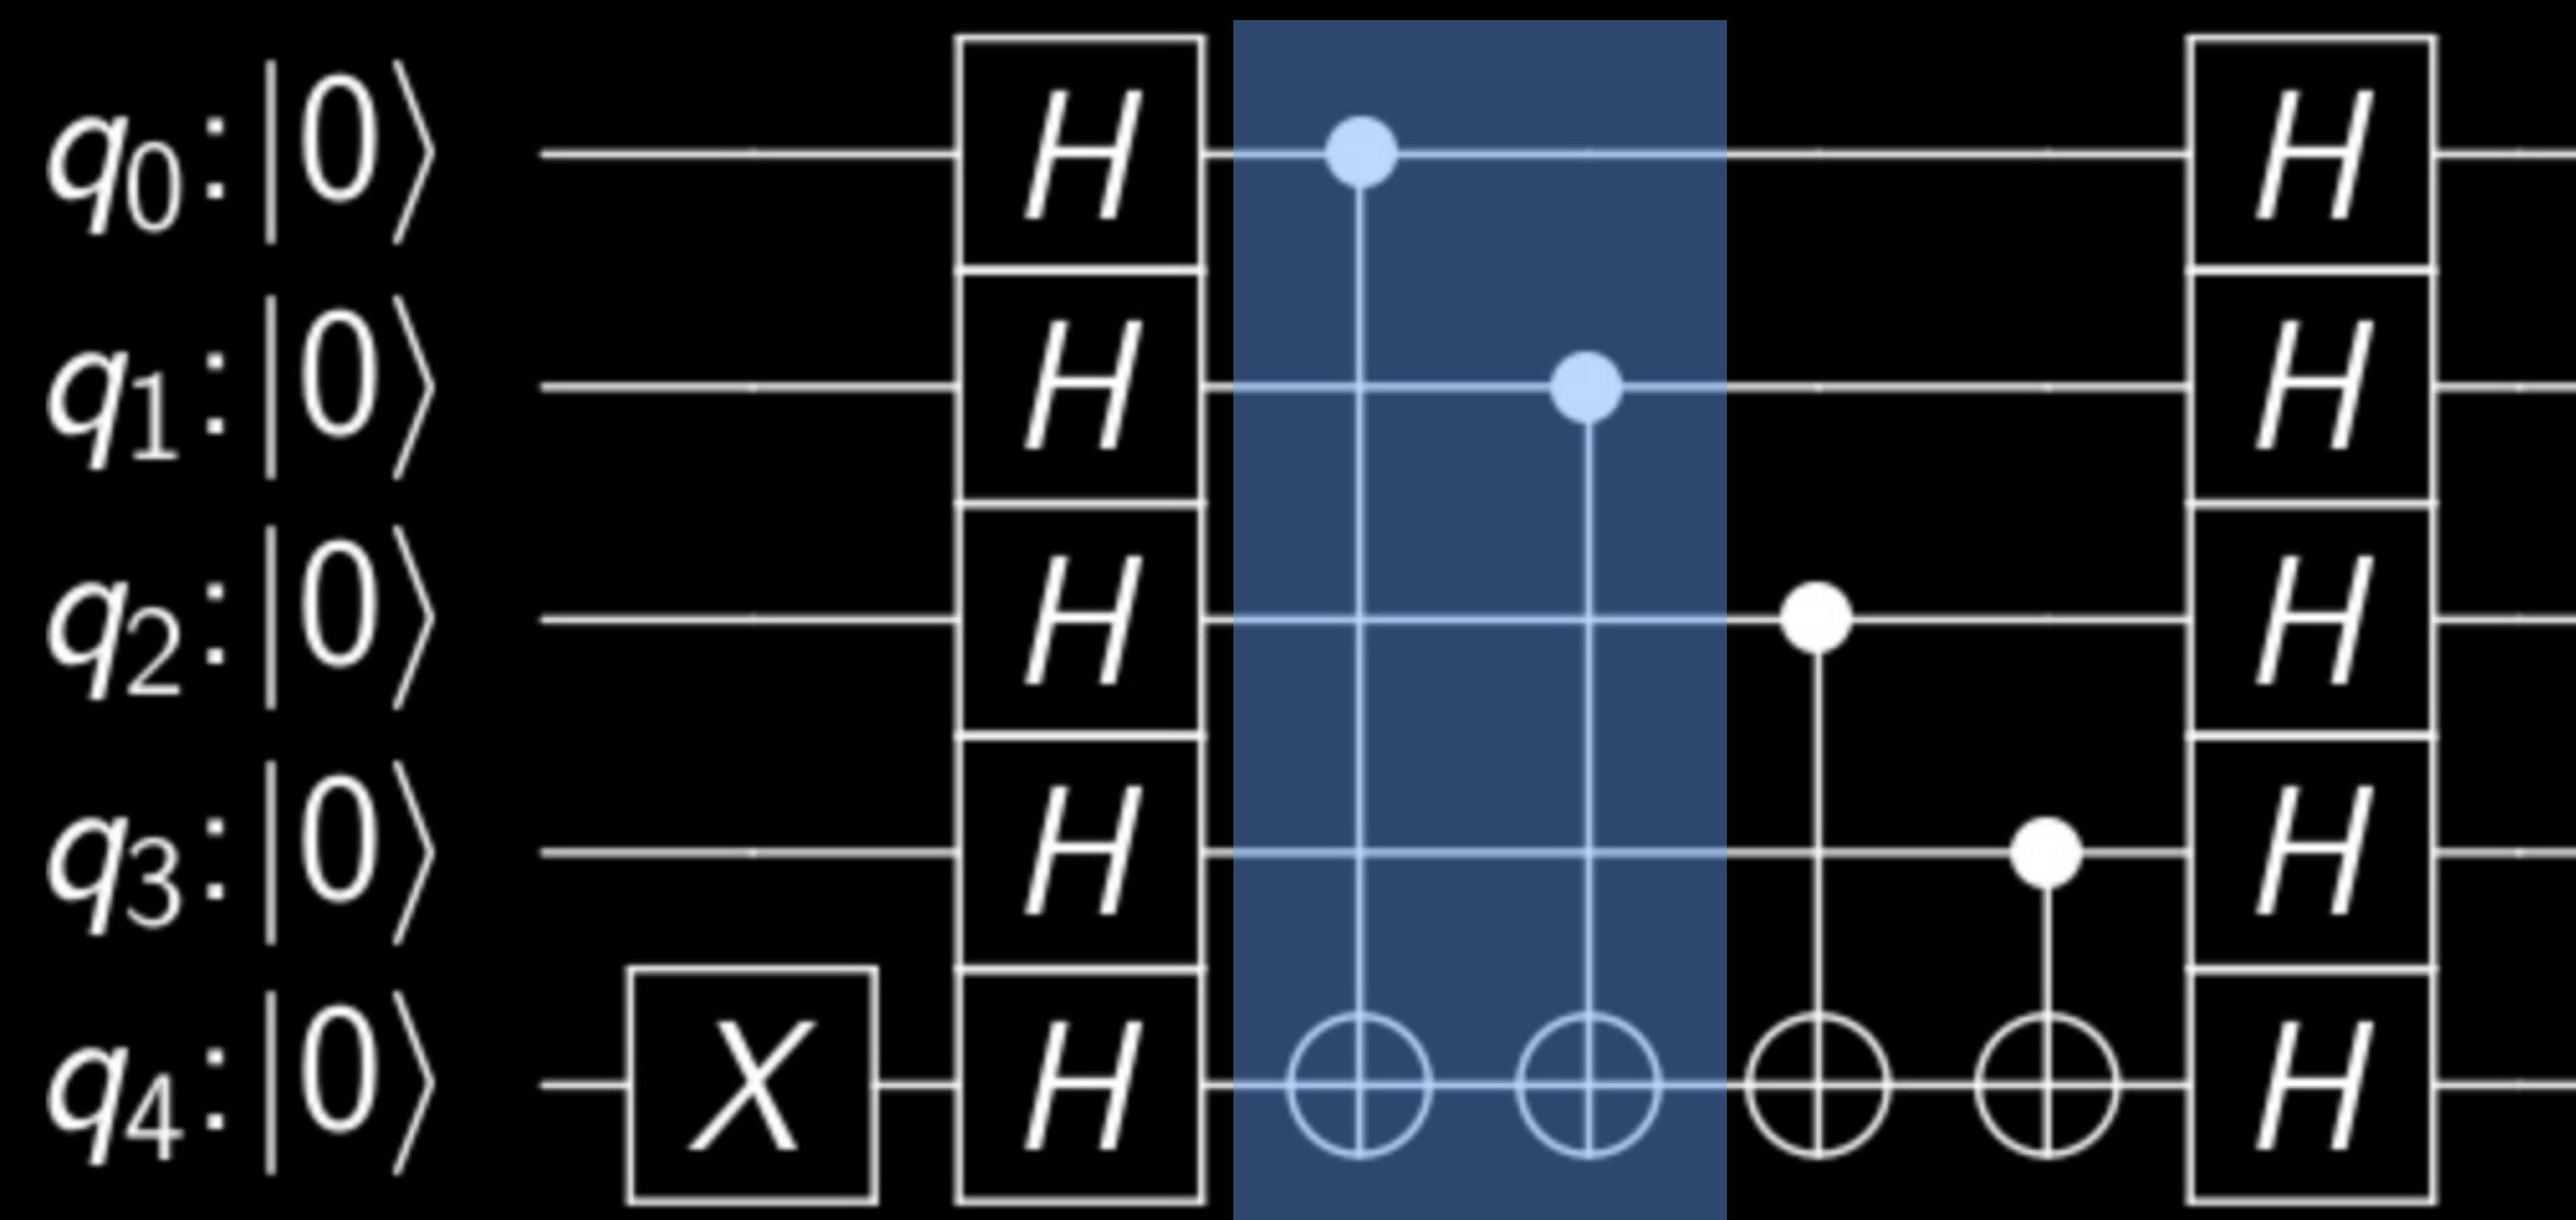
Why compiling?

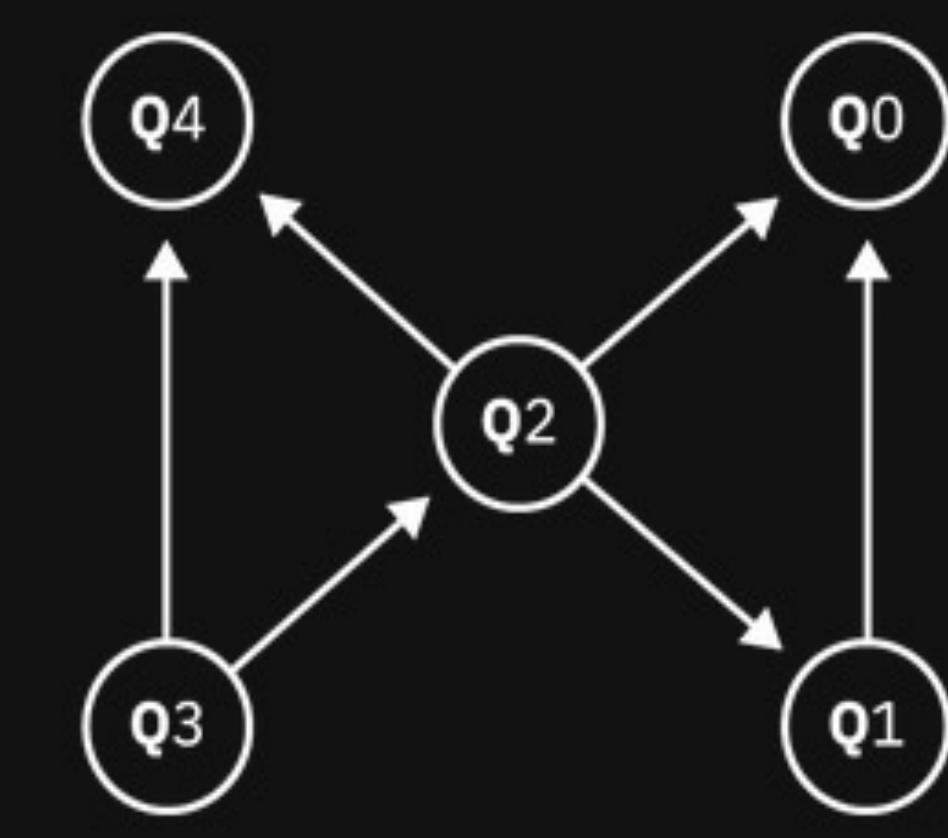
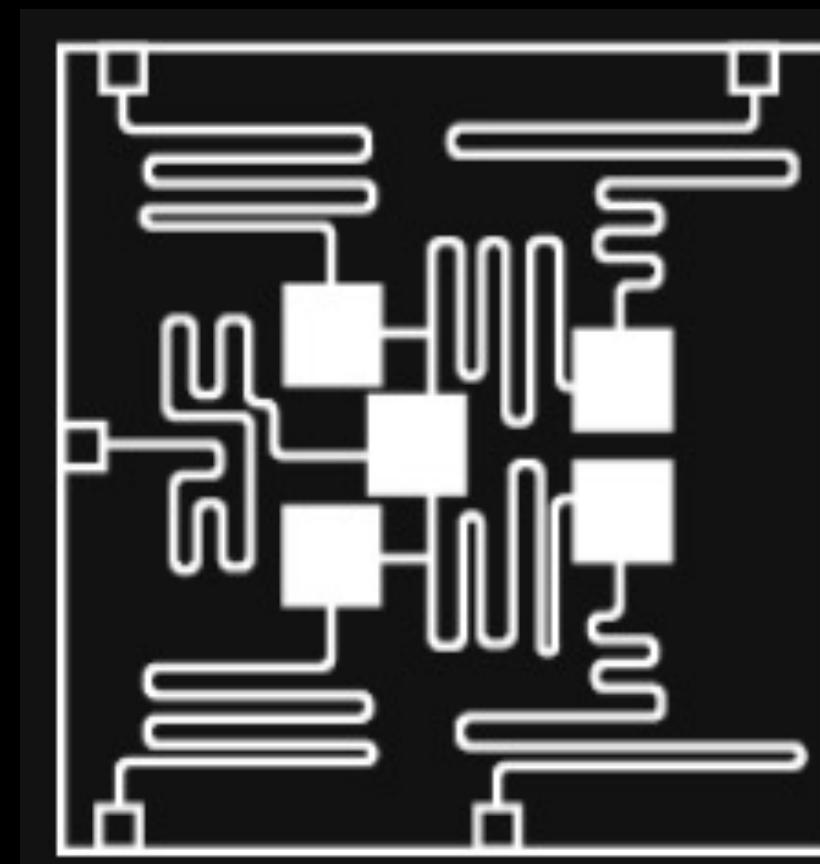
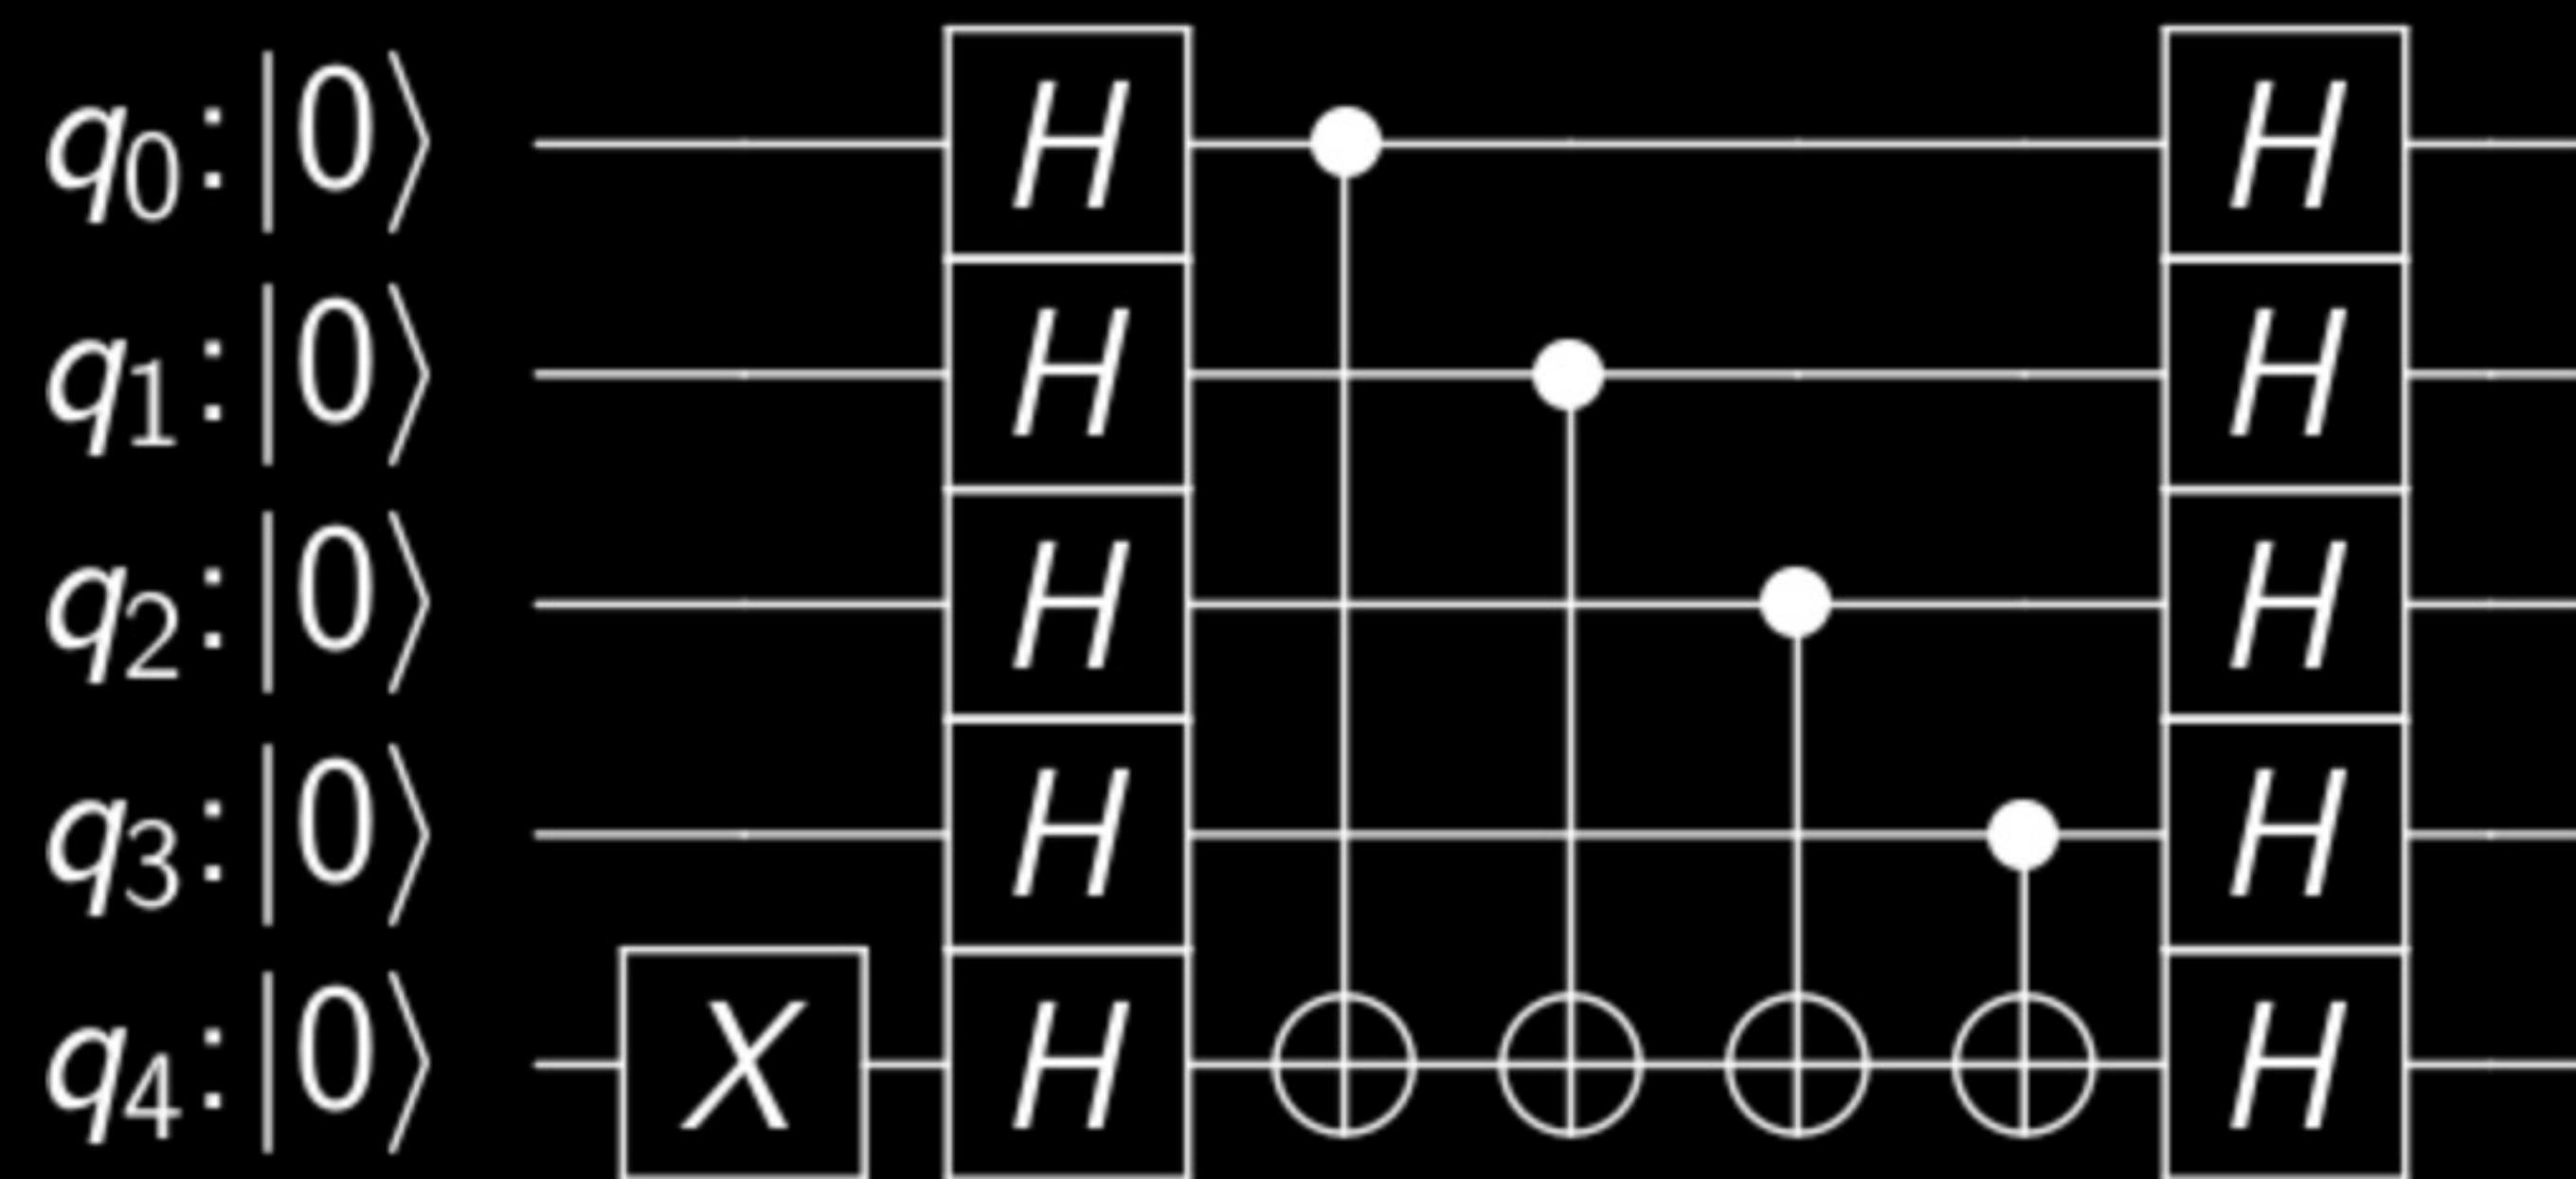
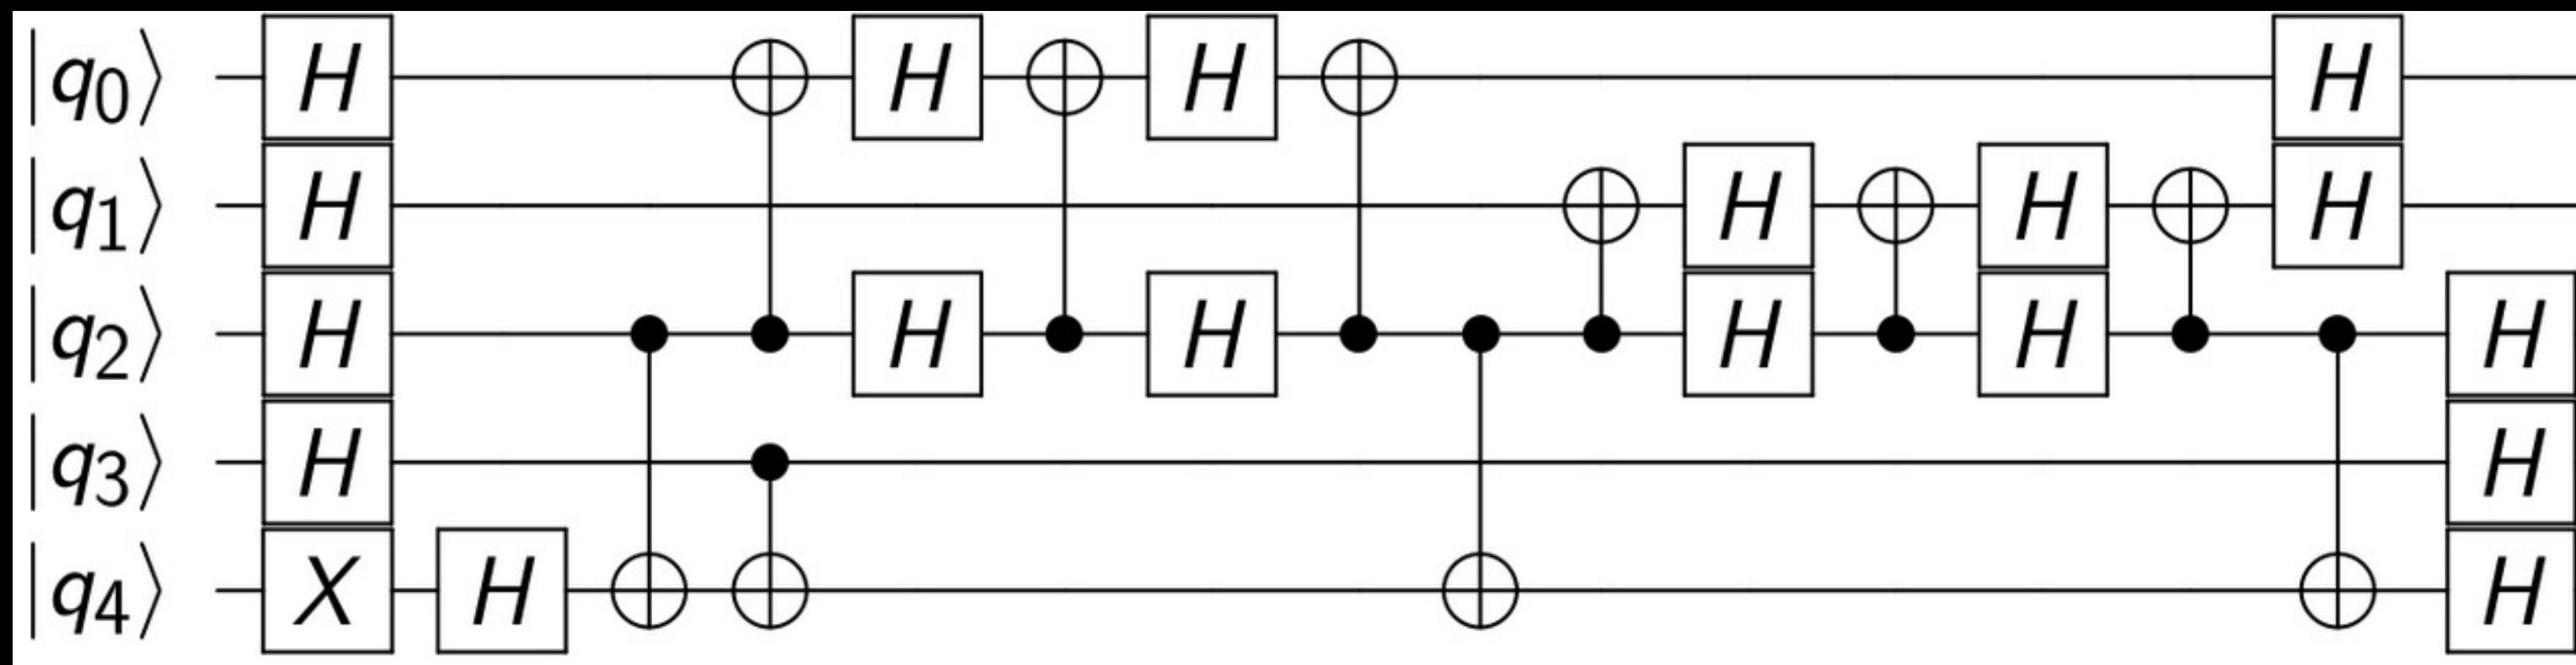
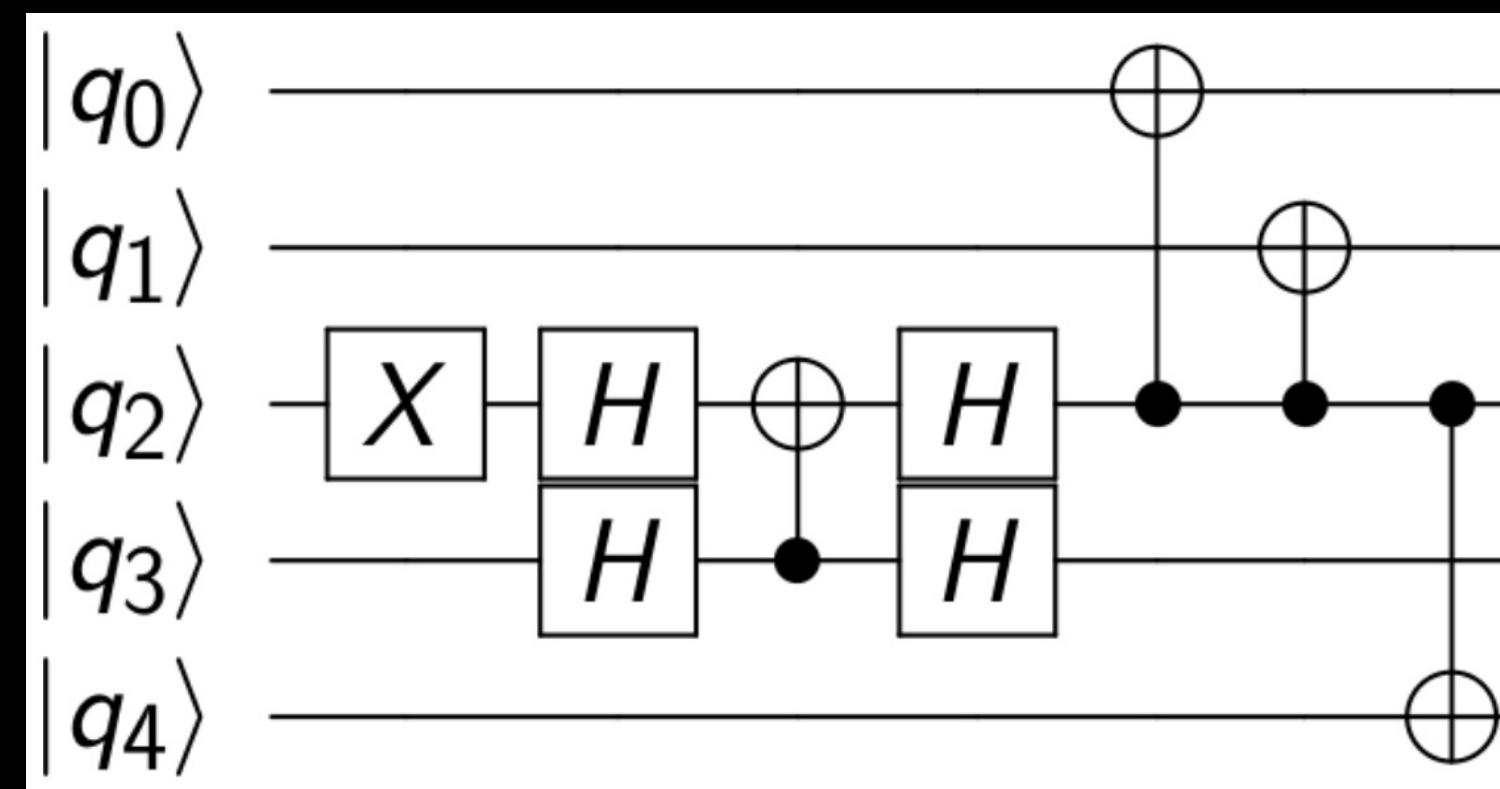
Optimizations

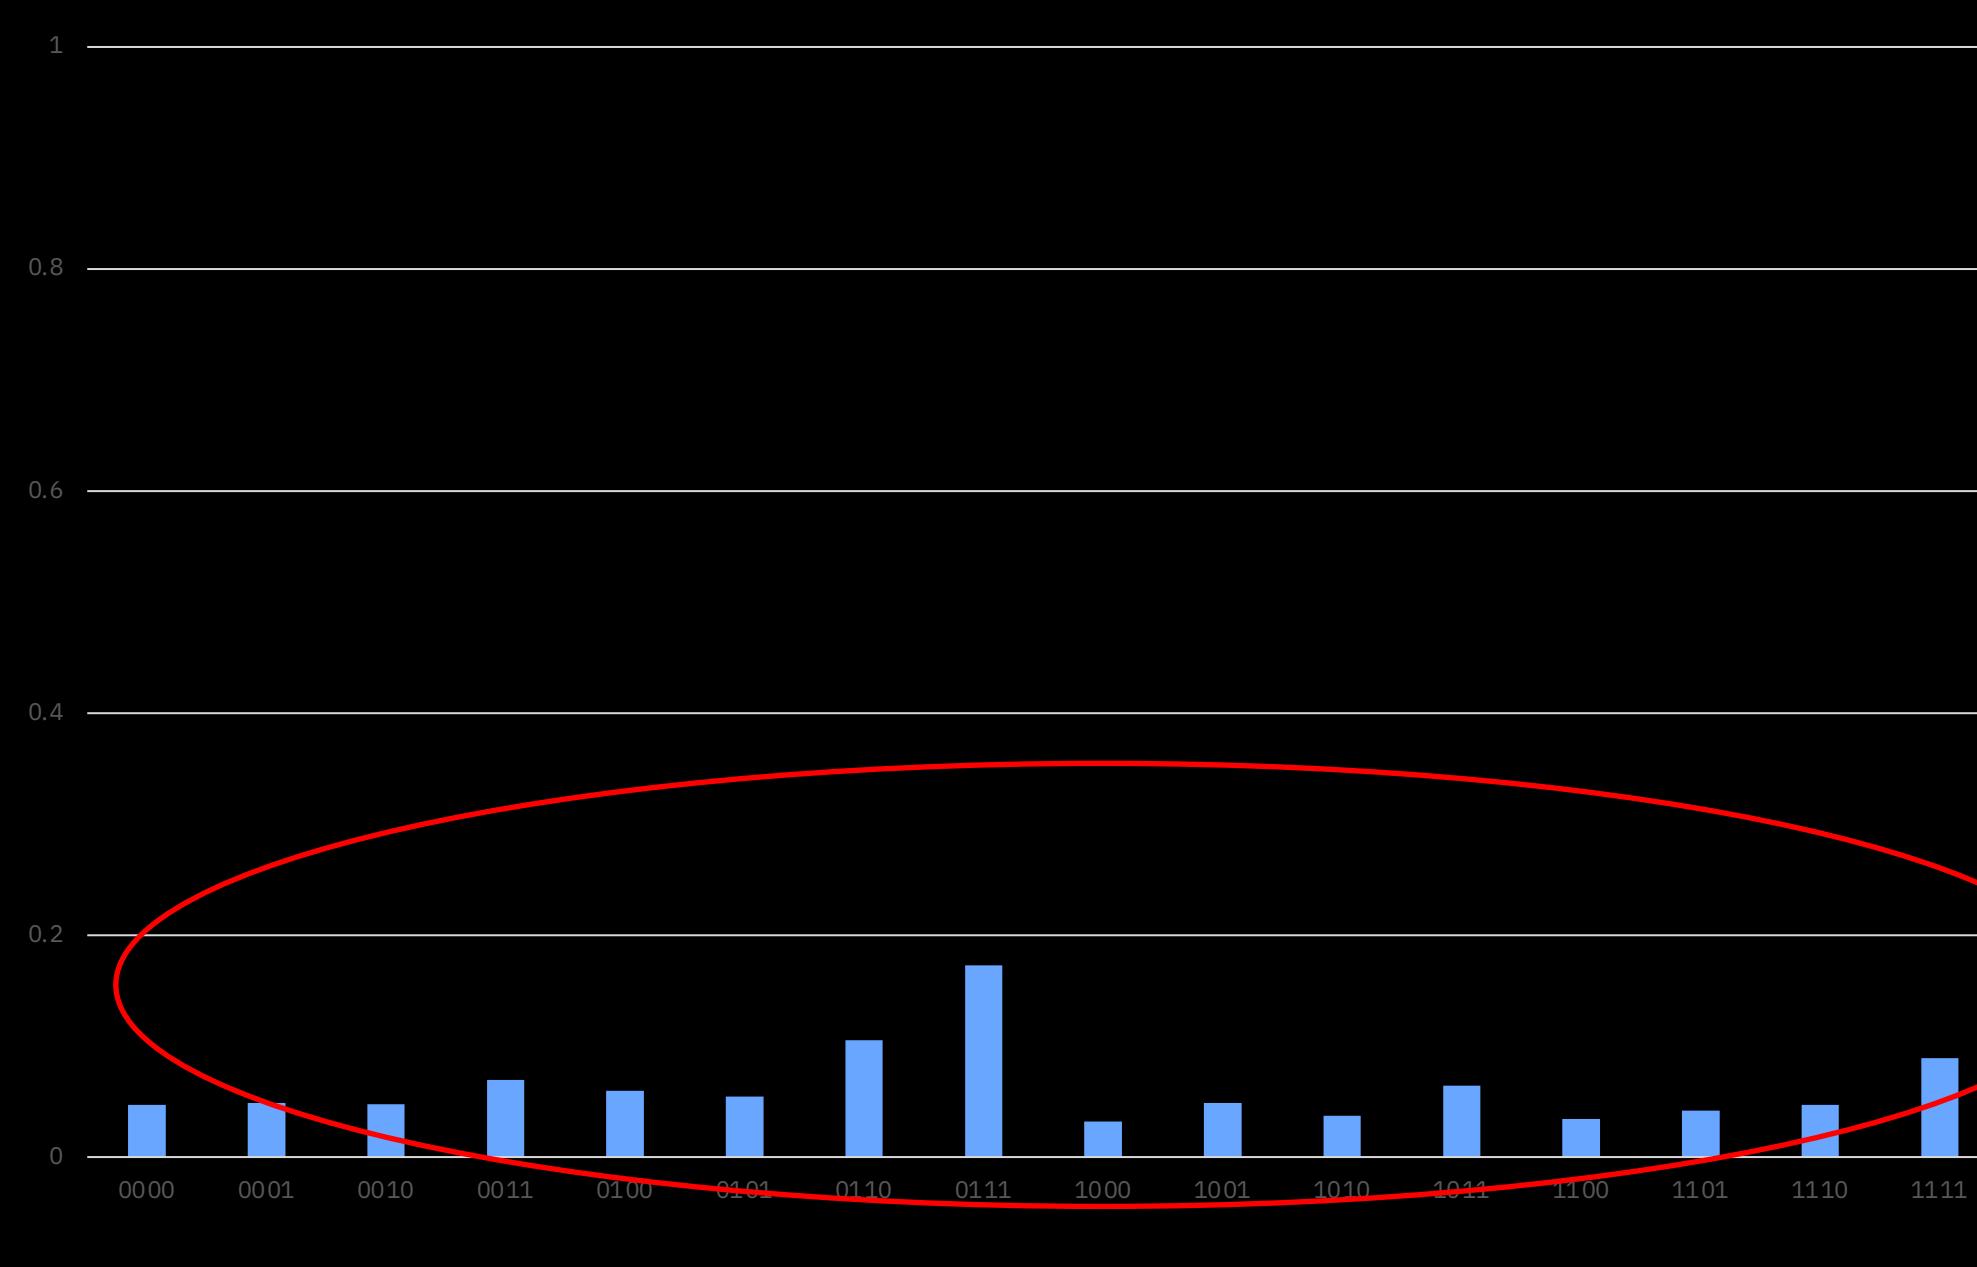
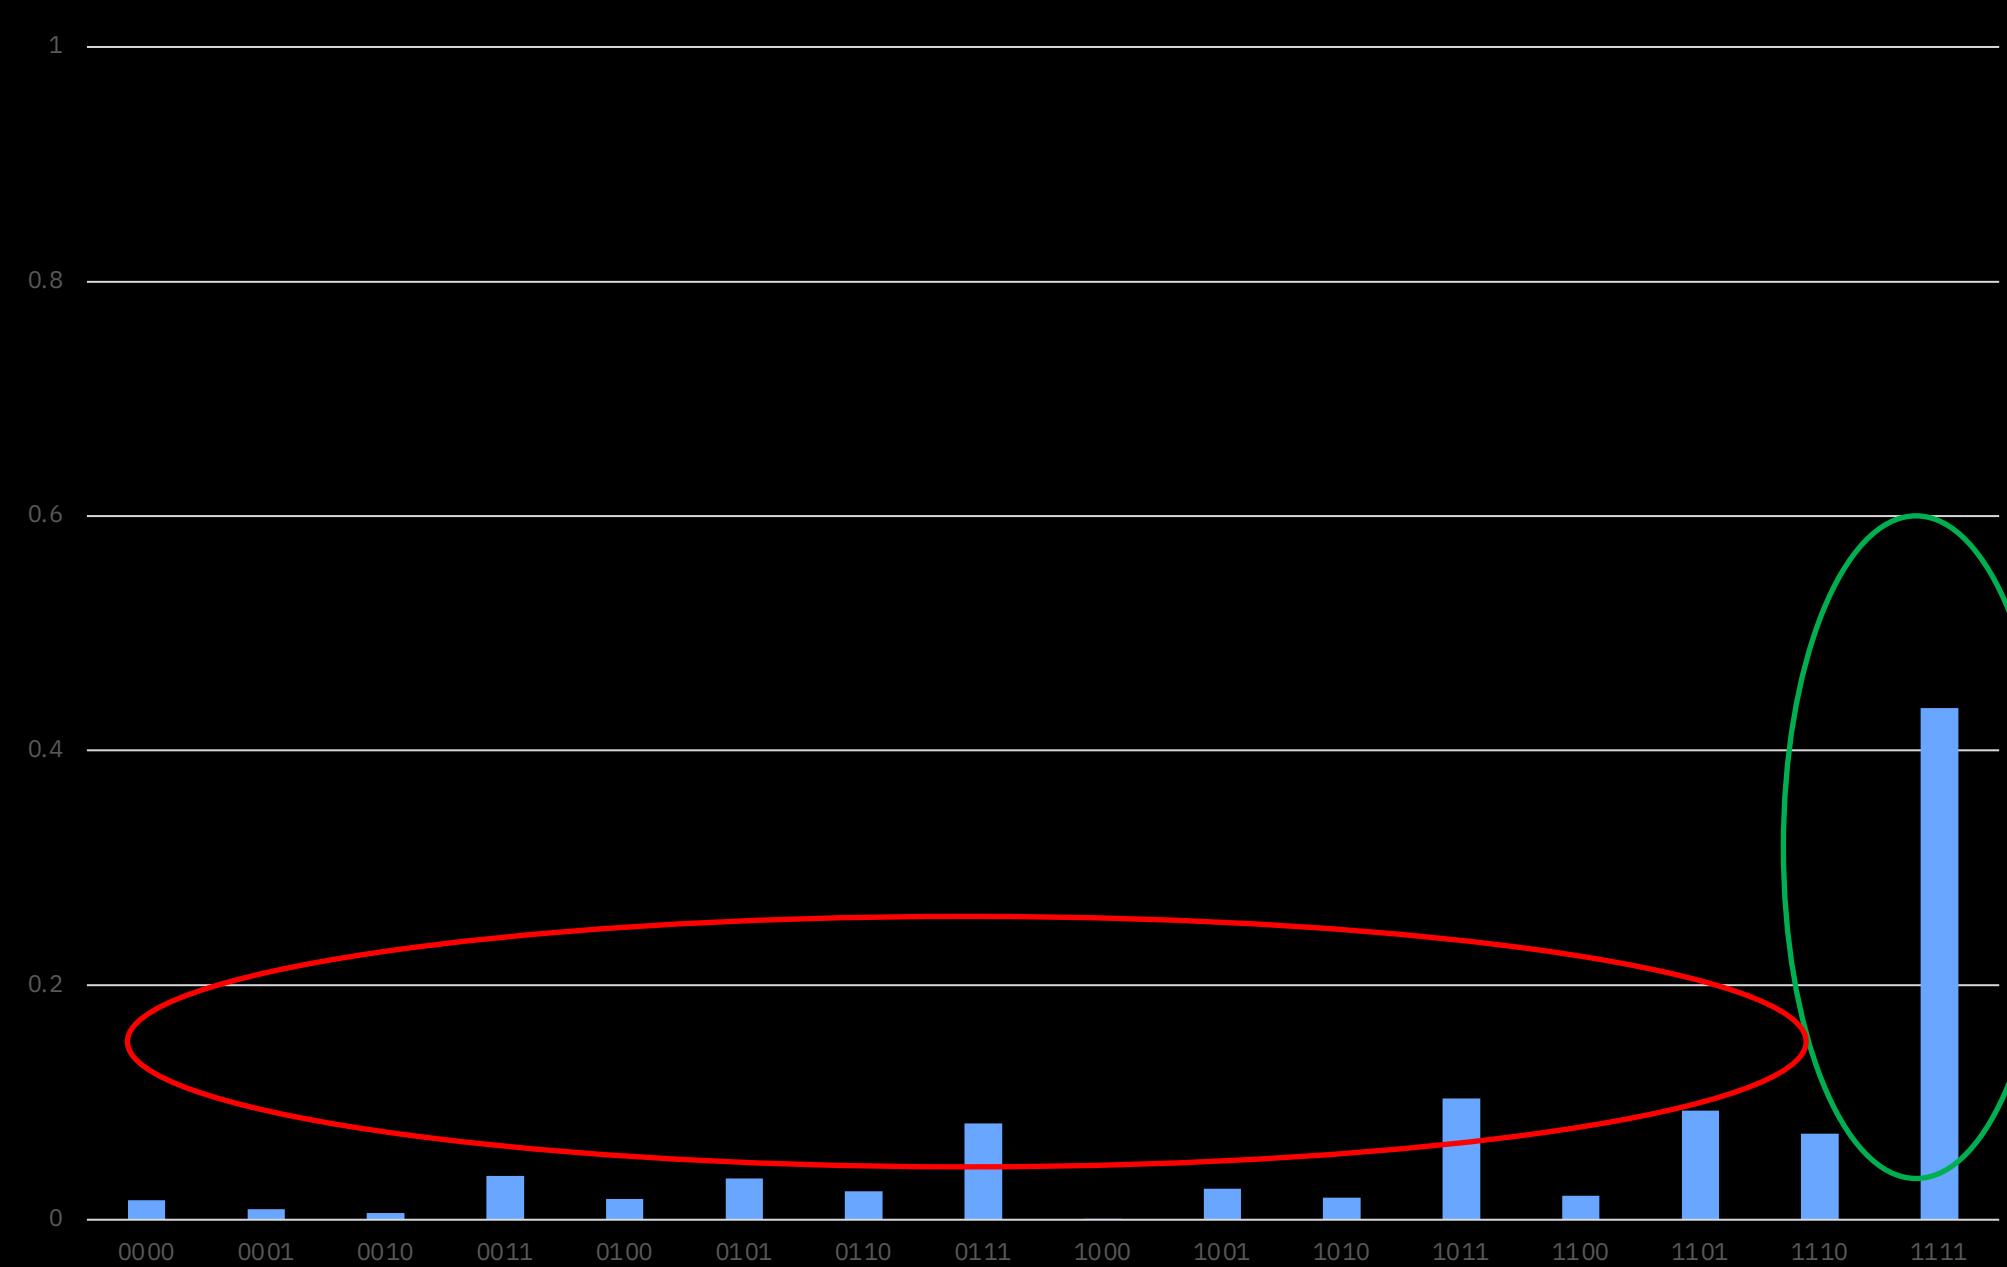
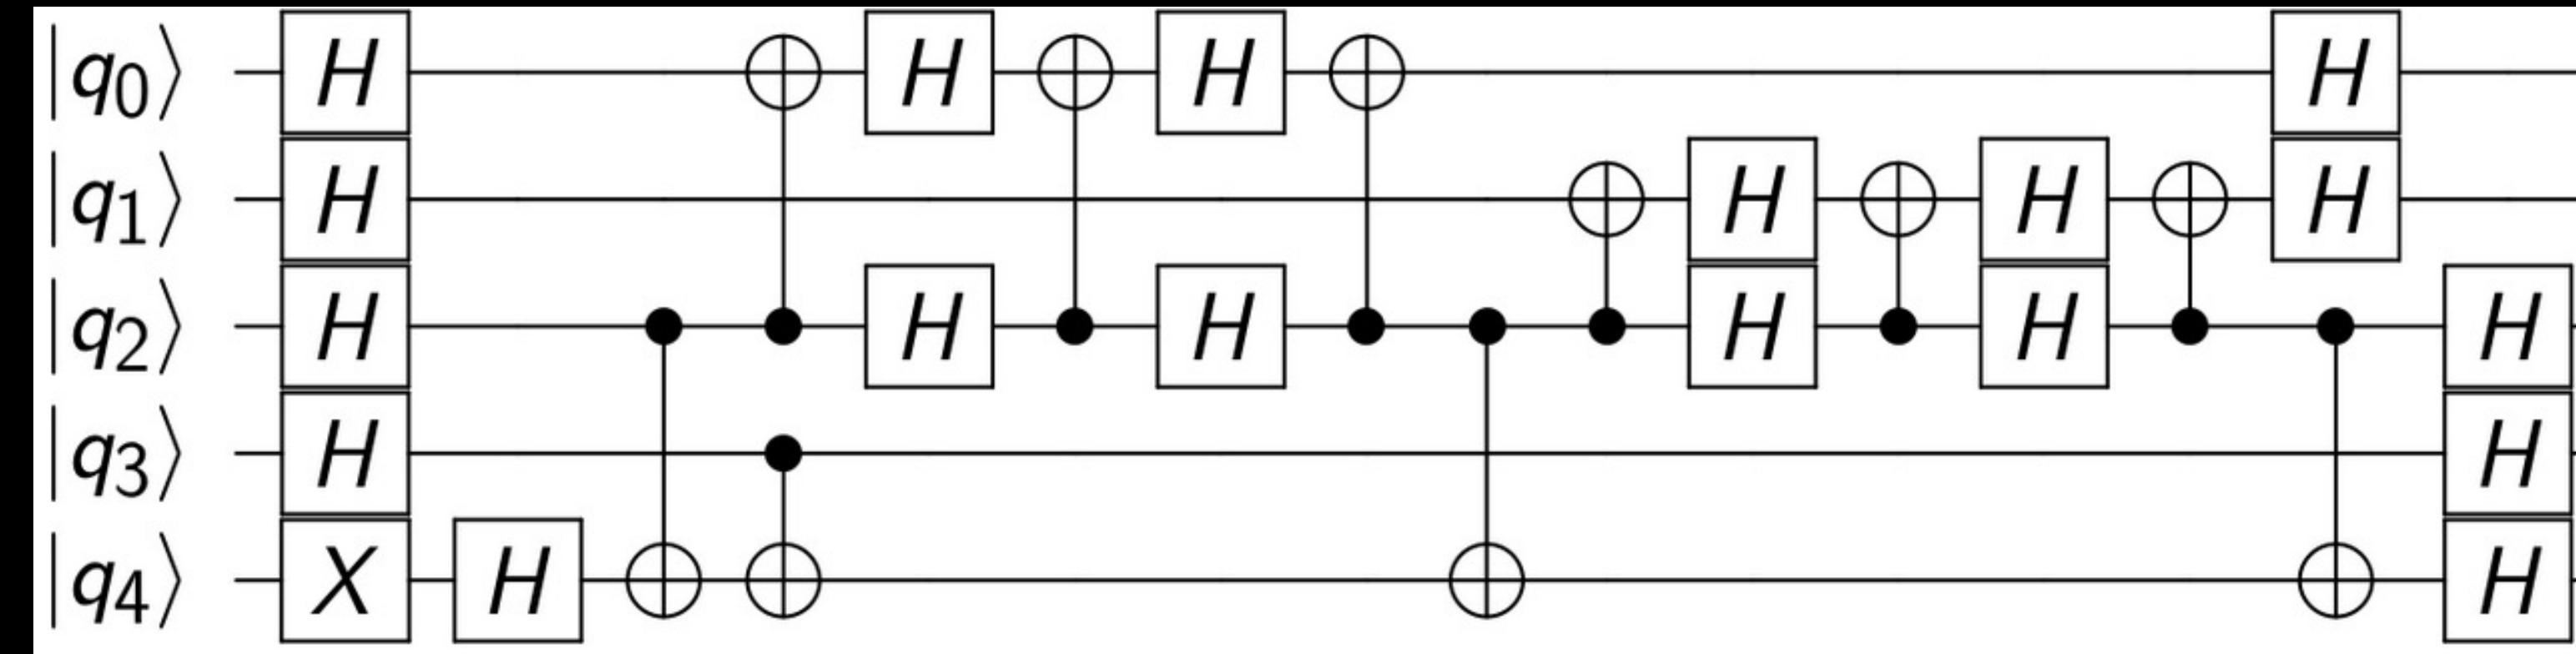
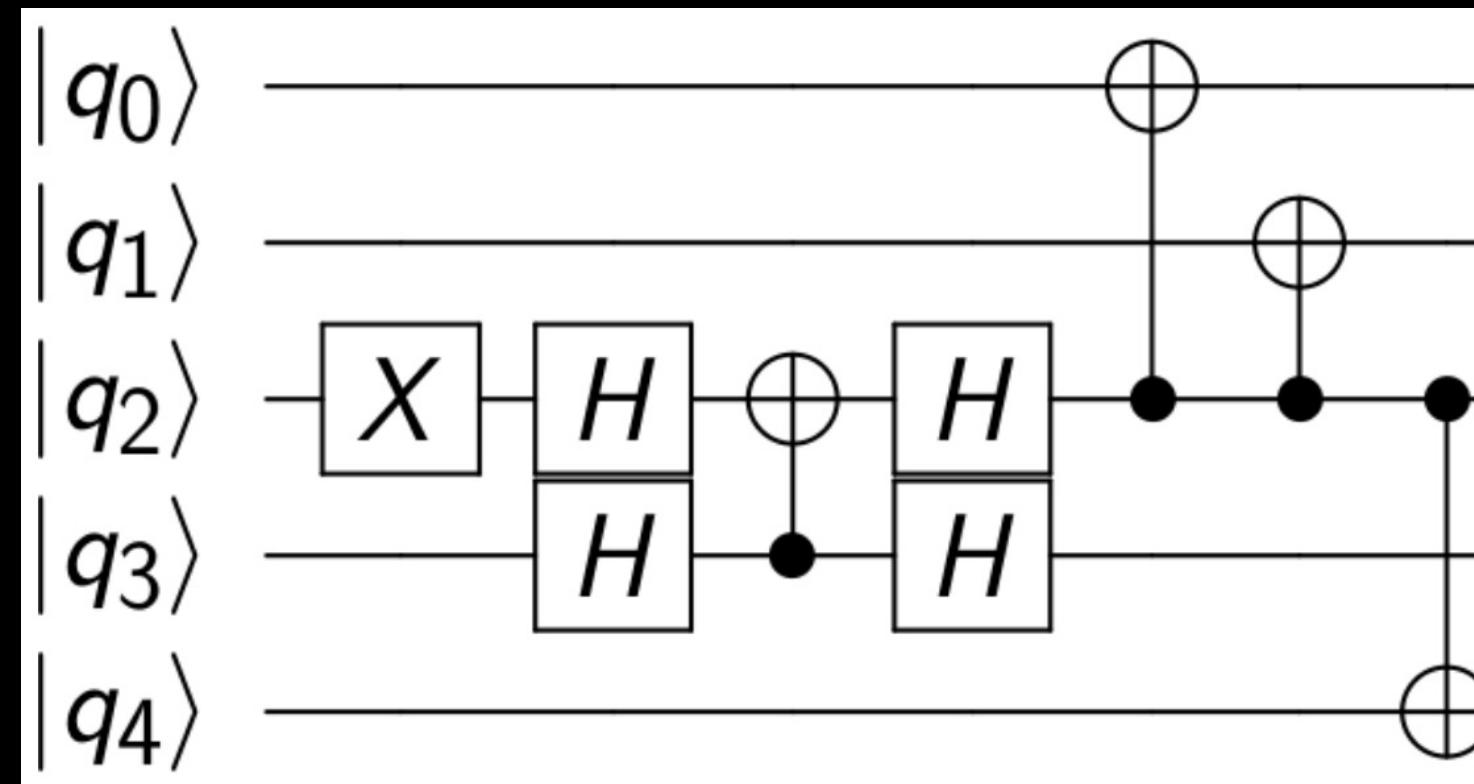
Instruction set

Hardware connectivity

The Routing Problem

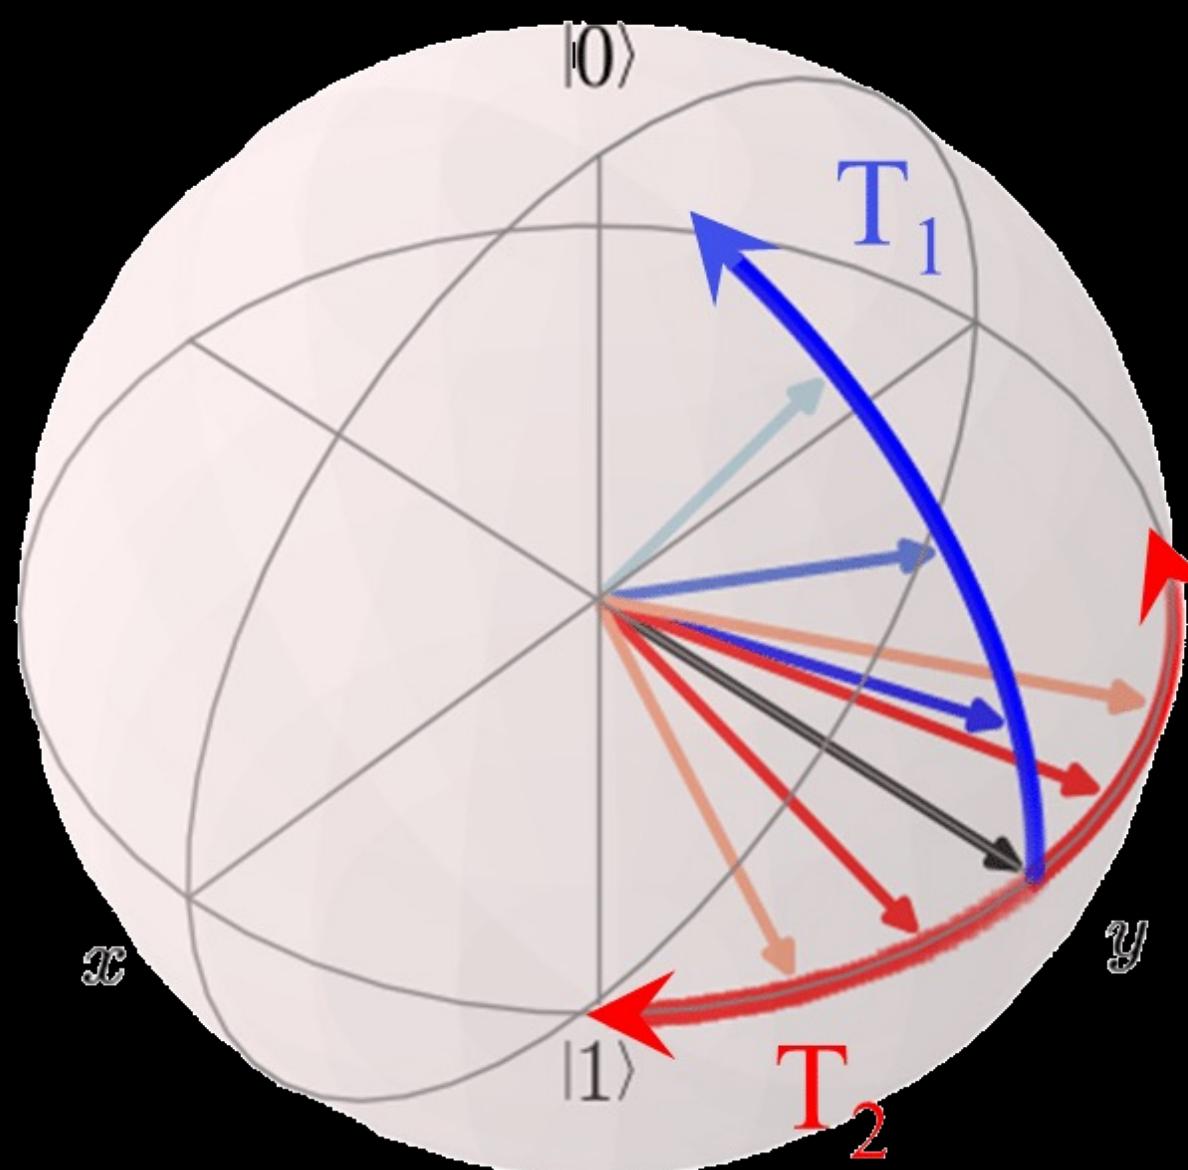
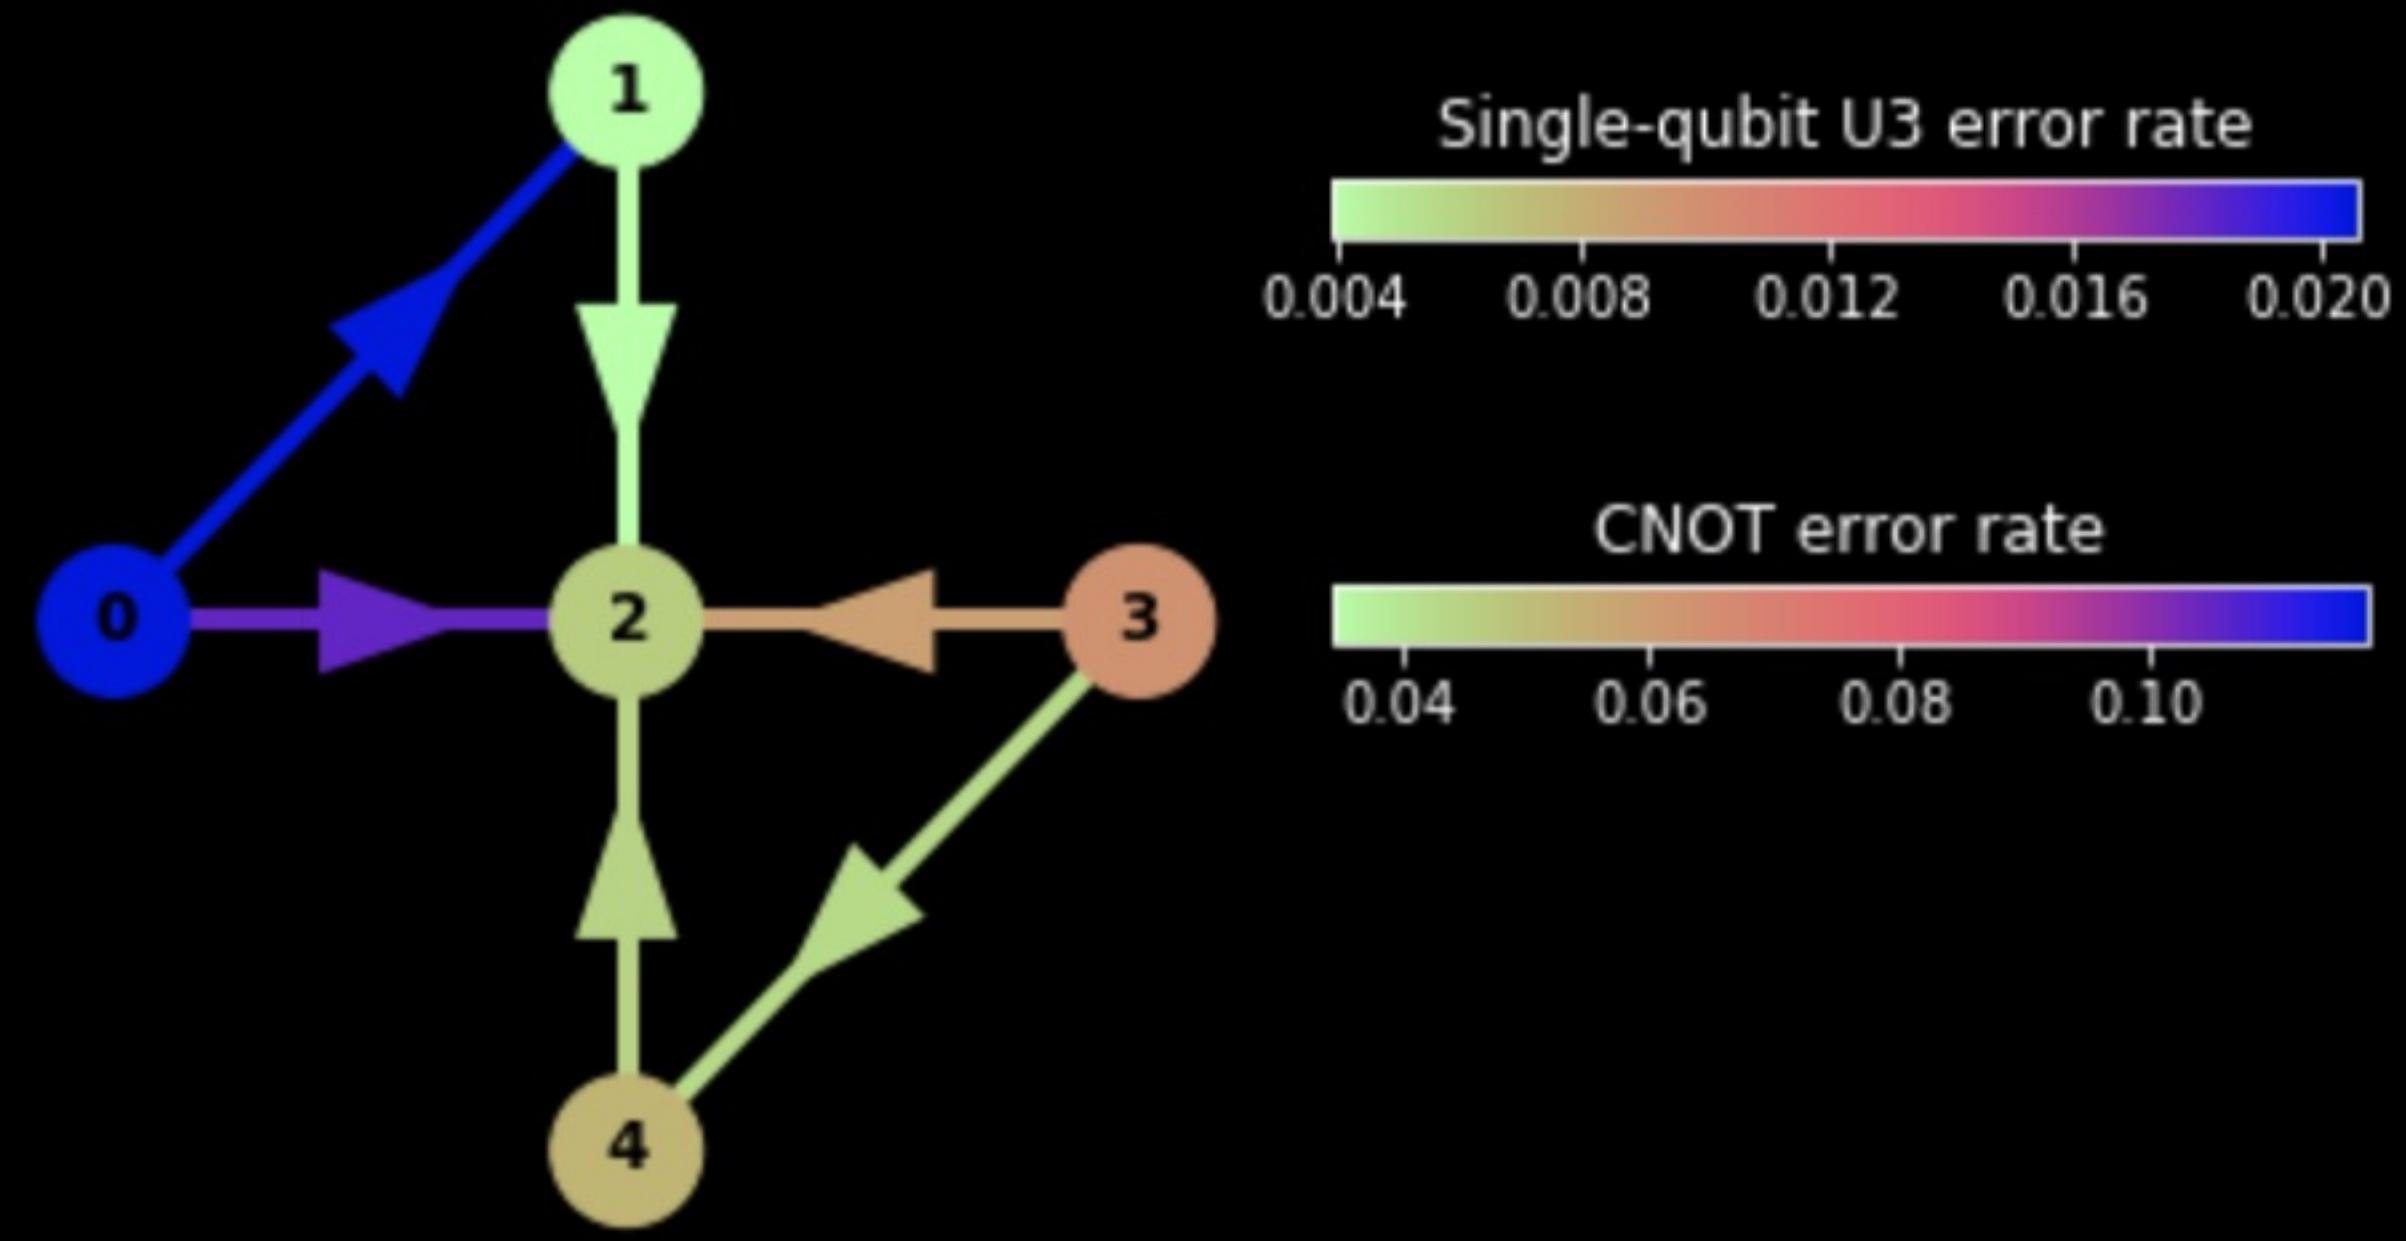






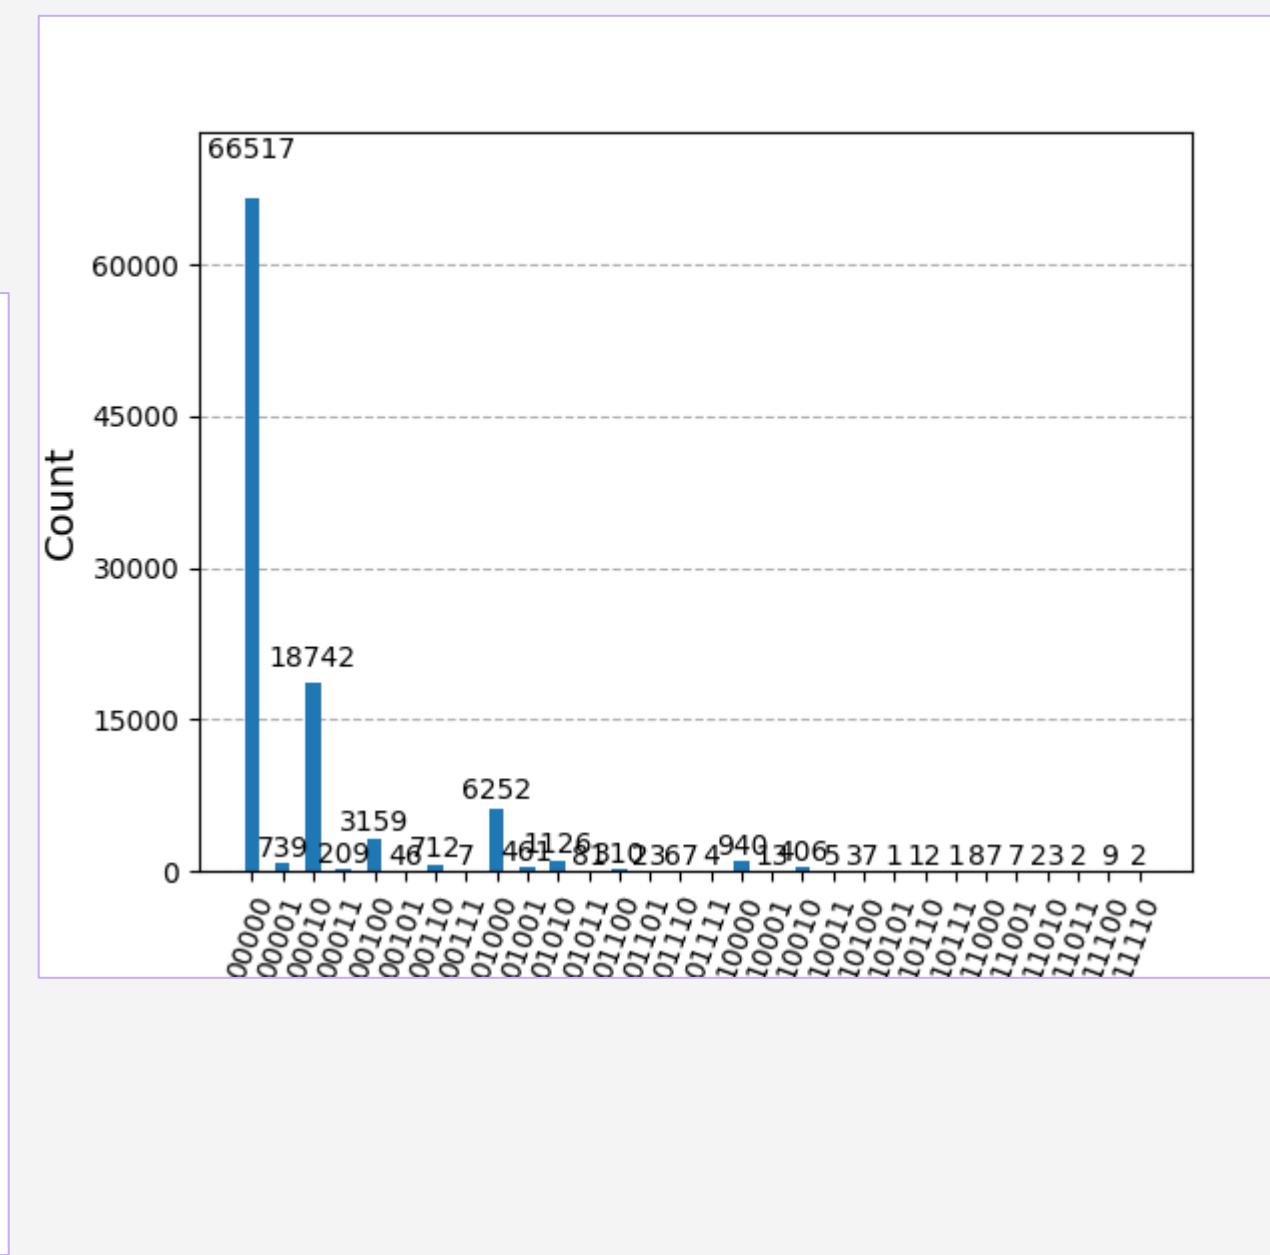
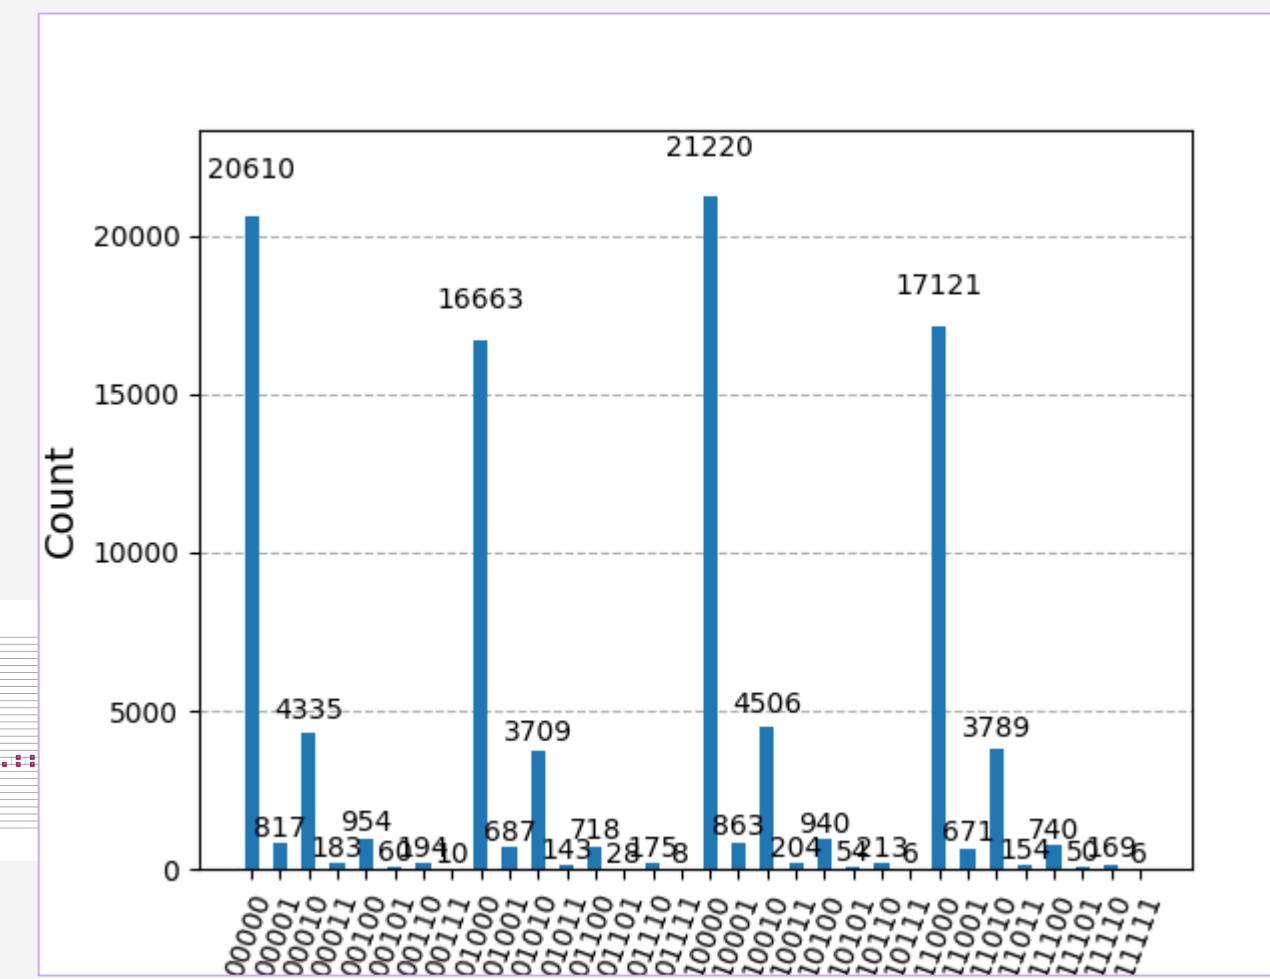
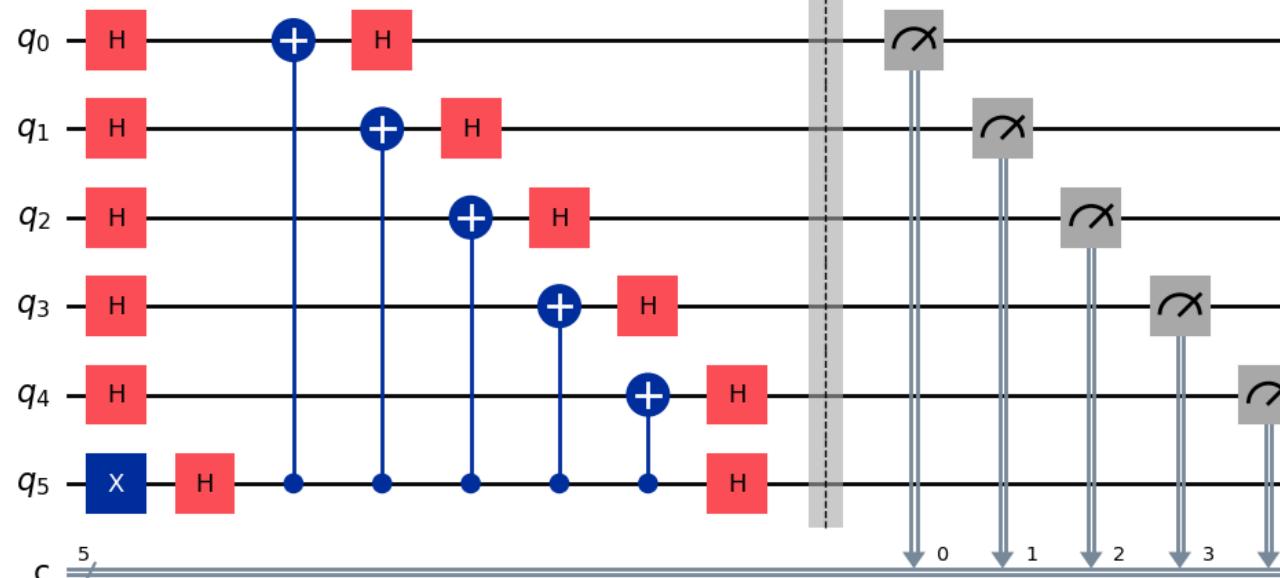
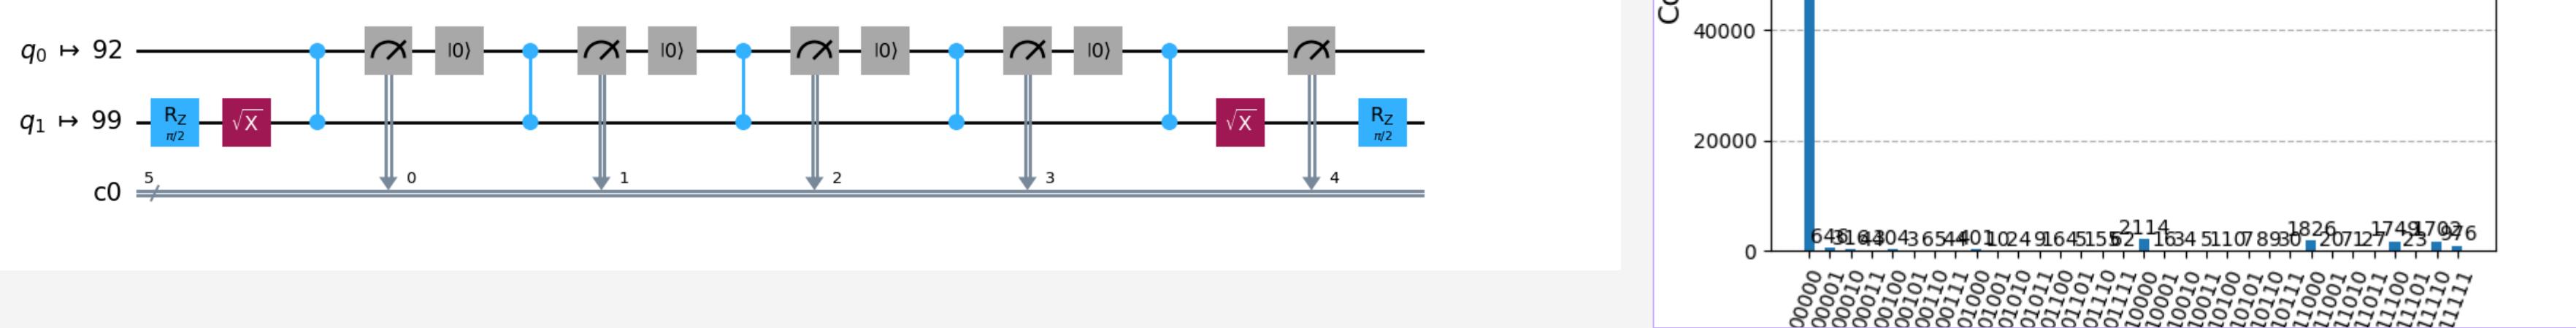
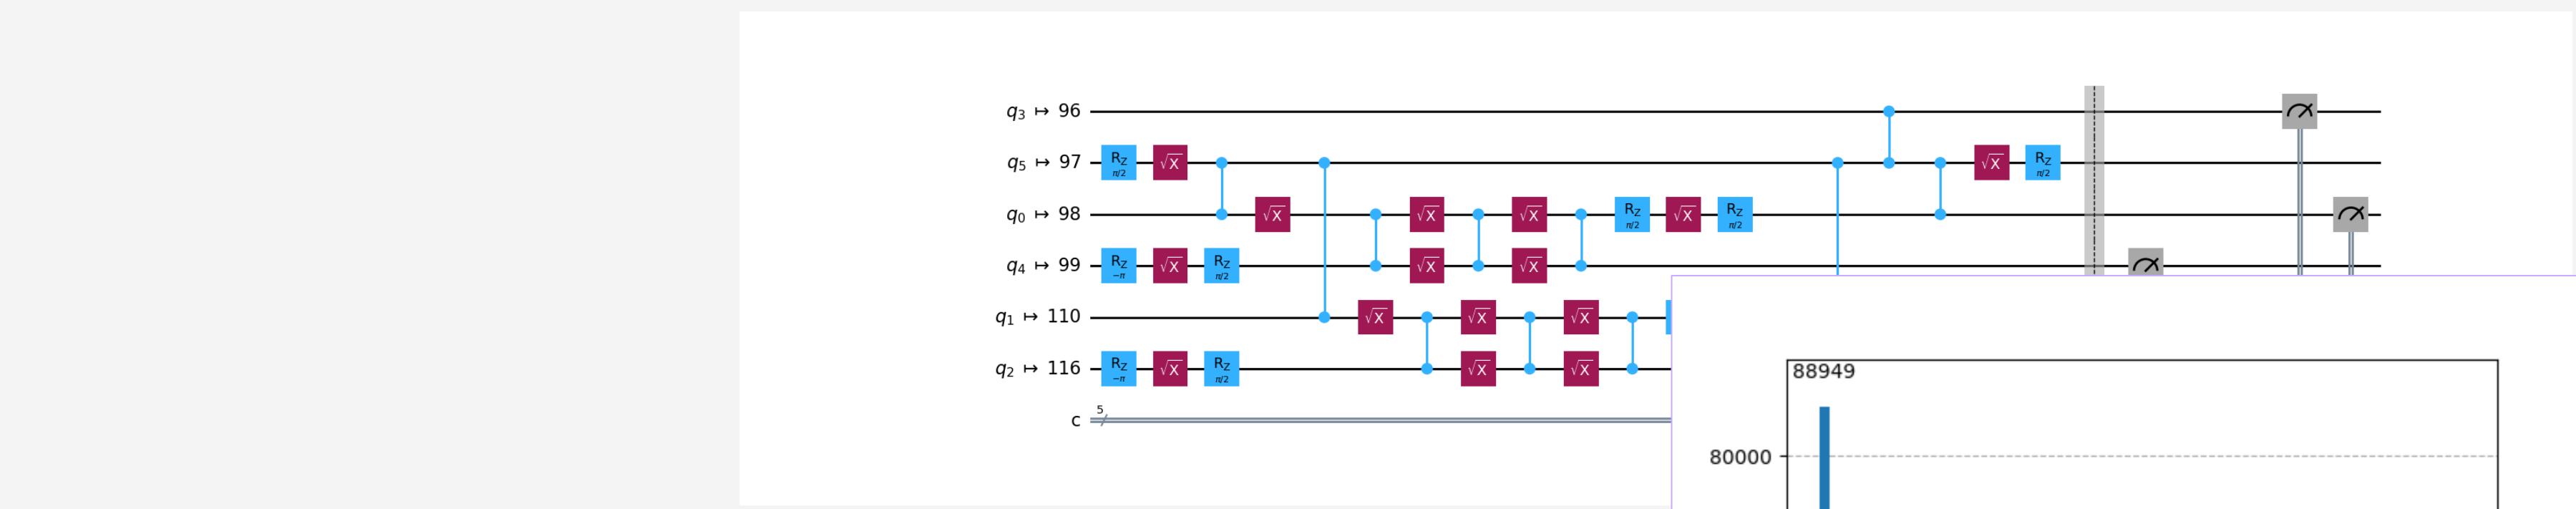
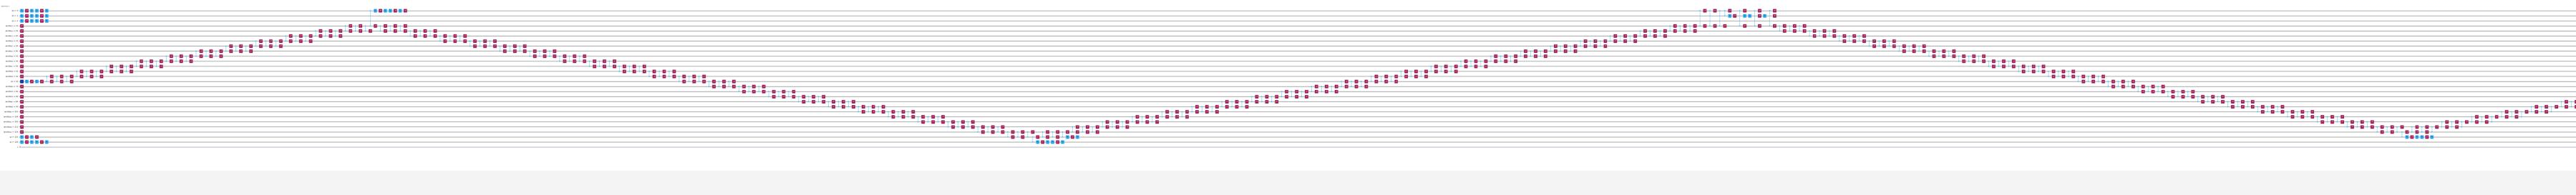
Noise everywhere!

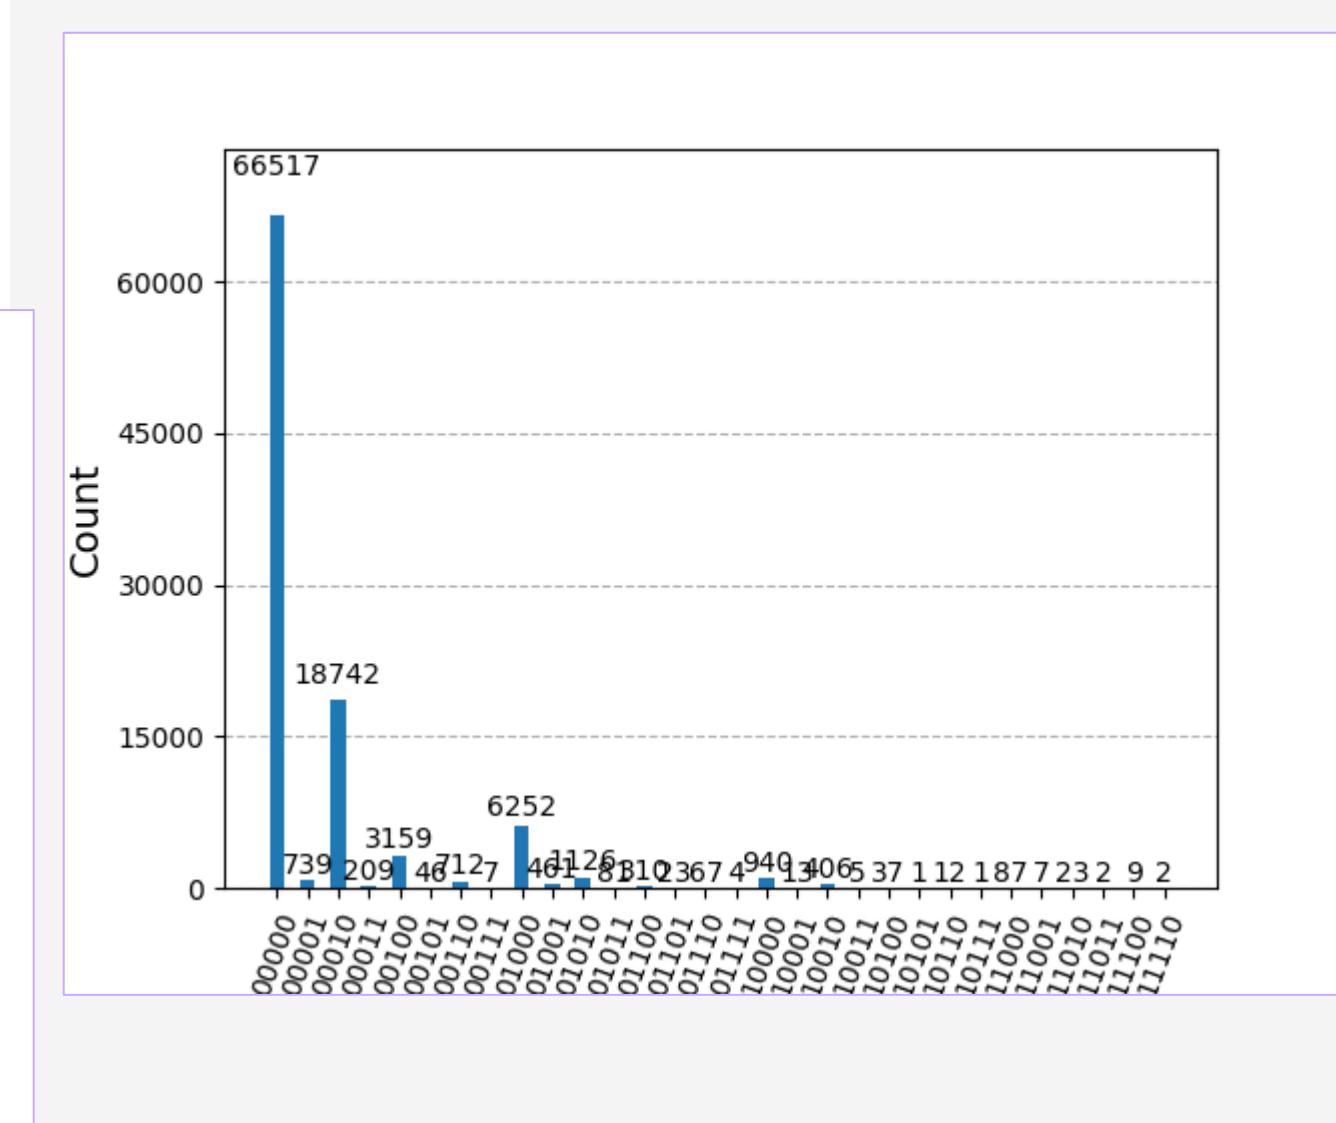
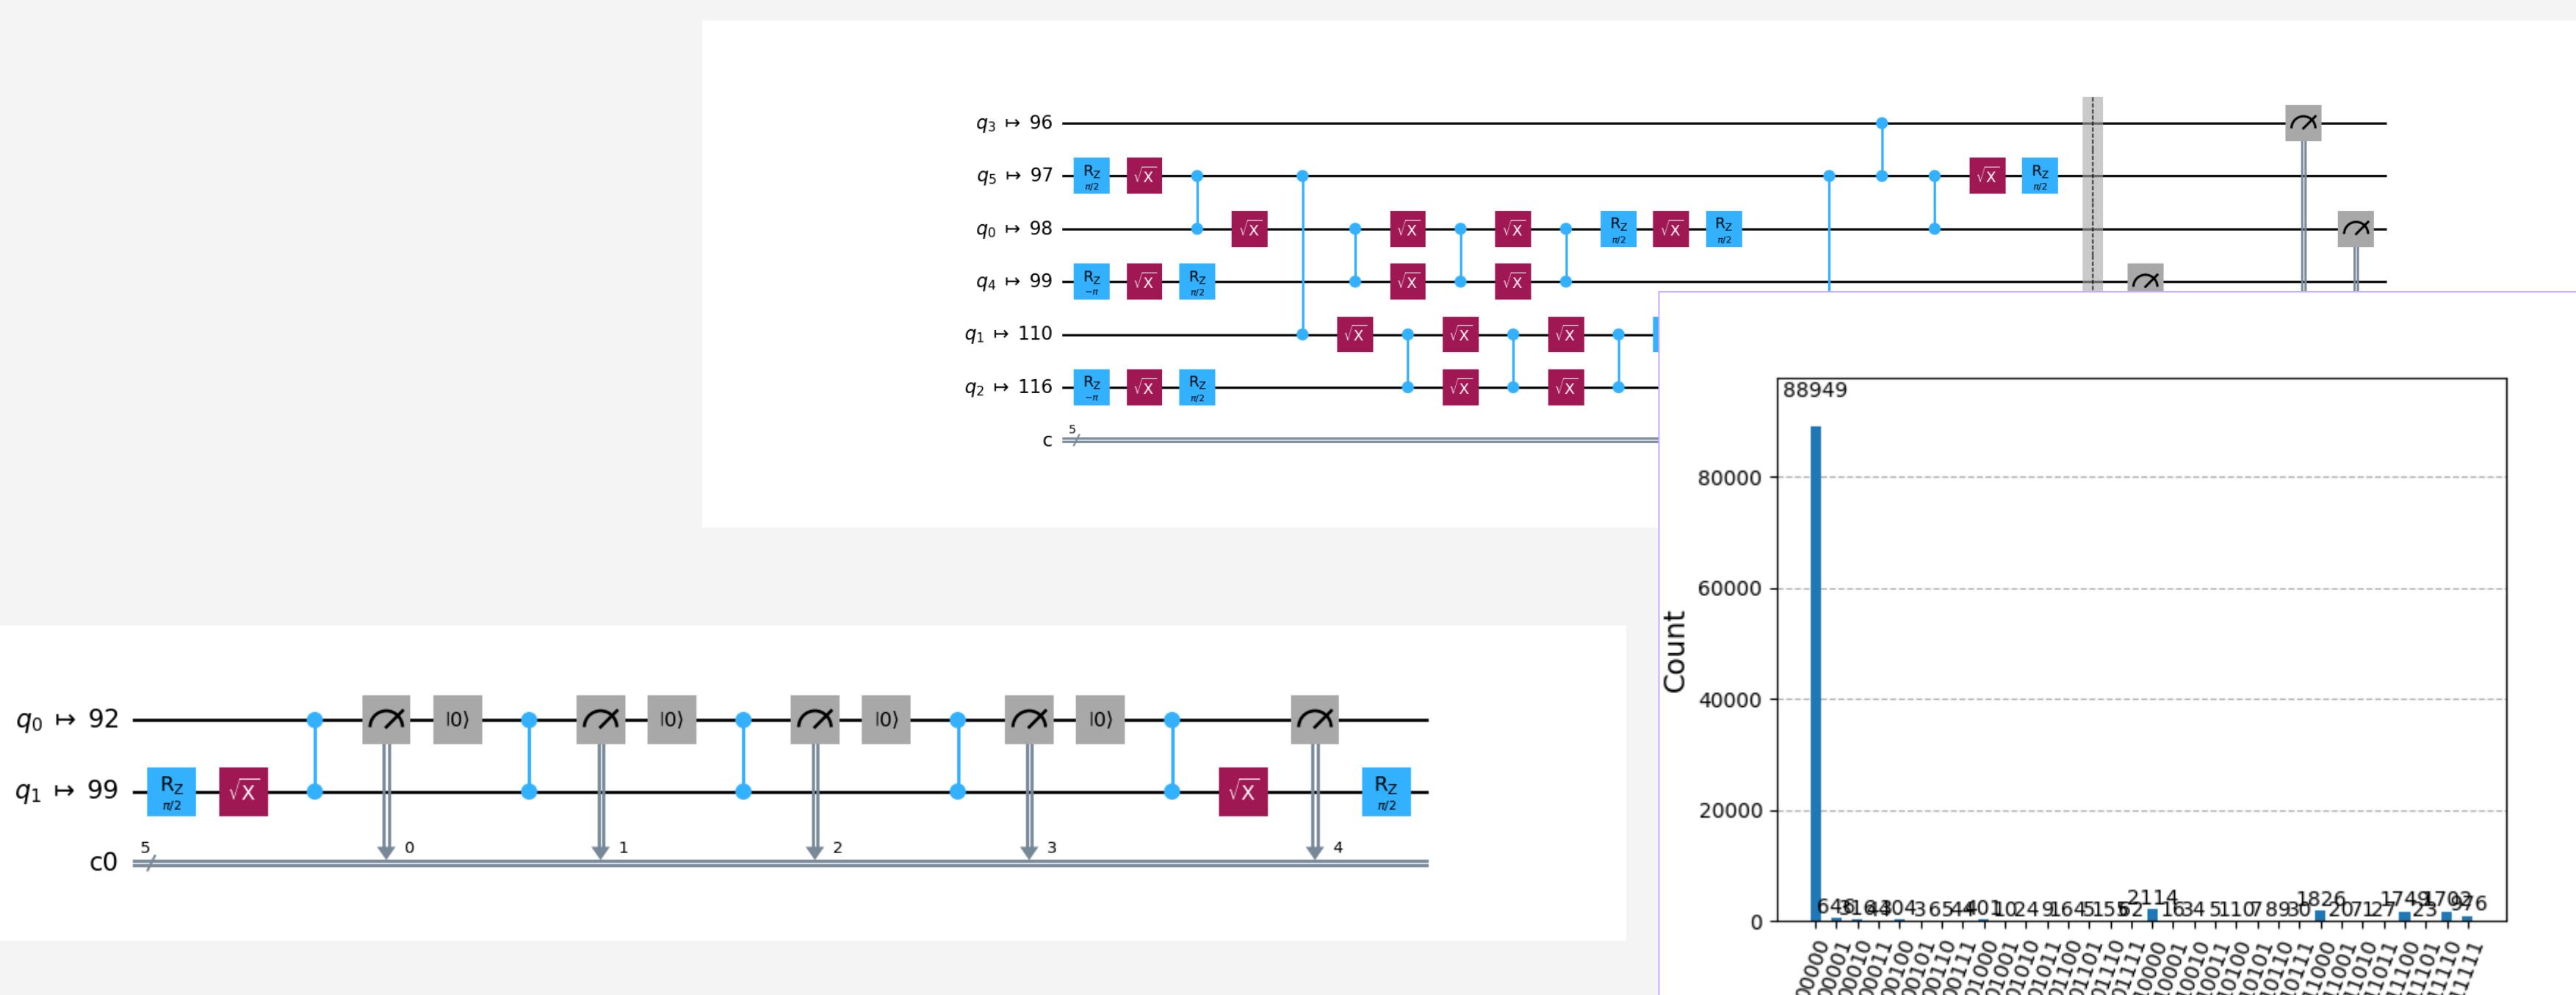
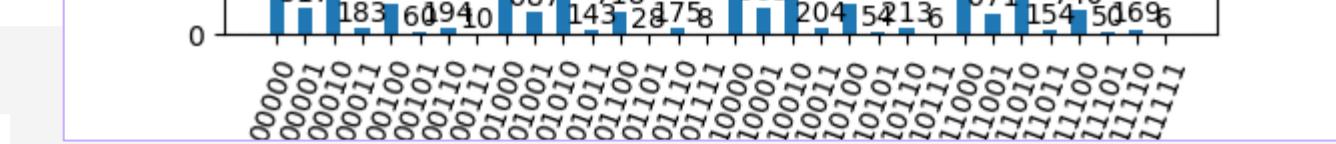
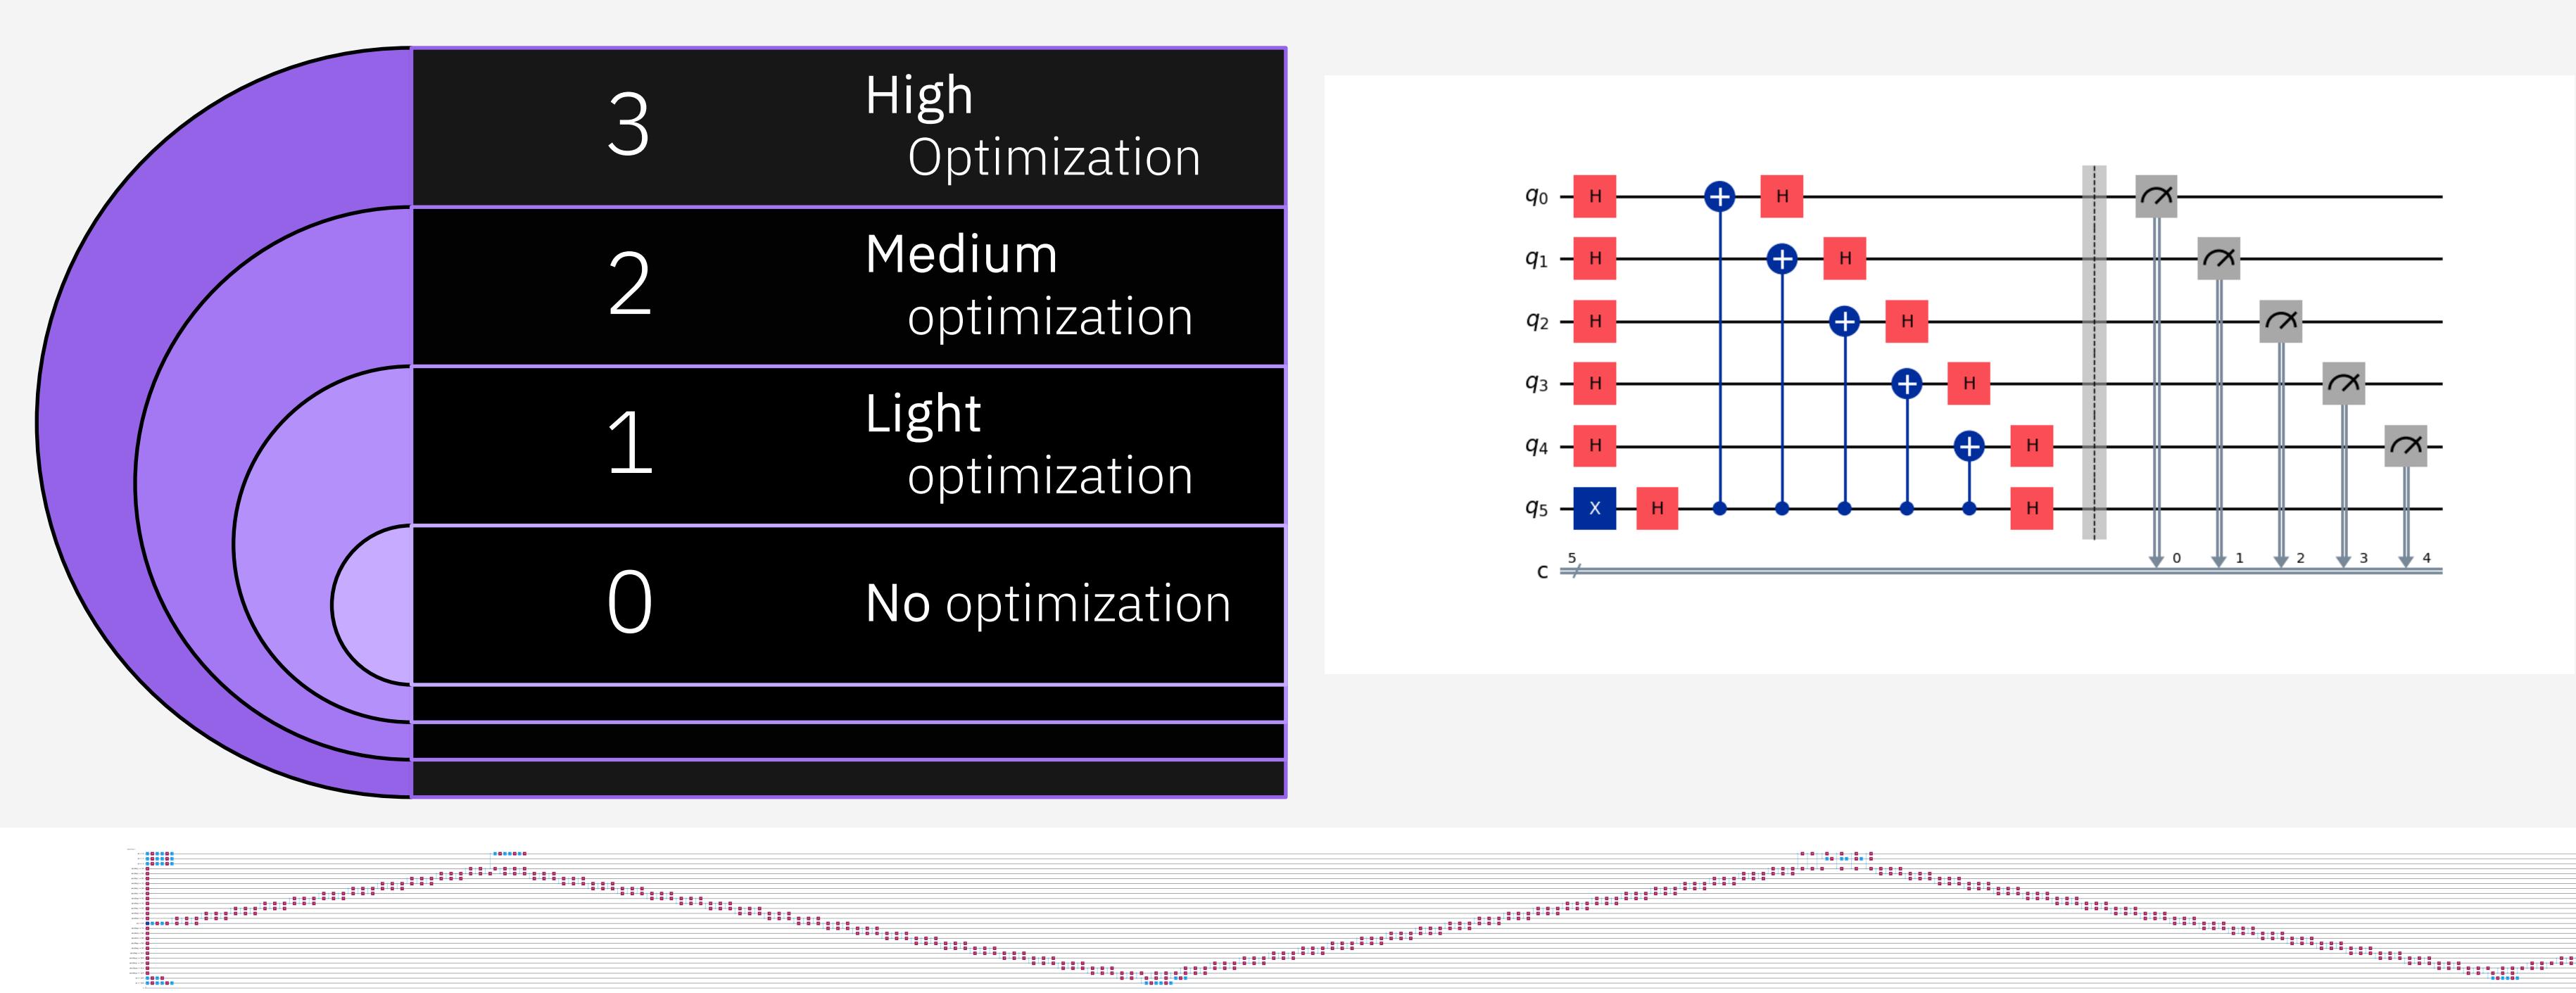
- Gate error
 - Single qubit gate
 - Multiple qubit gates
- Readout error
- Decoherence time

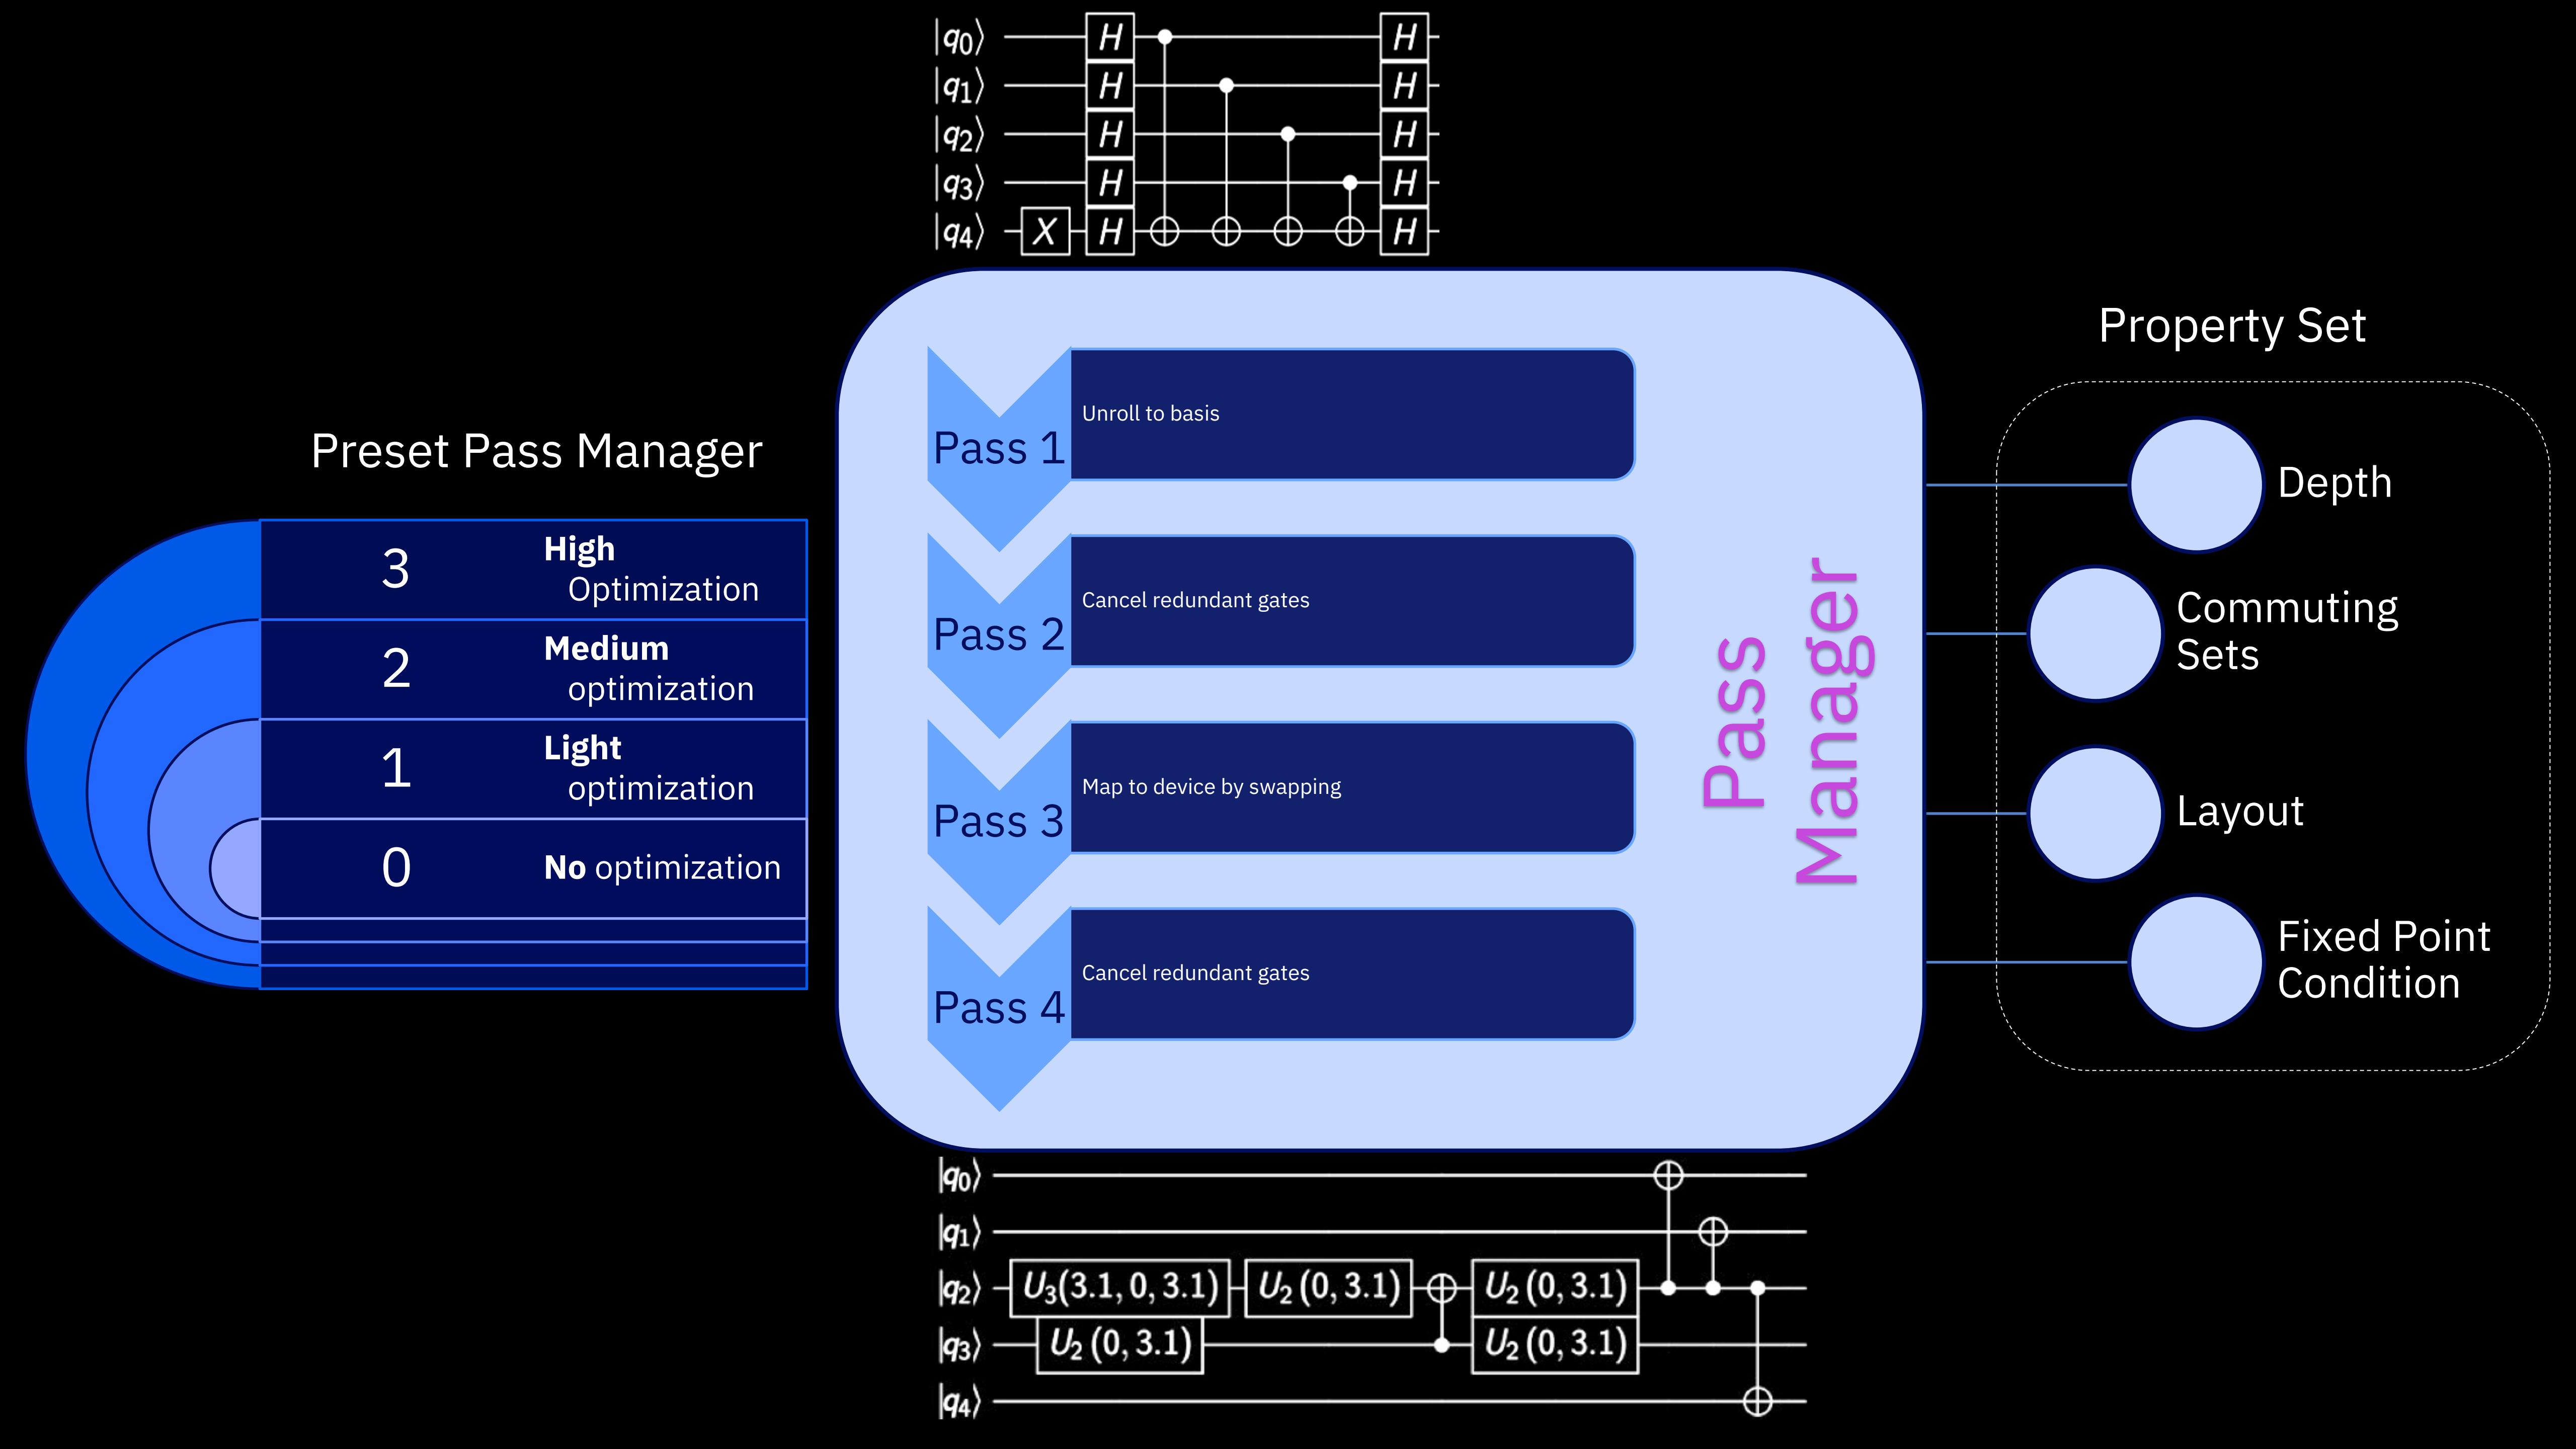


Why compiling?

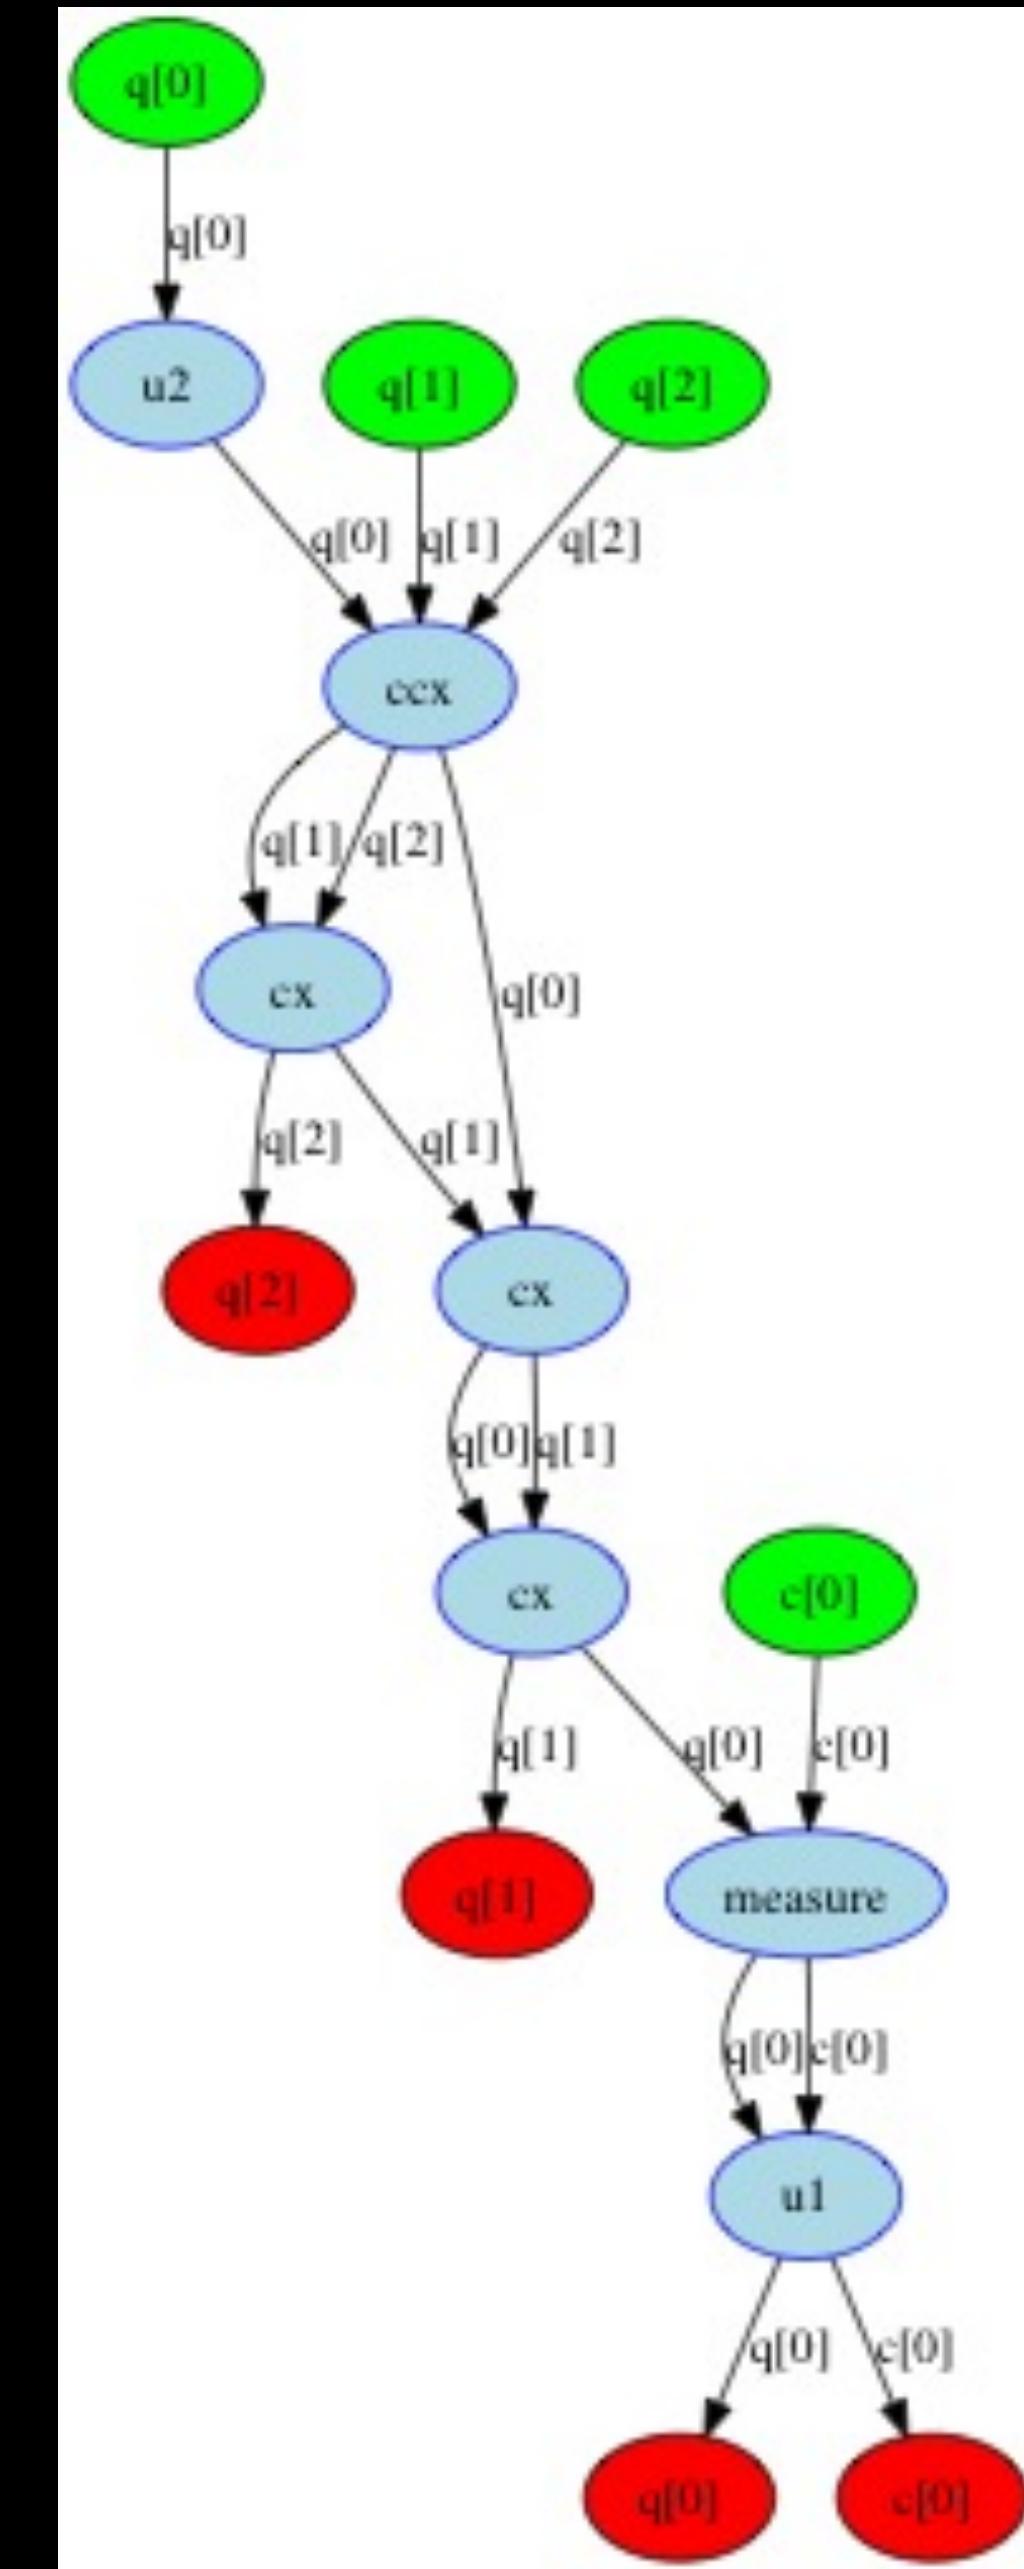
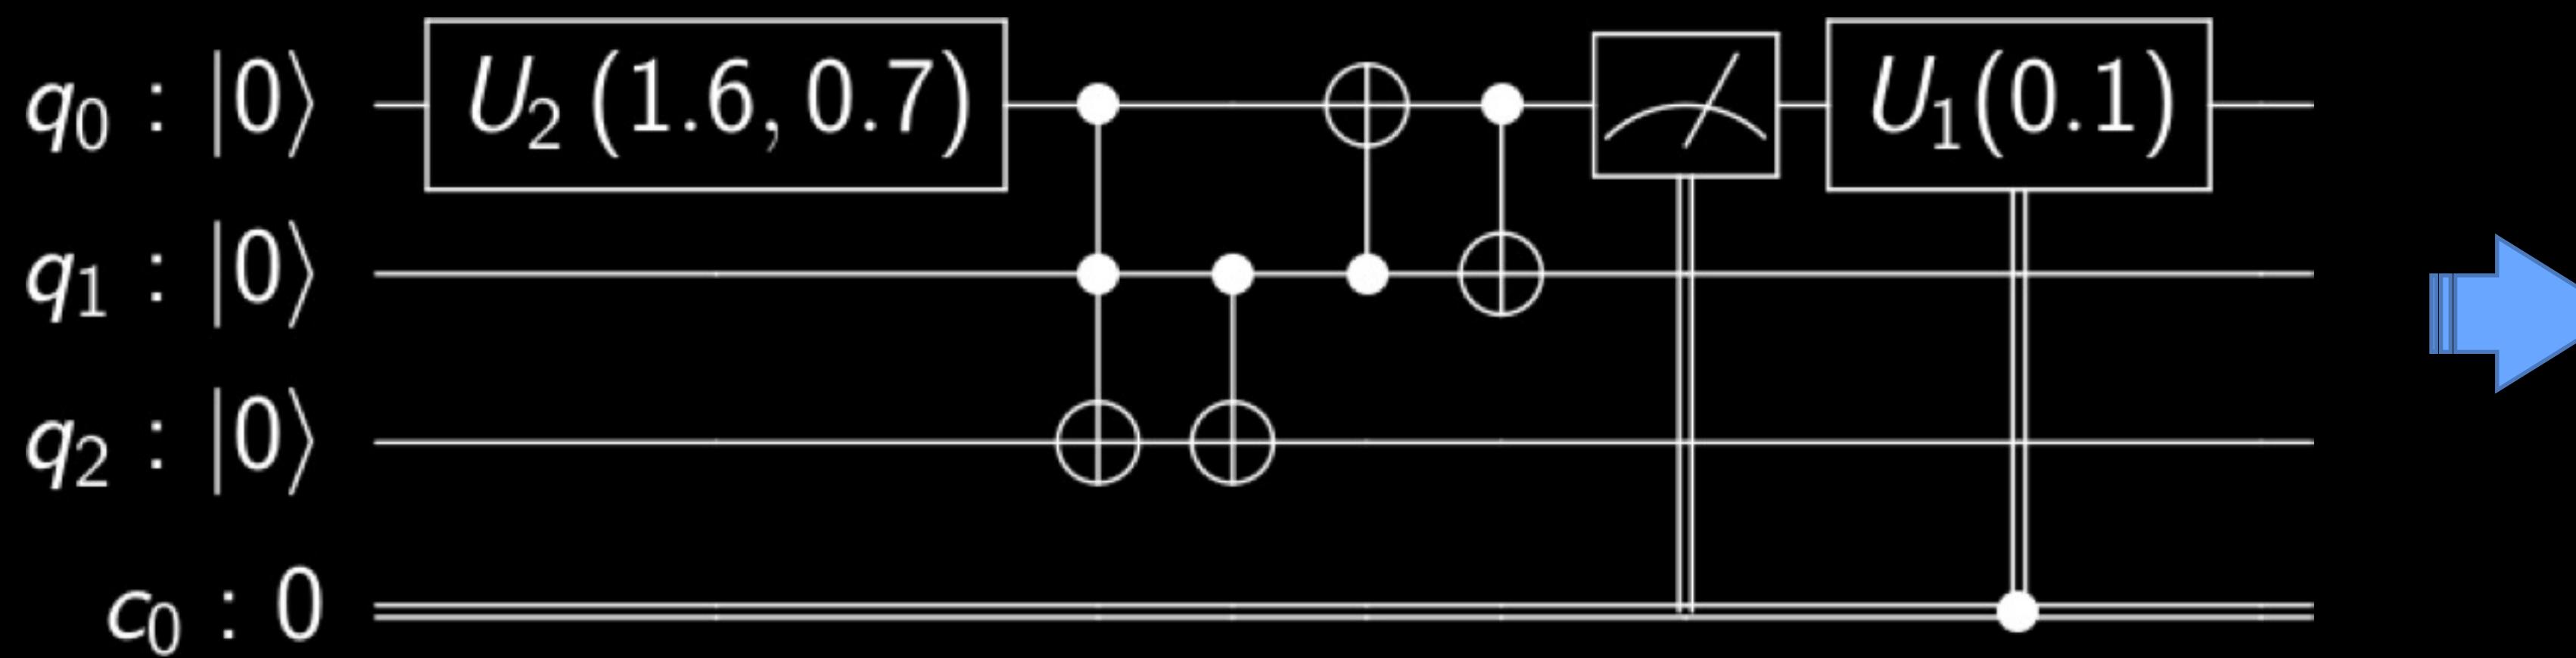
- Optimizations
- Instruction set
- Hardware connectivity
- Noise awareness



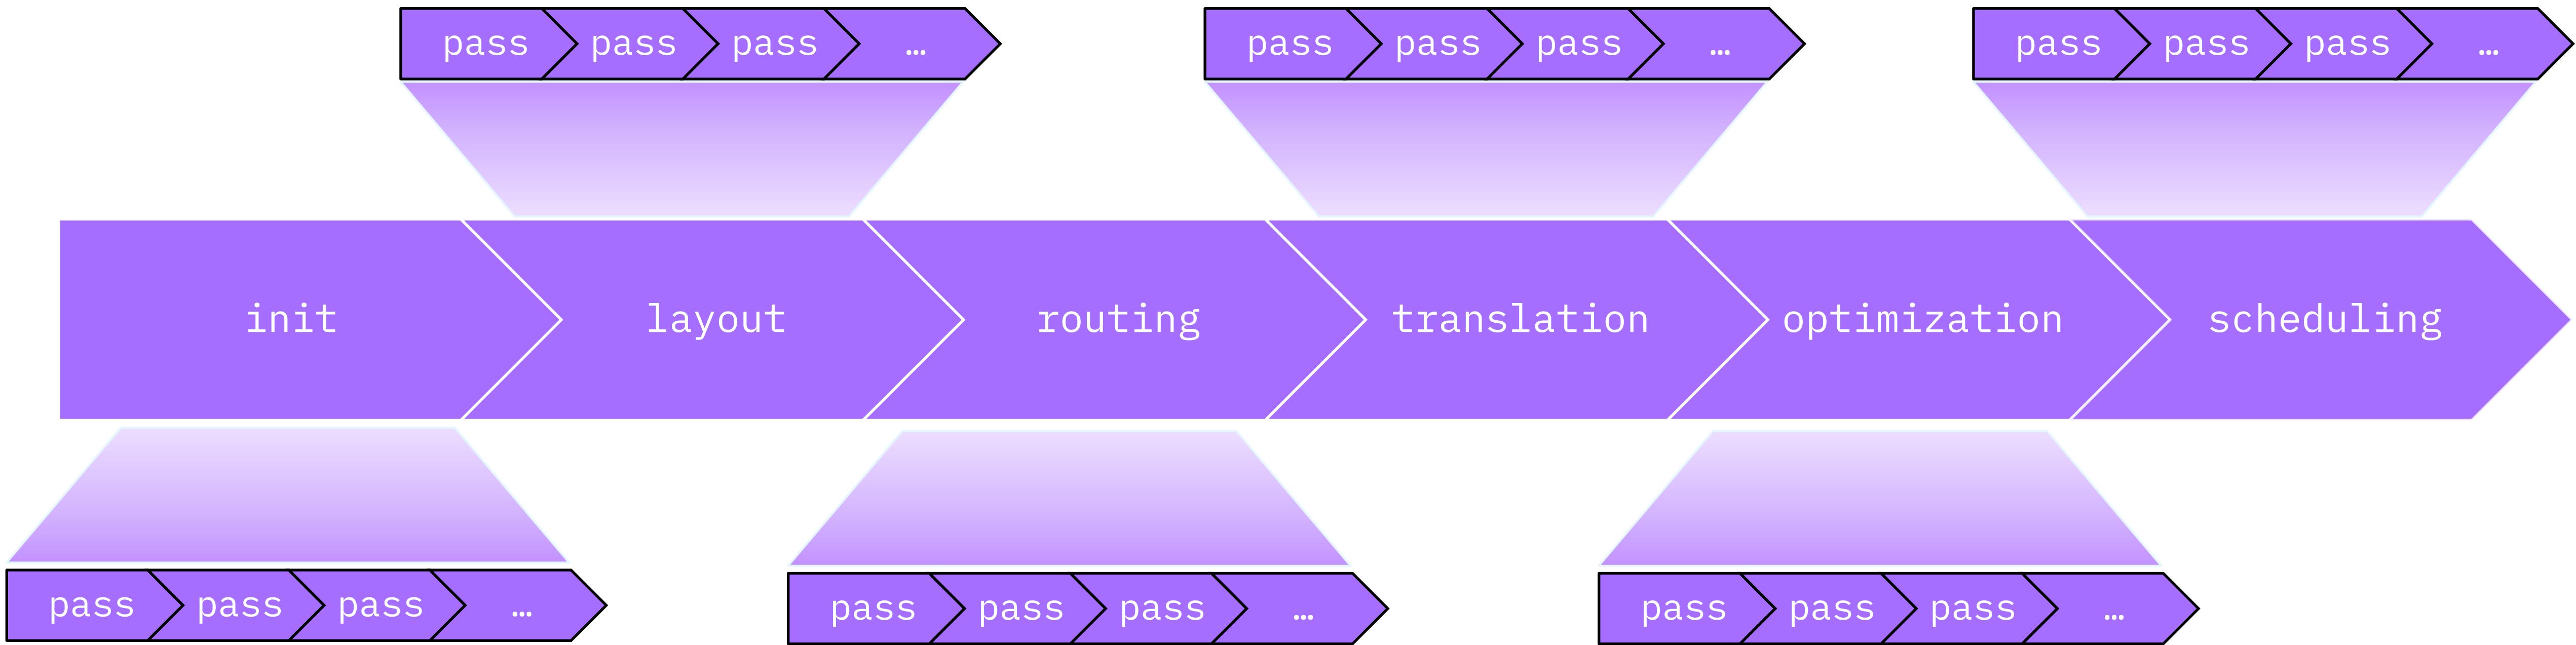




Directed Acyclic Graph

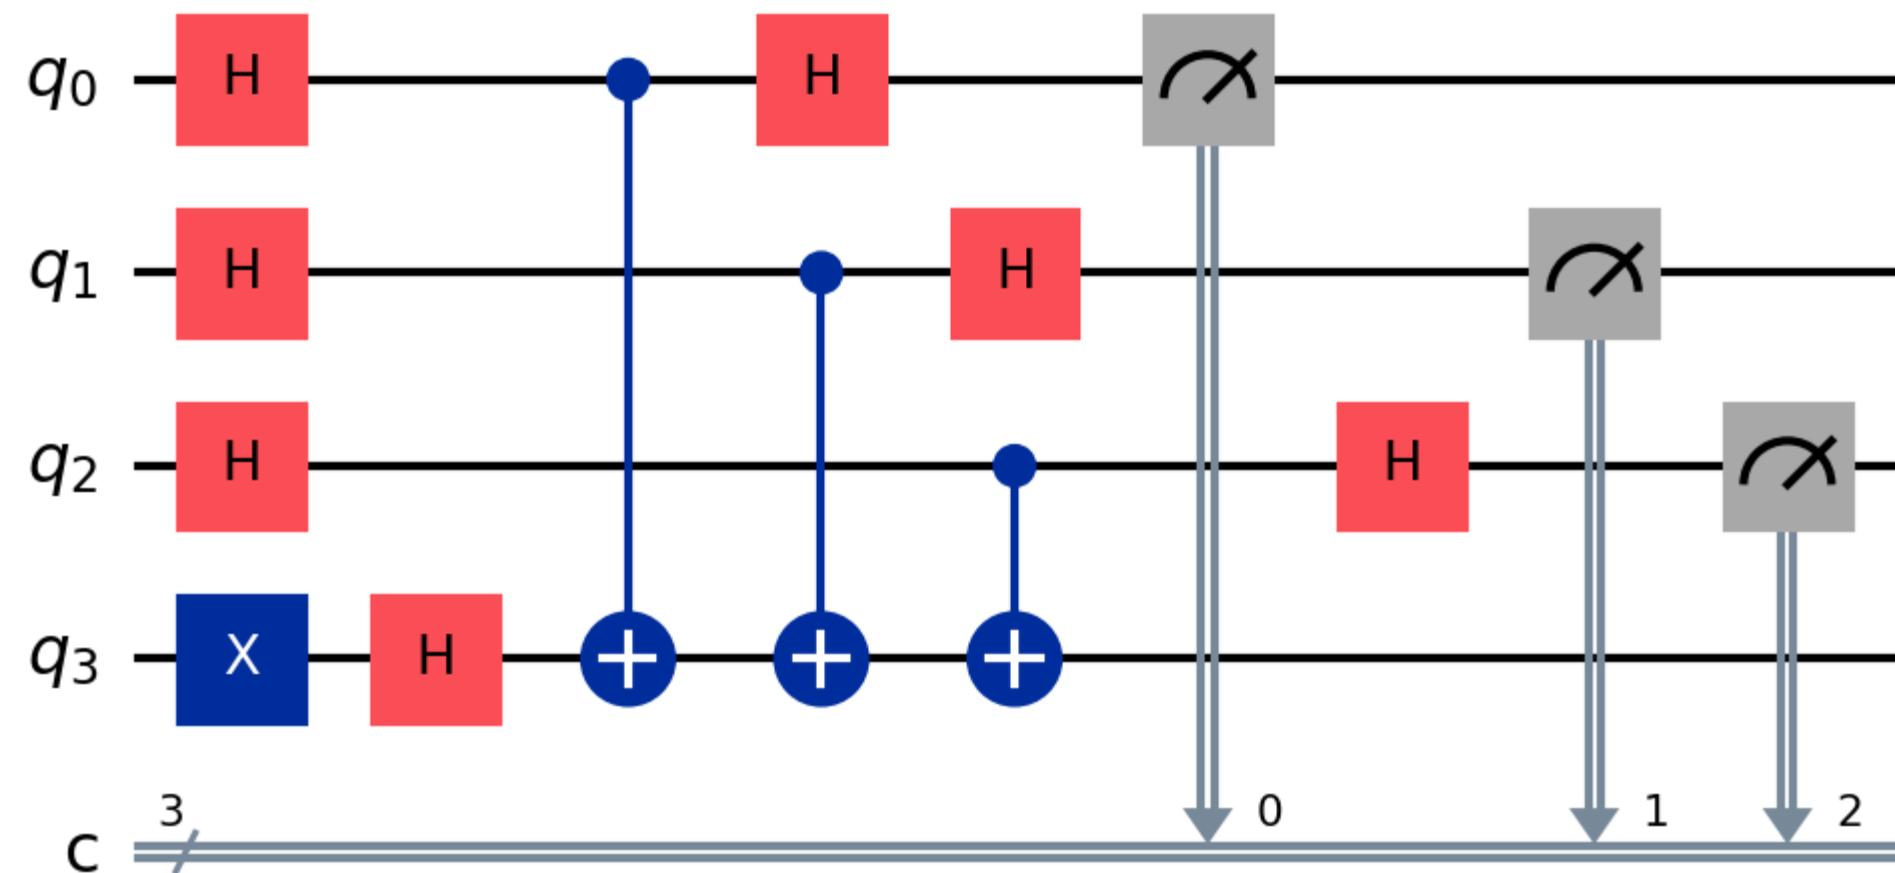


Stages of a Pass Manager

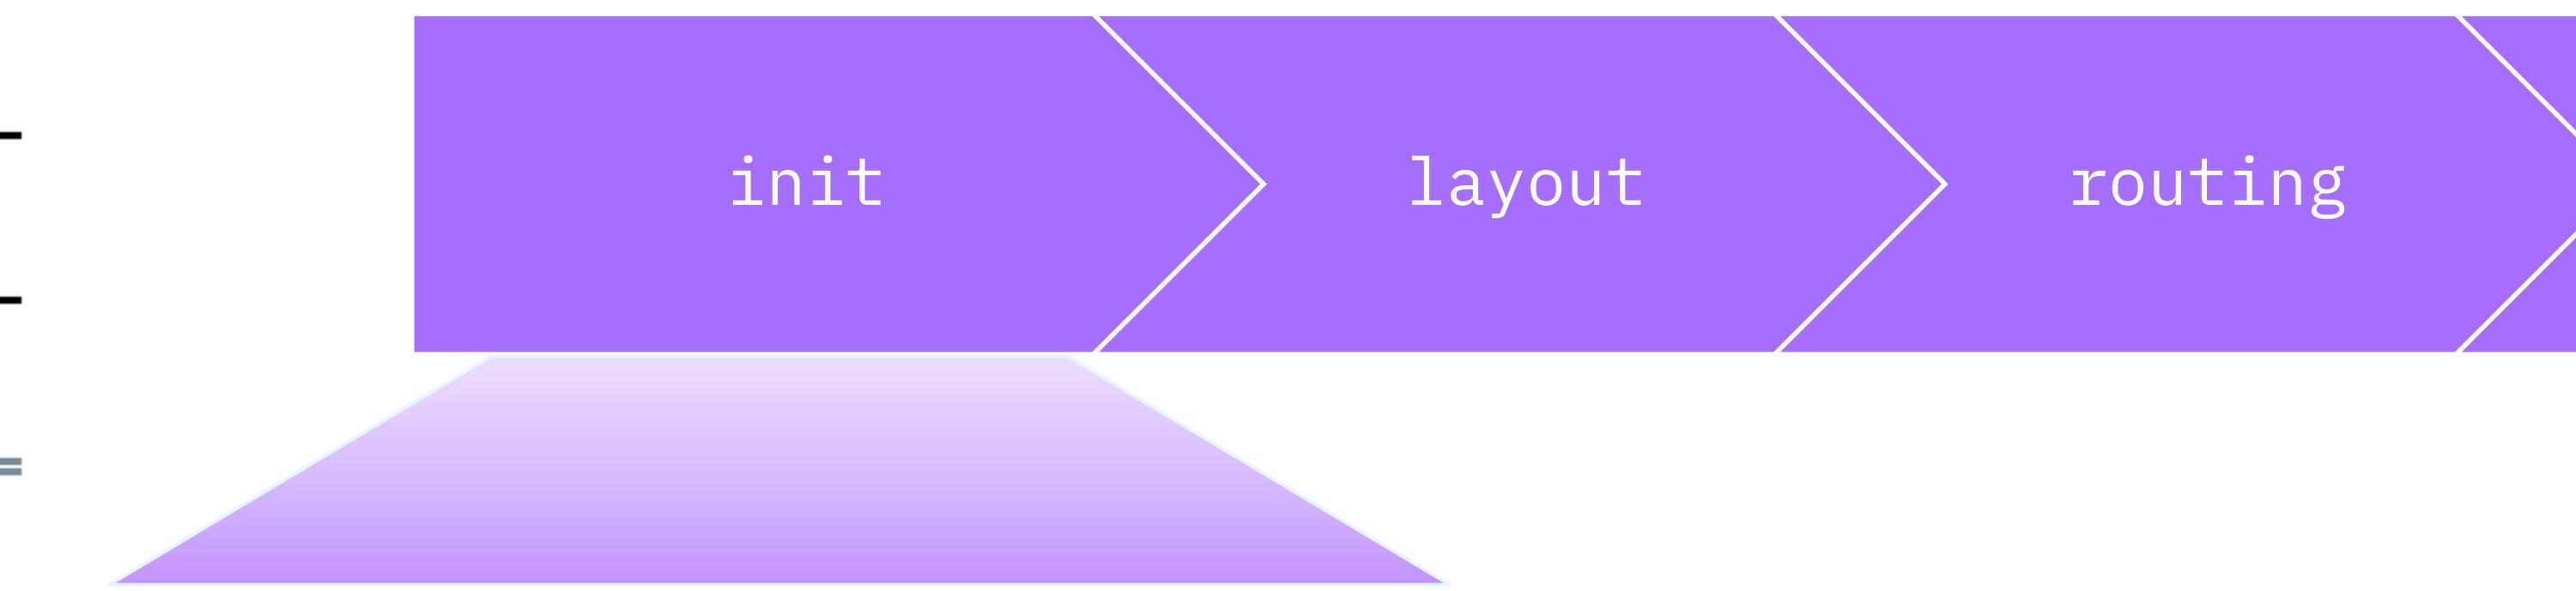


Stages of a Pass Manager - Plugins

```
pip install qiskit-qubit-reuse
```

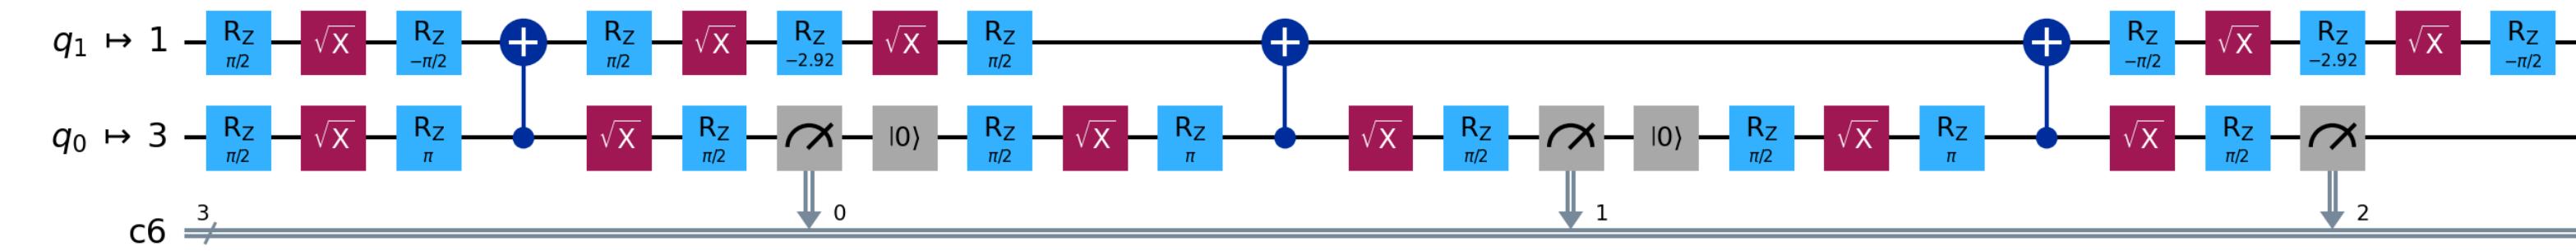


```
transpile(qc, target, init_method="qubit_reuse")
```



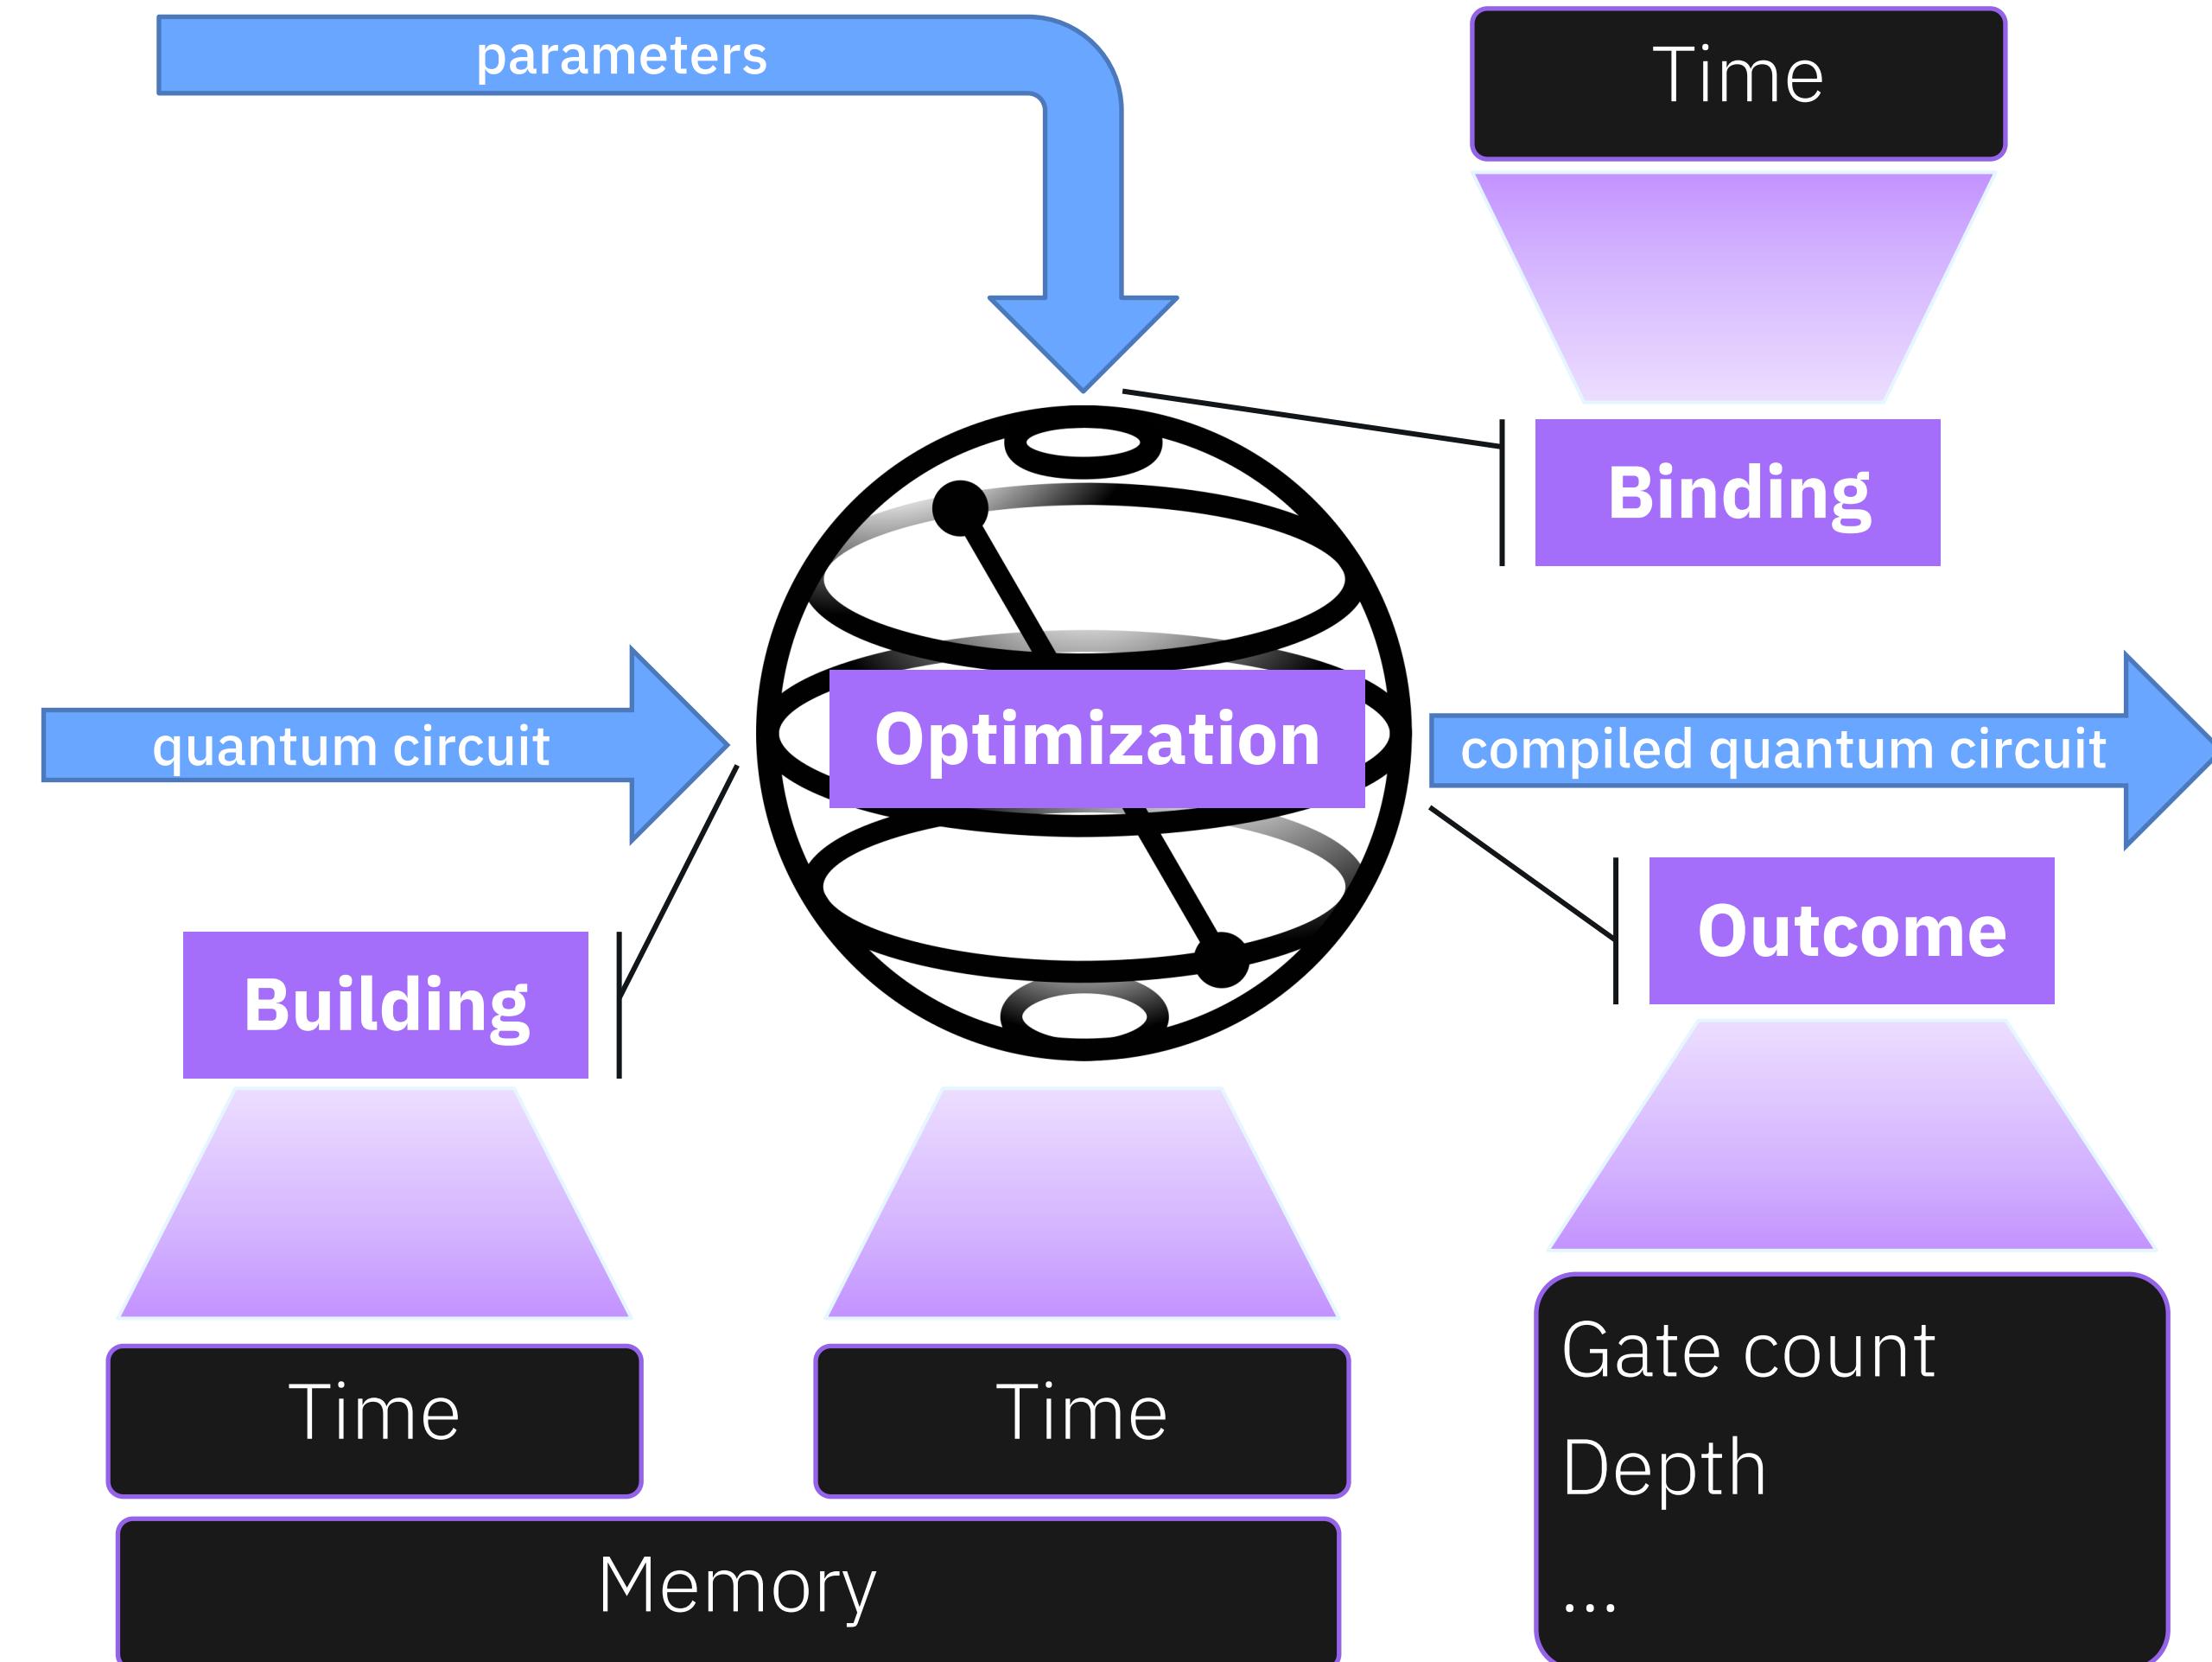
Stages of a Pass Manager - Plugins

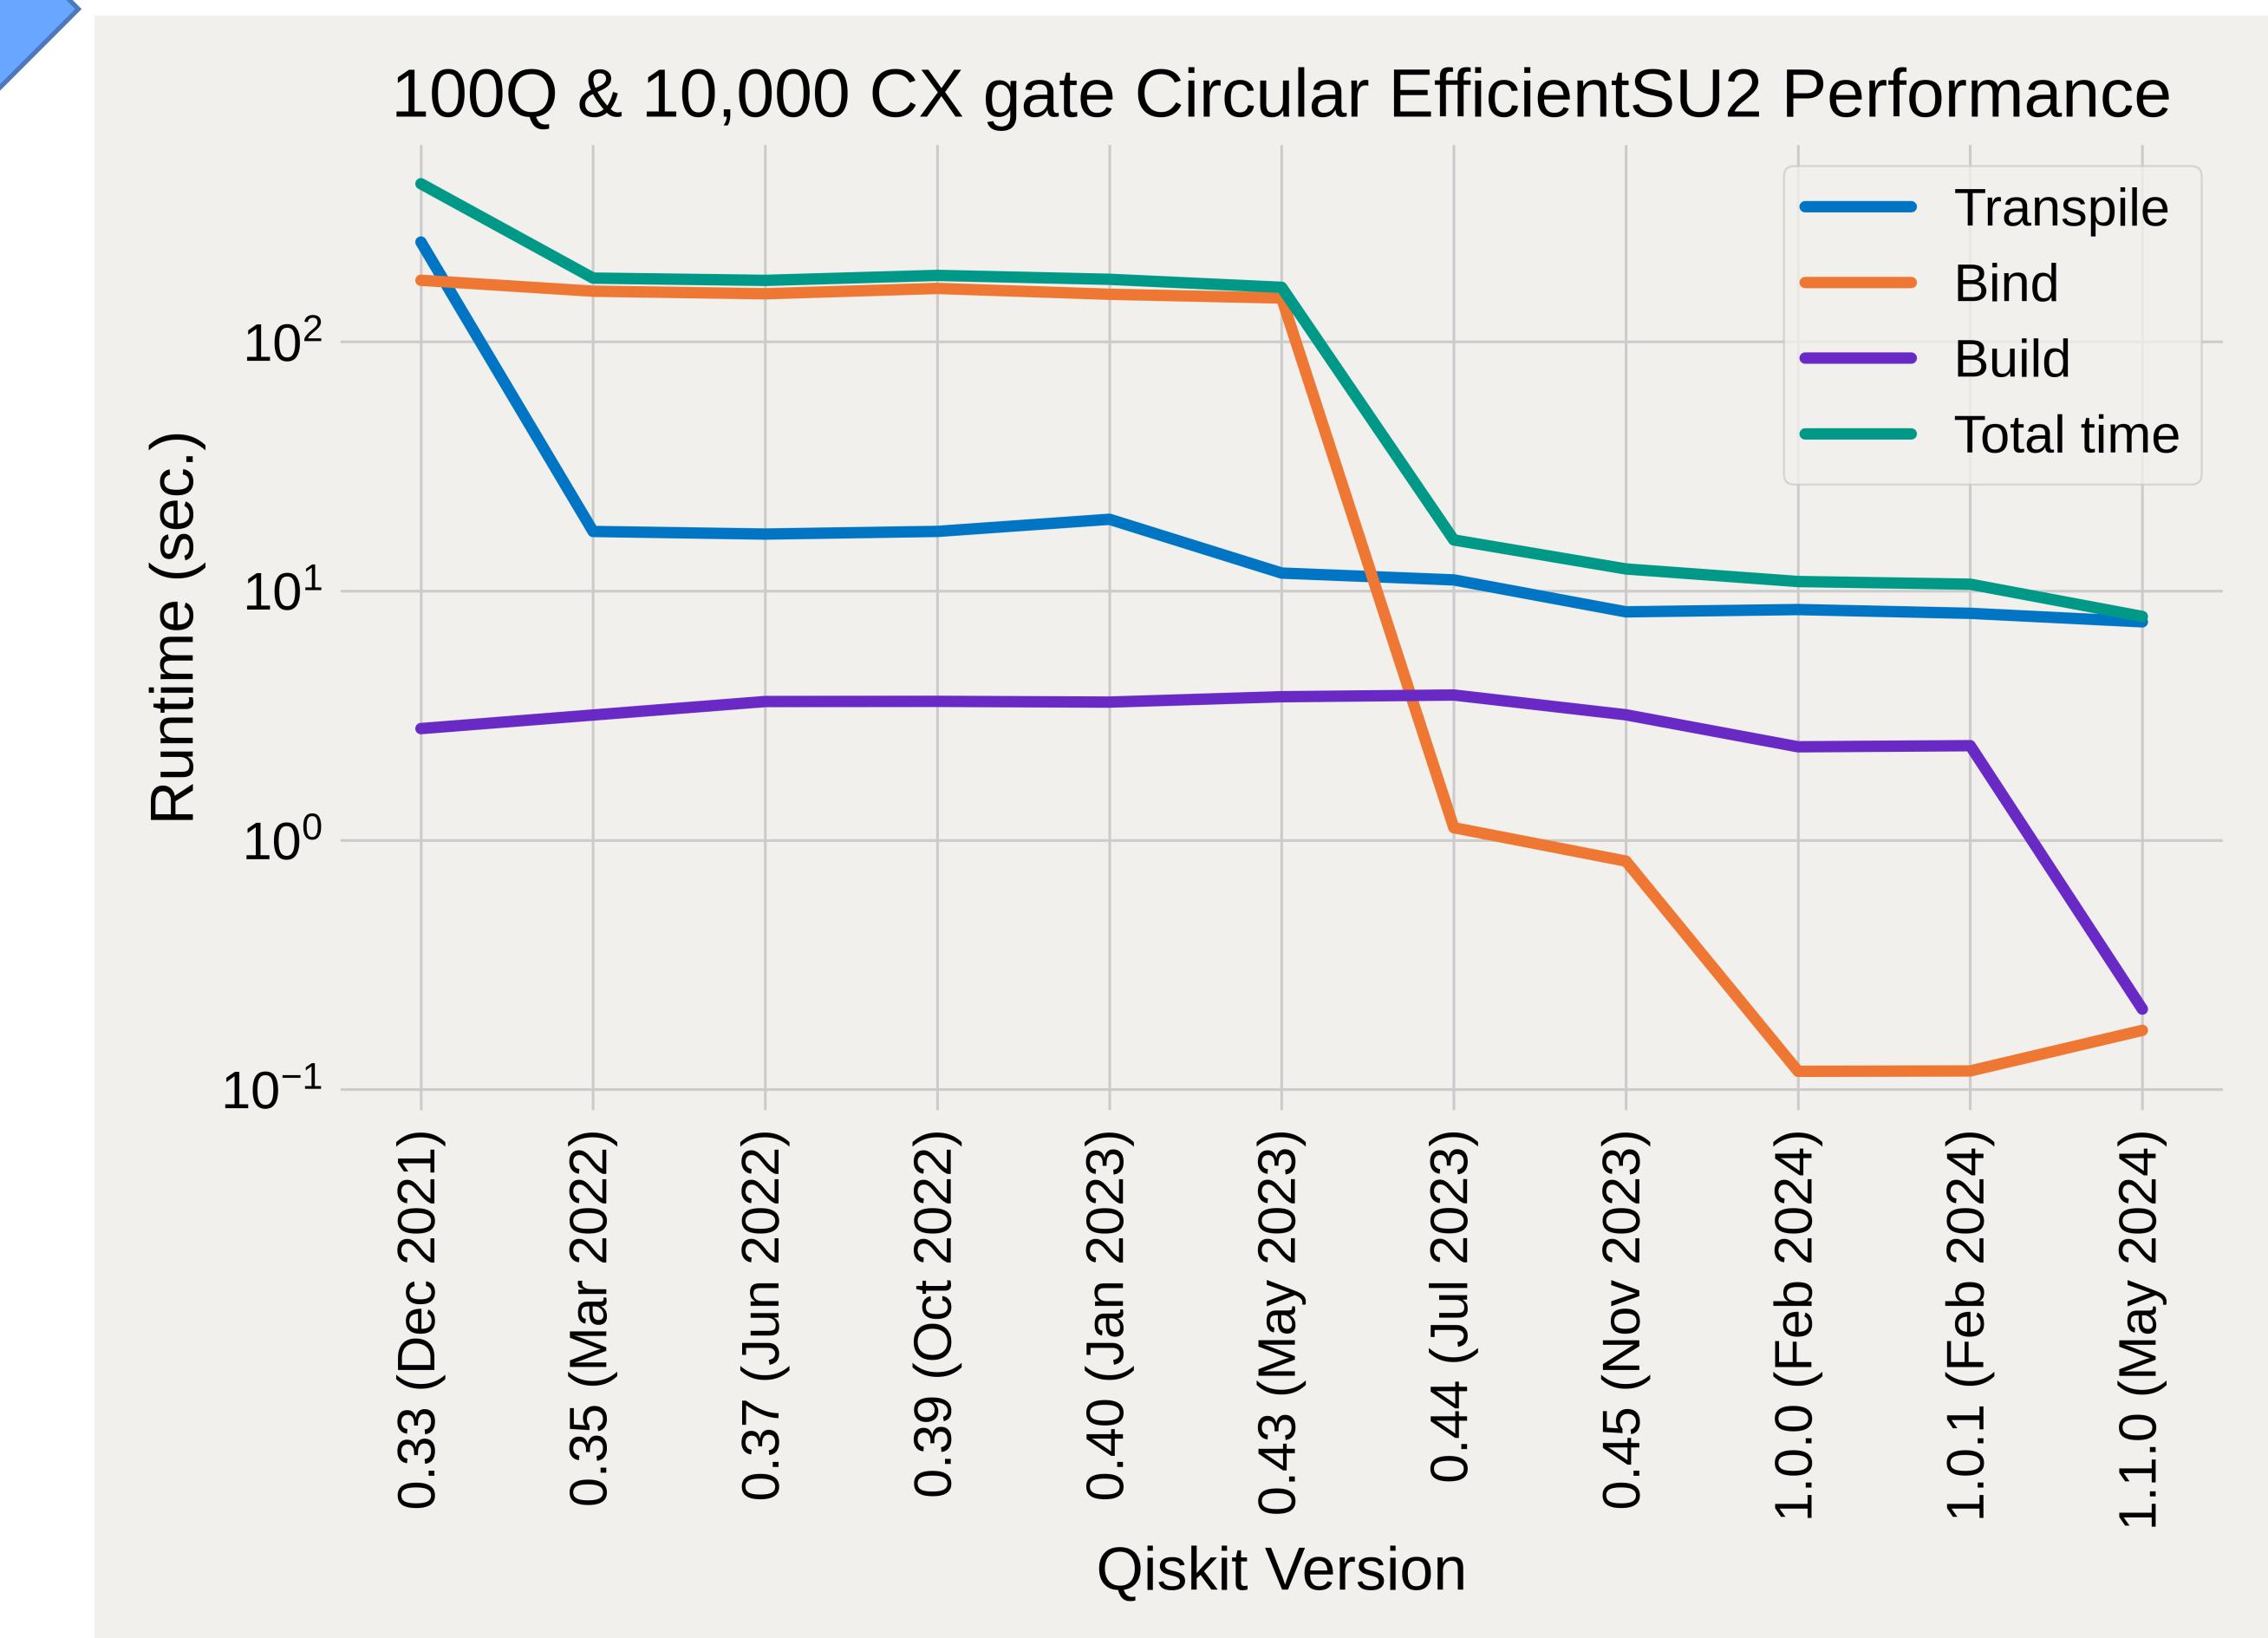
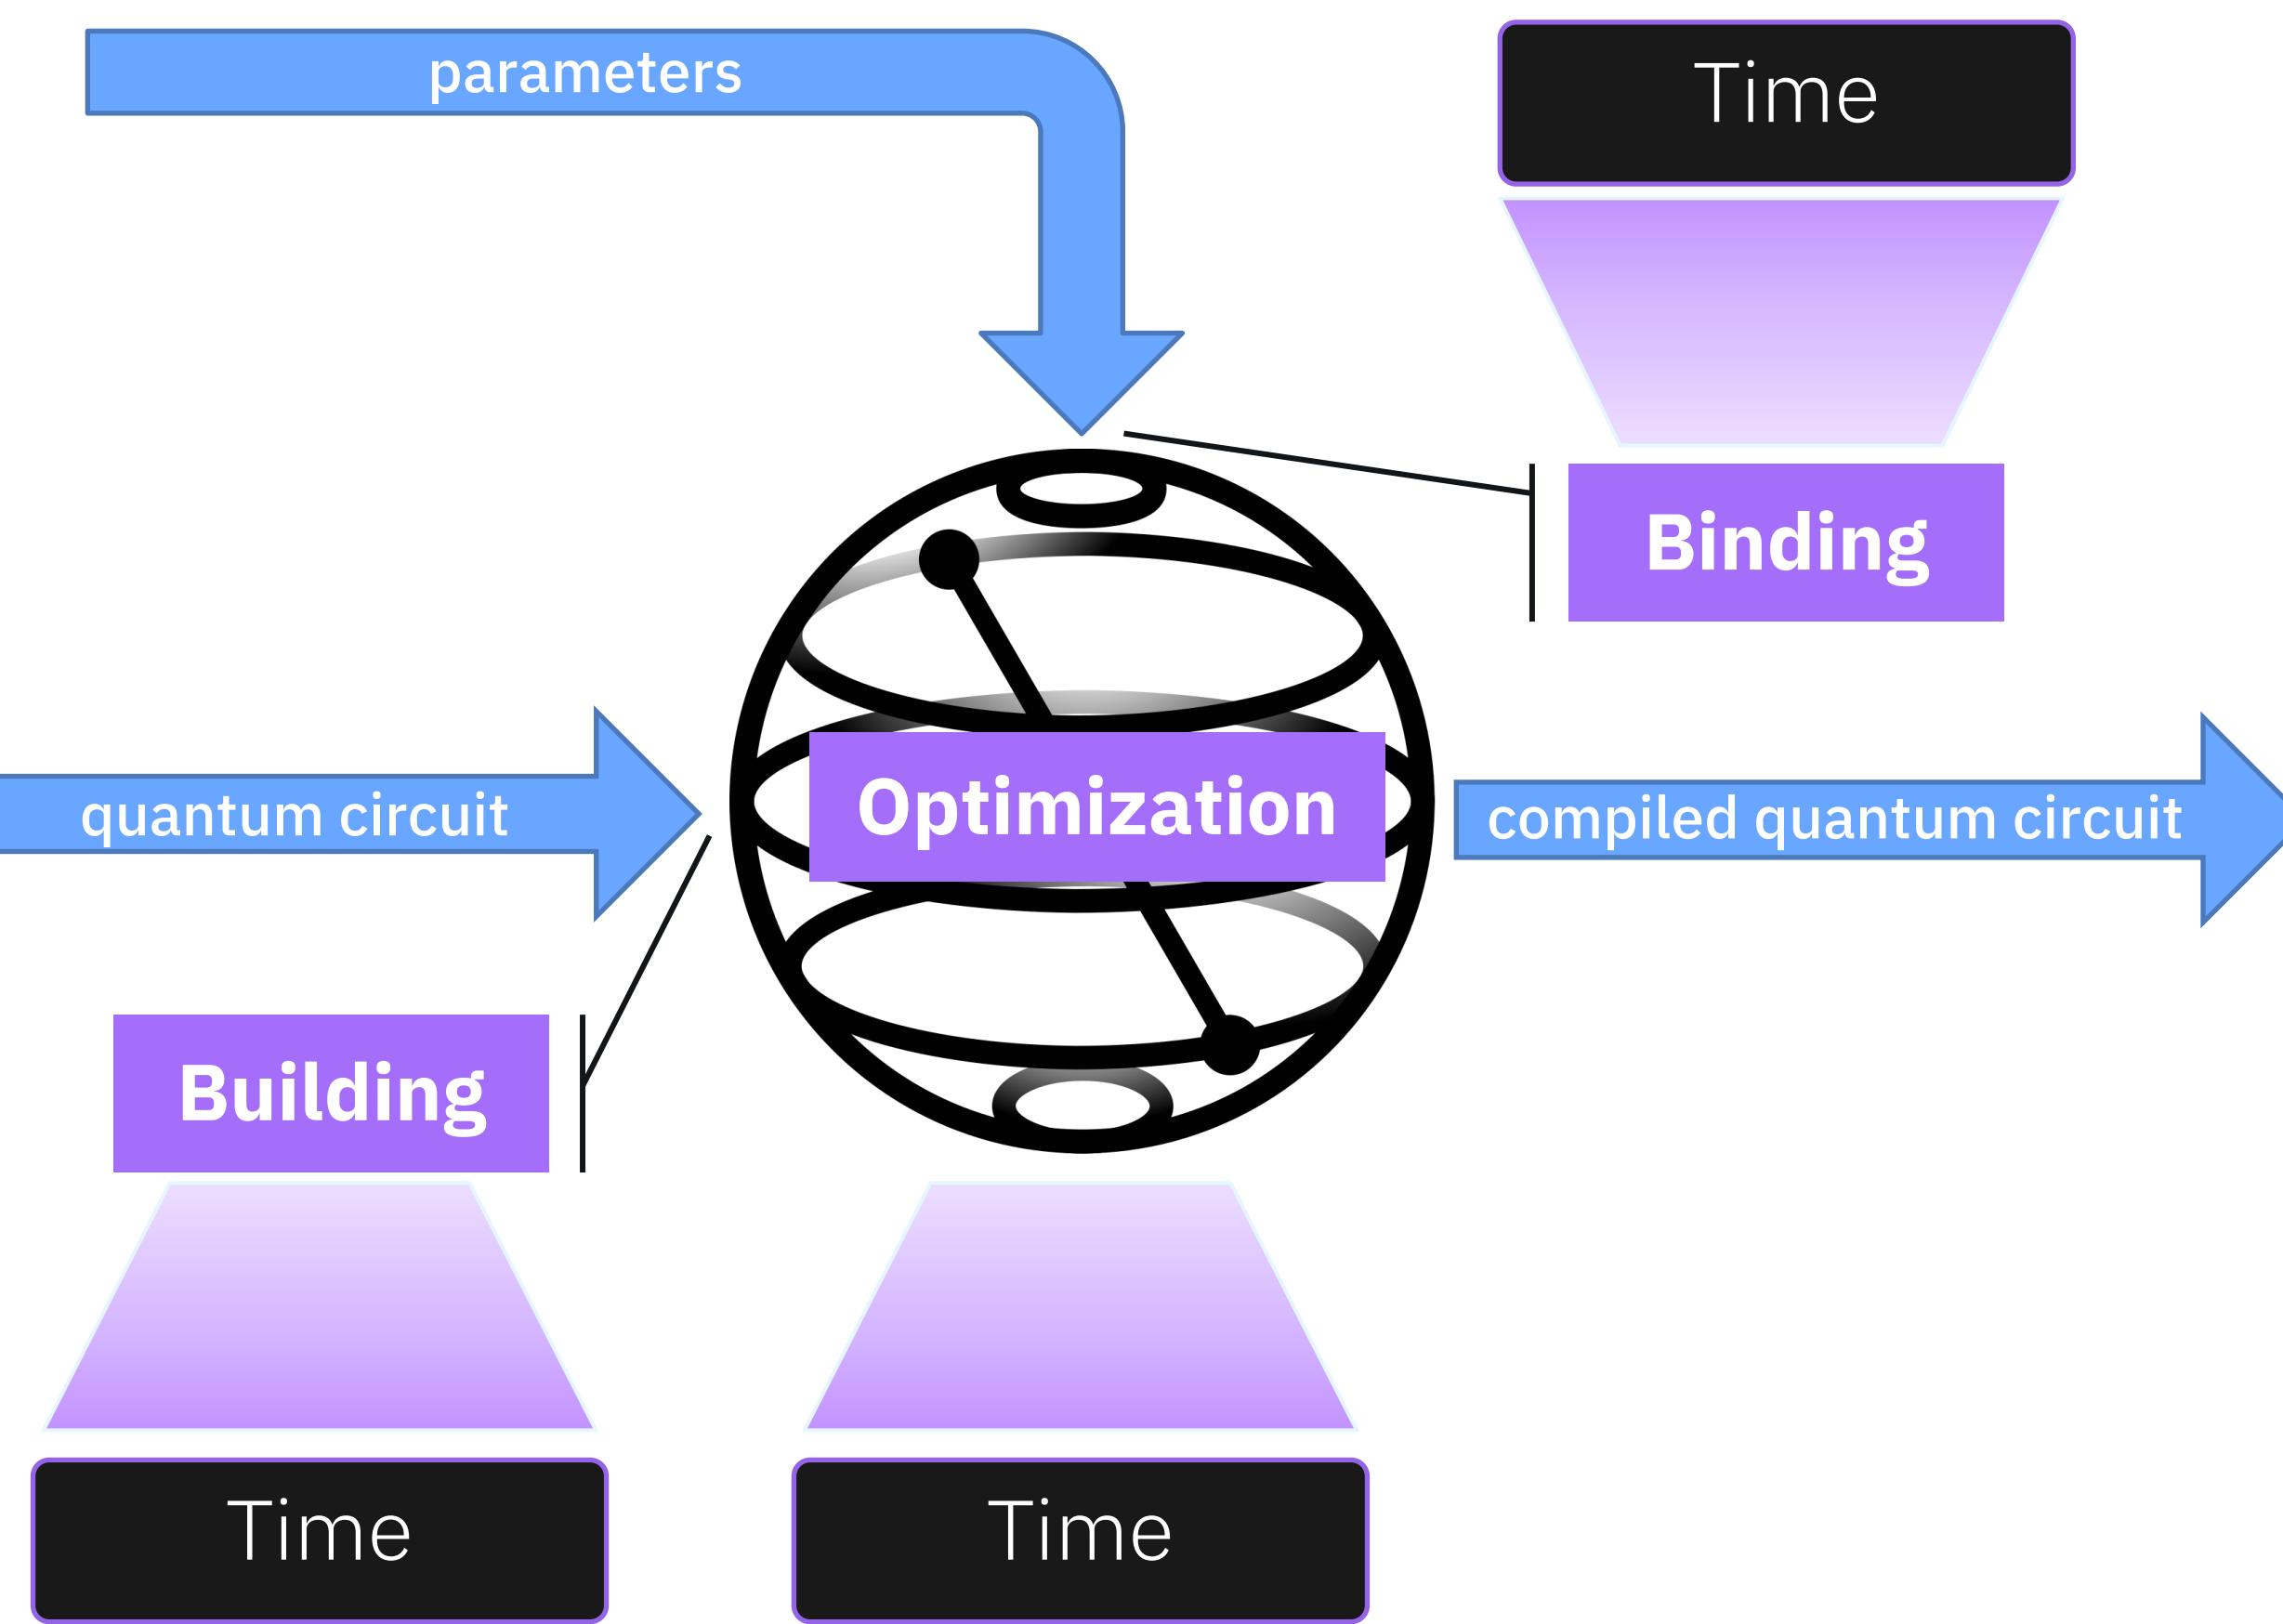
pip install qiskit-qubit-reuse



transpile(qc, target, init_method="qubit_reuse")

Performance





Building

Optimization

Time

Time

Benchmark performance (shorter is better)

Circuit Construction

test_QV100_build

test_DTC100_set_build

test_multi_control_circuit

test_clifford_build

test_param_circSU2_100_build

test_param_circSU2_100_bind

test_QV100_qasm2_import

FAILED

test_bigint_qasm2_import

FAILED

FAILED



Circuit Manipulation

test_DTC100_twirling

test_multi_control_decompose

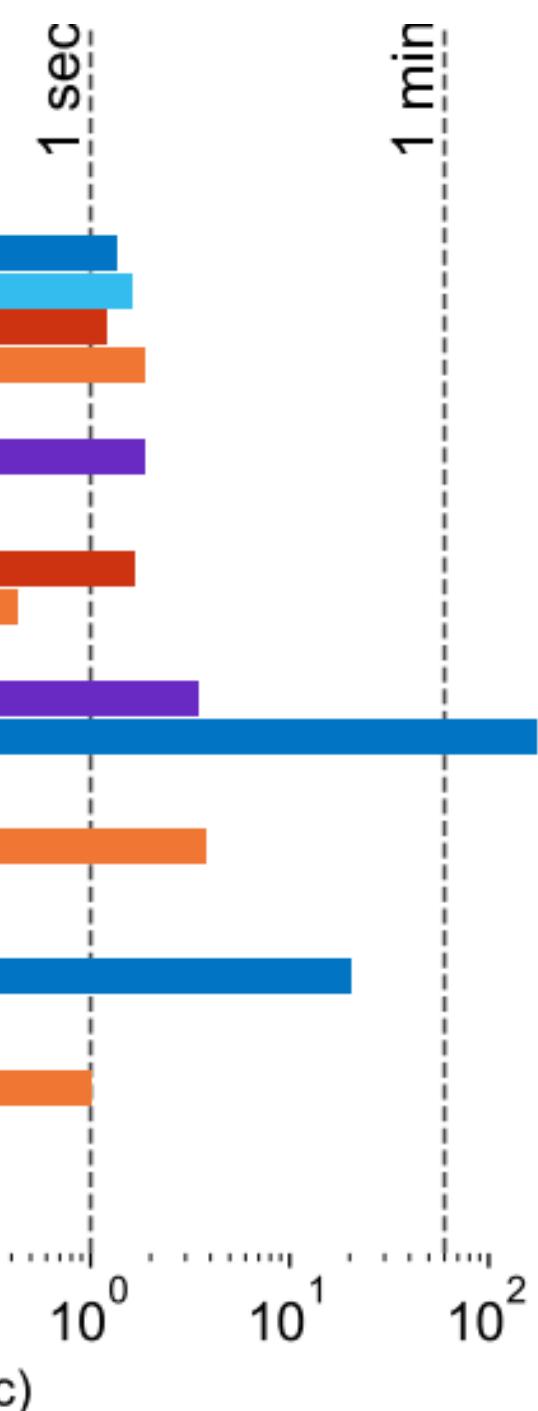
XFAIL
SKIPPED

test_QV100_basis_change

SKIPPED
FAILED

test_random_clifford_decompose

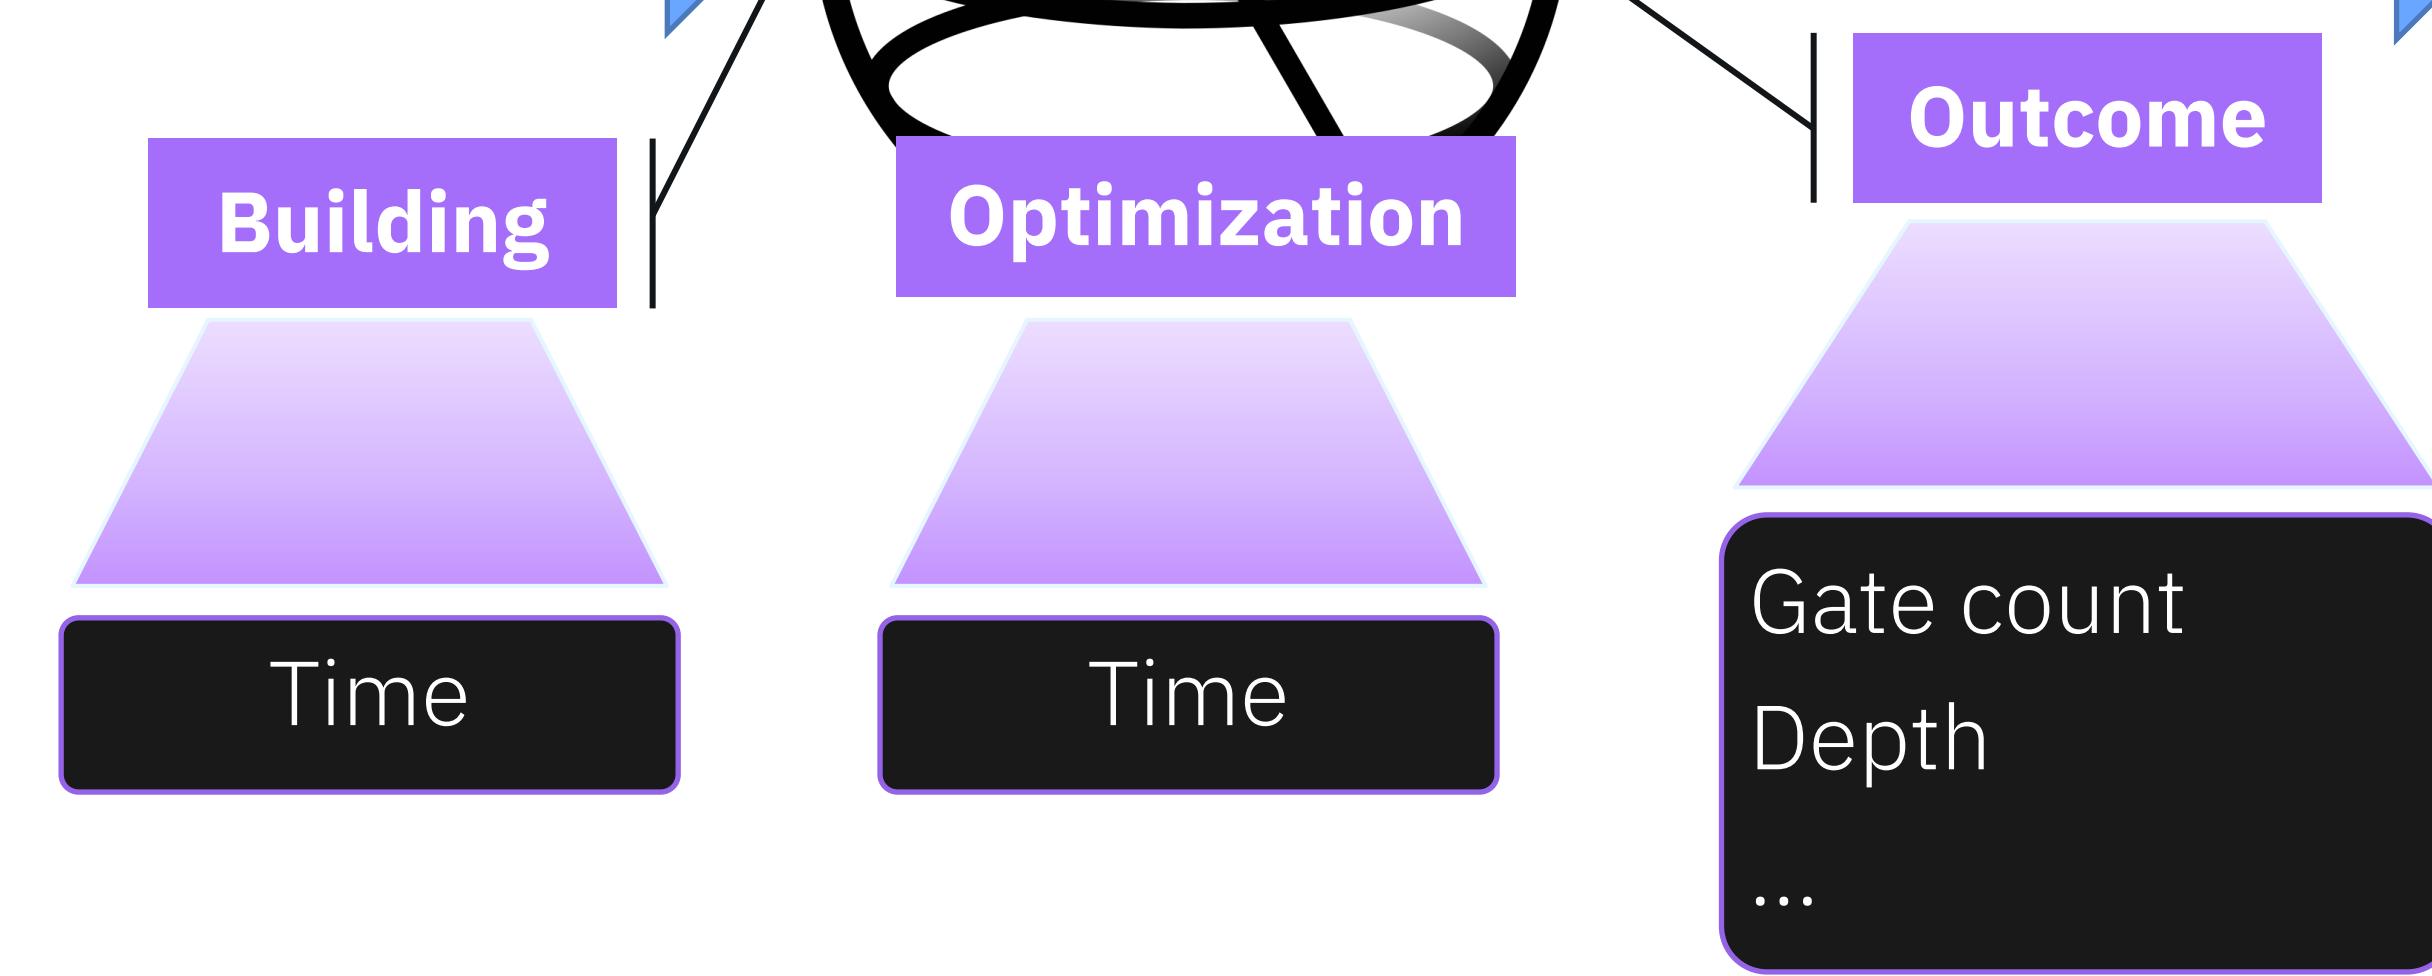
SKIPPED
FAILED



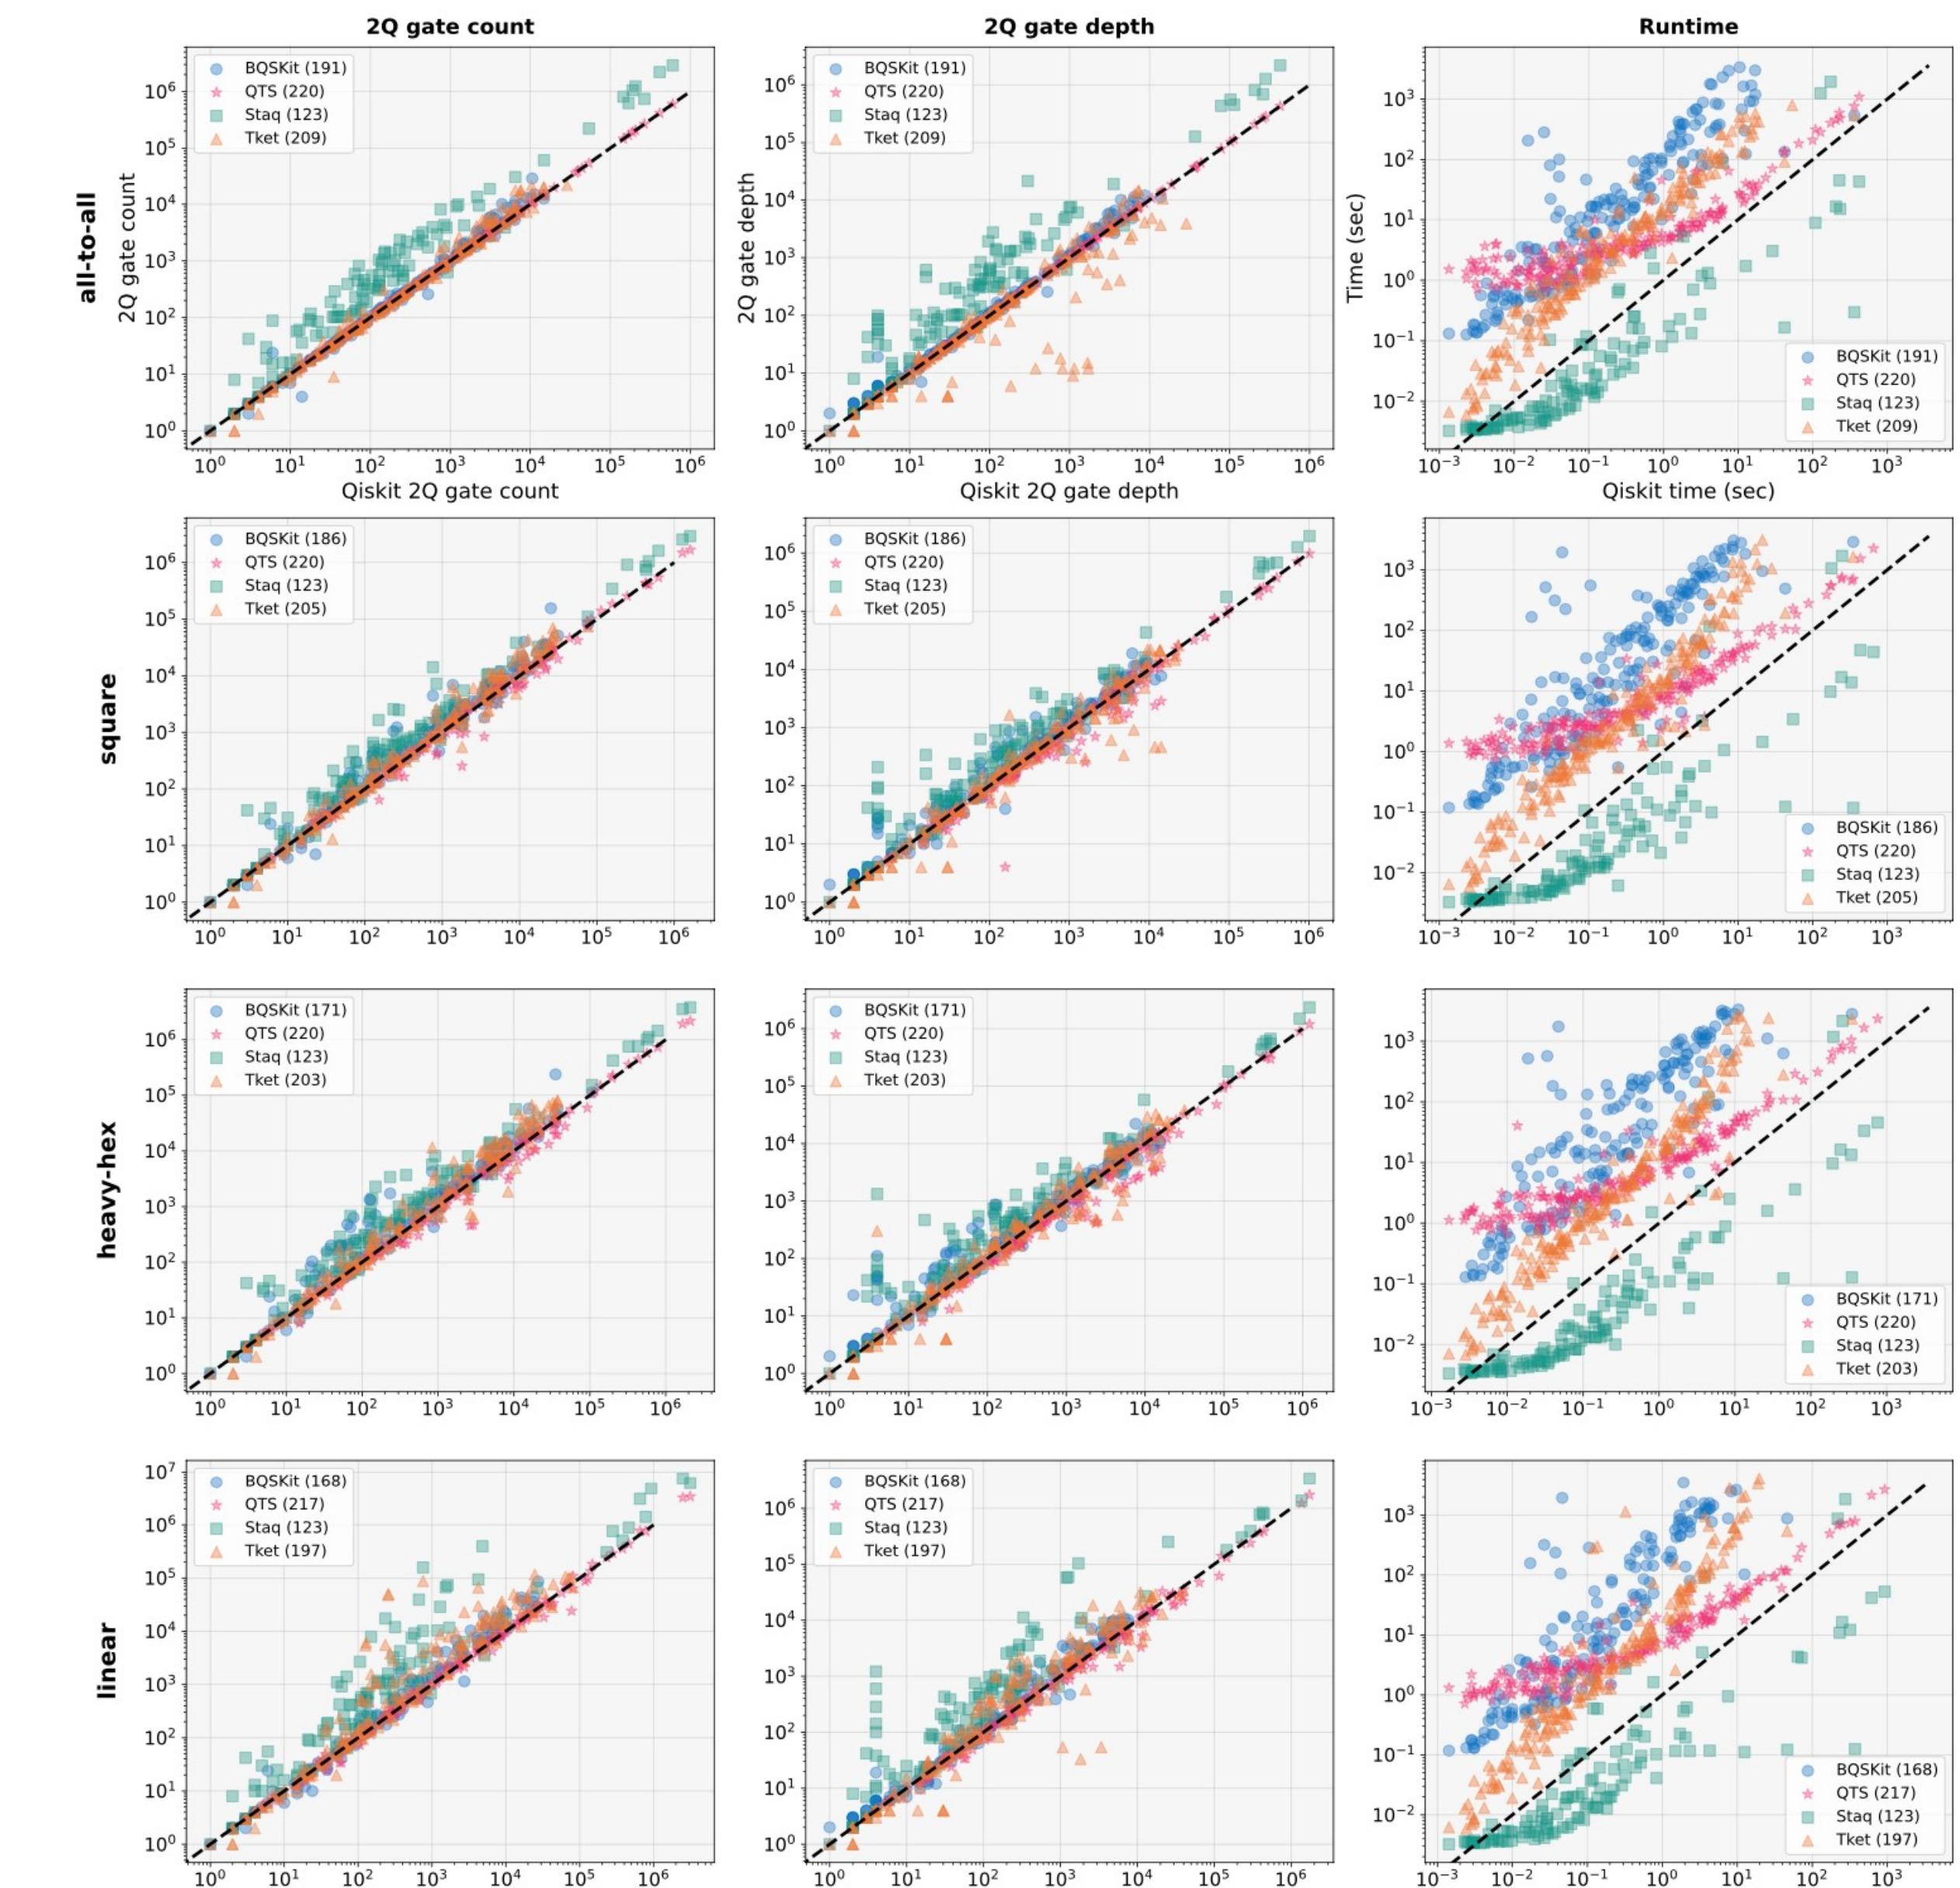
cond

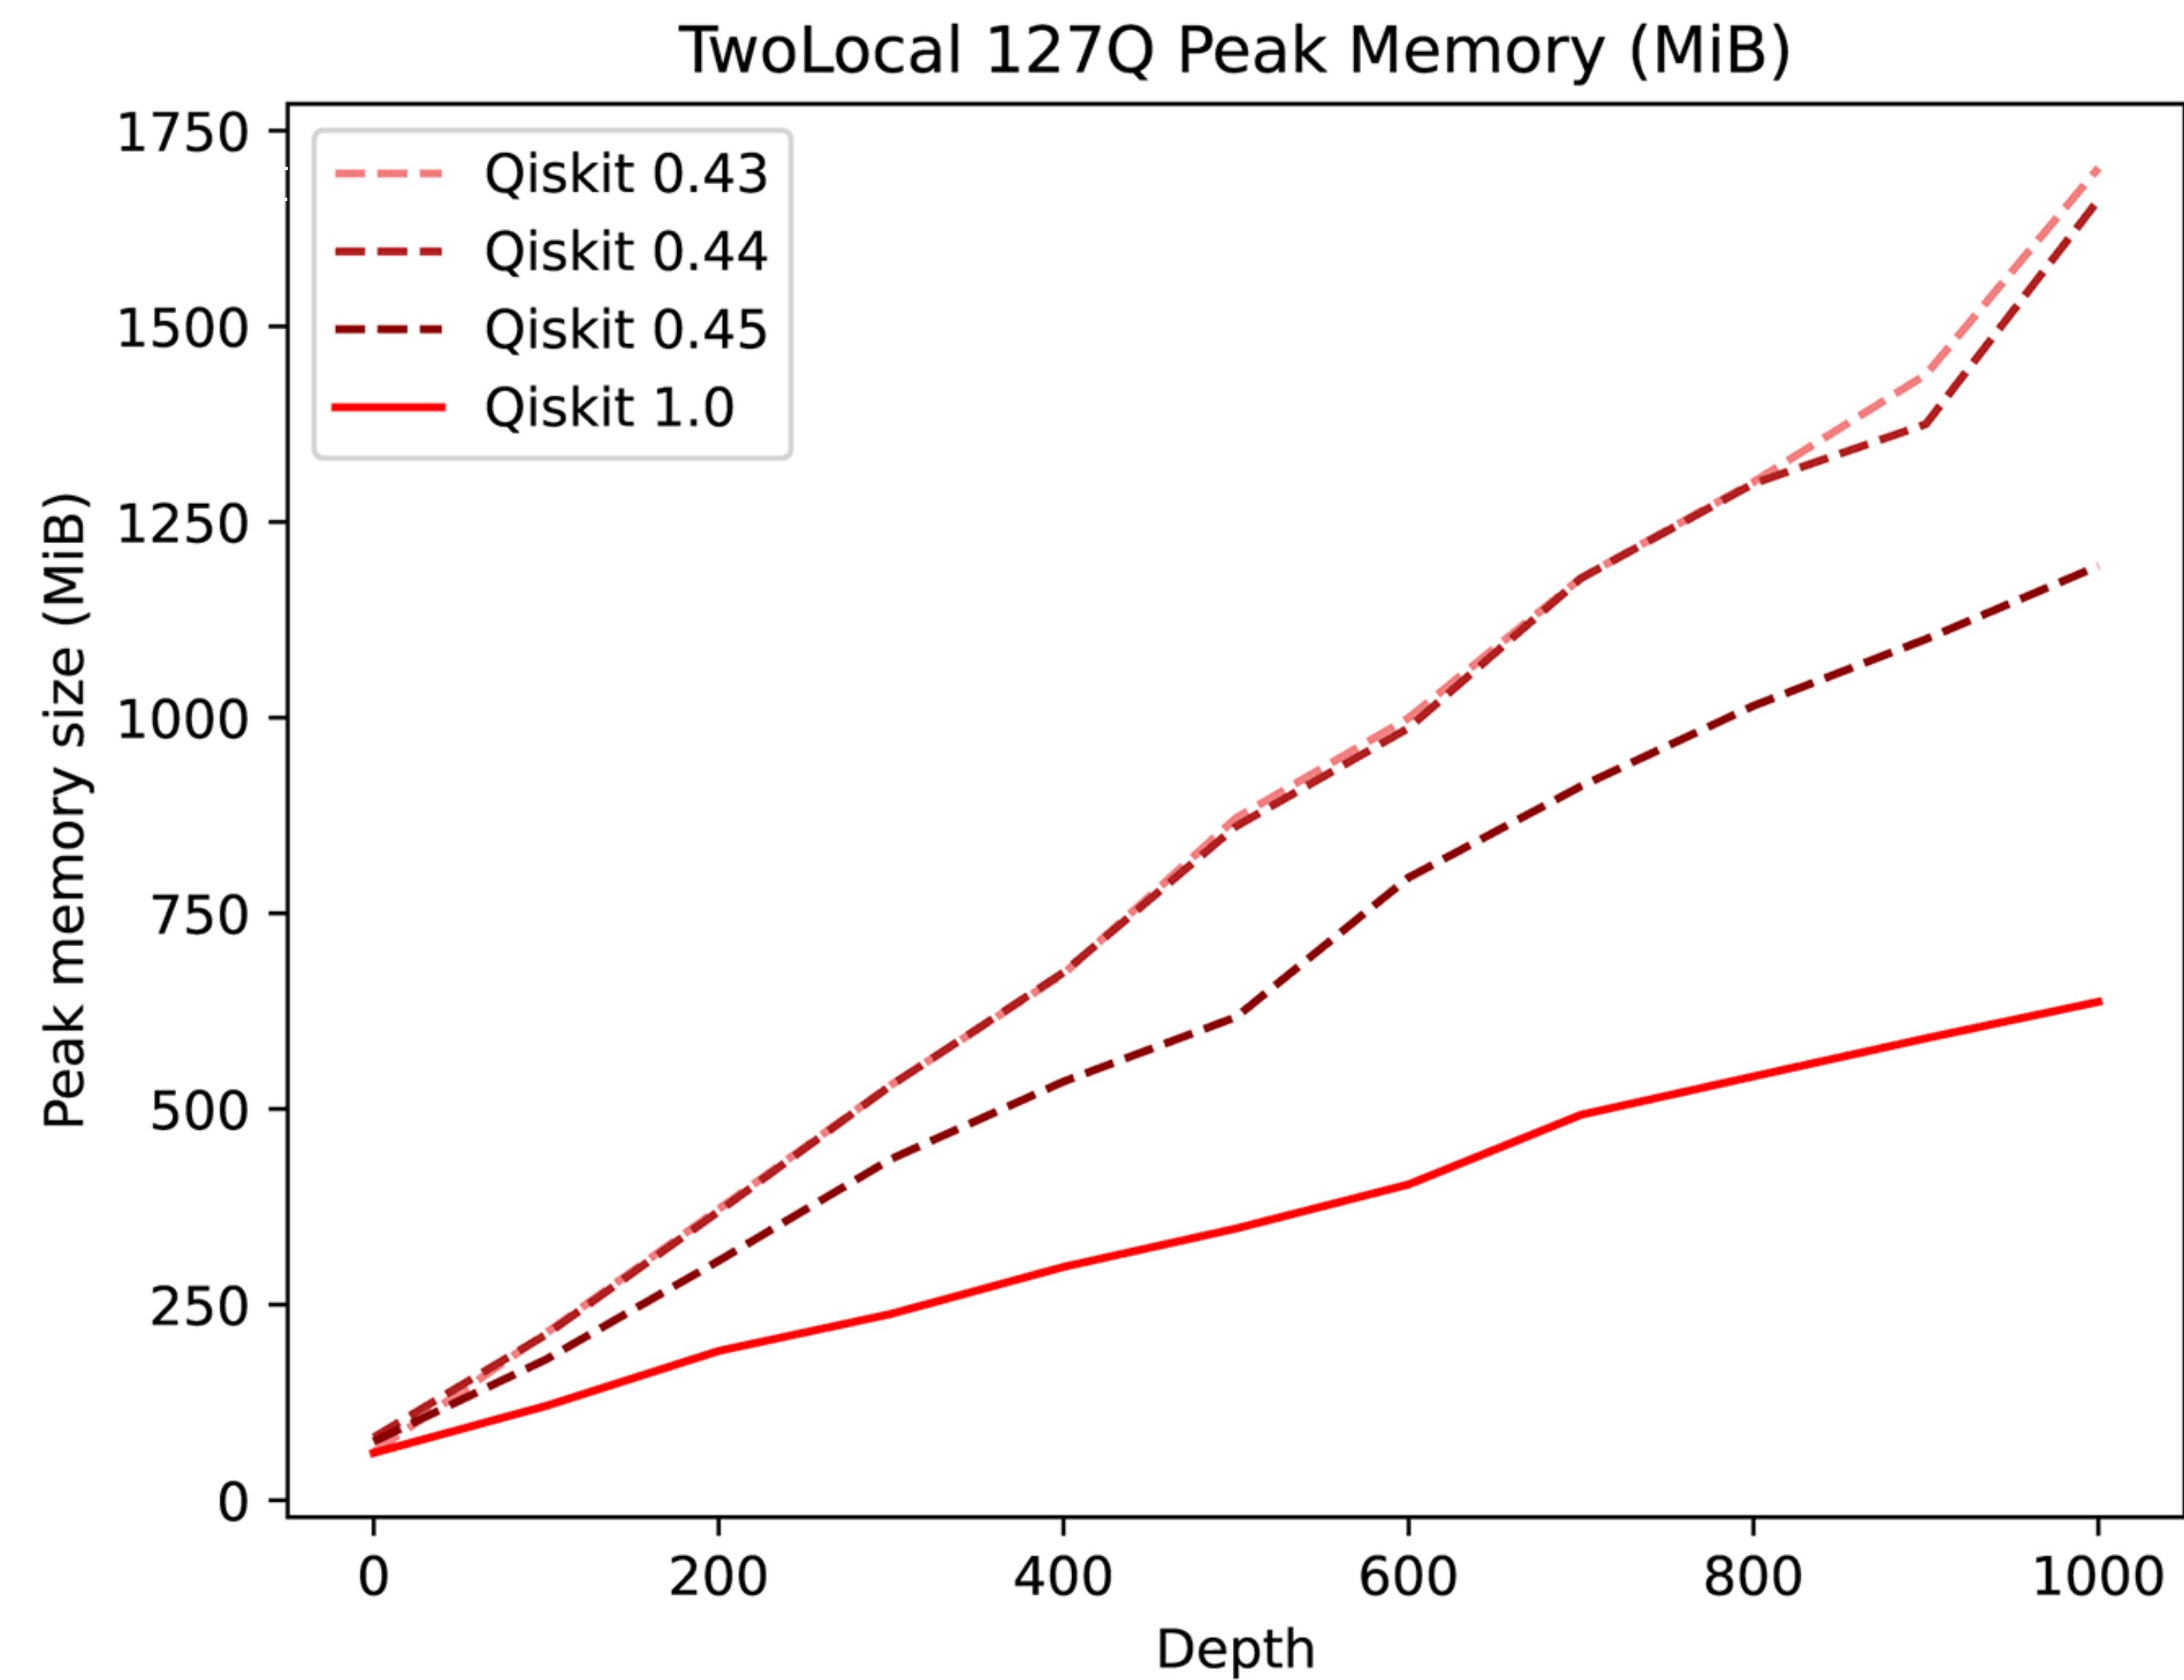
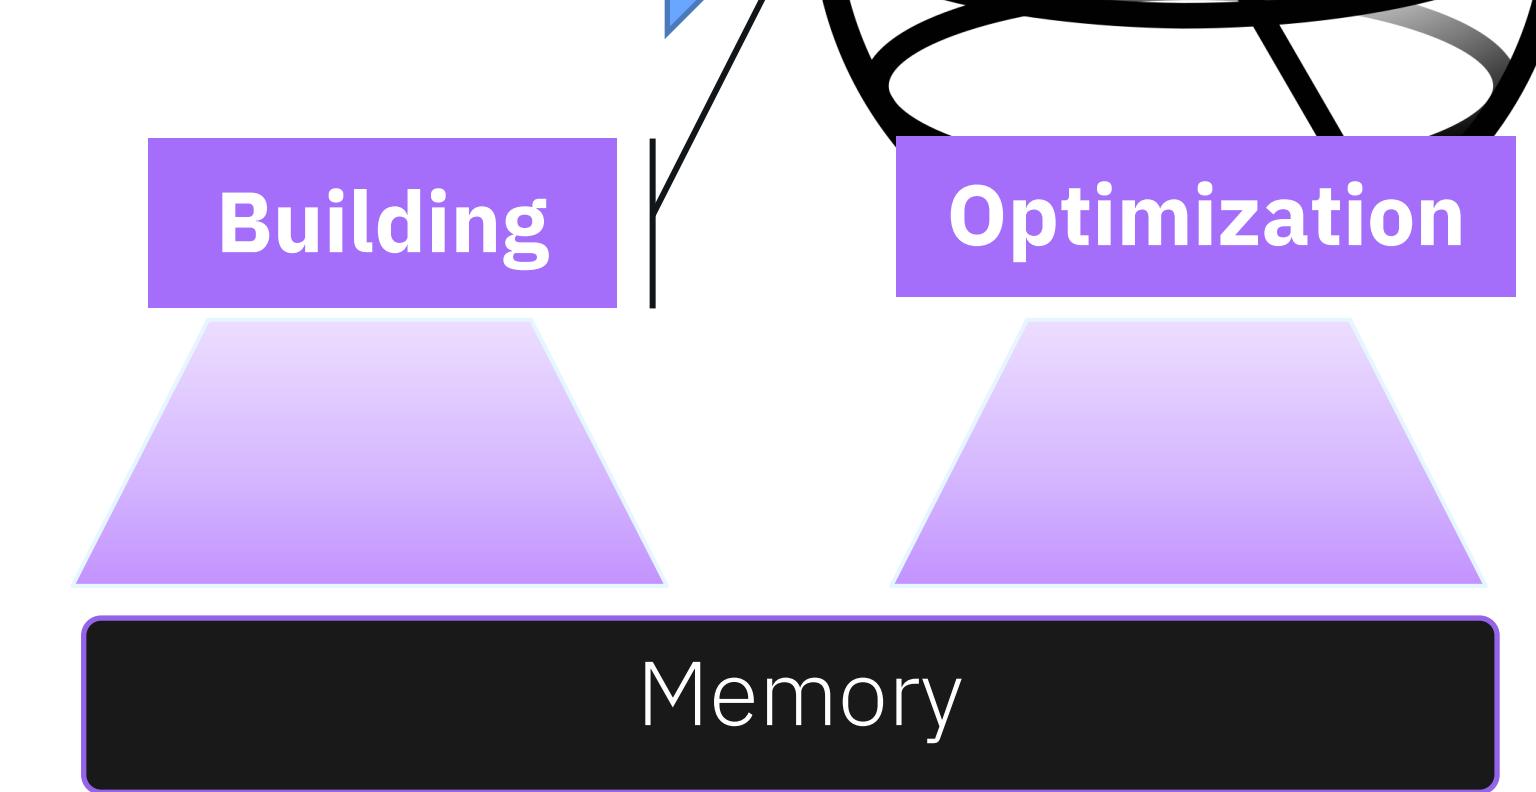
nute

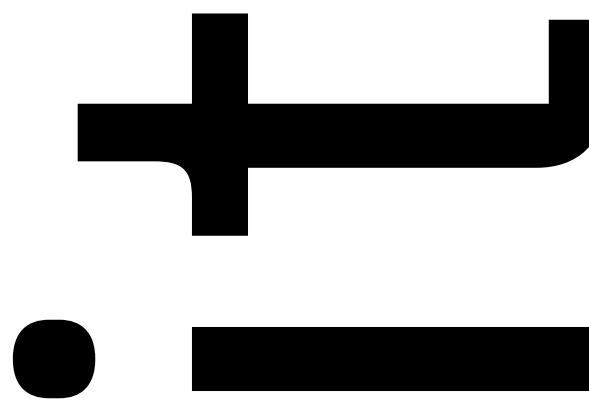
<https://arxiv.org/abs/2409.08844>



- Qiskit compared to Tket:
 - **53%** mean reduction on 2Q gate count
 - **29%** mean transpilation time reduction
- Qiskit compared to BQSKit:
 - **29%** mean reduction on 2Q gate count
 - **572%** mean transpilation time reduction







112

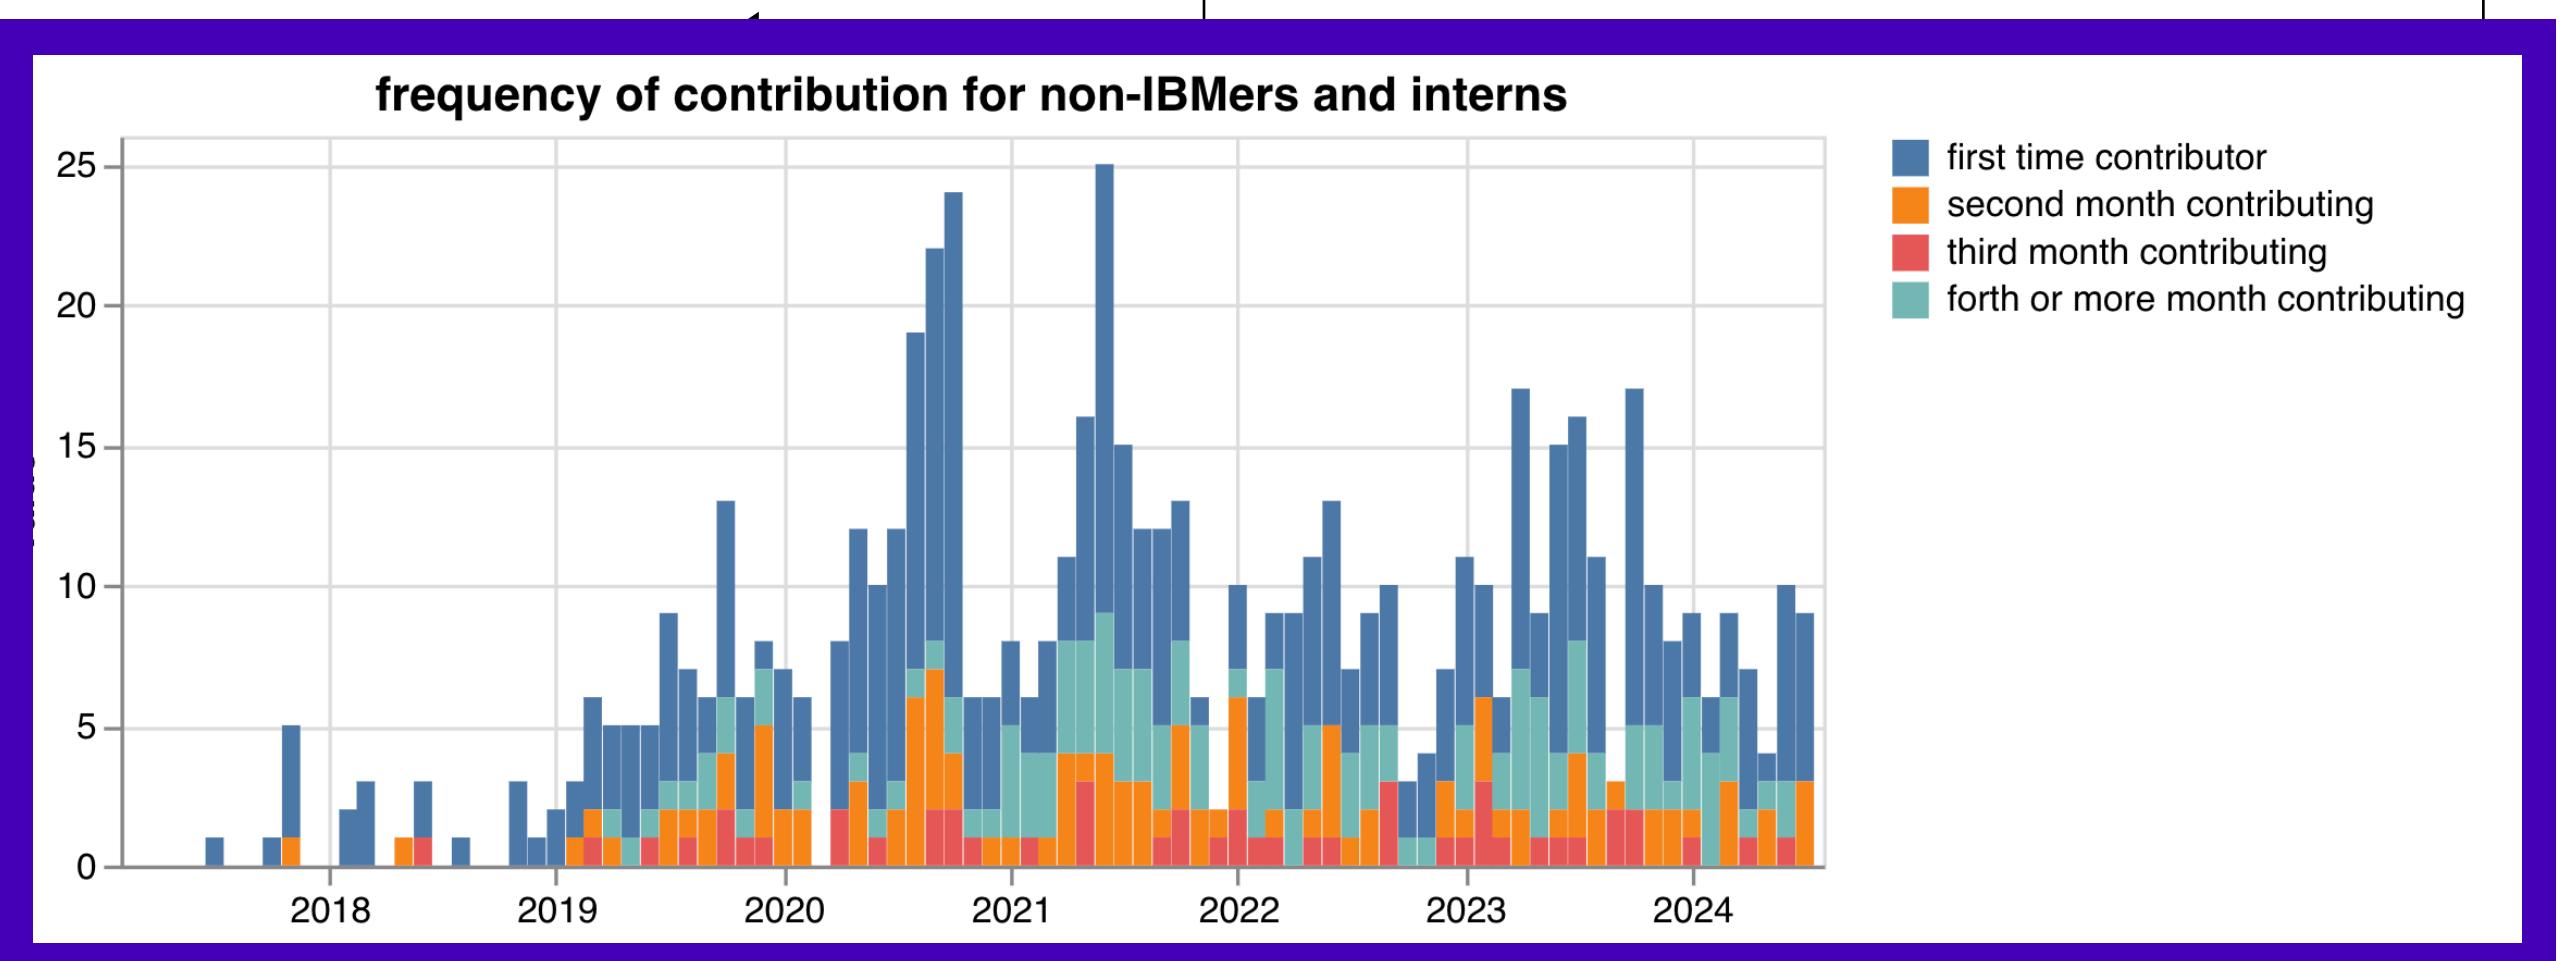
releases, including patch and pre-releases.

546

contributors, many first-OSS timers.
2.2k forks and 4.4k GitHub stars.

+1600

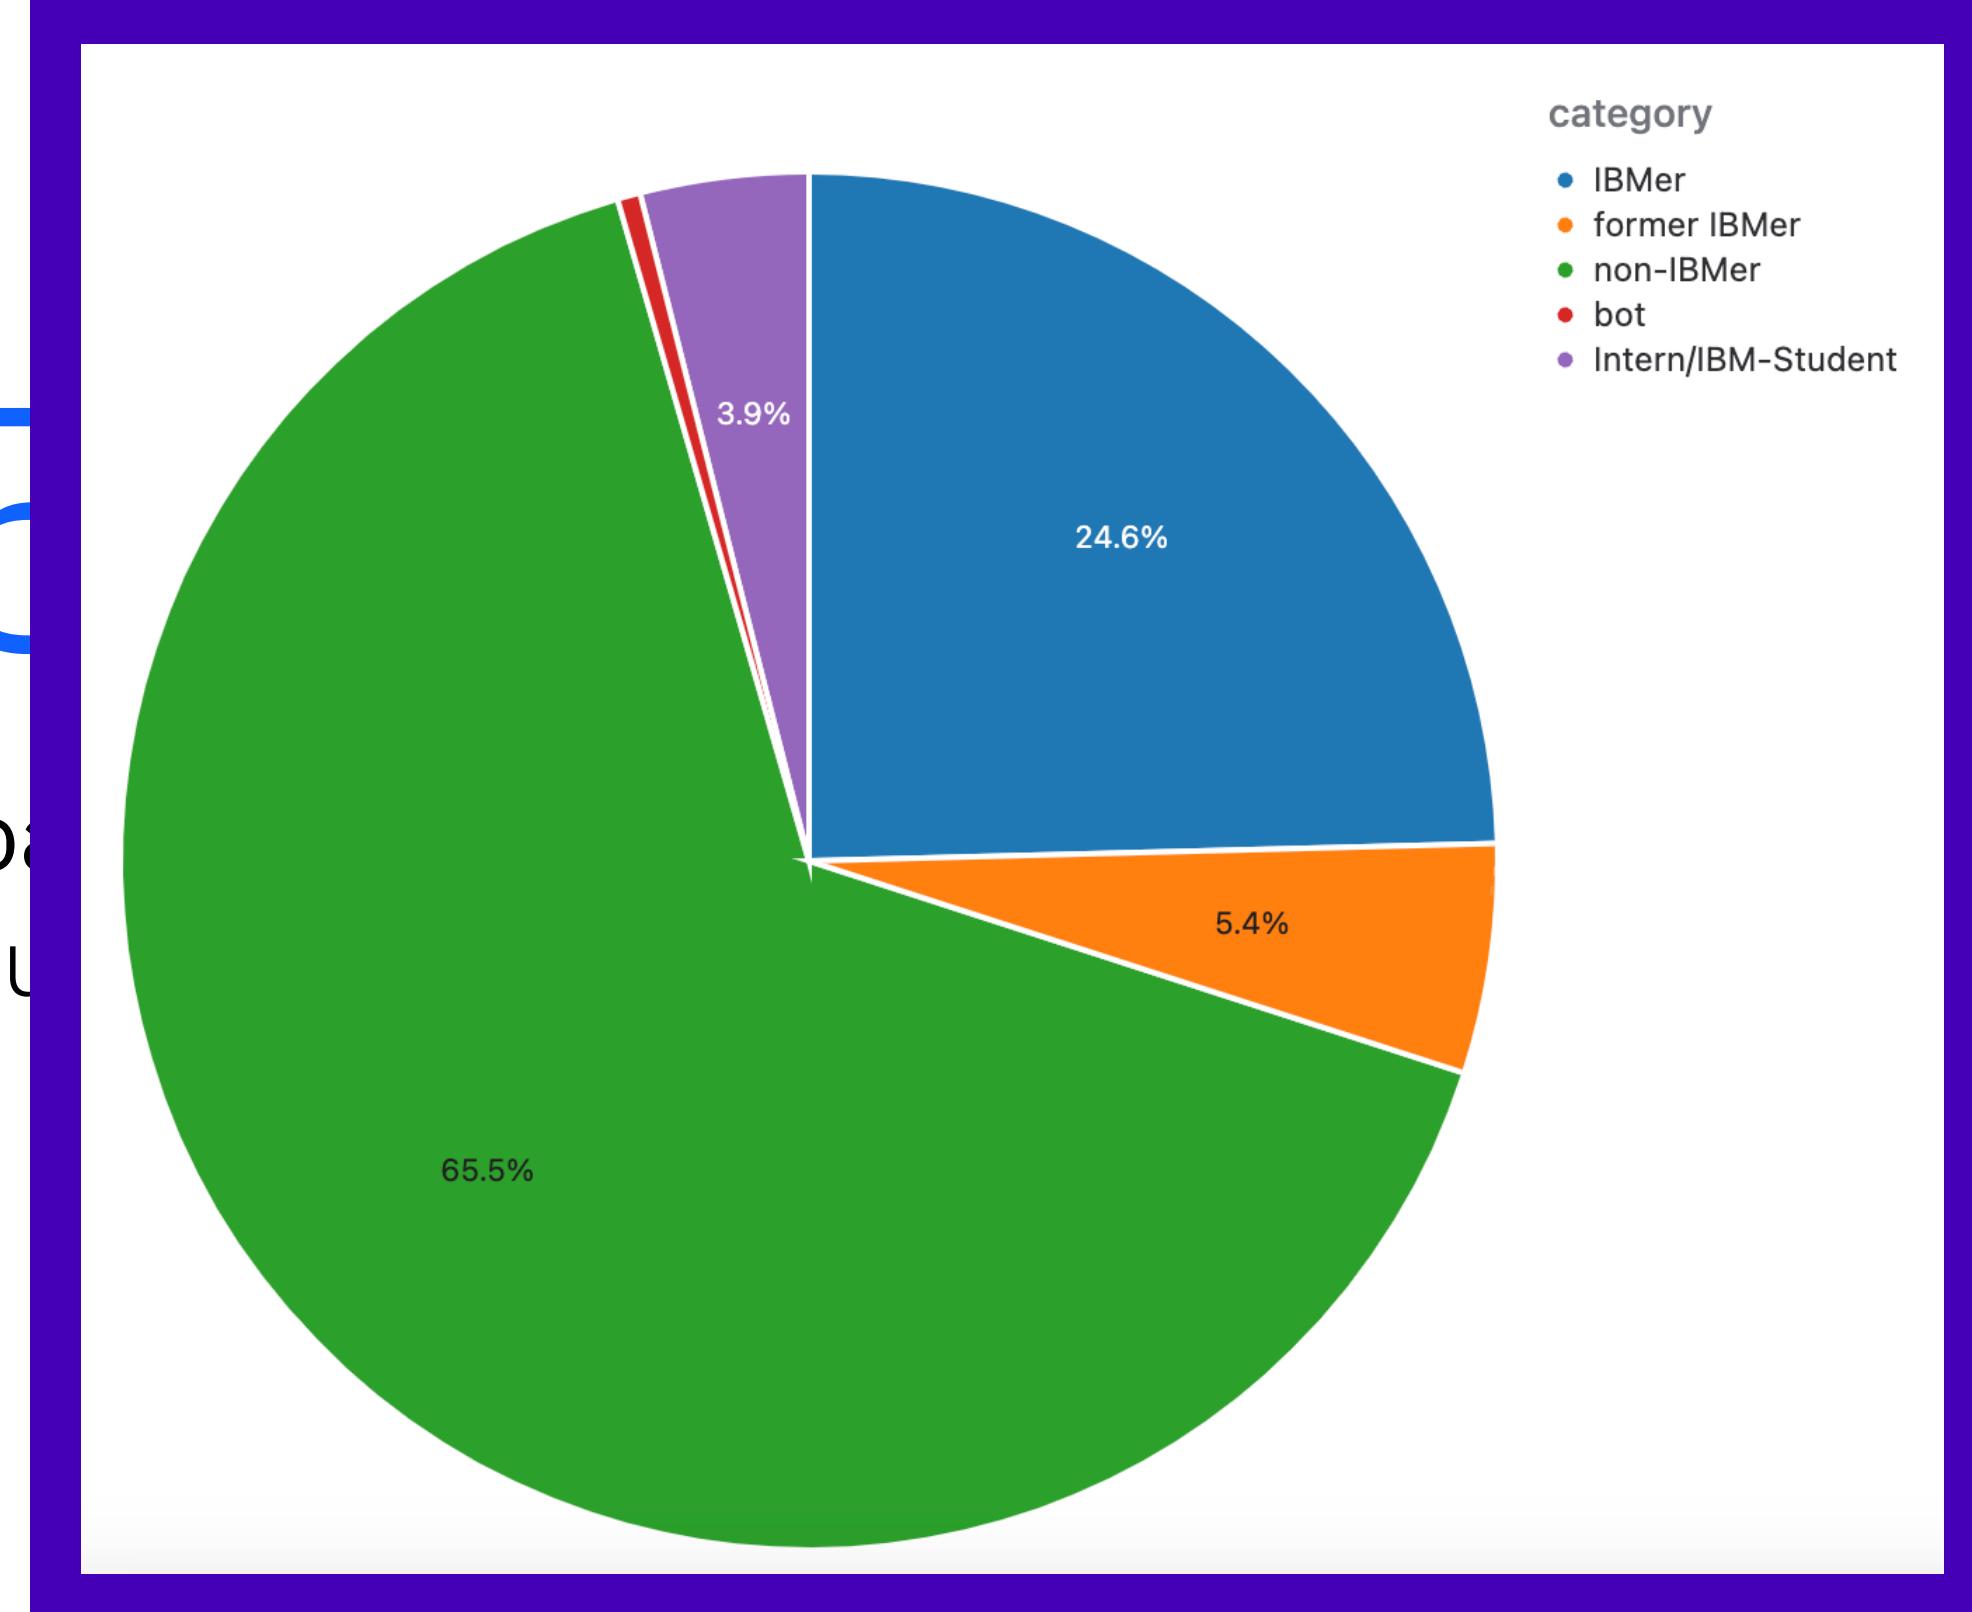
research papers extending, using, and benchmarking Qiskit

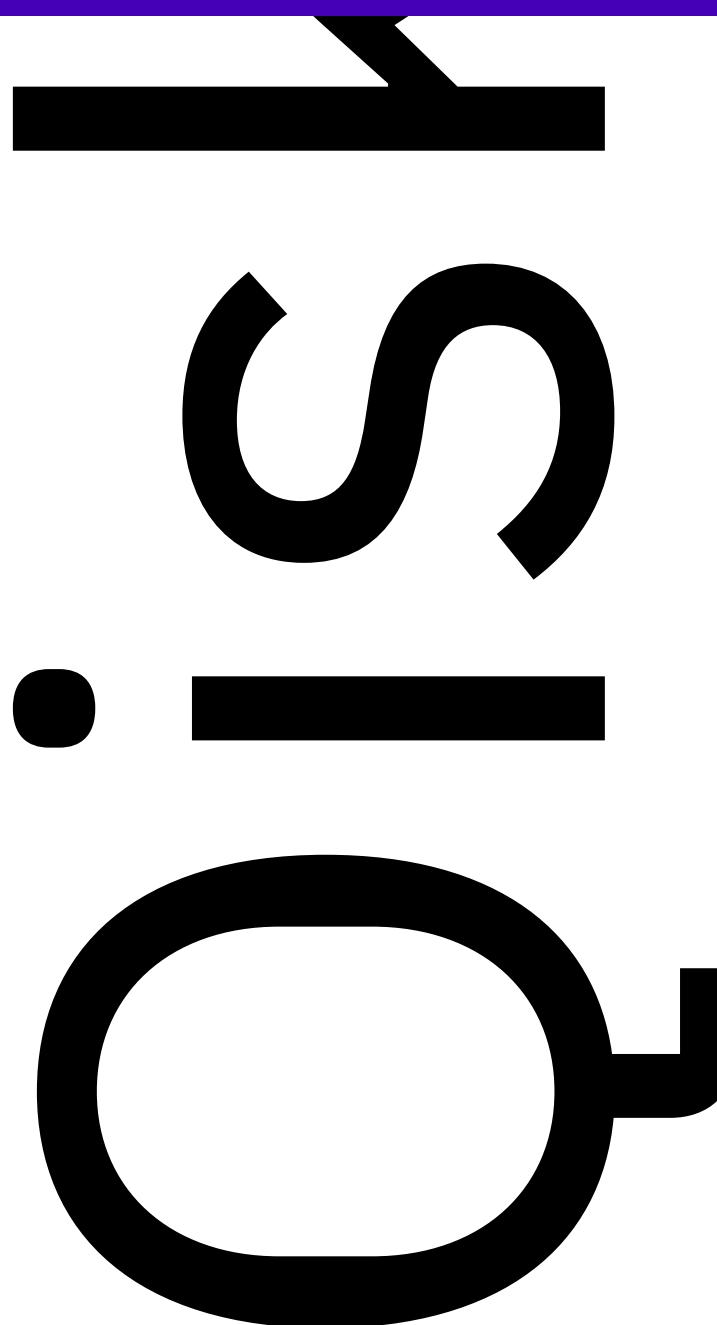
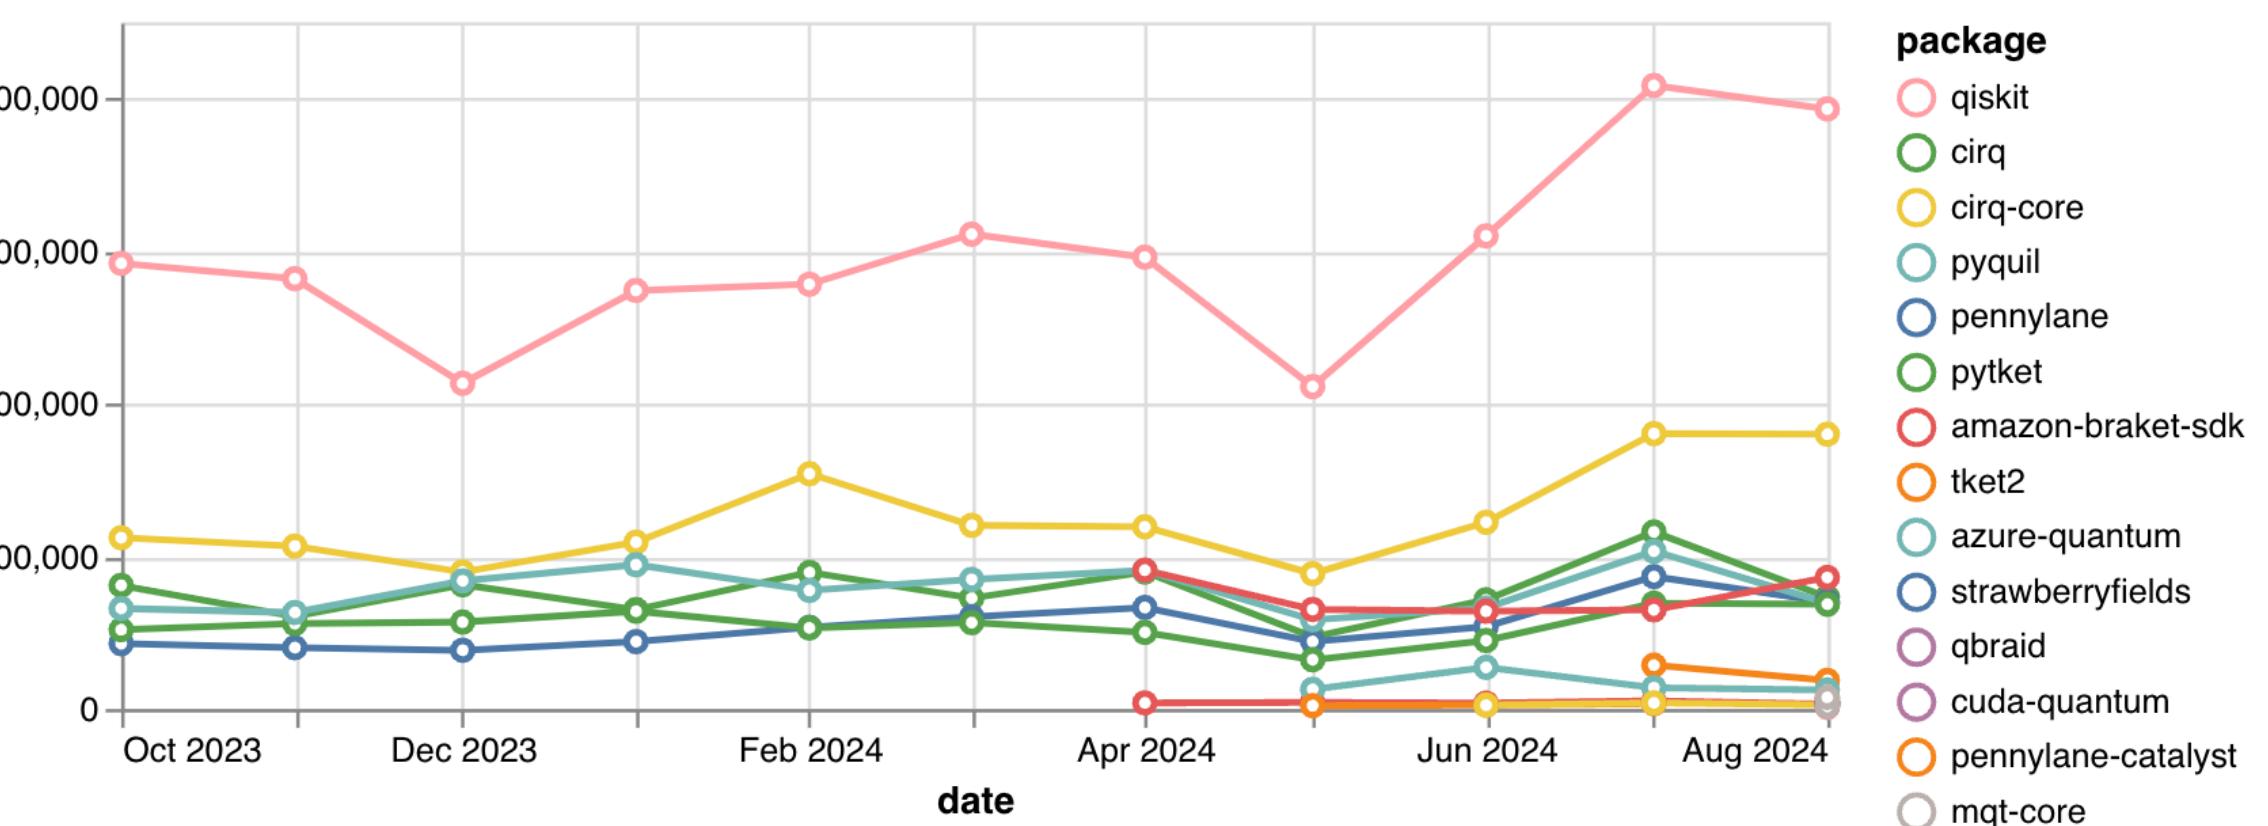


7.5

downloads
since Aug 2021

position in Unitary Fund
and
qisk.it/good-first-issues and
qisk.it/help-wanted
most used
Quantum Framework





1st

position in Unitary Fund
surveys (2022 and
2023) for most used
Quantum Framework

546

contributors, many
first-OSS timers.
2.2k forks and 4.4k
GitHub stars.

7.5 million

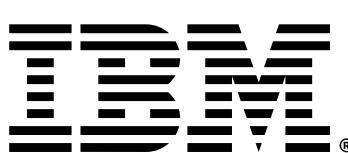
downloads via PyPI
since August 2017

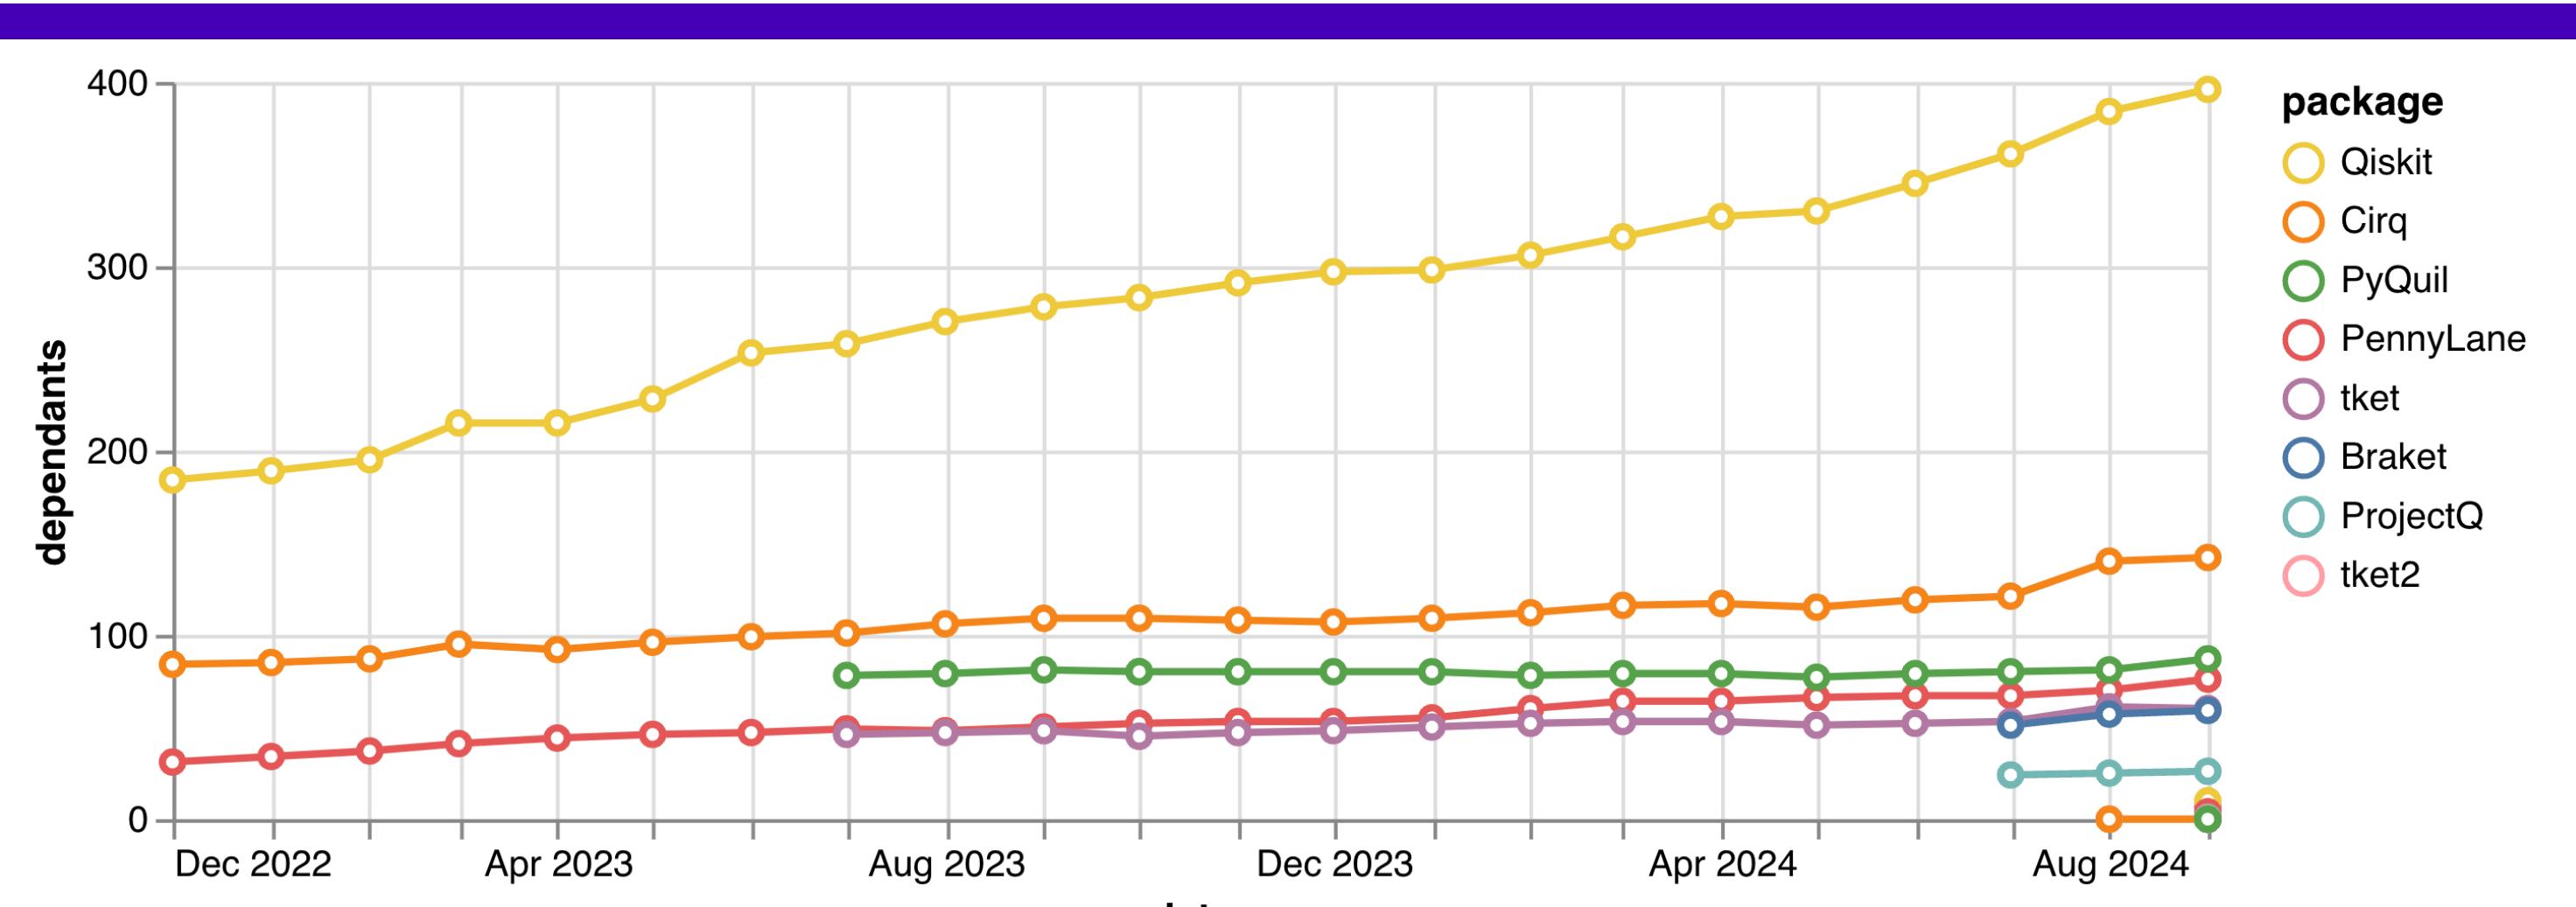
+1600

research papers
extending, using, and
benchmarking Qiskit

396

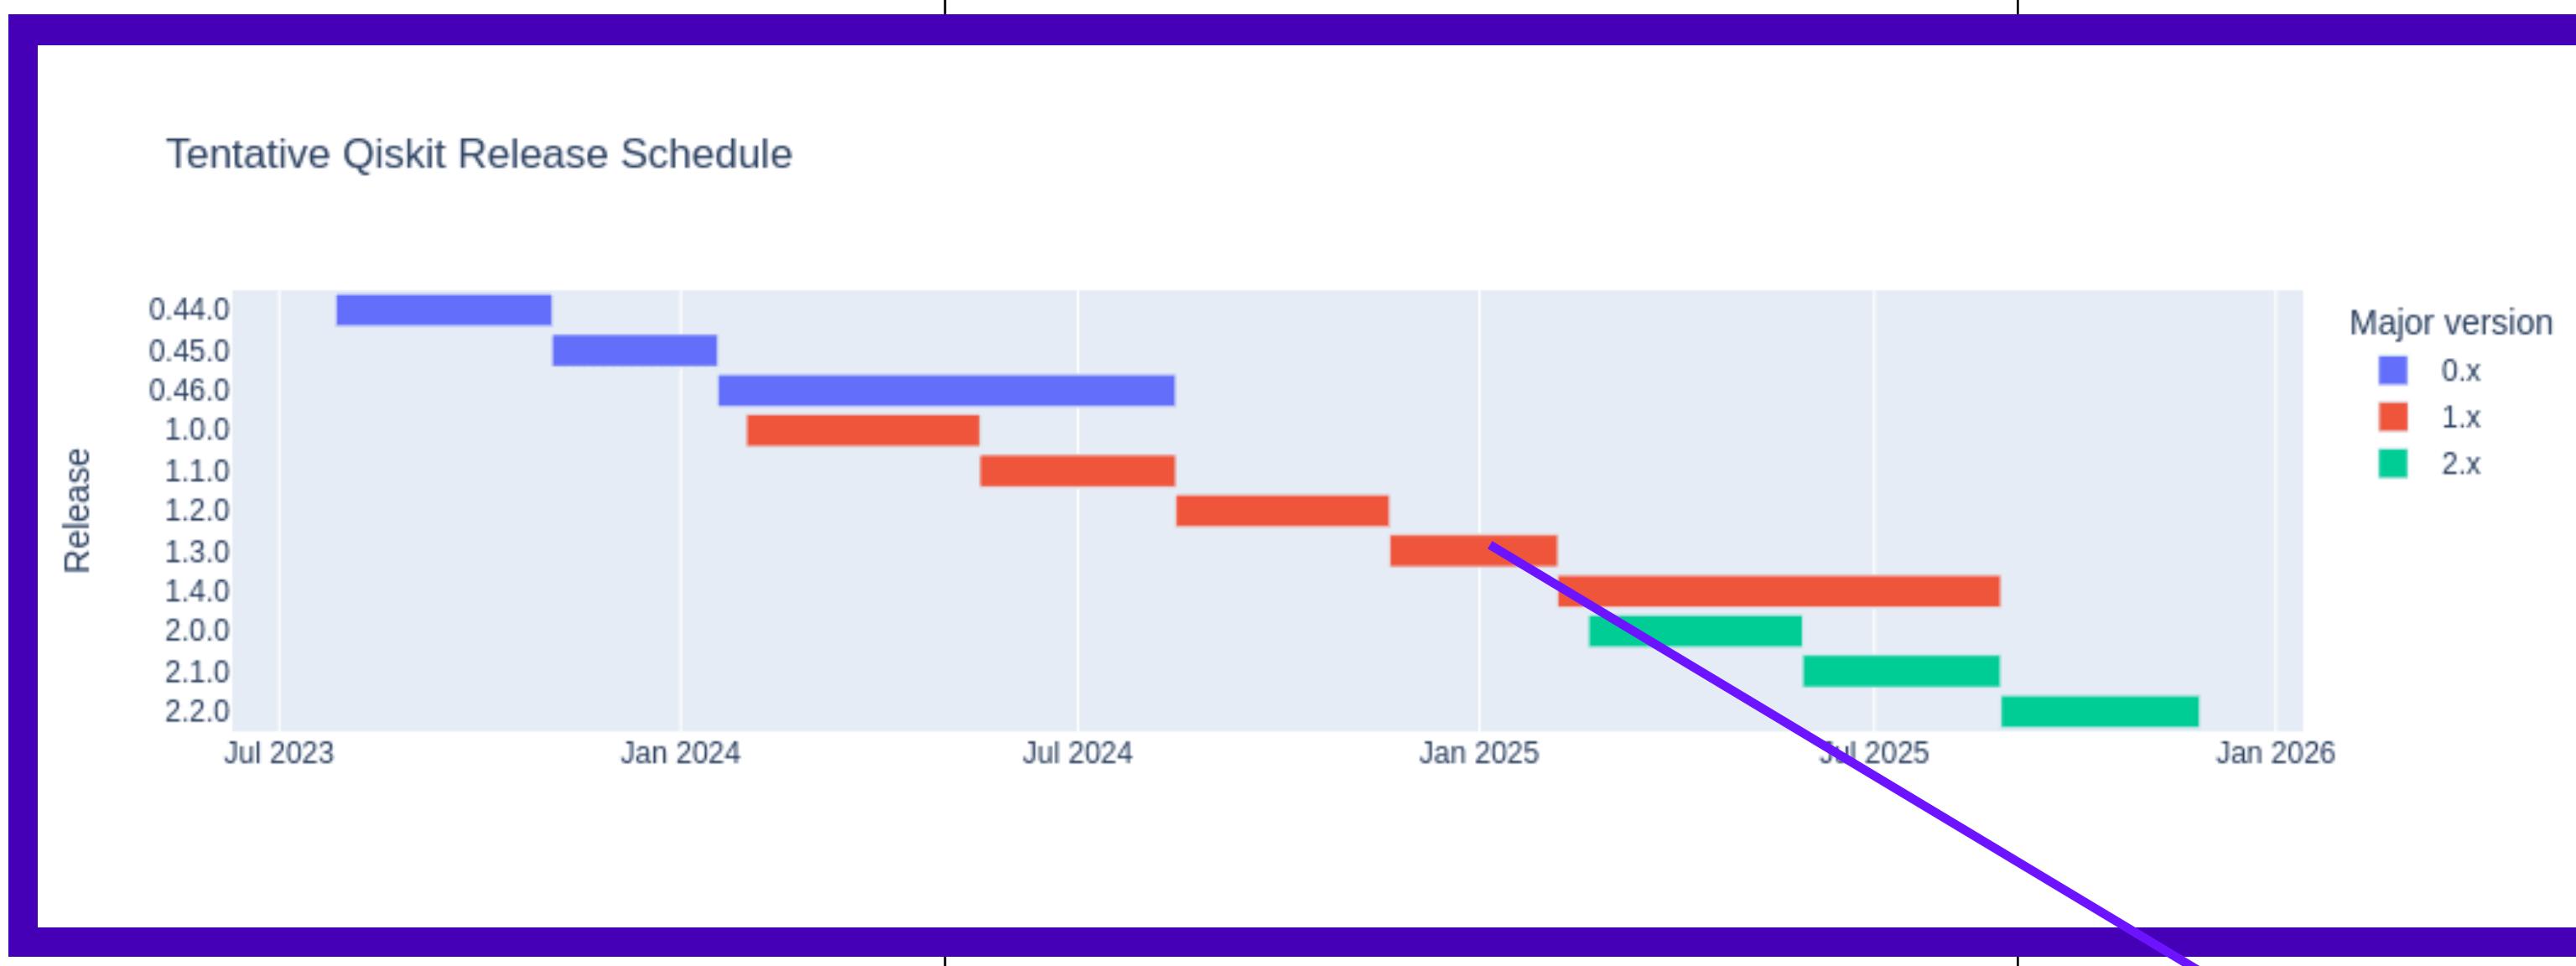
dependants, i.e
packages that directly
or indirectly depend on
Qiskit





+1600

research papers
extending, using, and
benchmarking Qiskit



million

via PyPI
2017

pip install "qiskit==1.3.0b1"

many
rs.
4.4k

396

dependants, i.e.
packages that directly
or indirectly depend on
Qiskit

IBM Quantum Learning – to learn, experiment, and prototype quantum algorithms and applications

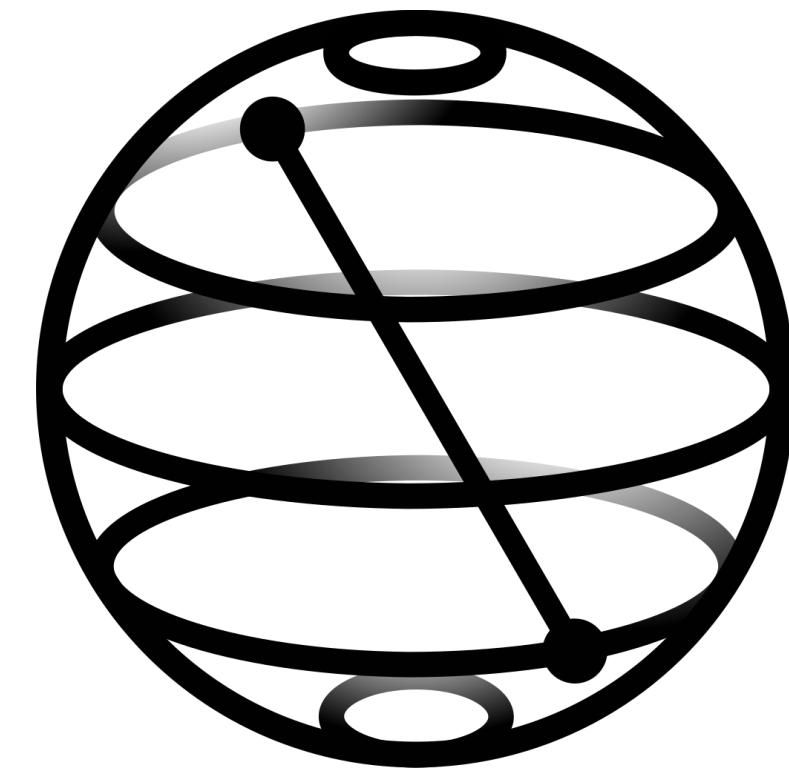
The screenshot shows the IBM Quantum Learning homepage. It features a large image of a quantum circuit with a person working on a computer. Below this, there's a section for the latest course, 'Quantum Computing in Practice', which has 2 lessons and 0% progress. The course description mentions learning about realistic potential use cases for quantum computing. There are also sections for 'Courses' with links to 'Basics of Quantum Information', 'Fundamentals of Quantum Algorithms', and 'Variational Algorithm Design'.

learning.quantum.ibm.com

IBM Quantum Blog – for news, announcements, and release reviews

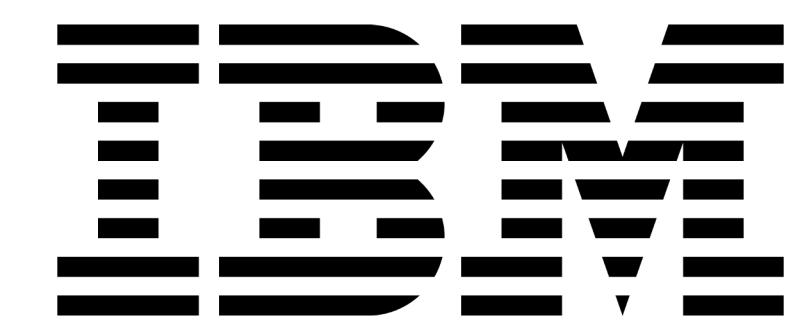
The screenshot shows the IBM Quantum Blog page. The main headline is 'Qiskit: The software for quantum performance'. Below it, there's a section for '24 results for Qiskit'. The first result is a post titled 'Defining – and citing – the Qiskit SDK' with a date of 26 Jun 2024. The second result is 'Charting the path to quantum utility at Qiskit Global Summer School 2024' with a date of 10 Jun 2024. The third result is 'Explore newly recommended notebook environments for Qiskit' with a date of 3 Jun 2024. The fourth result is 'Release news: Qiskit SDK v1.1.0 is here!' with a date of 29 May 2024. At the bottom, there's a section for 'What are ISA circuits?'.

www.ibm.com/quantum/blog



www.ibm.com/quantum/qiskit

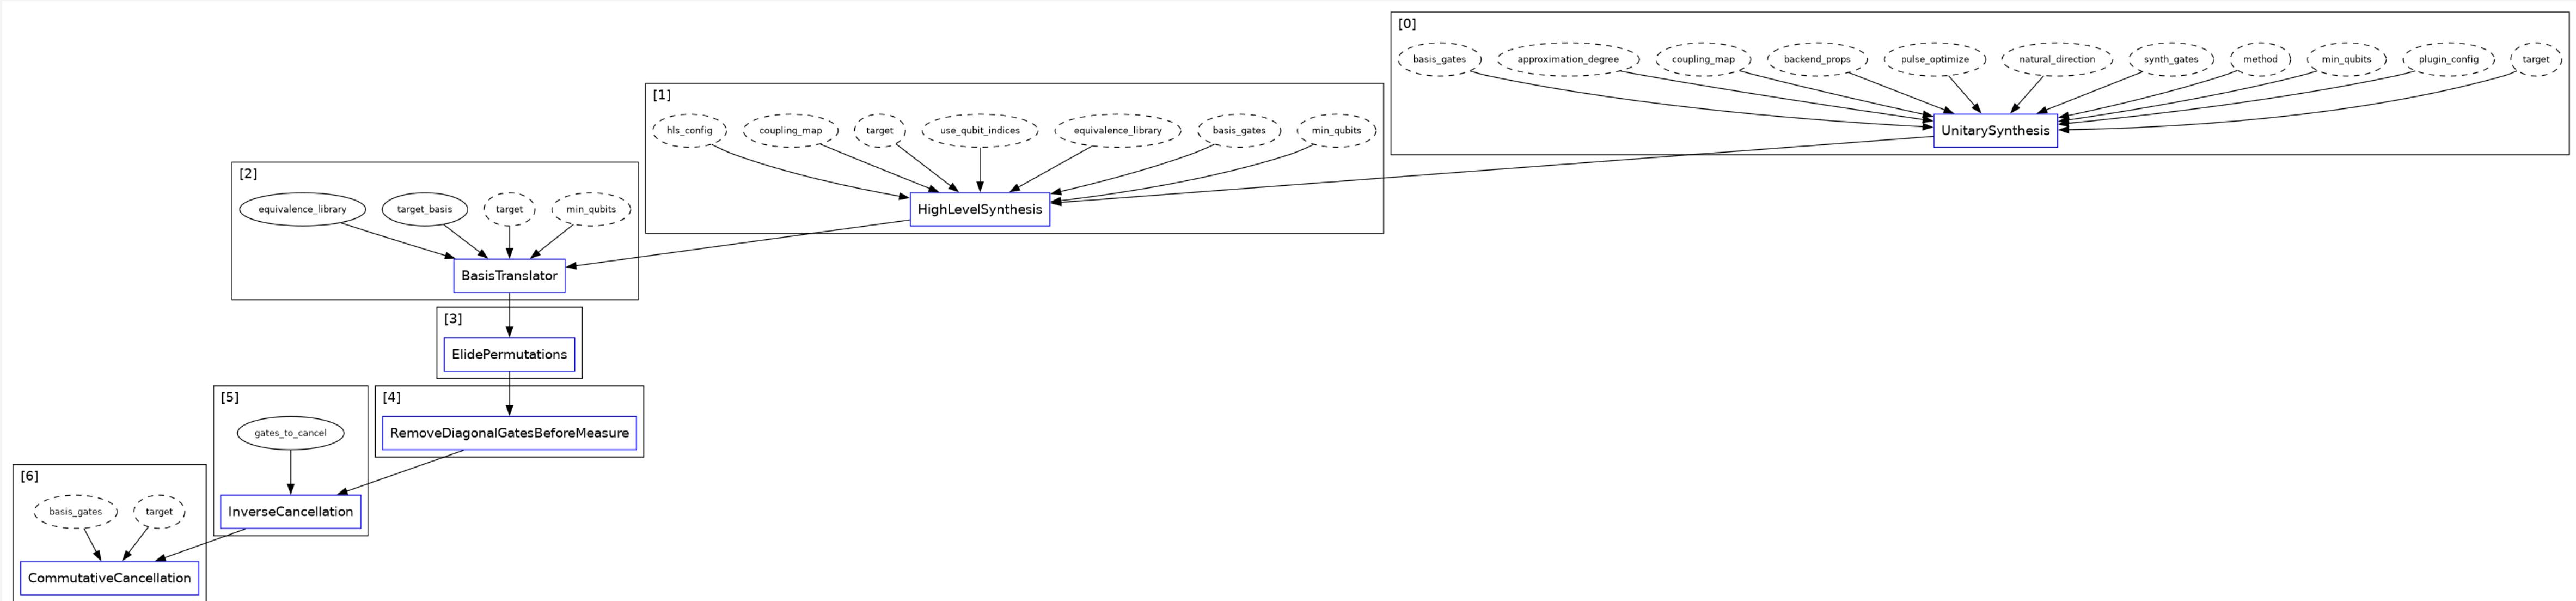
github.com/Qiskit/



Preset Pass Managers

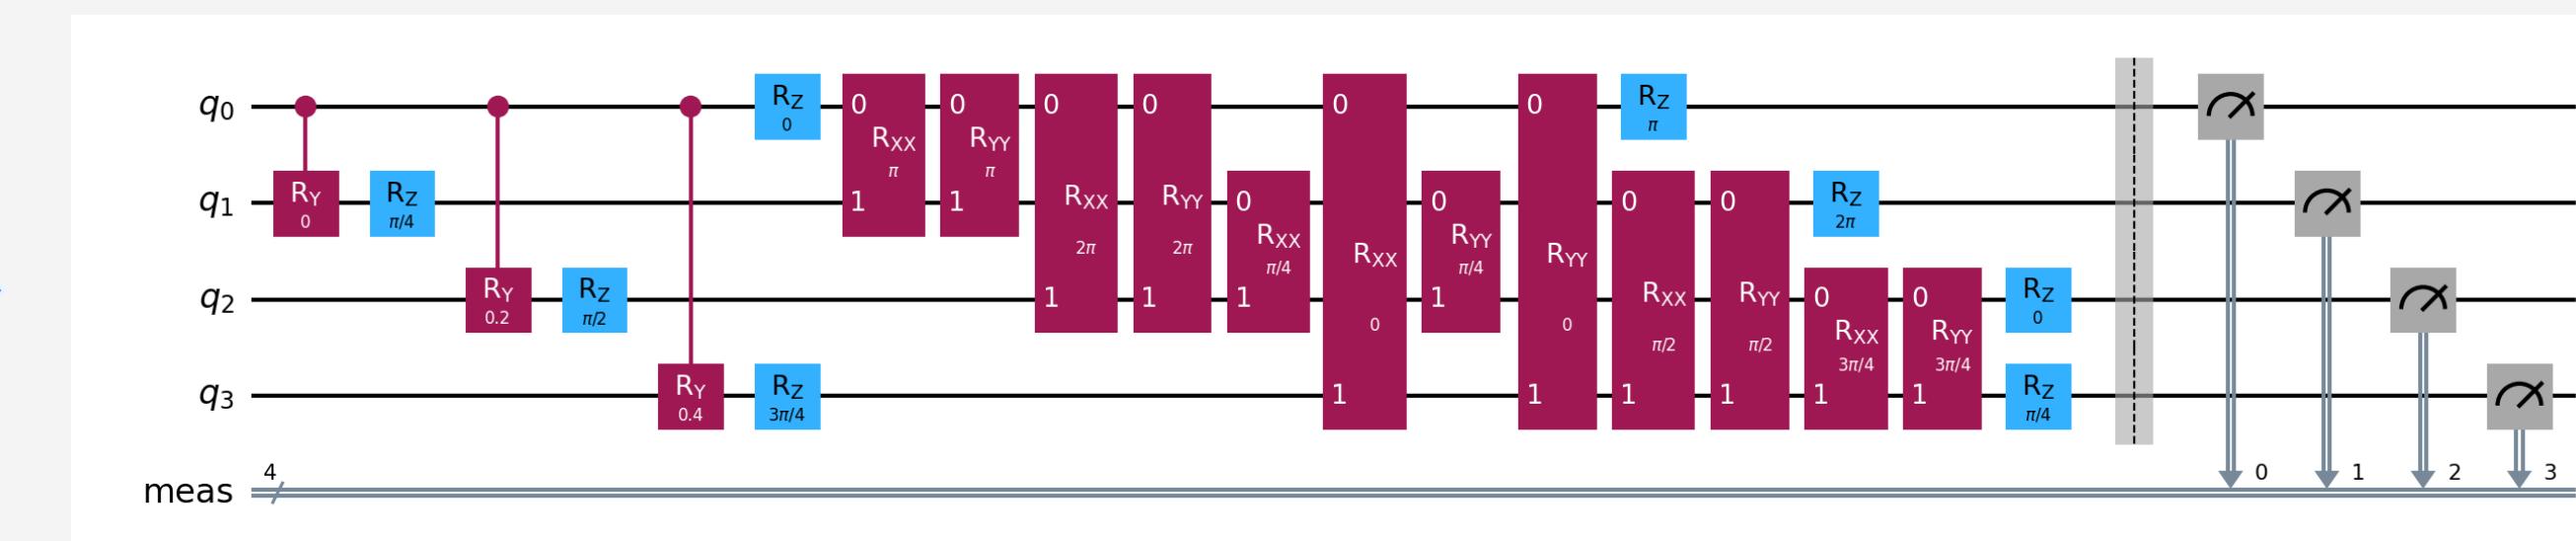
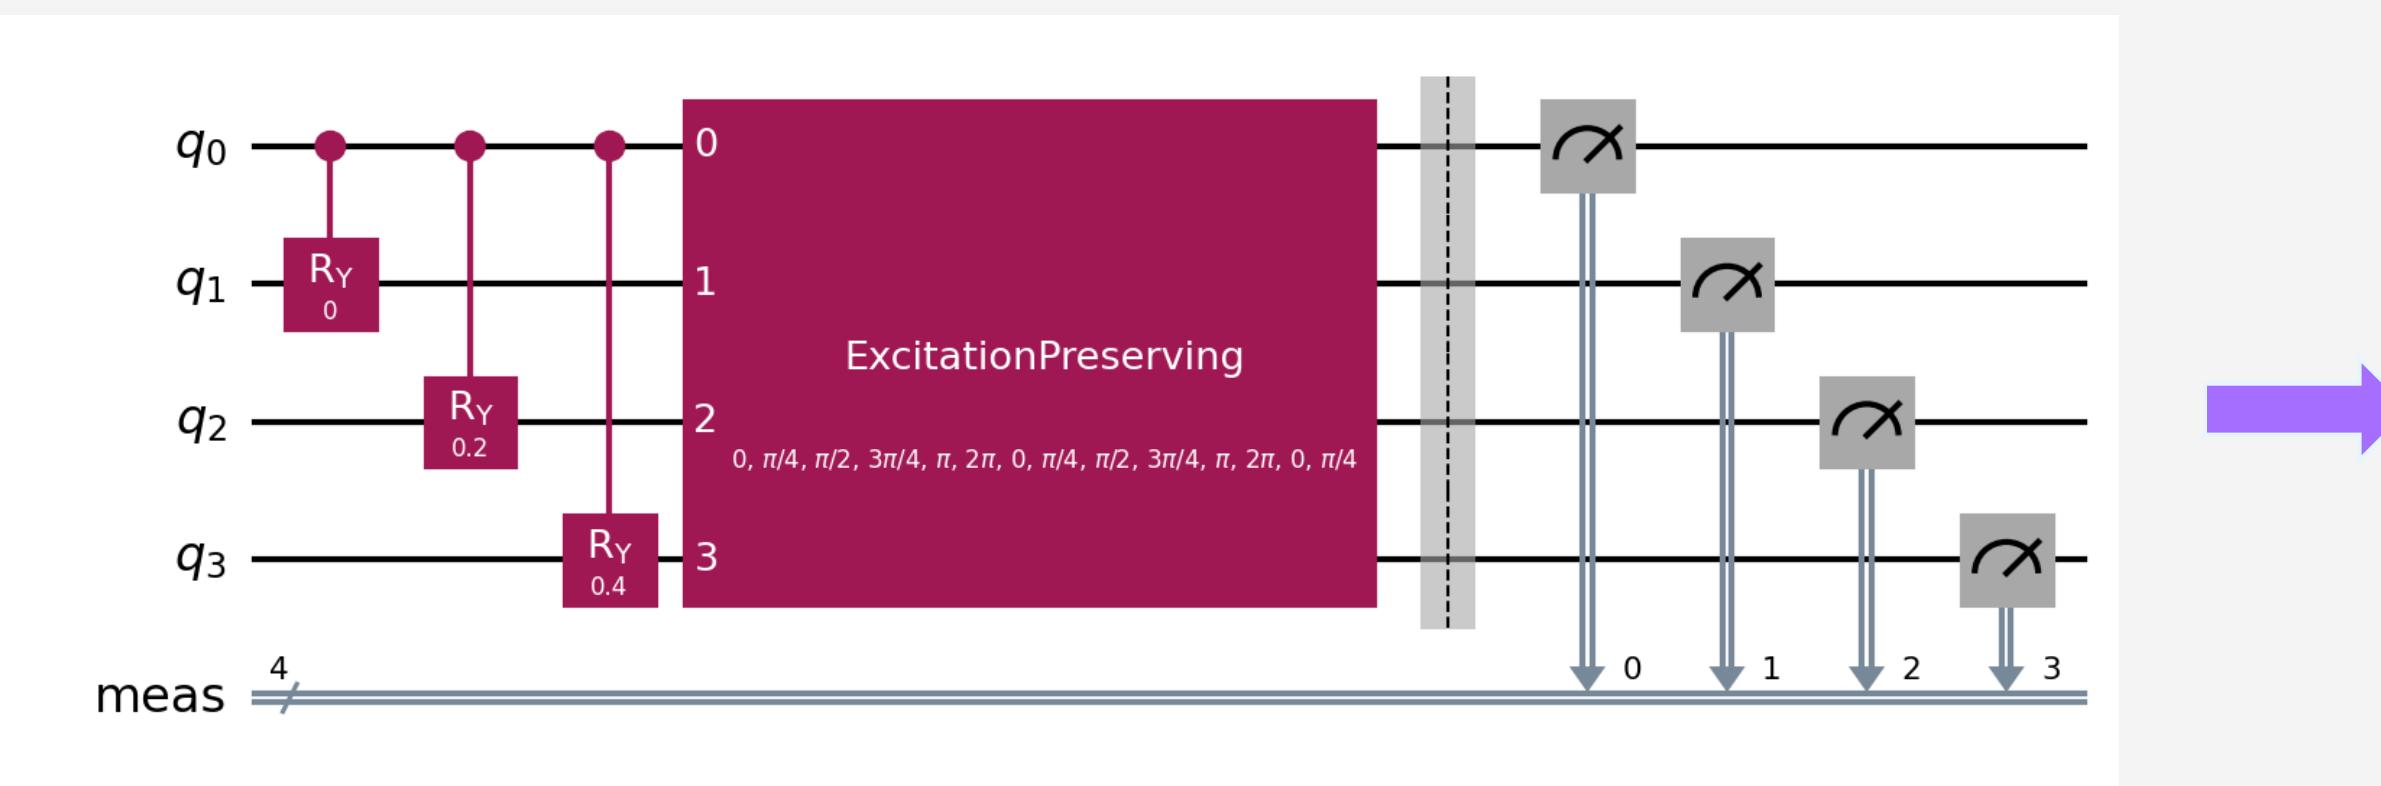
Init stage

- Stage is broken into two phases:
 - Decompose larger (≥ 3 qubit) gates to 1 and 2 qubit gates
 - Logical optimizations
- At the end of the stage, we are able to run layout (which models the circuit as a graph)



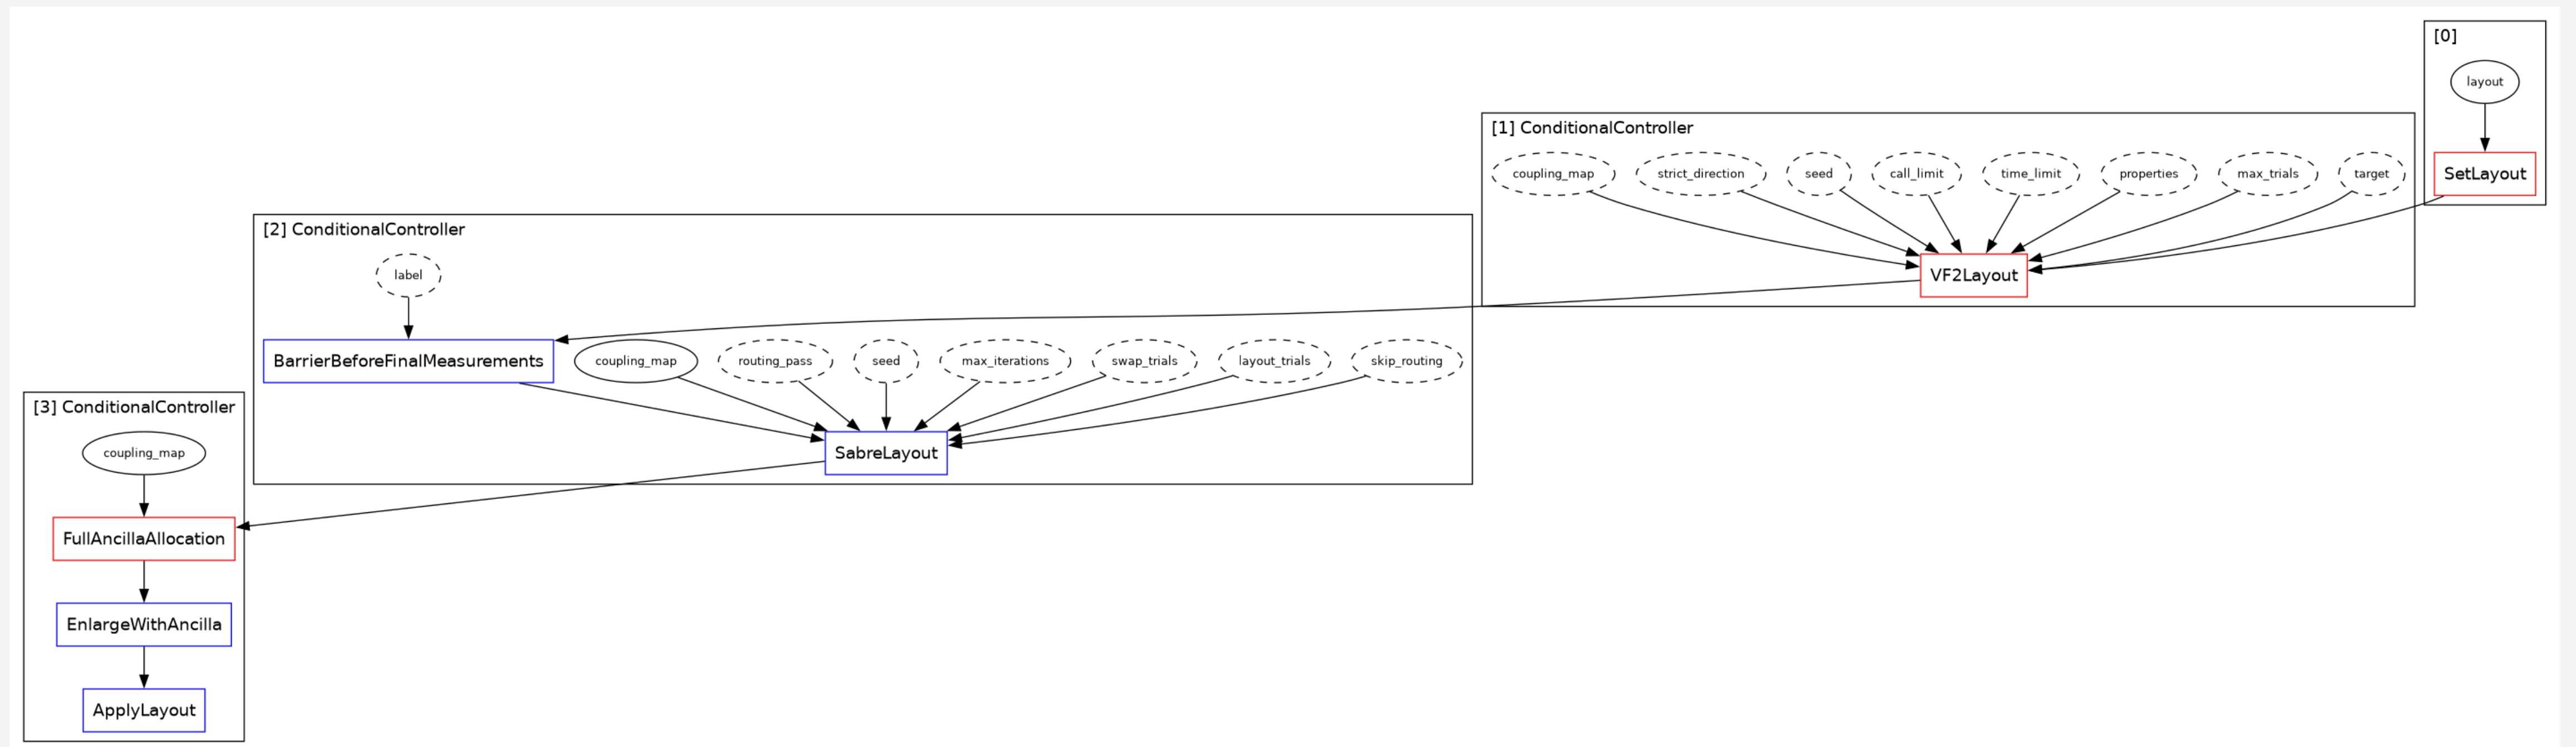
Init stage

- Stage is broken into two phases:
 - Decompose larger (≥ 3 qubit) gates to 1 and 2 qubit gates
 - Logical optimizations
- At the end of the stage, we are able to run layout (which models the circuit as a graph)

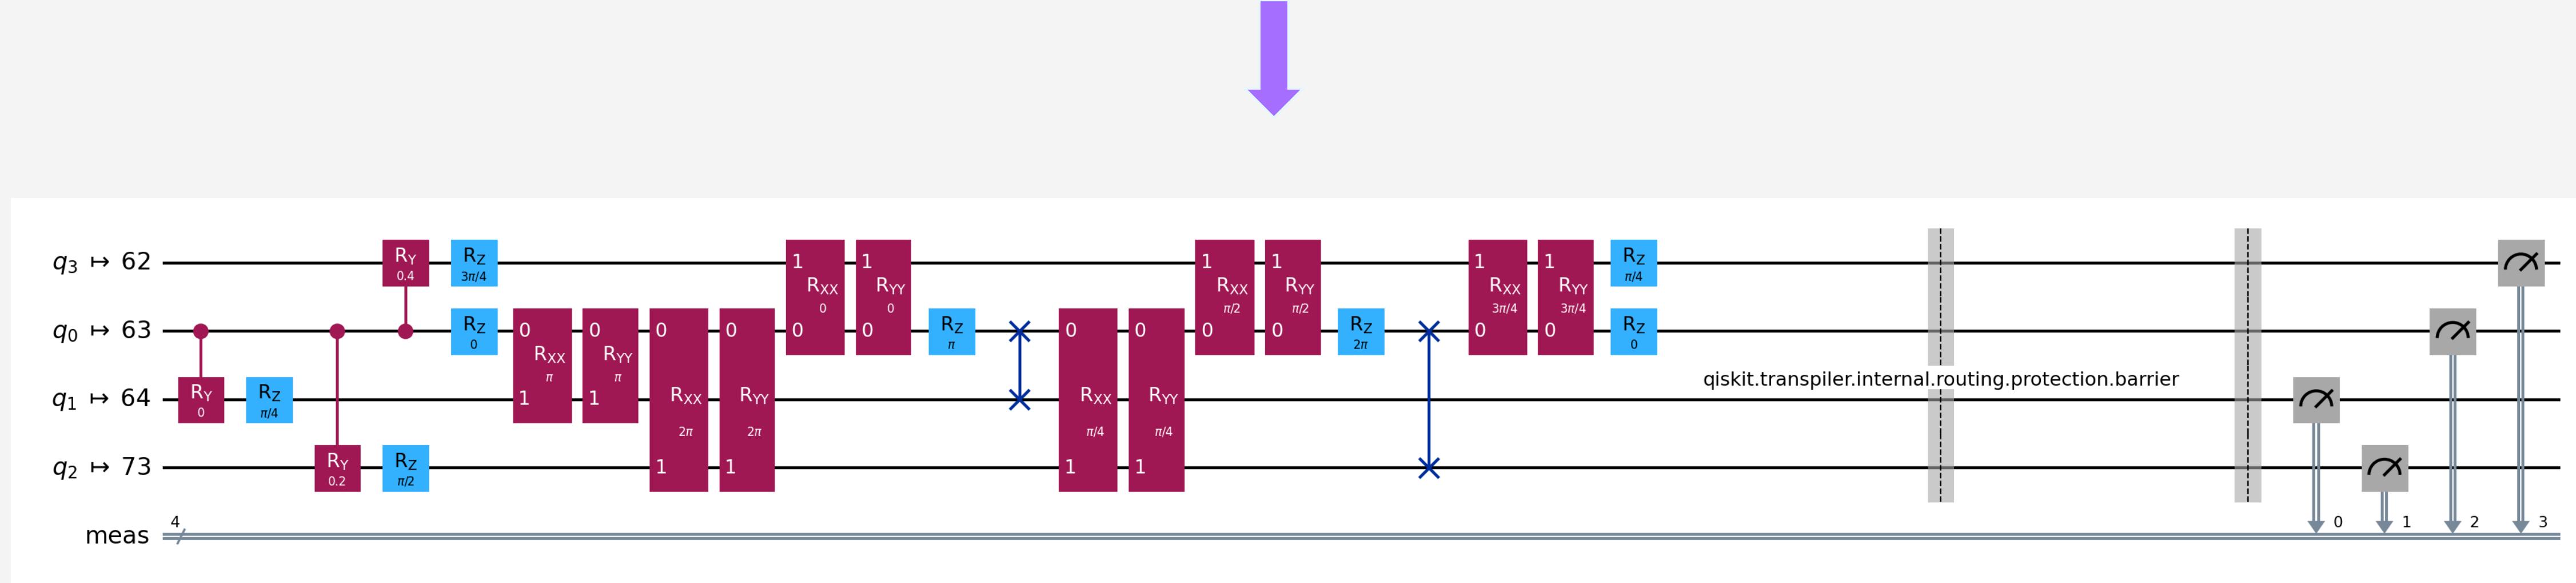
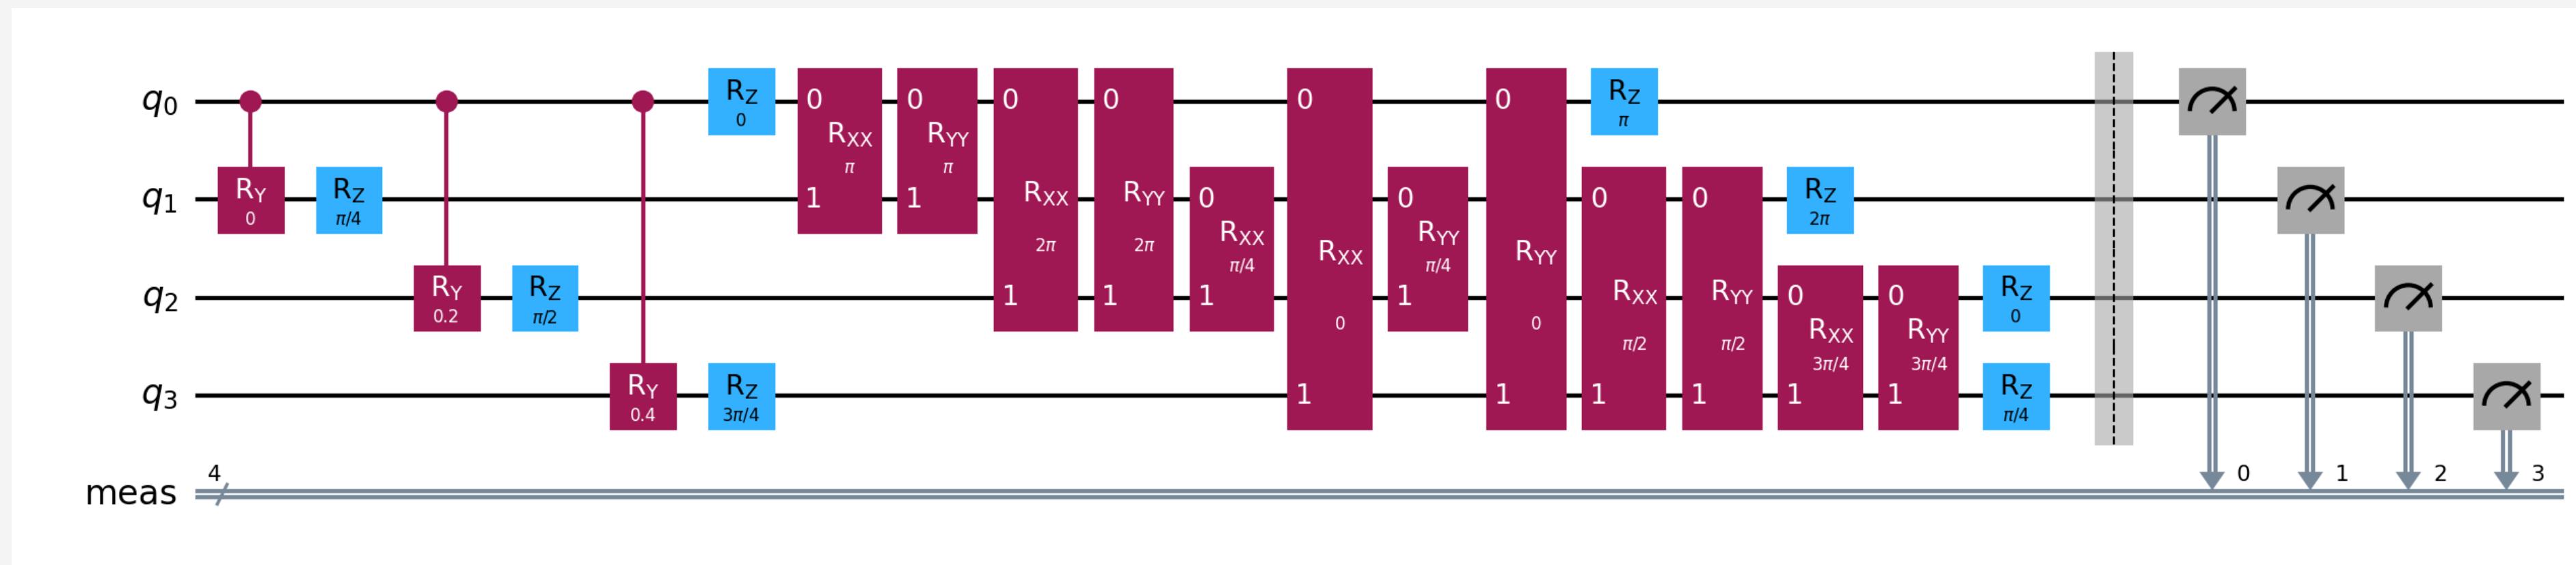


Layout stage

- This stage takes the virtual qubits from abstract circuit and maps them to a qubit as defined in the target
- Layout is of critical importance because not all qubits perform the same.
- This stage typically tries multiple techniques to find the best layout. In Qiskit we use two passes by default: **VF2Layout** and **Sabre**
- A poor layout can induce more Swaps because of the connectivity



Layout stage



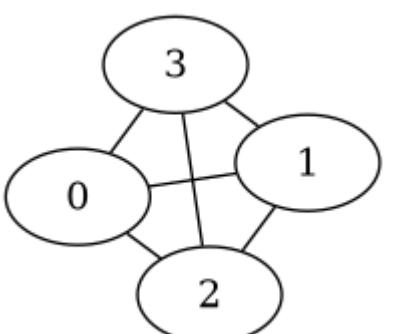
VF2Layout

This pass creates a graph from the 2 qubit interactions in the circuit and tries to find an isomorphic subgraph of the connectivity graph of the target. If a subgraph is found than that indicates a layout which does not require swaps can be used. If a layout is found, then the pass will search for additional layouts and pick the one with the lowest predicted error rate.

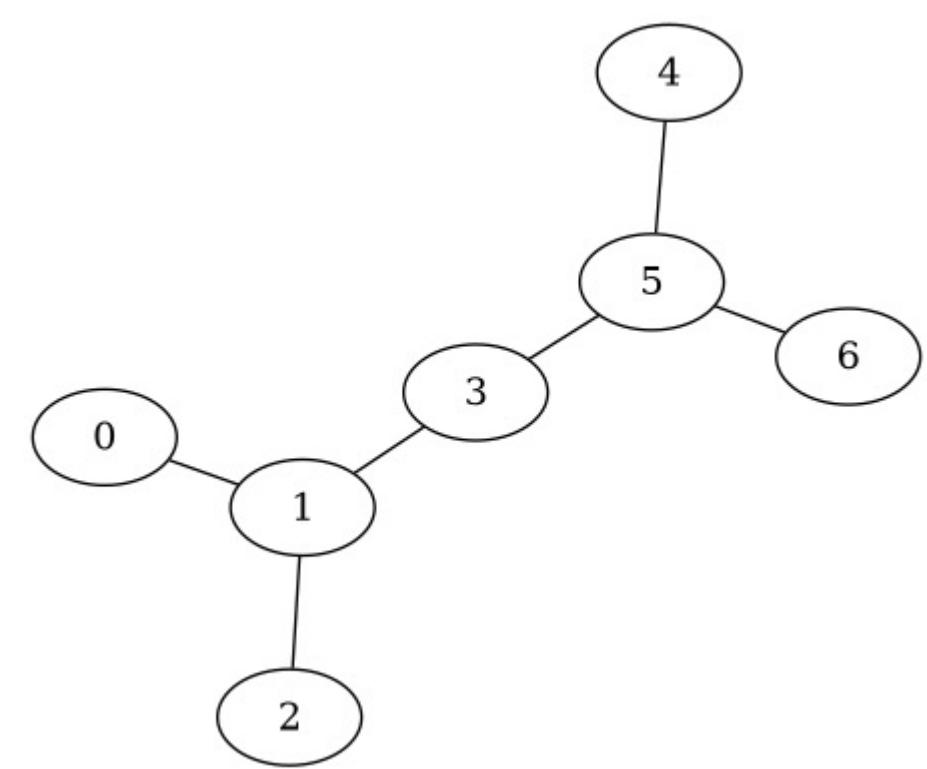
It leverages the [rustworkx's](#) library's VF2 algorithm [implementation](#) to efficiently search for an isomorphism.

<https://docs.quantum.ibm.com/api/qiskit/qiskit.transpiler.passes.VF2Layout>

Interaction Graph



Connectivity Graph



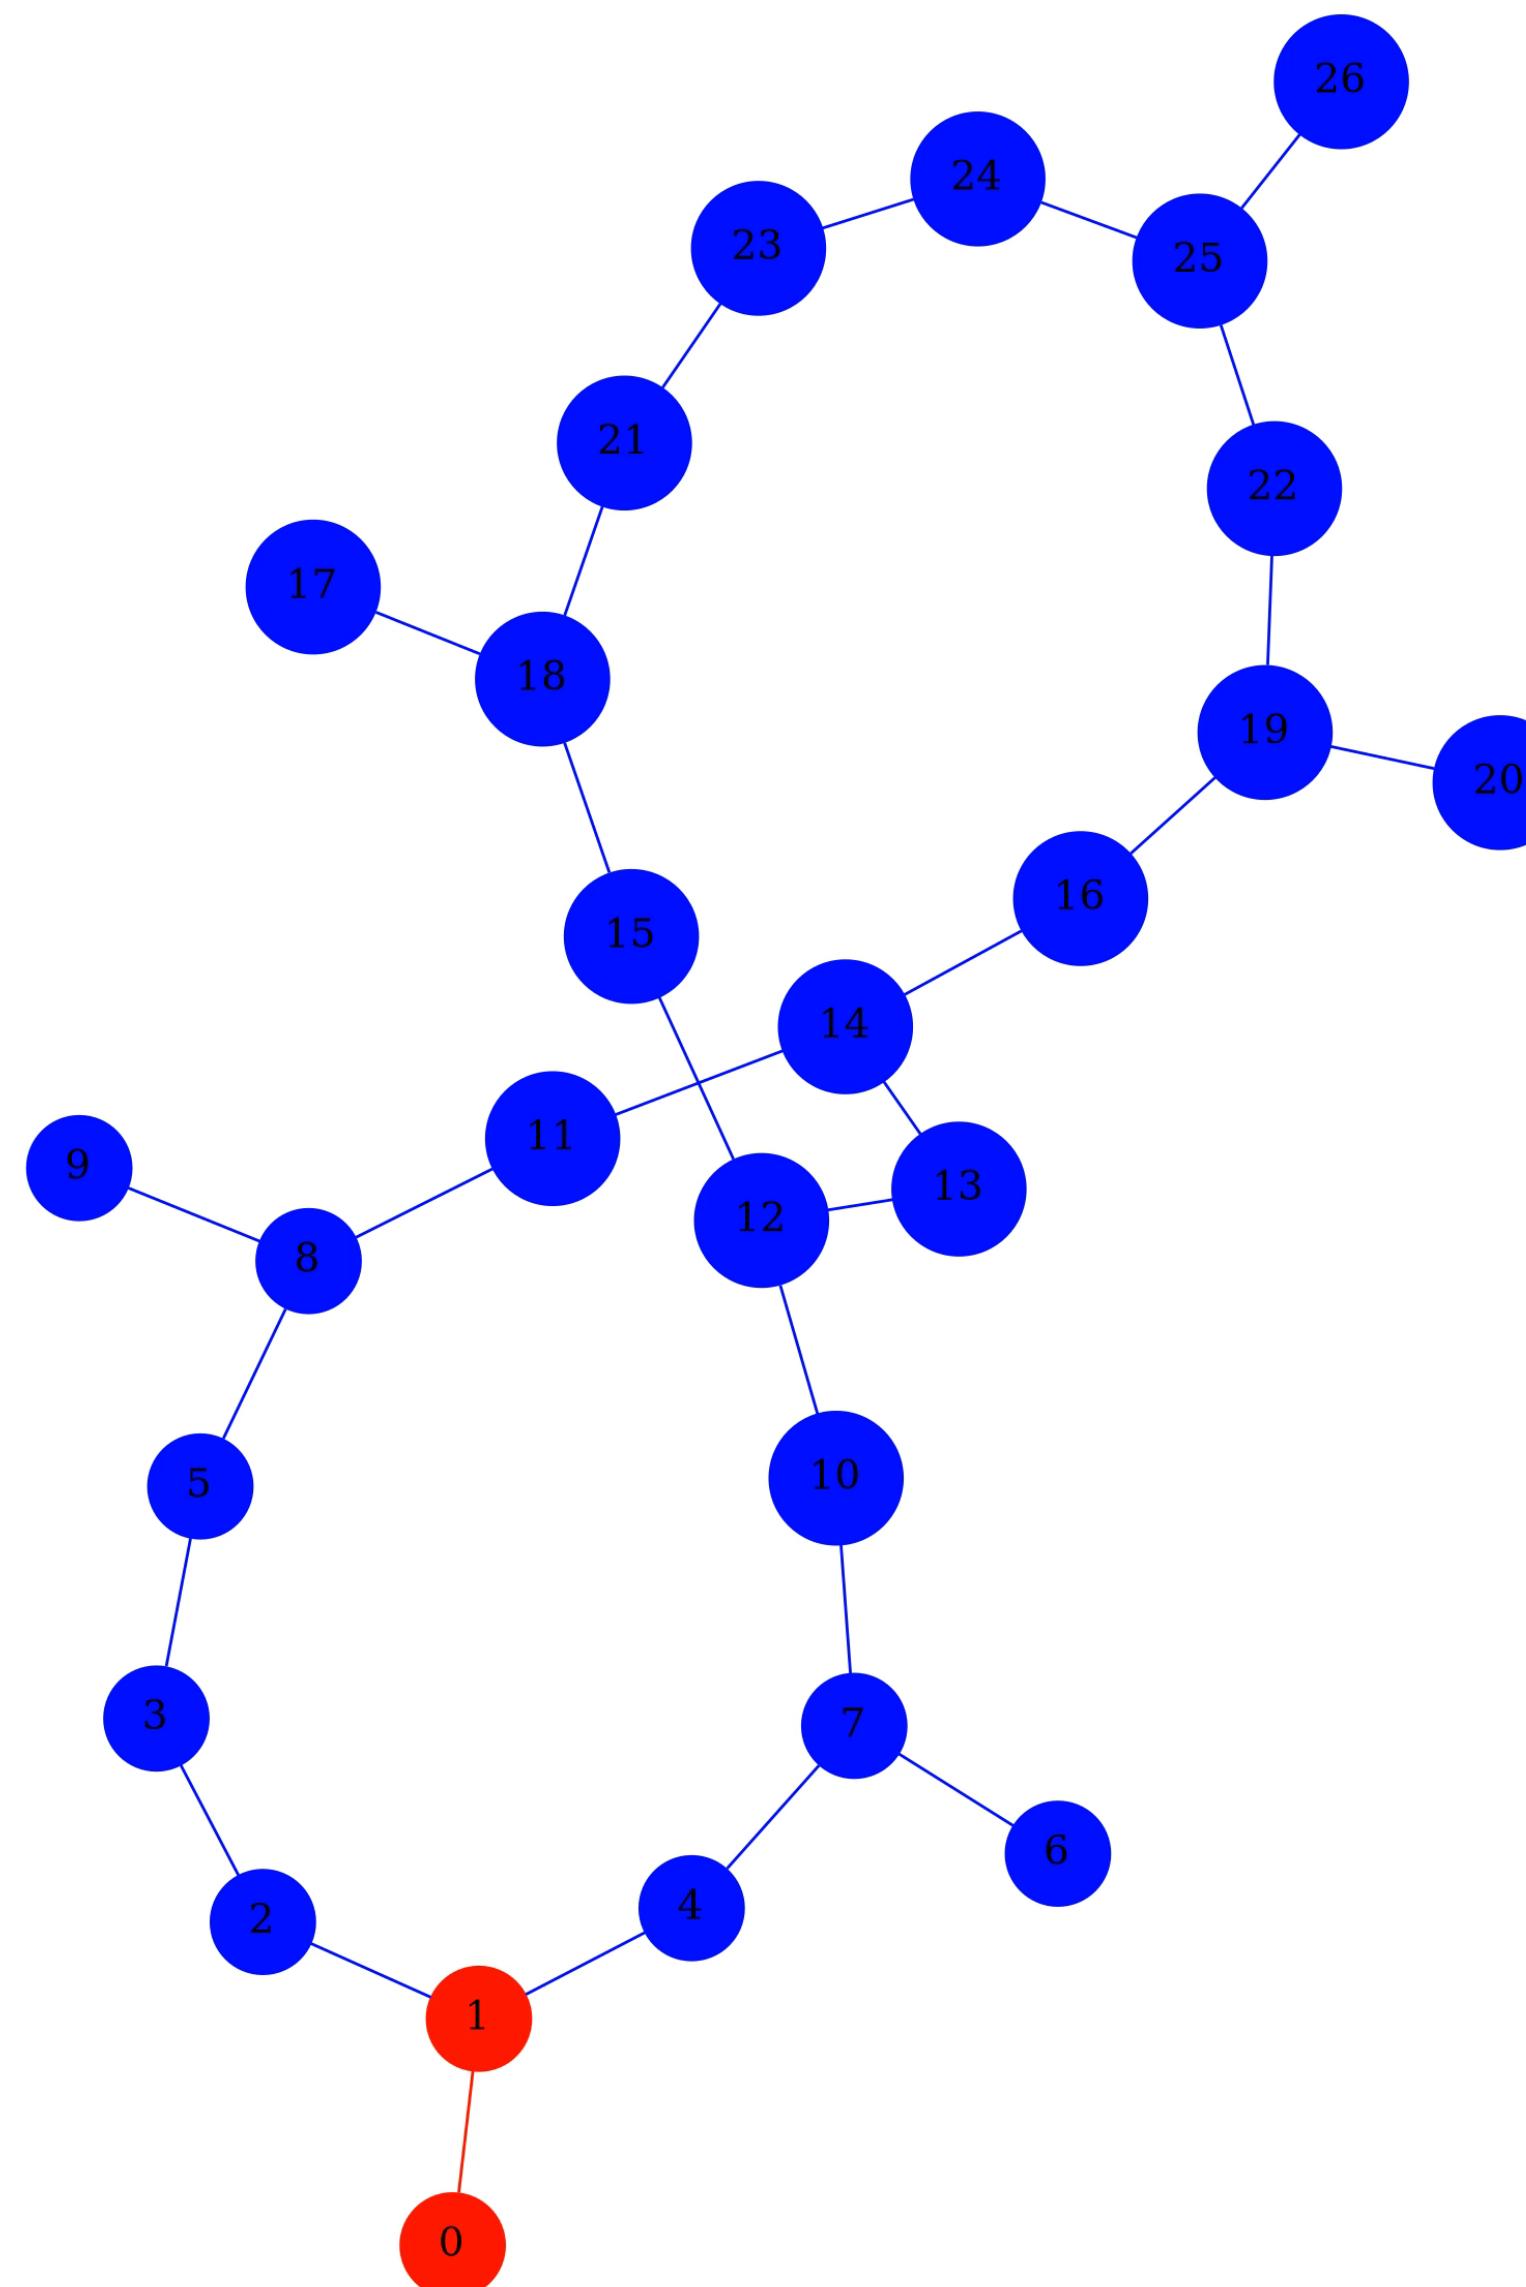
Sabre

If `VF2Layout` does not find a layout then we run `SabreLayout` pass. This pass is based on the SABRE (SWAP-based BidiREctional) algorithm originally detailed in: <https://arxiv.org/pdf/1809.02573.pdf>.

The SABRE algorithm works by starting with an initial random guess, then running a routing algorithm on it to insert swap gates and instead of inserting a swap it swaps the qubits in the layout. It fully "routes" the circuit then reverses the edge direction in the dag and repeats. This is performed multiple times, doing this will minimize the number of swaps needed.

The implementation in Qiskit has significantly changed and improved on the algorithm from the original paper.

<https://docs.quantum.ibm.com/api/qiskit/qiskit.transpiler.passes.SabreLayout>



Sabre

If `VF2Layout` does not find a layout then we run `SabreLayout` pass. This pass is based on the SABRE (SWAP-based BiDiREctional) algorithm originally detailed in: <https://arxiv.org/pdf/1809.02573.pdf>.

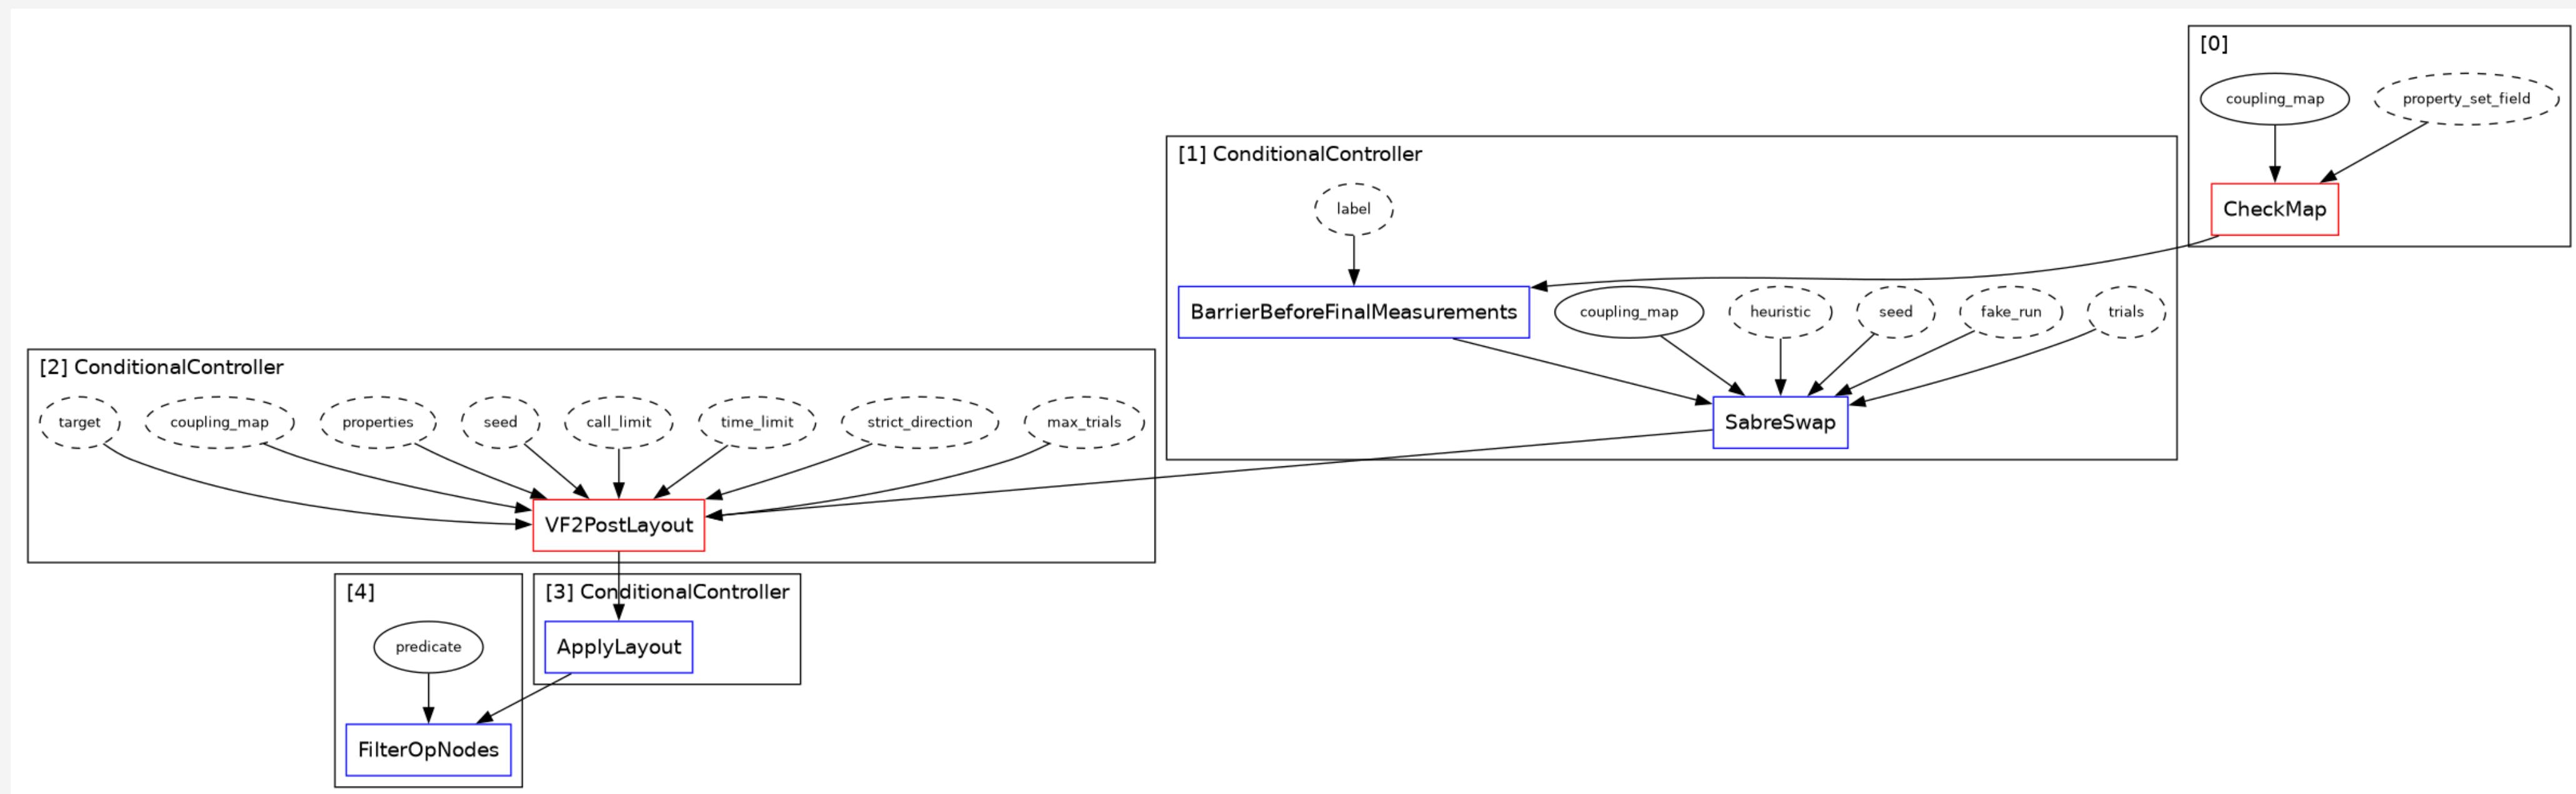
The SABRE algorithm works by starting with an initial random guess, then running a routing algorithm on it to insert swap gates and instead of inserting a swap it swaps the qubits in the layout. It fully "routes" the circuit then reverses the edge direction in the dag and repeats. This is performed multiple times, doing this will minimize the number of swaps needed.

The implementation in Qiskit has significantly changed and improved on the algorithm from the original paper.

<https://docs.quantum.ibm.com/api/qiskit/qiskit.transpiler.passes.SabreLayout>

Routing stage

- This stage takes the circuit with a layout set and inserts Swap gates where necessary.
- The default pass for this is **SabreSwap**.
- If running with defaults however this pass is skipped for efficiency, and it's run as part of **SabreLayout** in the Layout stage.
- We re-run layout with **VF2PostLayout** after swap insertion to search for a better layout



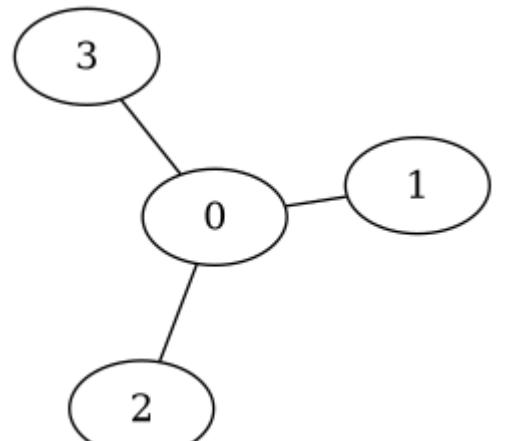
VF2PostLayout

Just as with VF2Layout this pass creates a graph from the 2 qubit interactions in the circuit and tries to find an isomorphic subgraph of the connectivity graph of the target. The difference is by running after routing we know that there is at least one isomorphic subgraph. This pass then searches for other layouts and picks the one found with the lowest predicted error rate.

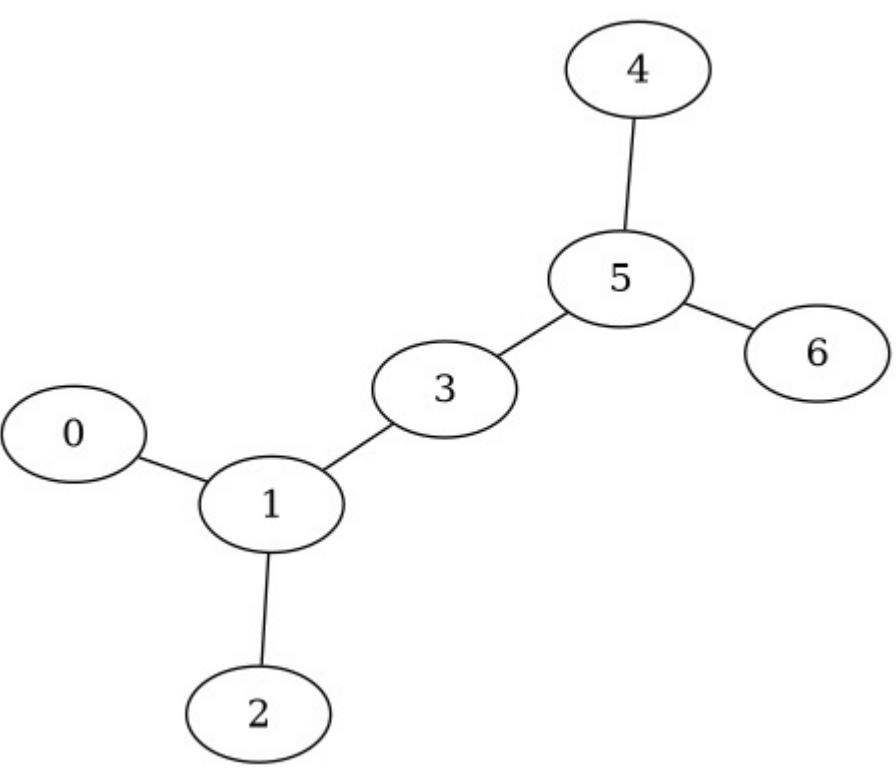
<https://journals.aps.org/prxquantum/abstract/10.1103/PRXQuantum.4.010327>

<https://docs.quantum.ibm.com/api/qiskit/qiskit.transpiler.passes.VF2PostLayout>

Interaction Graph



Connectivity Graph

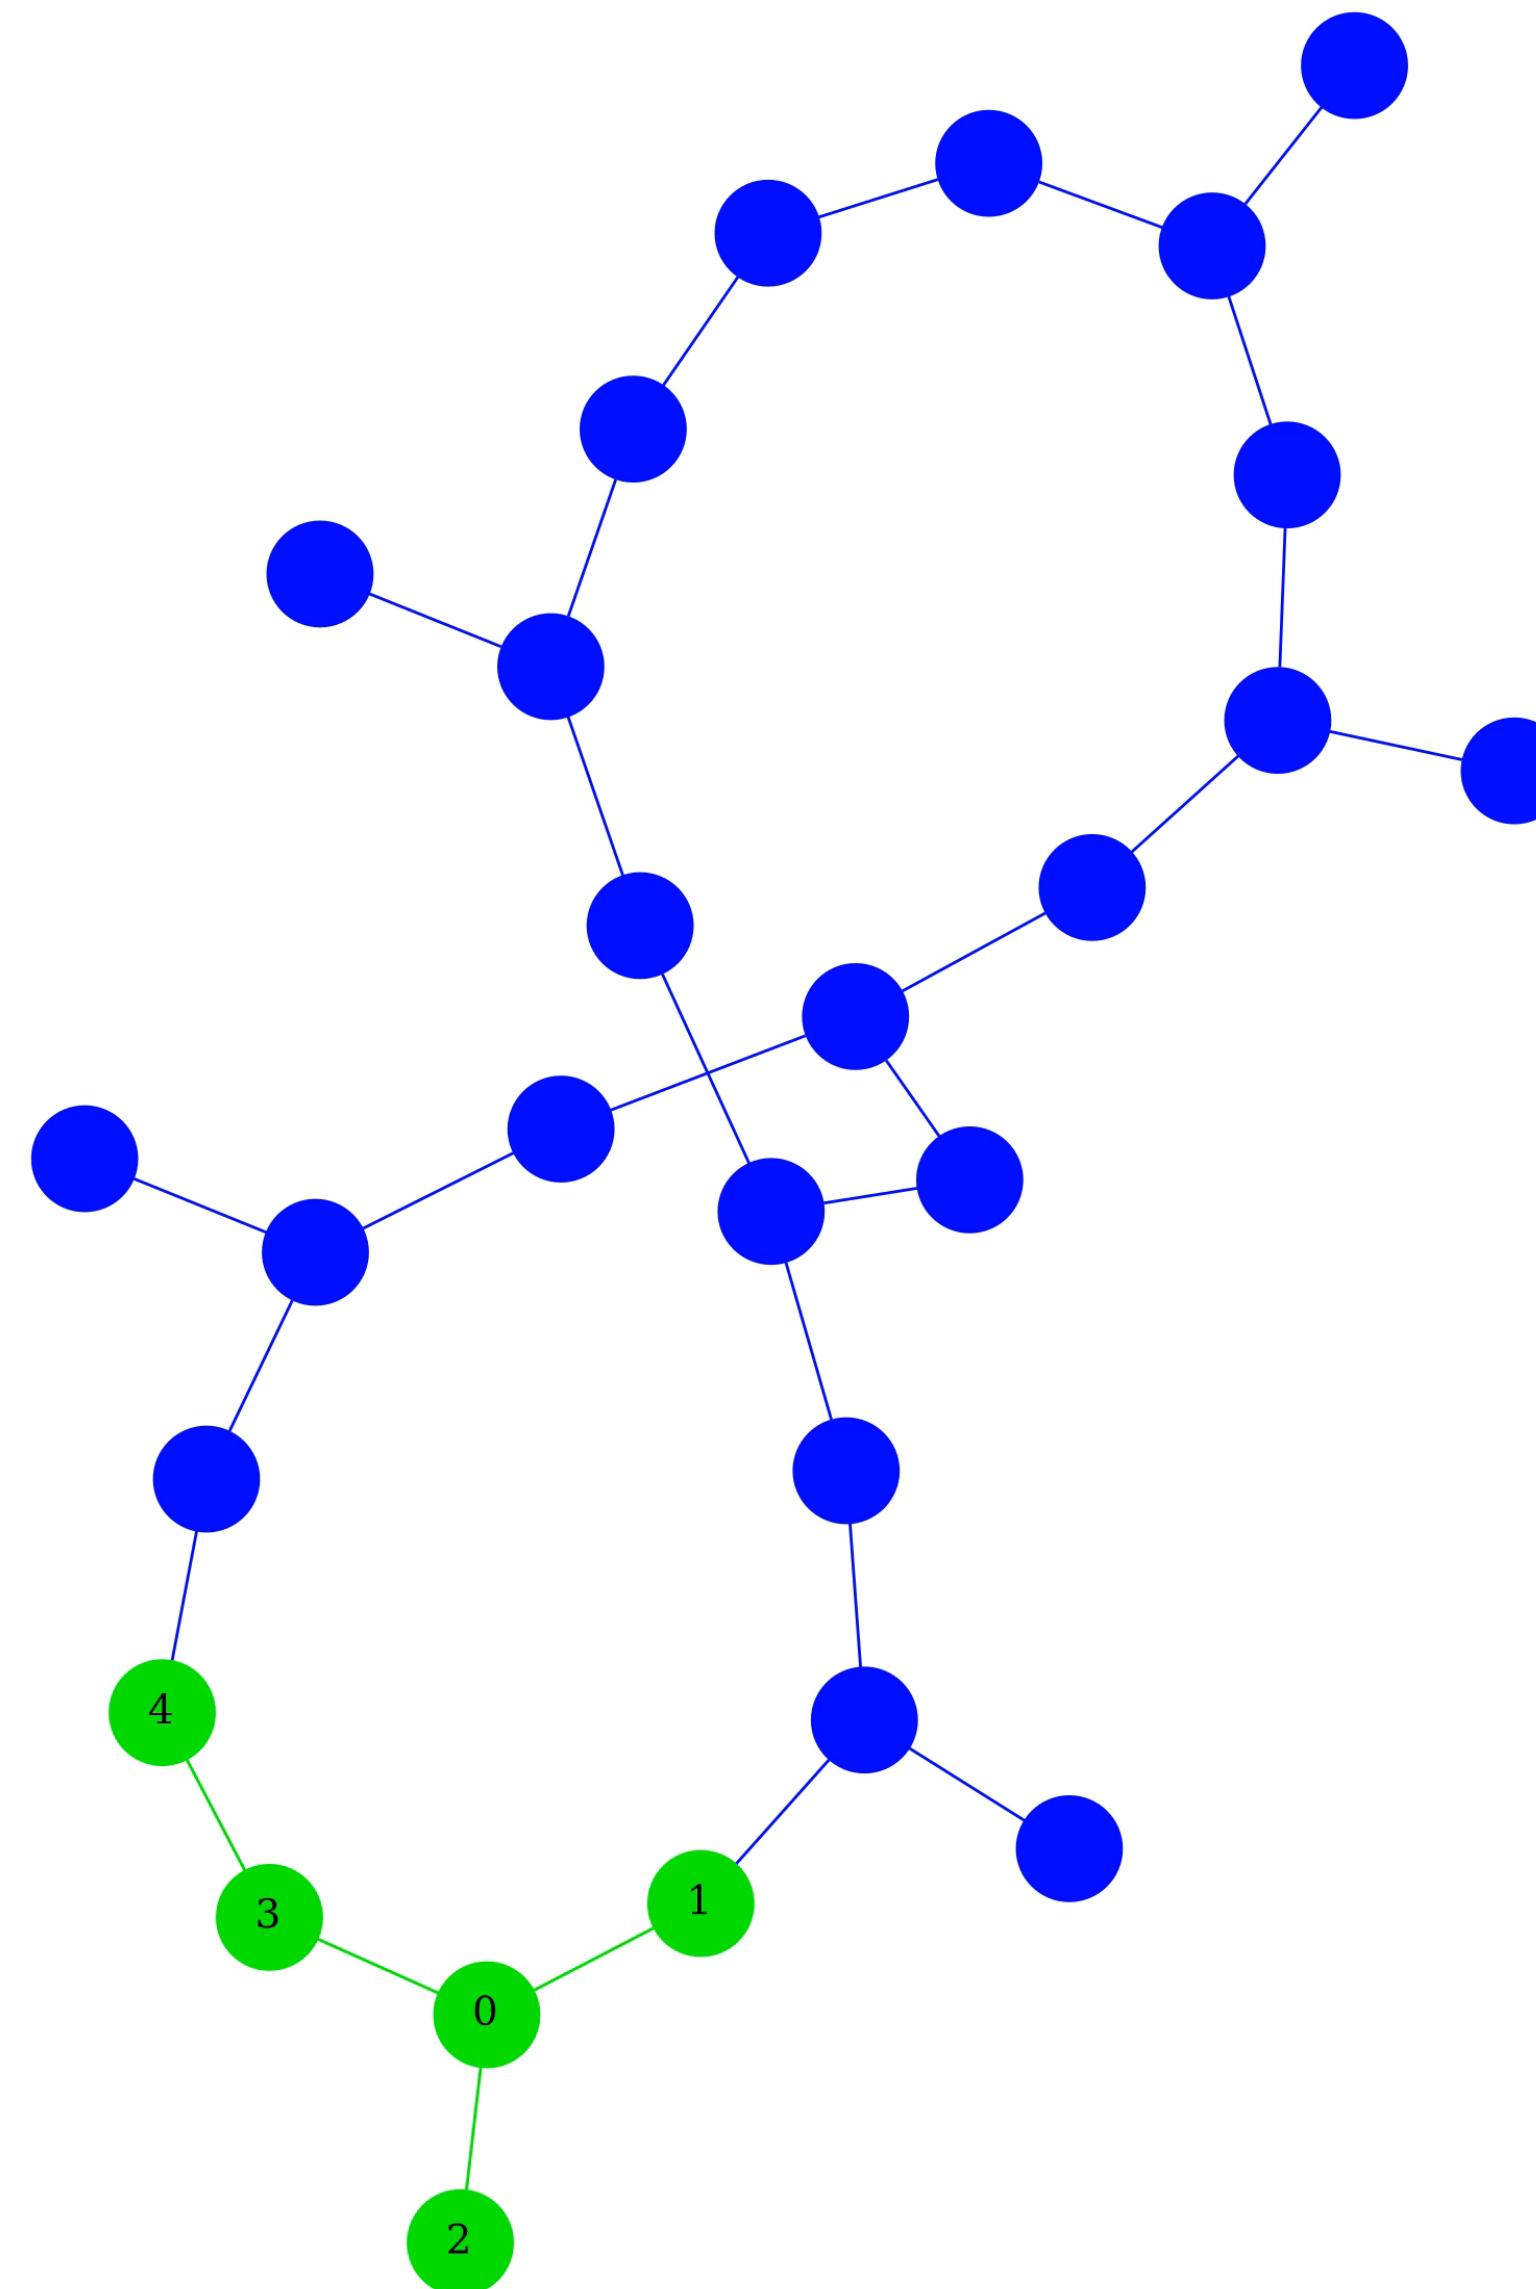


VF2PostLayout

Just as with VF2Layout this pass creates a graph from the 2 qubit interactions in the circuit and tries to find an isomorphic subgraph of the connectivity graph of the target. The difference is by running after routing we know that there is at least one isomorphic subgraph. This pass then searches for other layouts and picks the one found with the lowest predicted error rate.

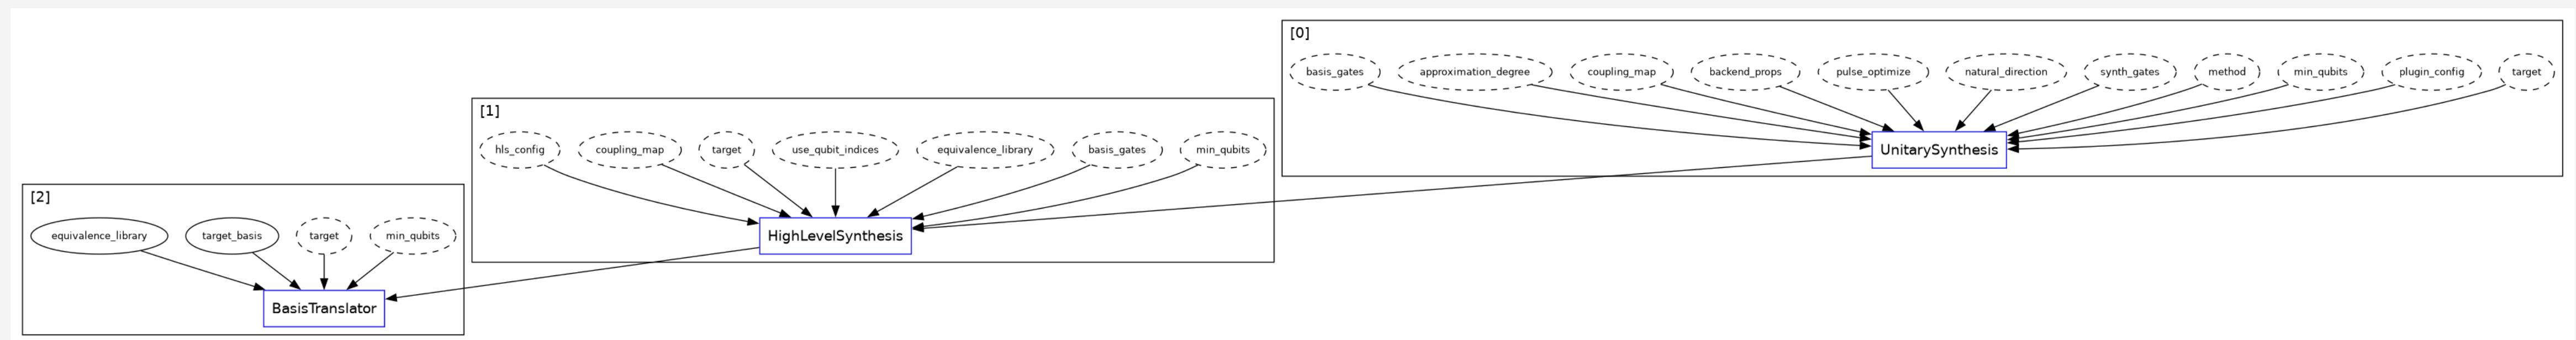
<https://journals.aps.org/prxquantum/abstract/10.1103/PRXQuantum.4.010327>

<https://docs.quantum.ibm.com/api/qiskit/qiskit.transpiler.passes.VF2PostLayout>



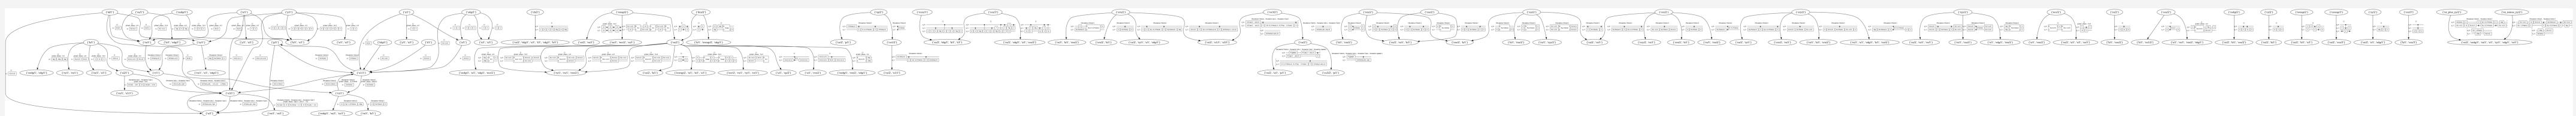
Translation stage

- This stage takes the routed circuit which meets the connectivity constraints of the target and is responsible for making sure that all the operations are supported on the target
- This stage runs the same passes we used in the first phase of the init stage but for all operations instead of just those with ≥ 3 qubits



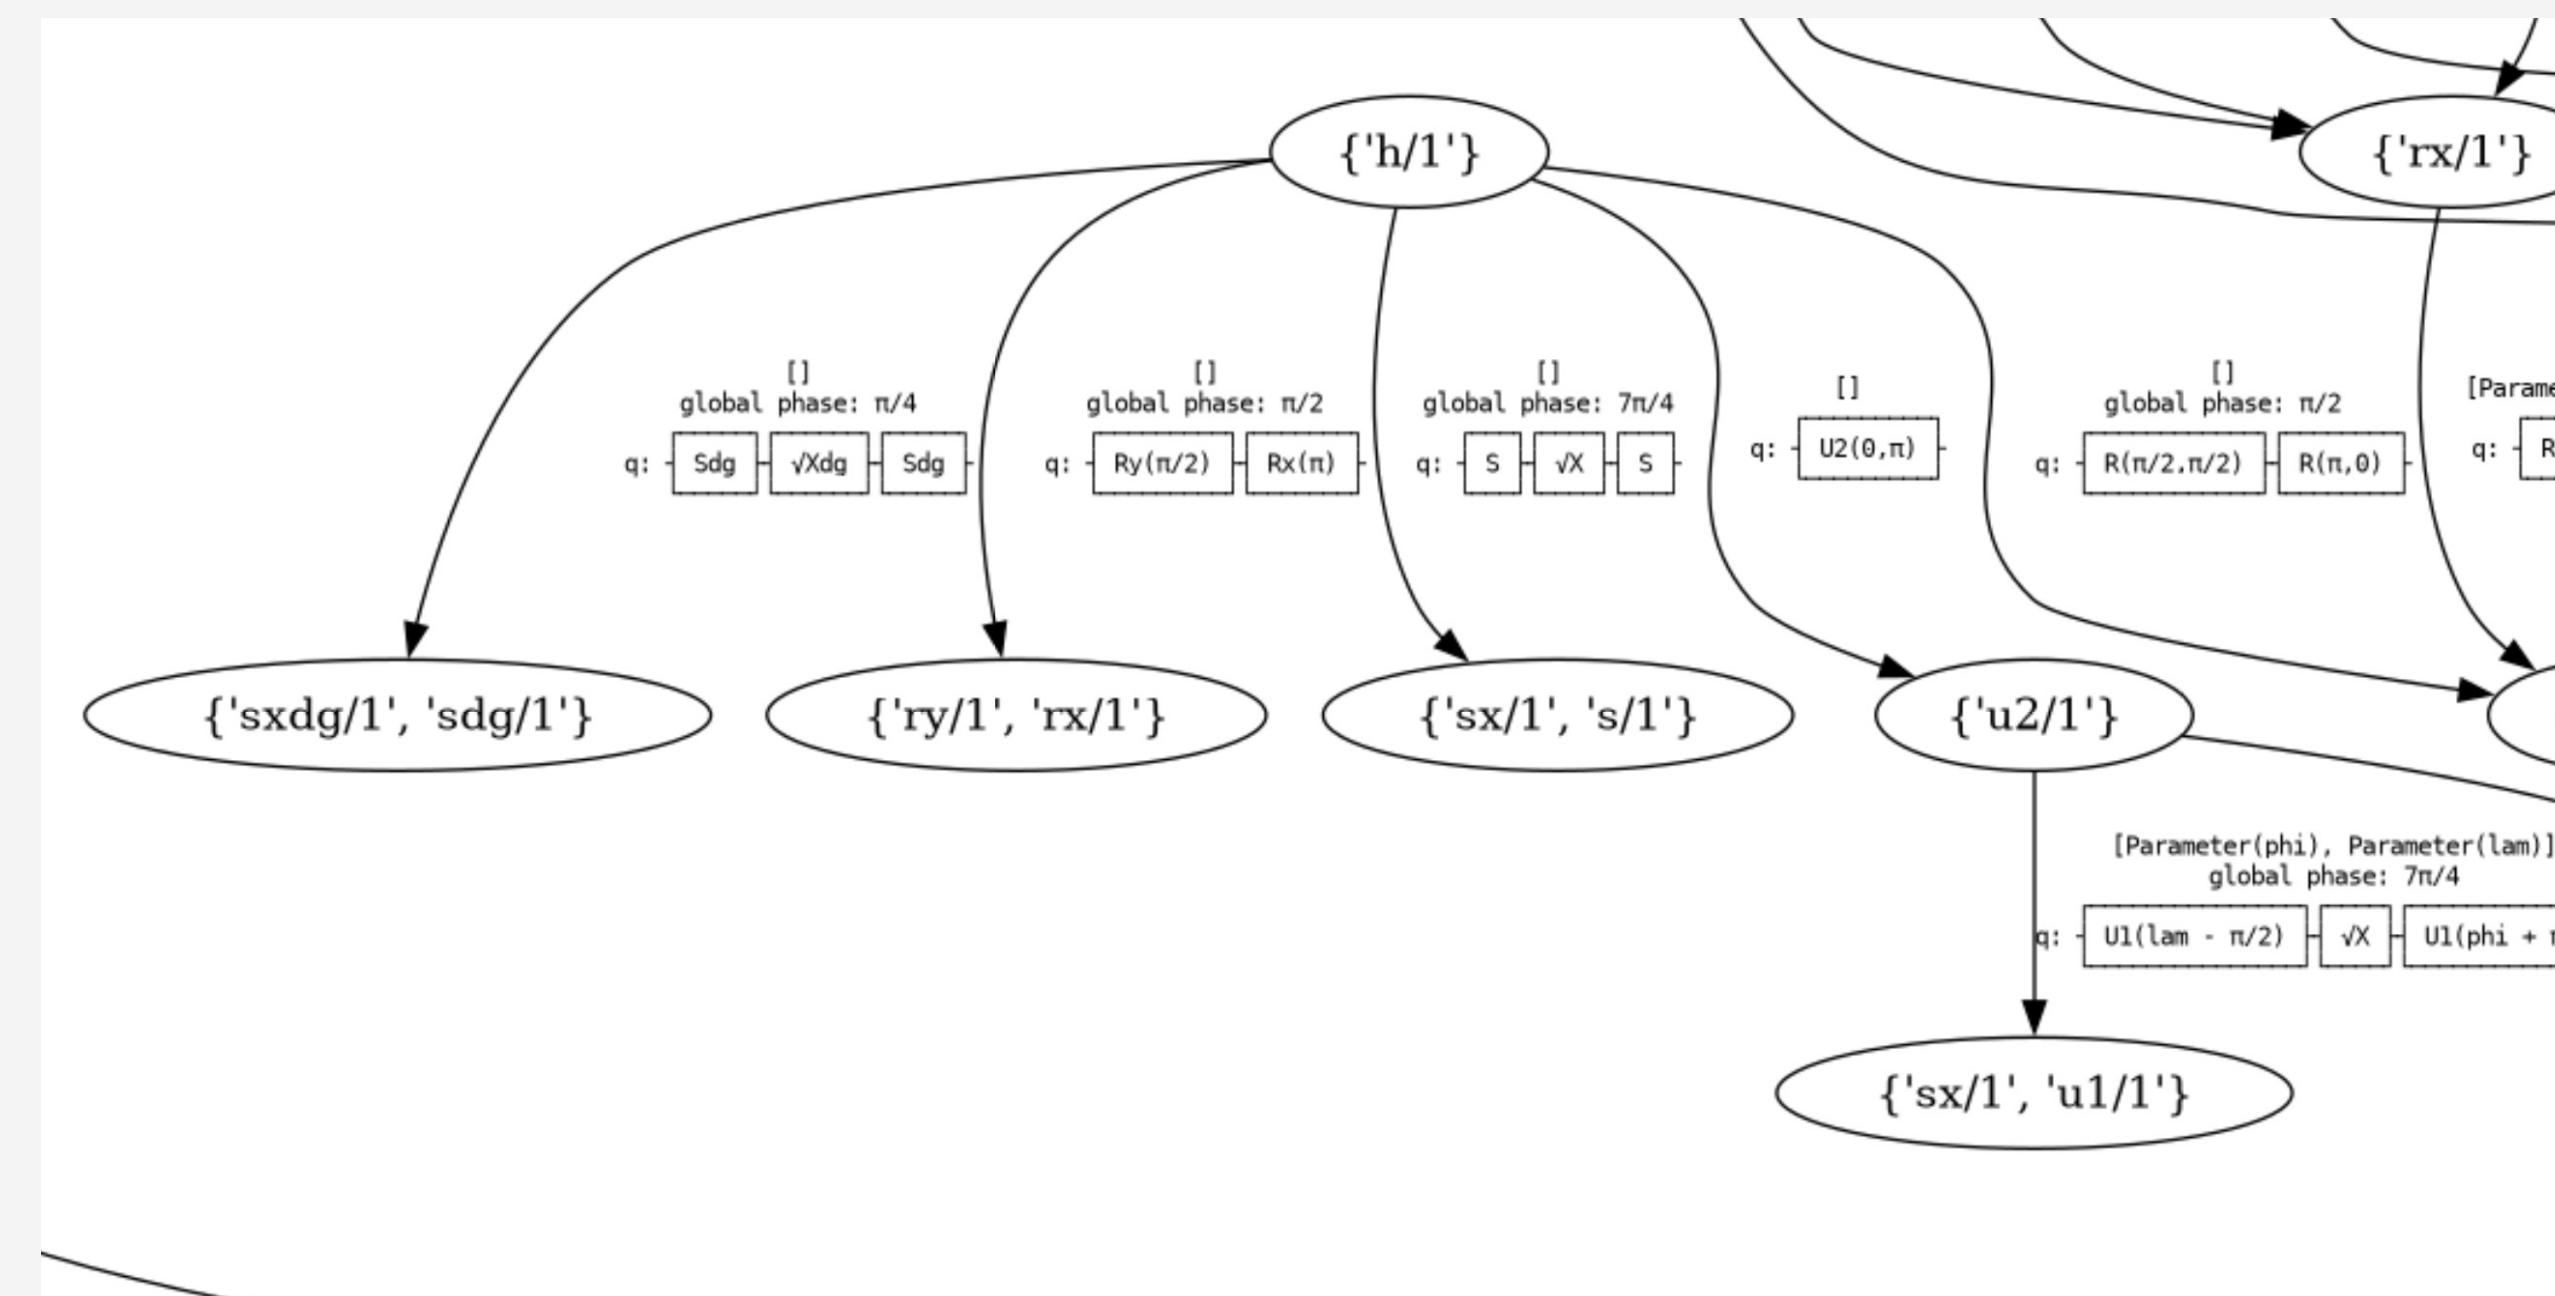
Basis Translator

- The basis translator uses an equivalence library that builds a graph of equivalent circuits to a gate
- Dijkstra's algorithm is used to find a path from a given gate to gates in the target
- The gates in the circuits are recursively substituted along the path to build an equivalent circuit that conforms to the target
- Each gate is then replaced in the DAG by that equivalent circuit

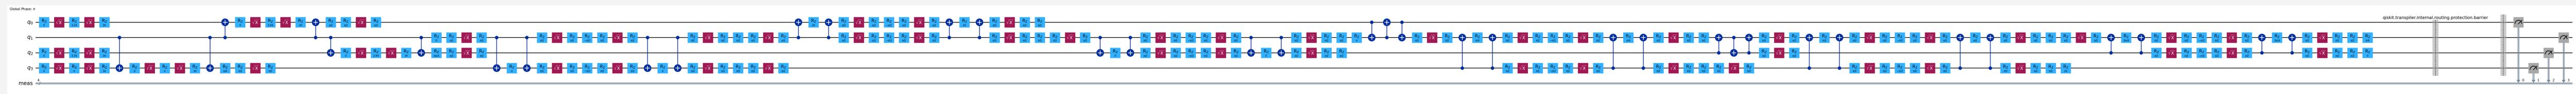
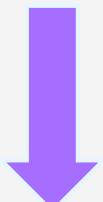
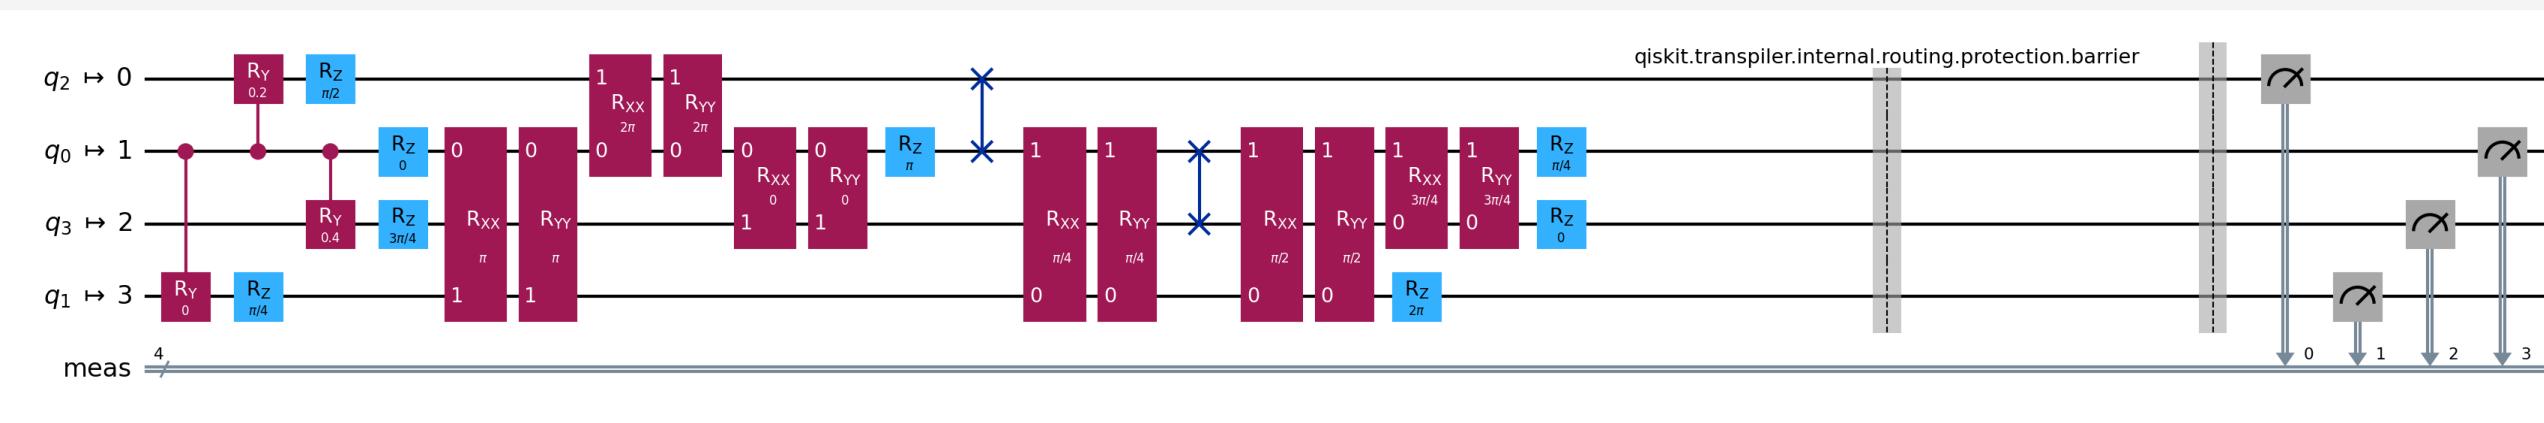


Basis Translator

- The basis translator uses an equivalence library that builds a graph of equivalent circuits to a gate
- Dijkstra's algorithm is used to find a path from a given gate to gates in the target
- The gates in the circuits are recursively substituted along the path to build an equivalent circuit that conforms to the target
- Each gate is then replaced in the DAG by that equivalent circuit

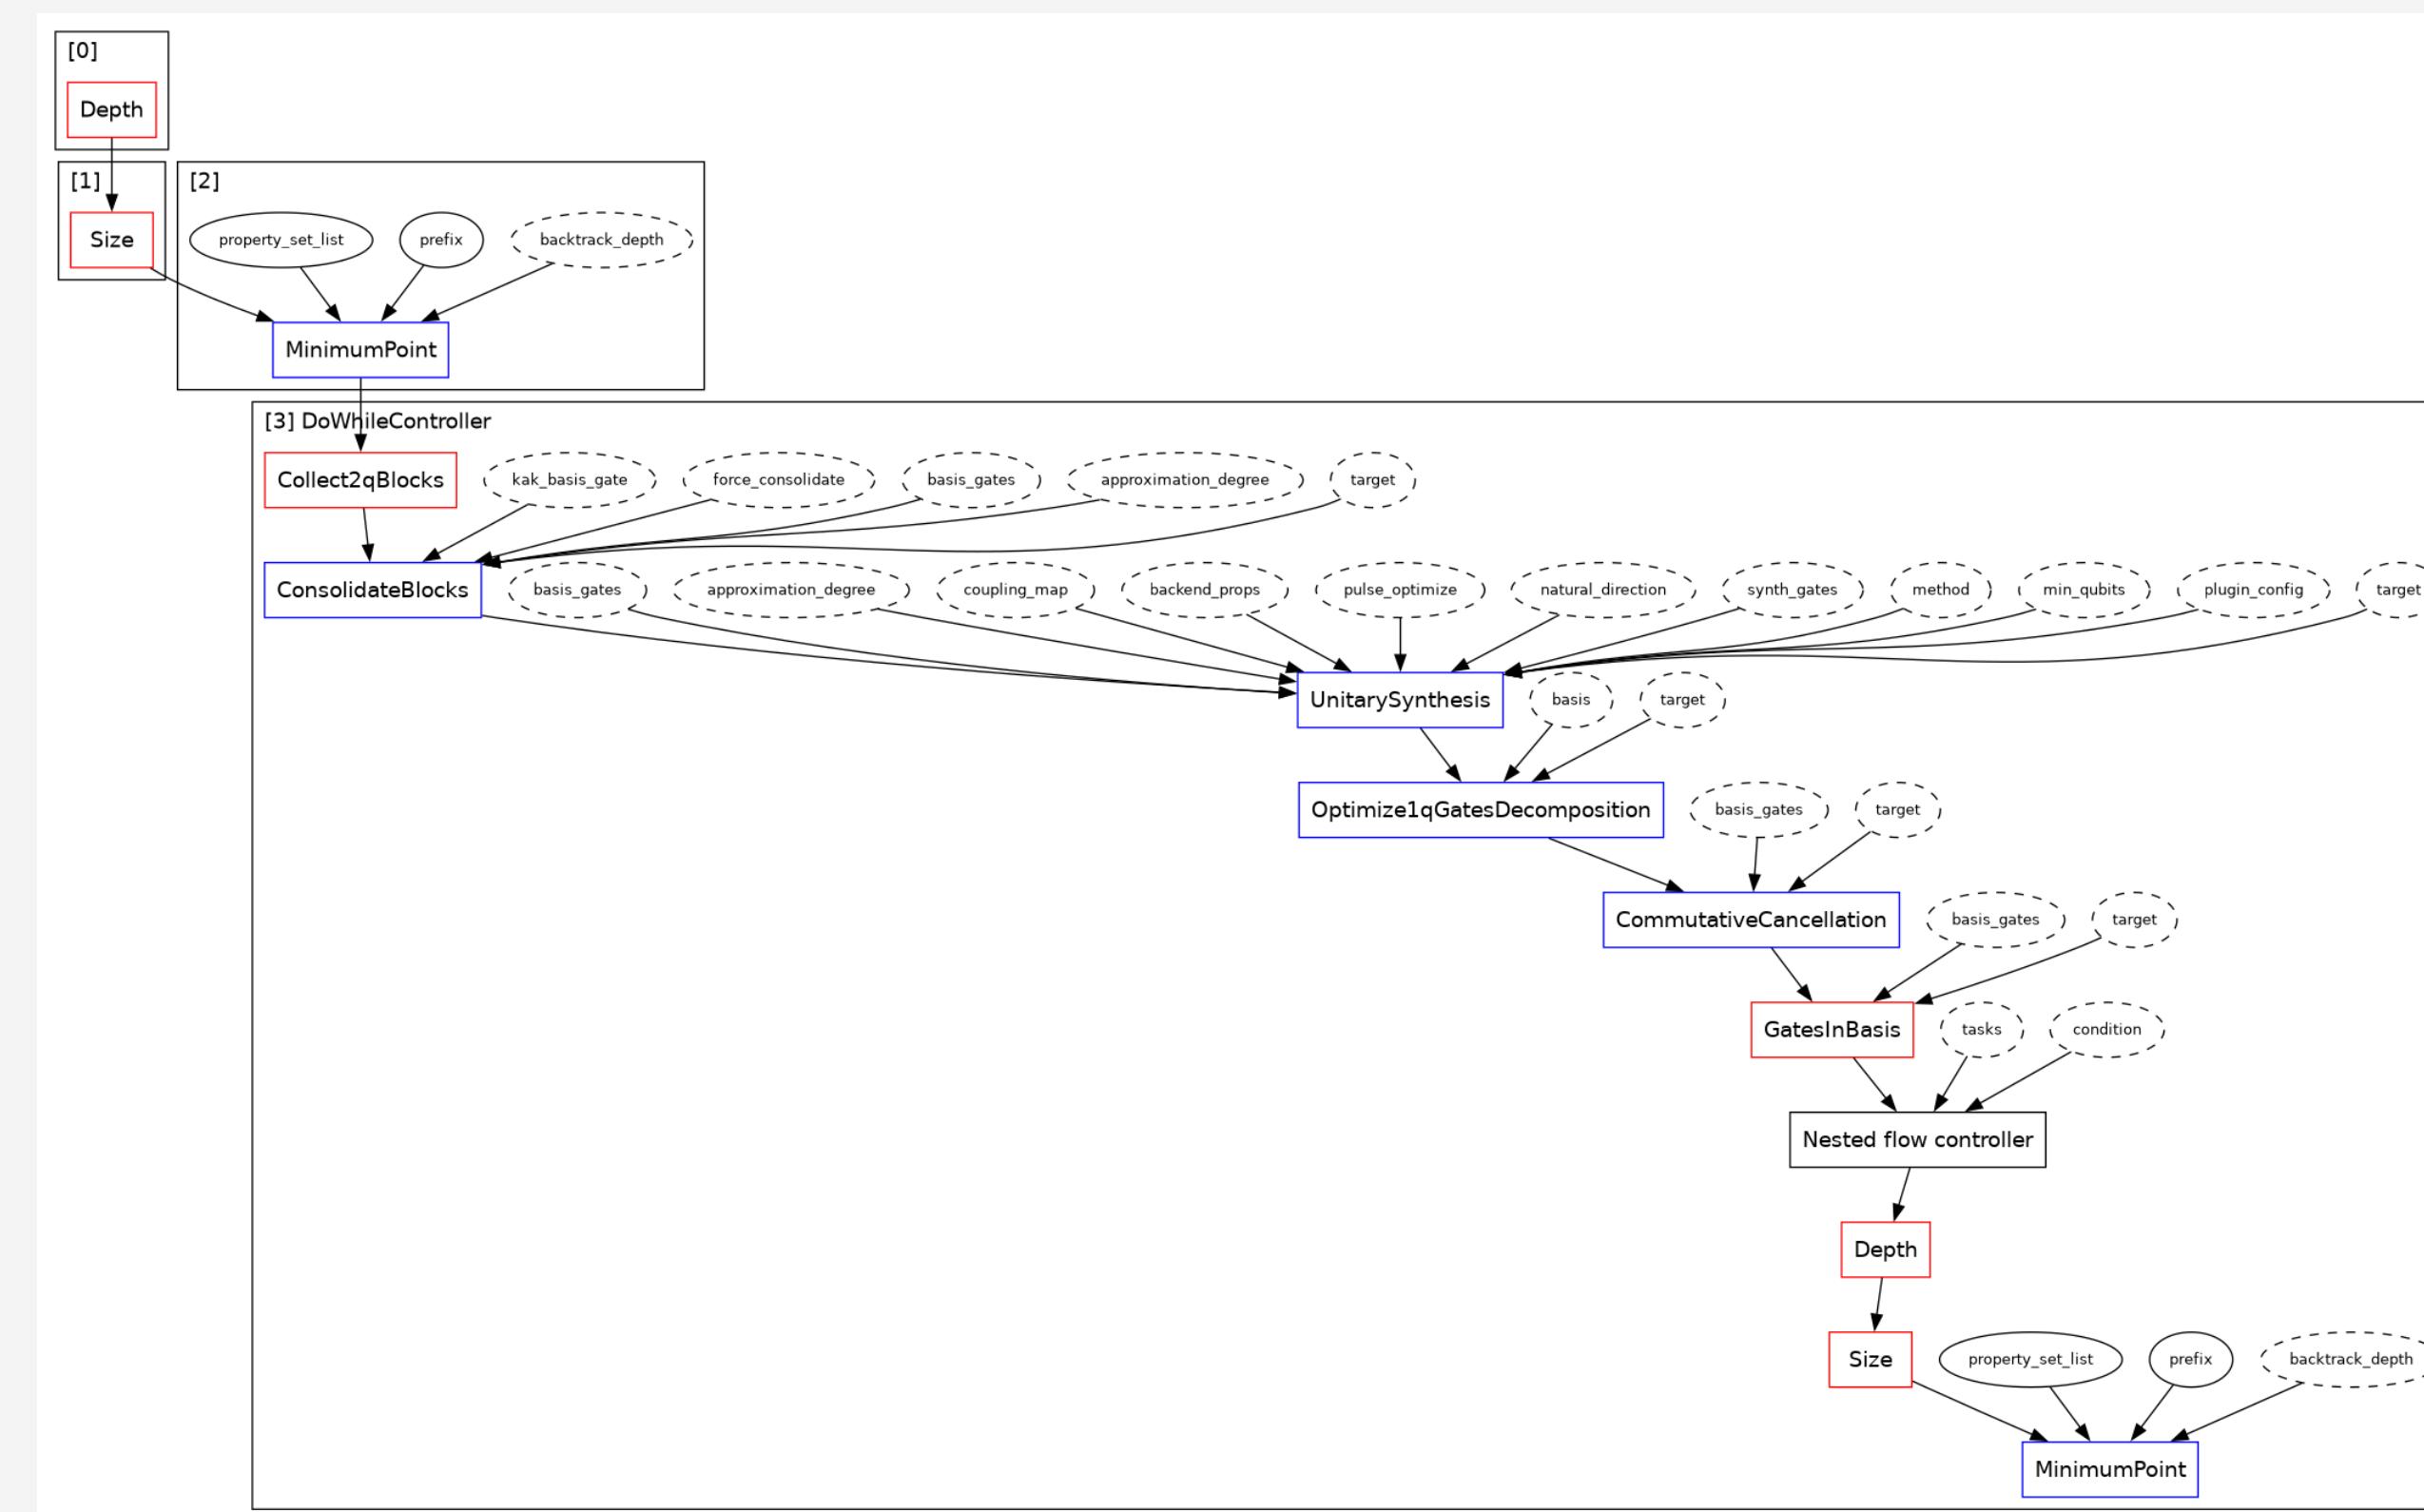


Basis Translator



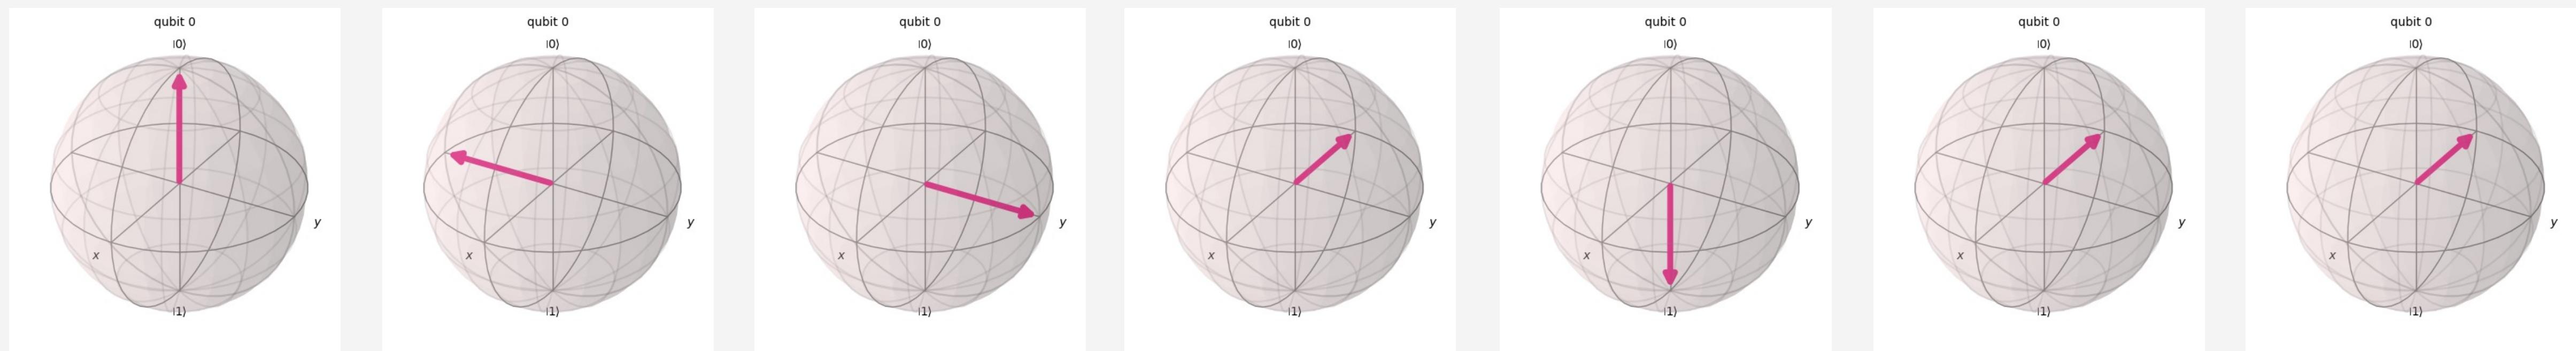
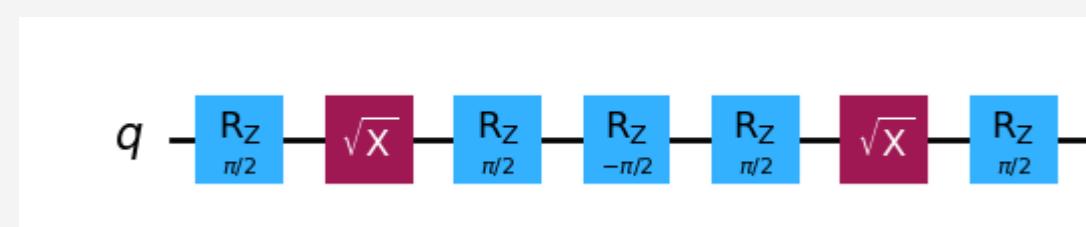
Optimization stage

- This stage takes the mapped circuit and optimizes it to reduce unnecessary operations
- This stage runs in a loop to repeatedly run the optimization passes until a minimum point in depth and size is found

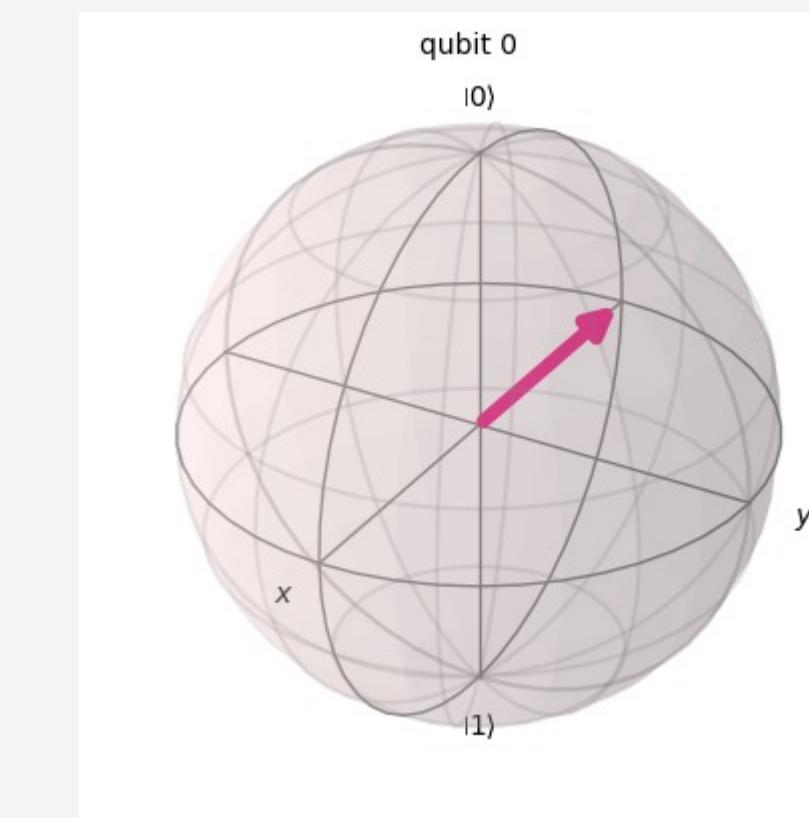
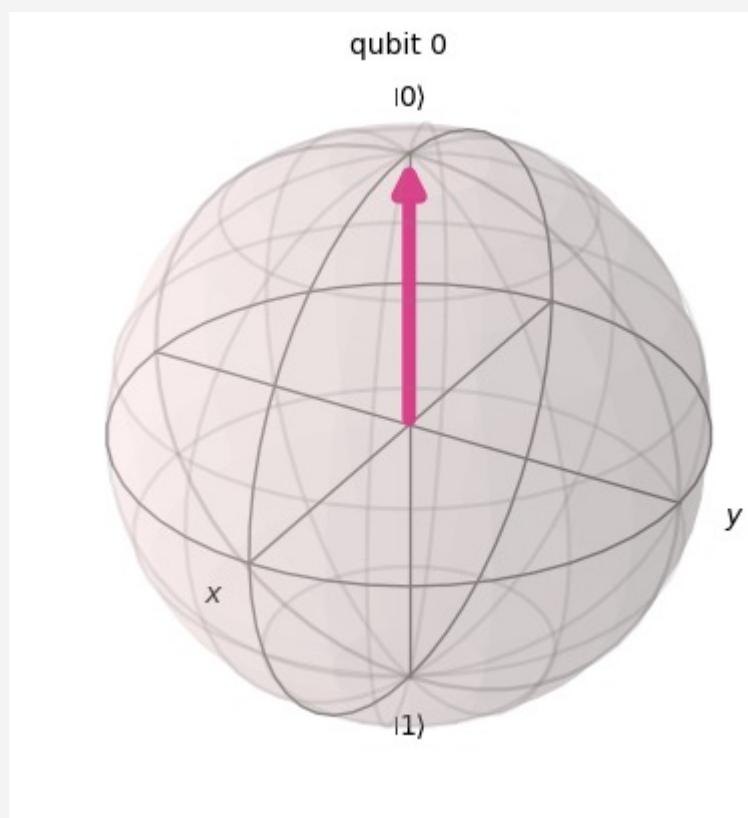
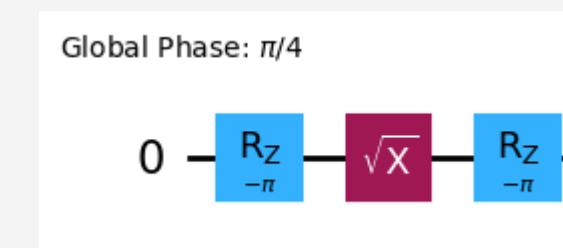


1 Qubit Unitary Peephole Optimization

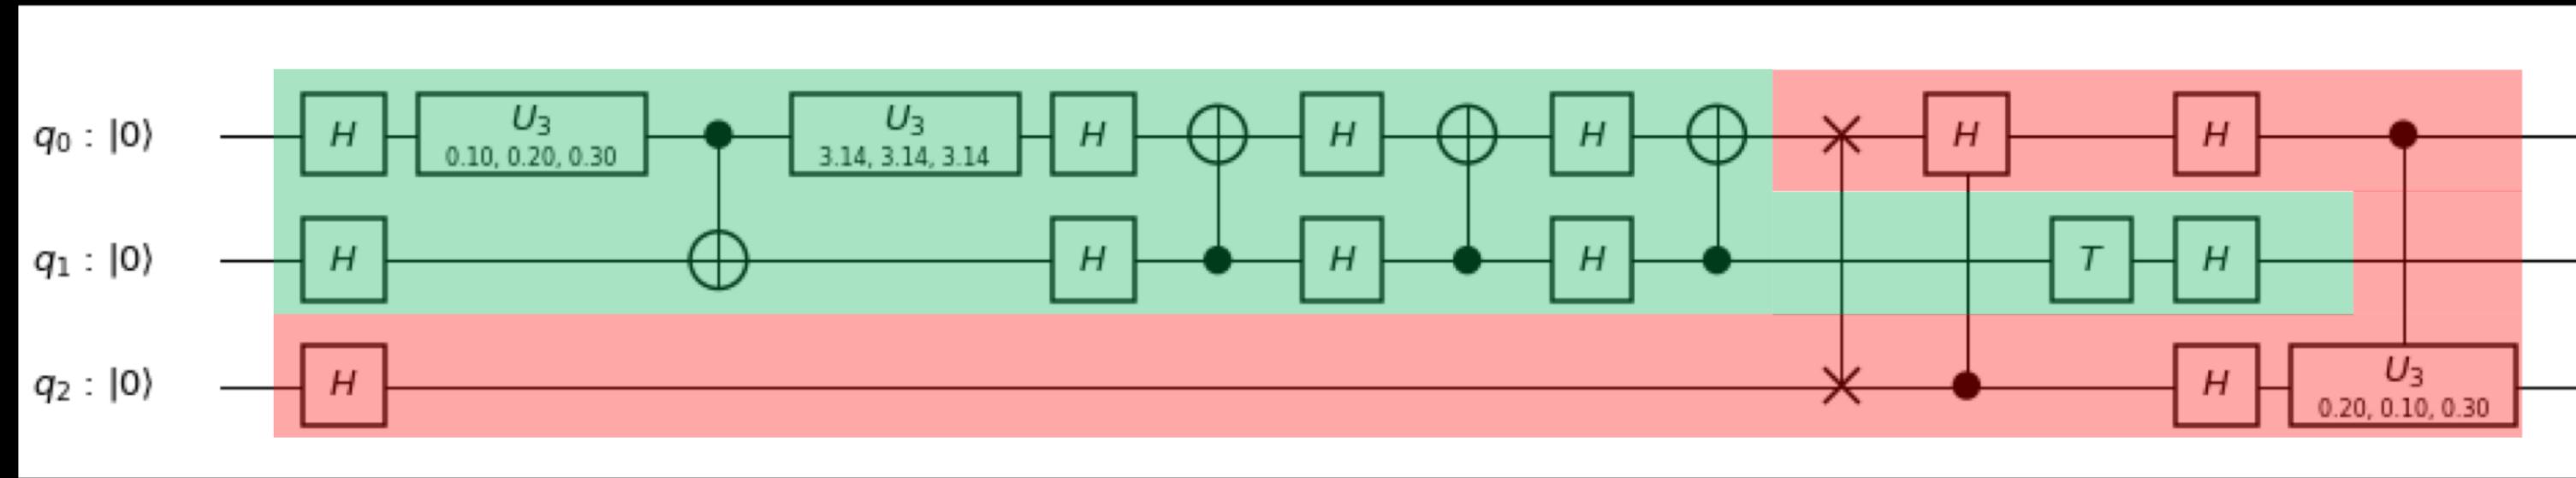
- This pass looks for sequences of single qubit gates



1 Qubit Unitary Peephole Optimization

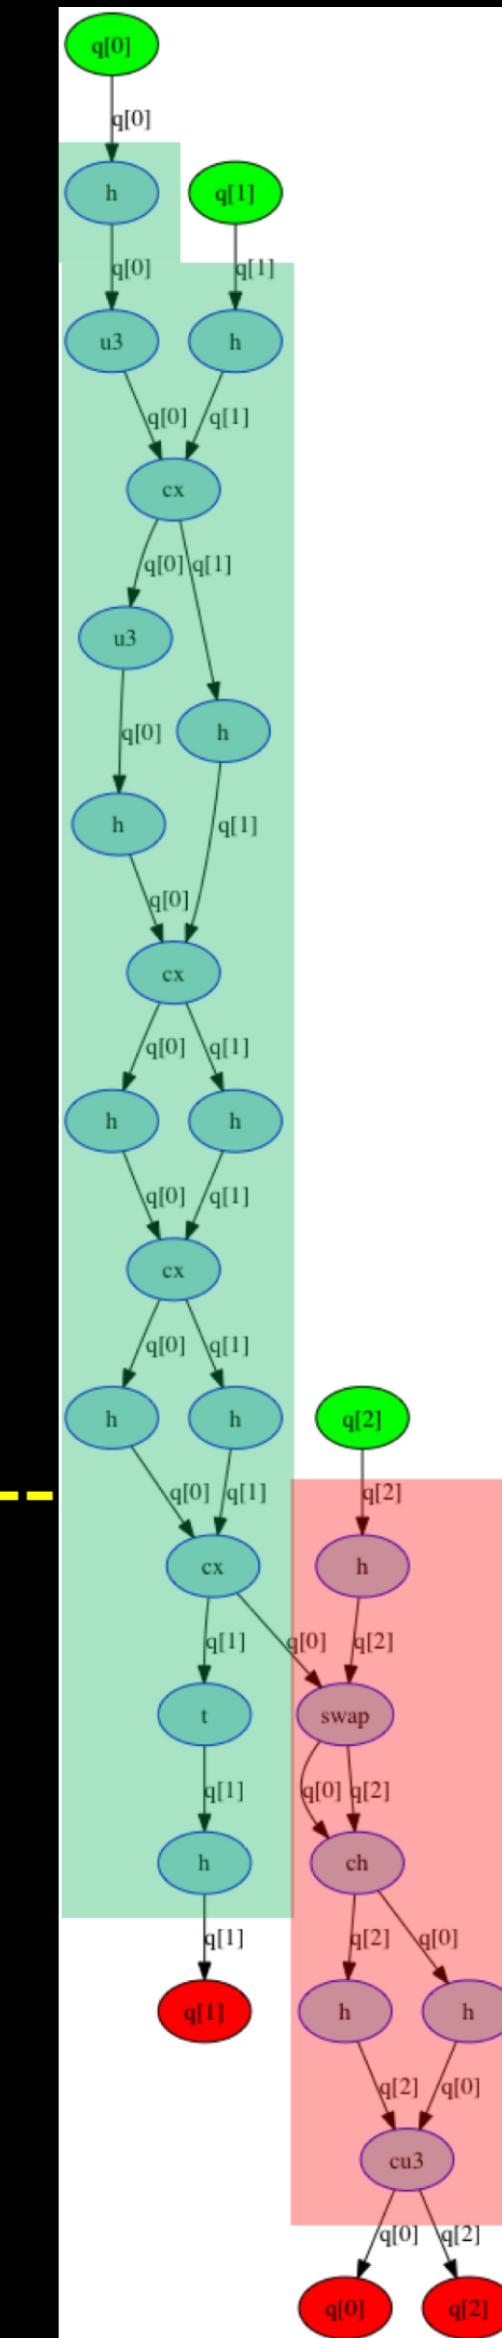


2-qubit Block Collection

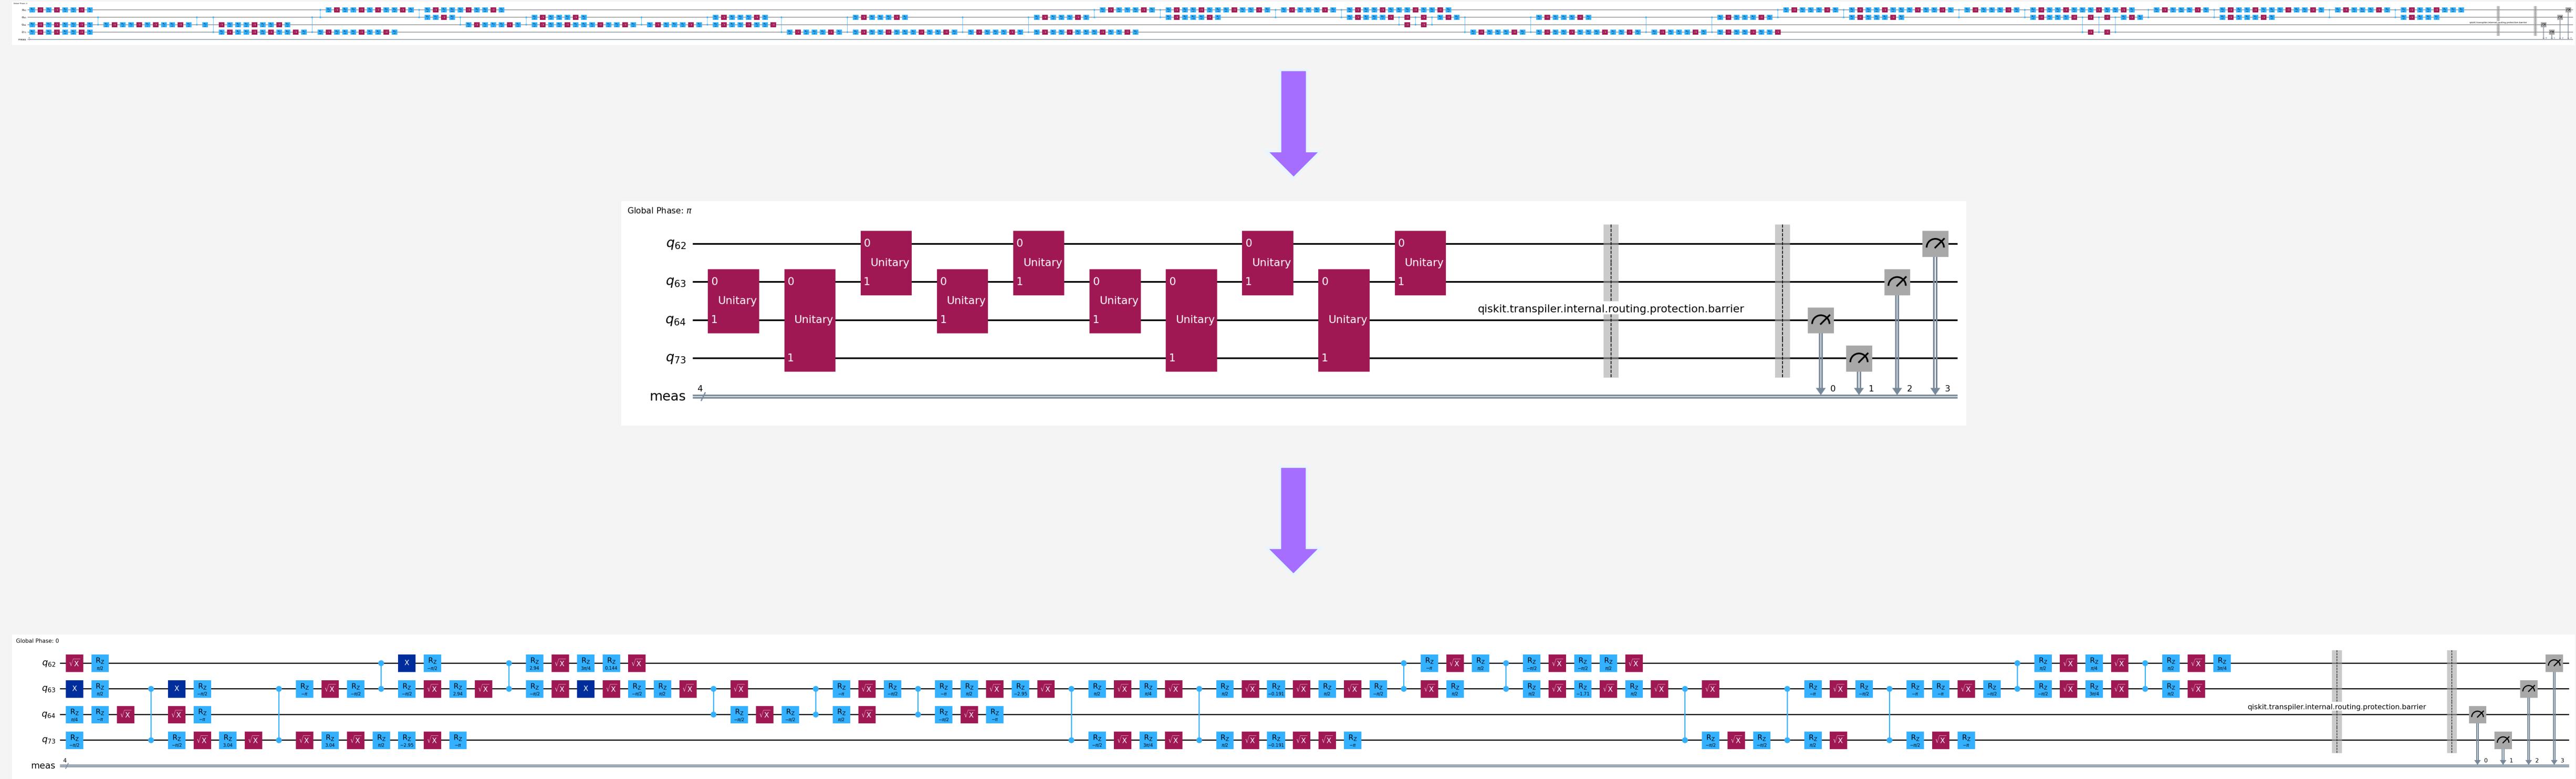


Efficient graph traversal on the DAG.

For each CNOT, collect ancestors and predecessors until a branch.

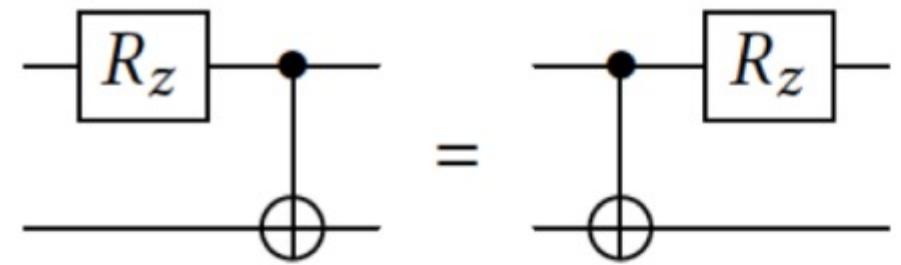


2 Qubit Unitary Peephole Optimization

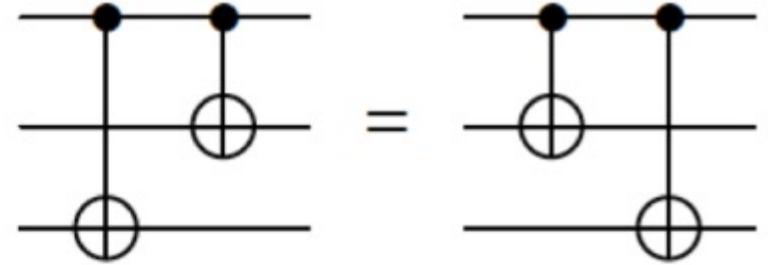


Commutative Cancellation

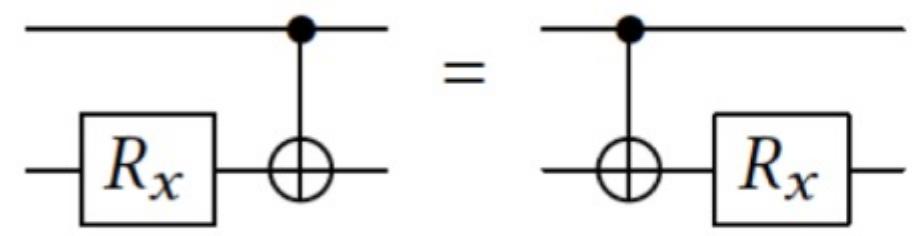
Examples of commuting gates:



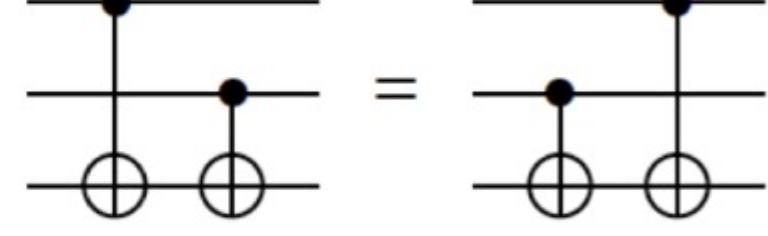
(a) R_z -control



(b) Control-control



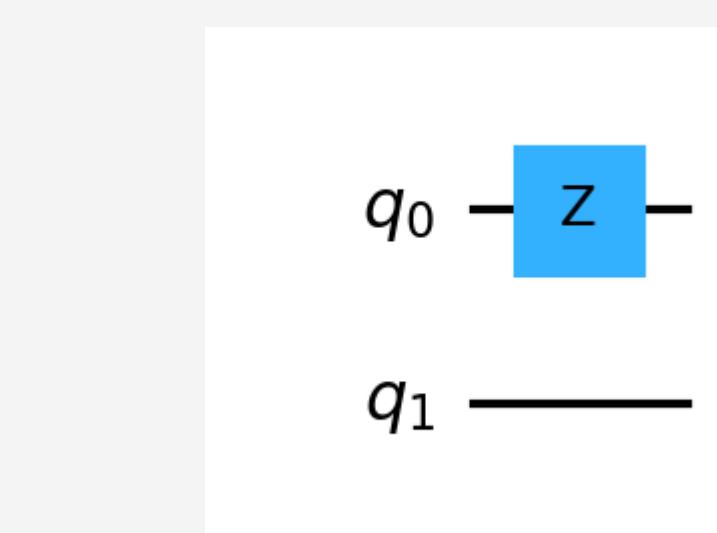
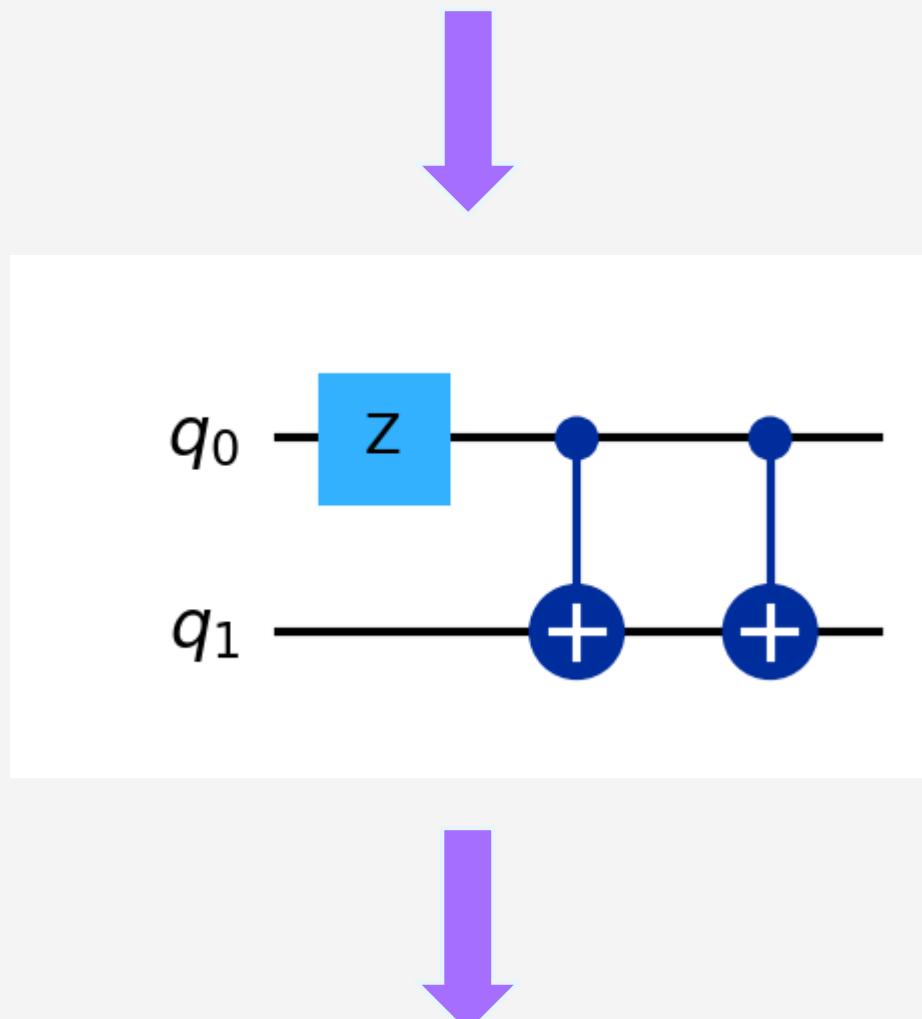
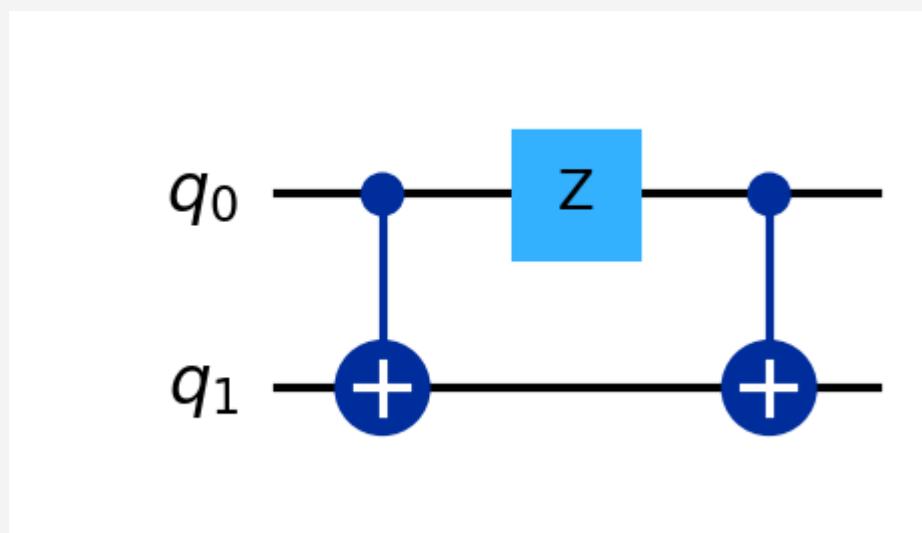
(c) R_x -target



(d) Target-target

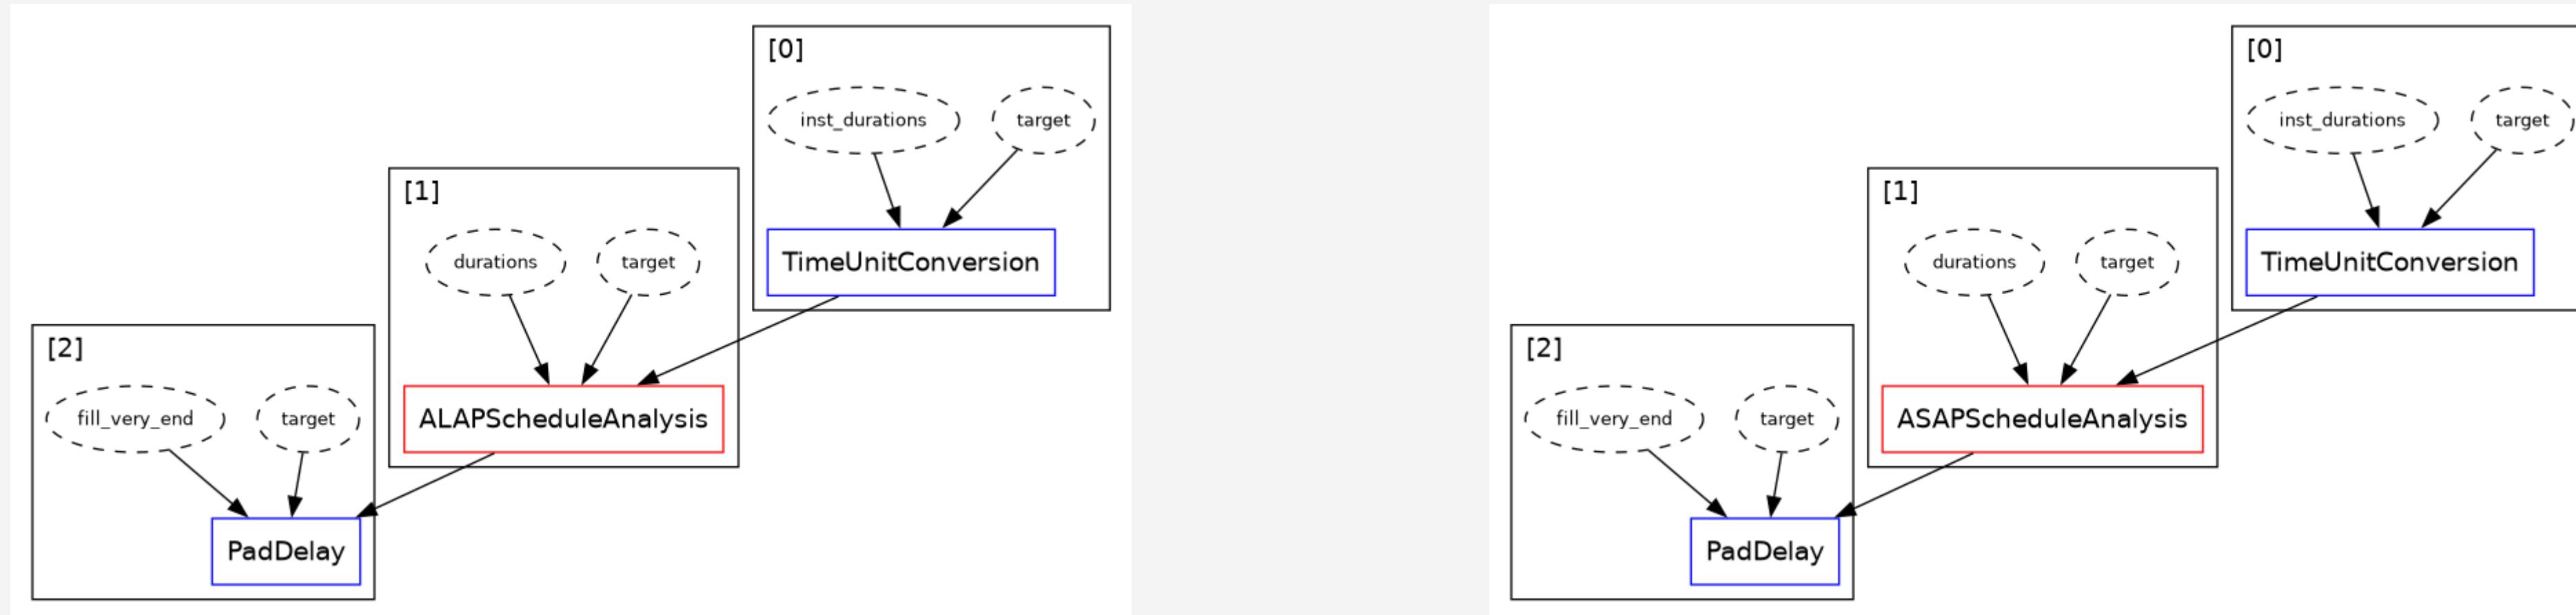
Two consecutive gates are commutative \Leftrightarrow

They can be exchanged without changing what they compute.



Scheduling stage

- This stage takes the mapped and optimized circuit and analyzes the runtime of the operations in the circuit and identify idle periods on each qubit
- It will then insert explicit operations to account for all the time in the circuit



Scheduling stage

- This stage takes the mapped and optimized circuit and analyzes the runtime of the operations in the circuit and identify idle periods on each qubit
- It will then insert explicit operations to account for all the time in the circuit

